



# Fundamentals of Data Engineering

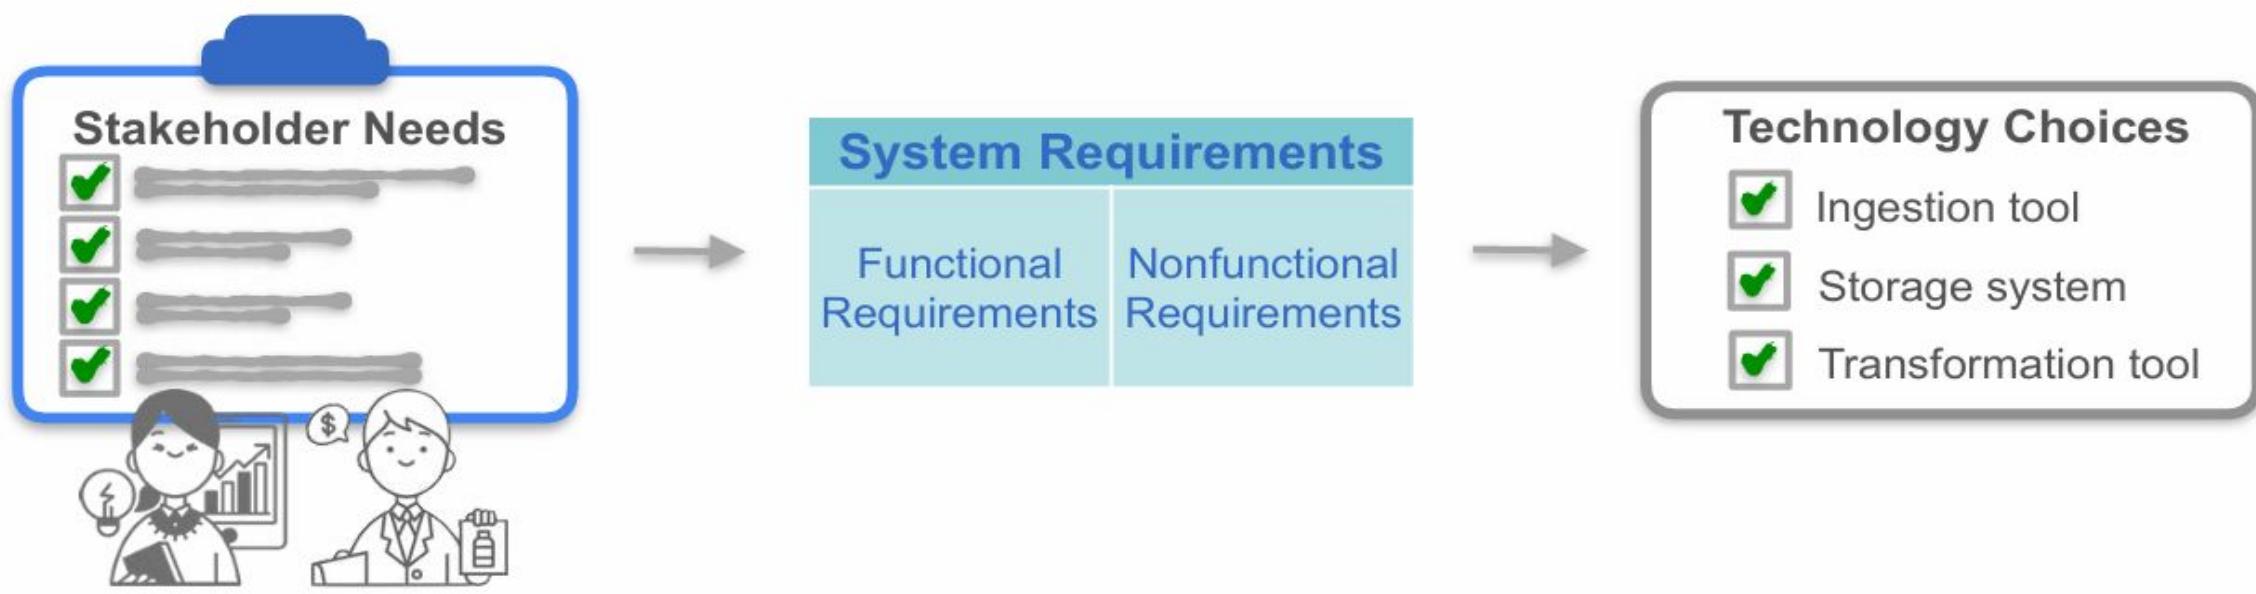
# Scenario

Every stakeholder interacts with the data ecosystem differently:  
 Business executives want results and trends. Data scientists need structured data to build models. Developers need usable interfaces (APIs). Customers want value and privacy. You, the data engineer, ensure the whole data flow works efficiently and ethically.



**Data Engineer**

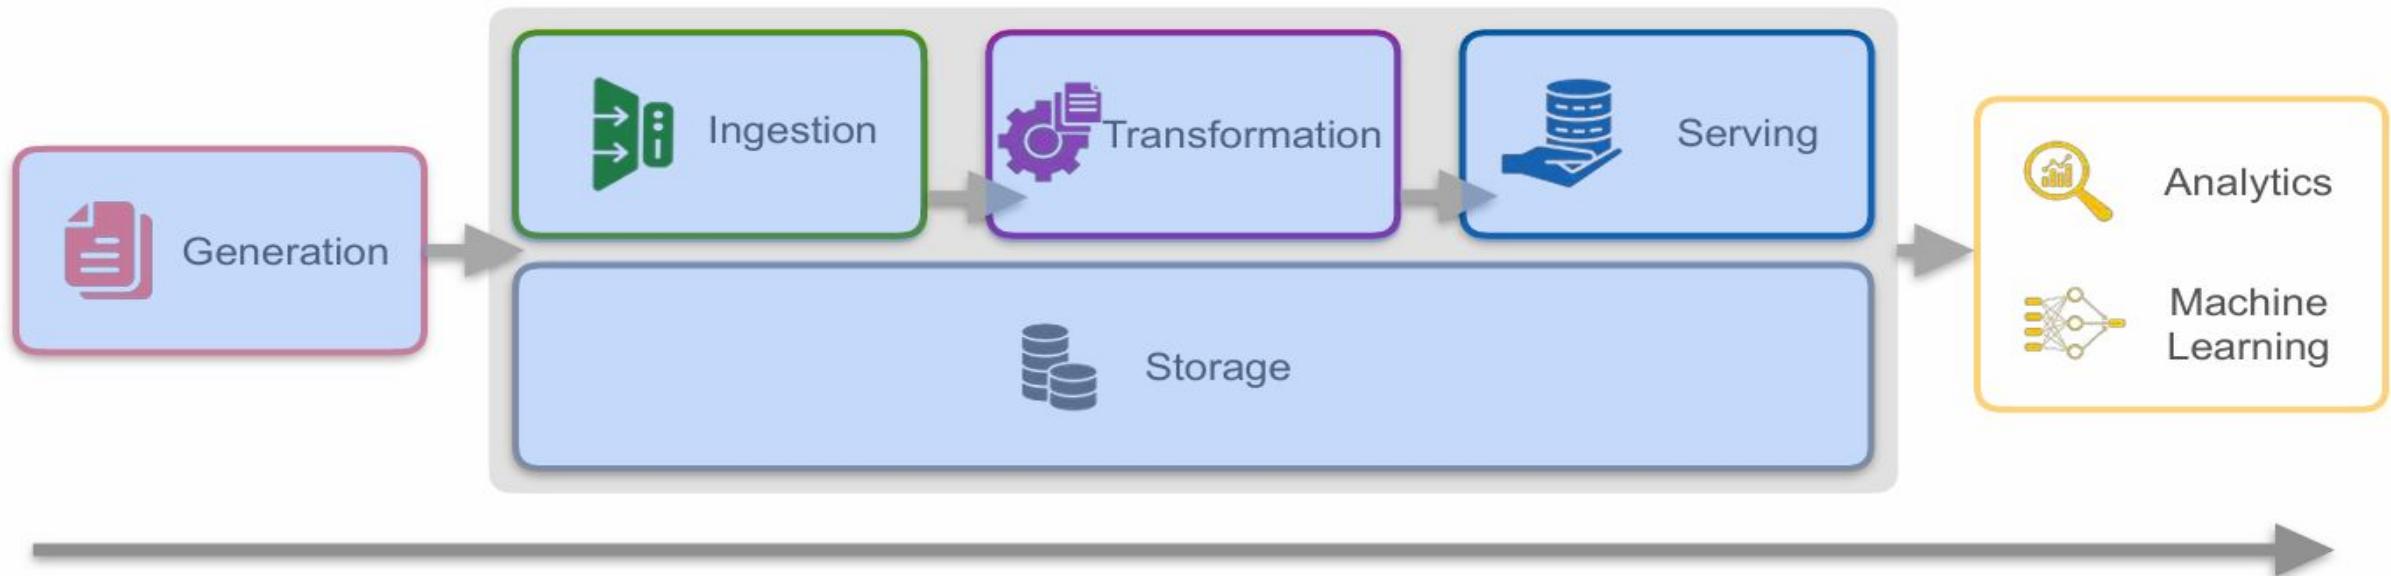
**Wasting time & resources!**



# Module-4

1. Storage Systems,
2. Cloud Storage Options: Block, Object and File storage,
3. Storage Tiers, Distributed Storage Systems,
4. Row vs Column Storage,
5. Graph Databases,
6. Vector Databases,
7. Neo4j Graph Database,
8. Storage Abstractions,
9. Data Warehouse,
10. Modern Cloud Data Warehouses,
11. Data Lakes, Data Lakehouse Implementations,

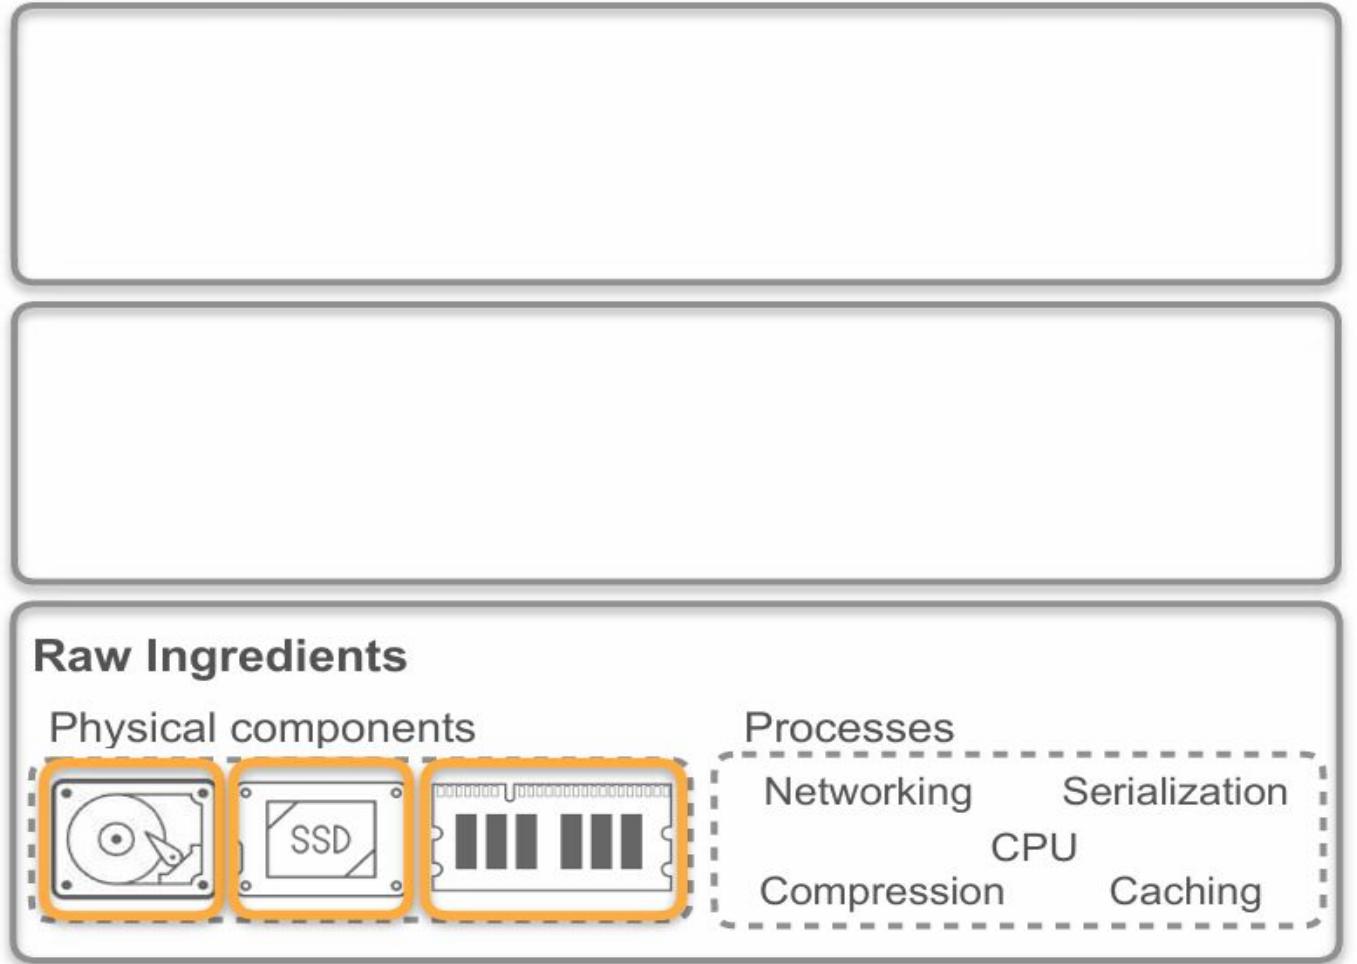
# Storage



Storage solution considerations:

- Data type
- Data format
- Data size
- Access and update pattern

# Storage Hierarchy



# Storage Hierarchy

## Management system:

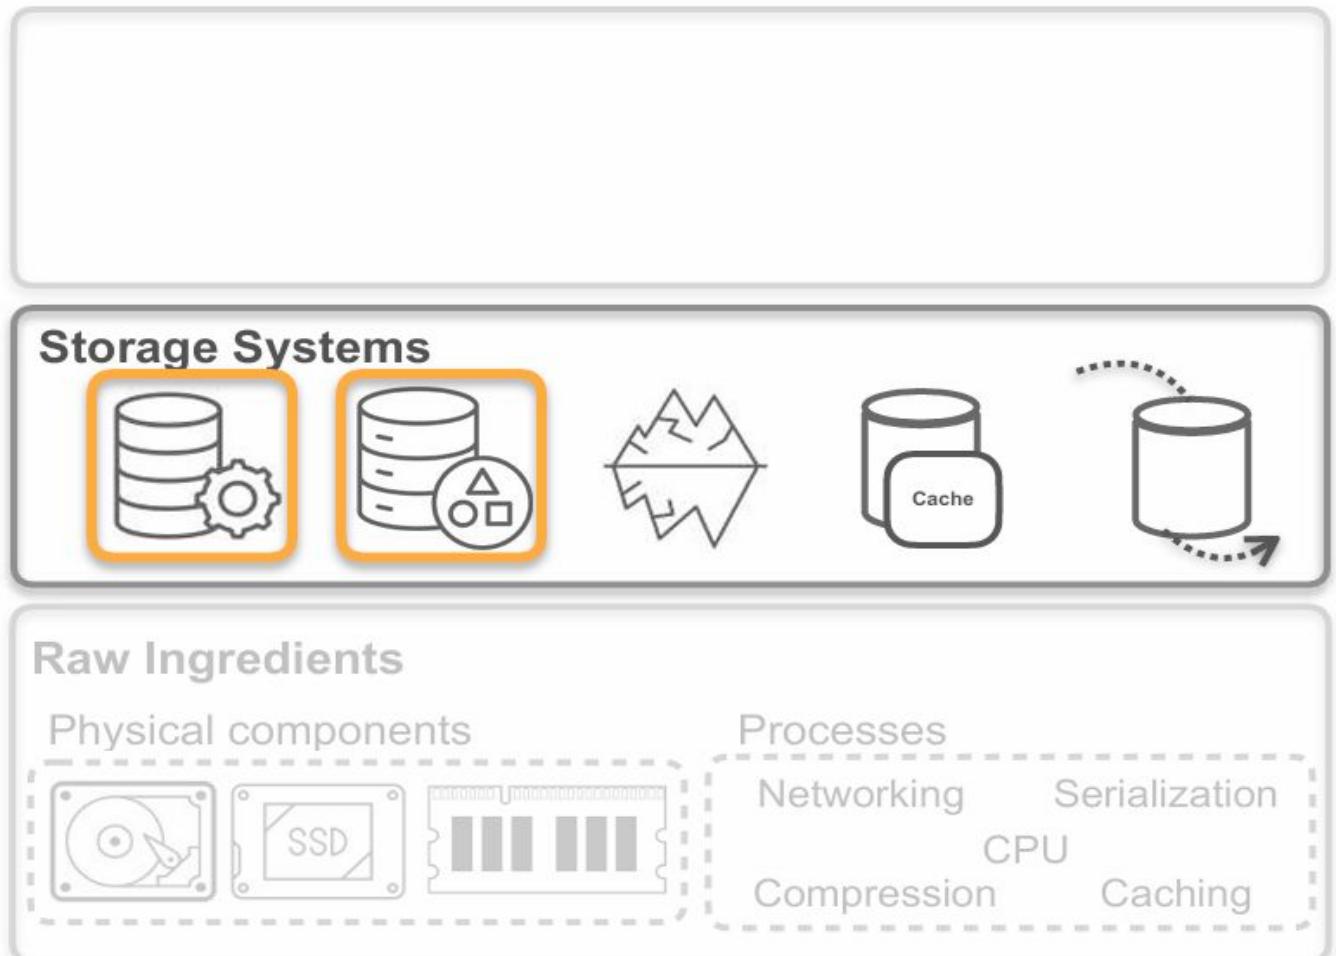
Organizes data in the raw components and allows you to interact with the stored data

### OLTP Systems

Online Transactional Processing Systems  
Focus on performing read and write queries with low latency

### OLAP Systems

Online Analytical Processing Systems  
Focus on applying analytical activities on data (e.g. aggregation, summarization)



# Storage Hierarchy

## Management system:

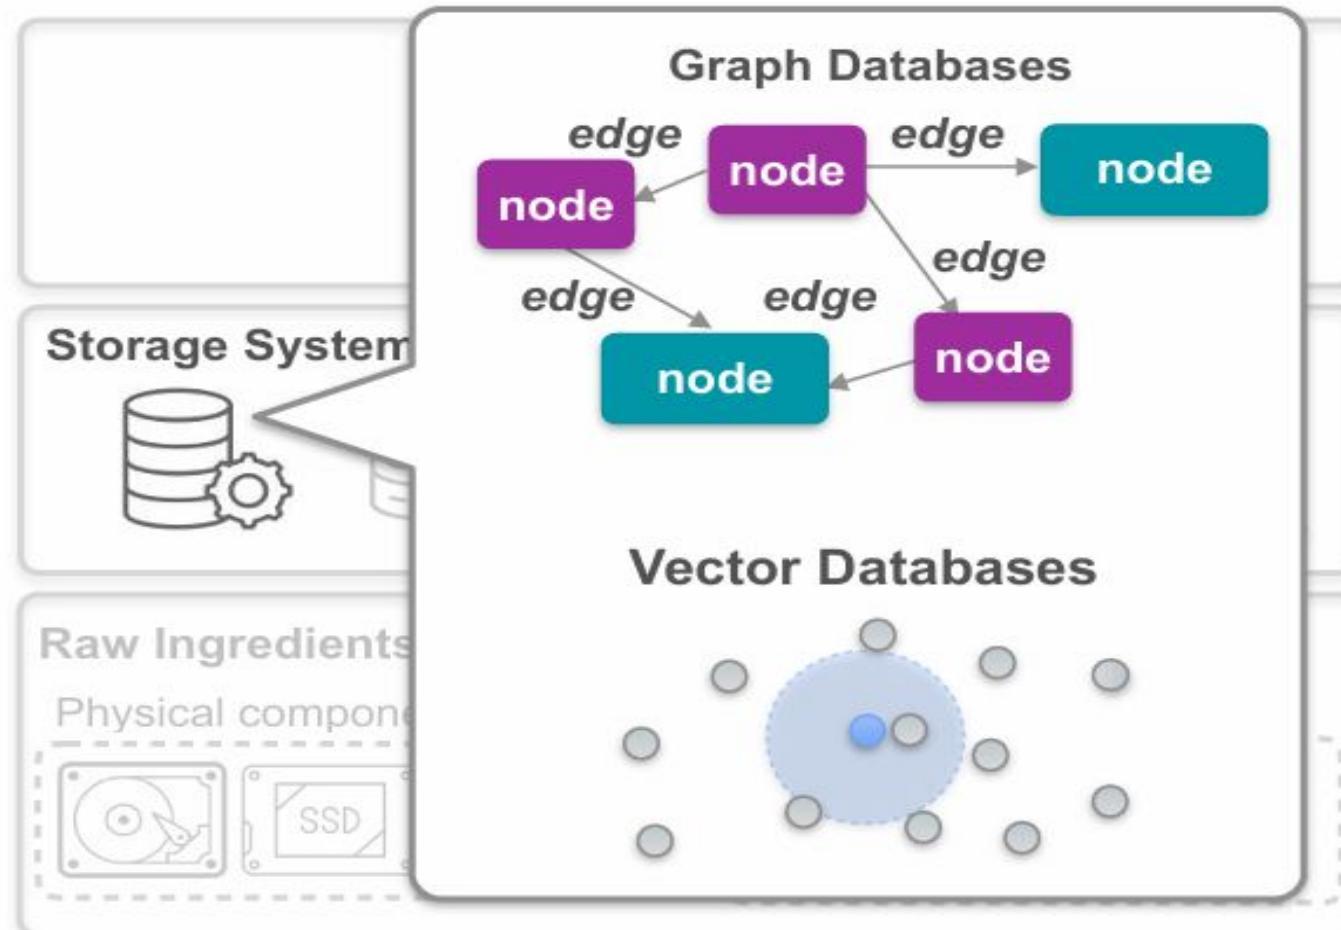
Organizes data in the raw components and allows you to interact with the stored data

### OLTP Systems

Online Transactional Processing Systems  
Focus on performing read and write queries with low latency

### OLAP Systems

Online Analytical Processing Systems  
Focus on applying analytical activities on data (e.g. aggregation, summarization)



# Storage

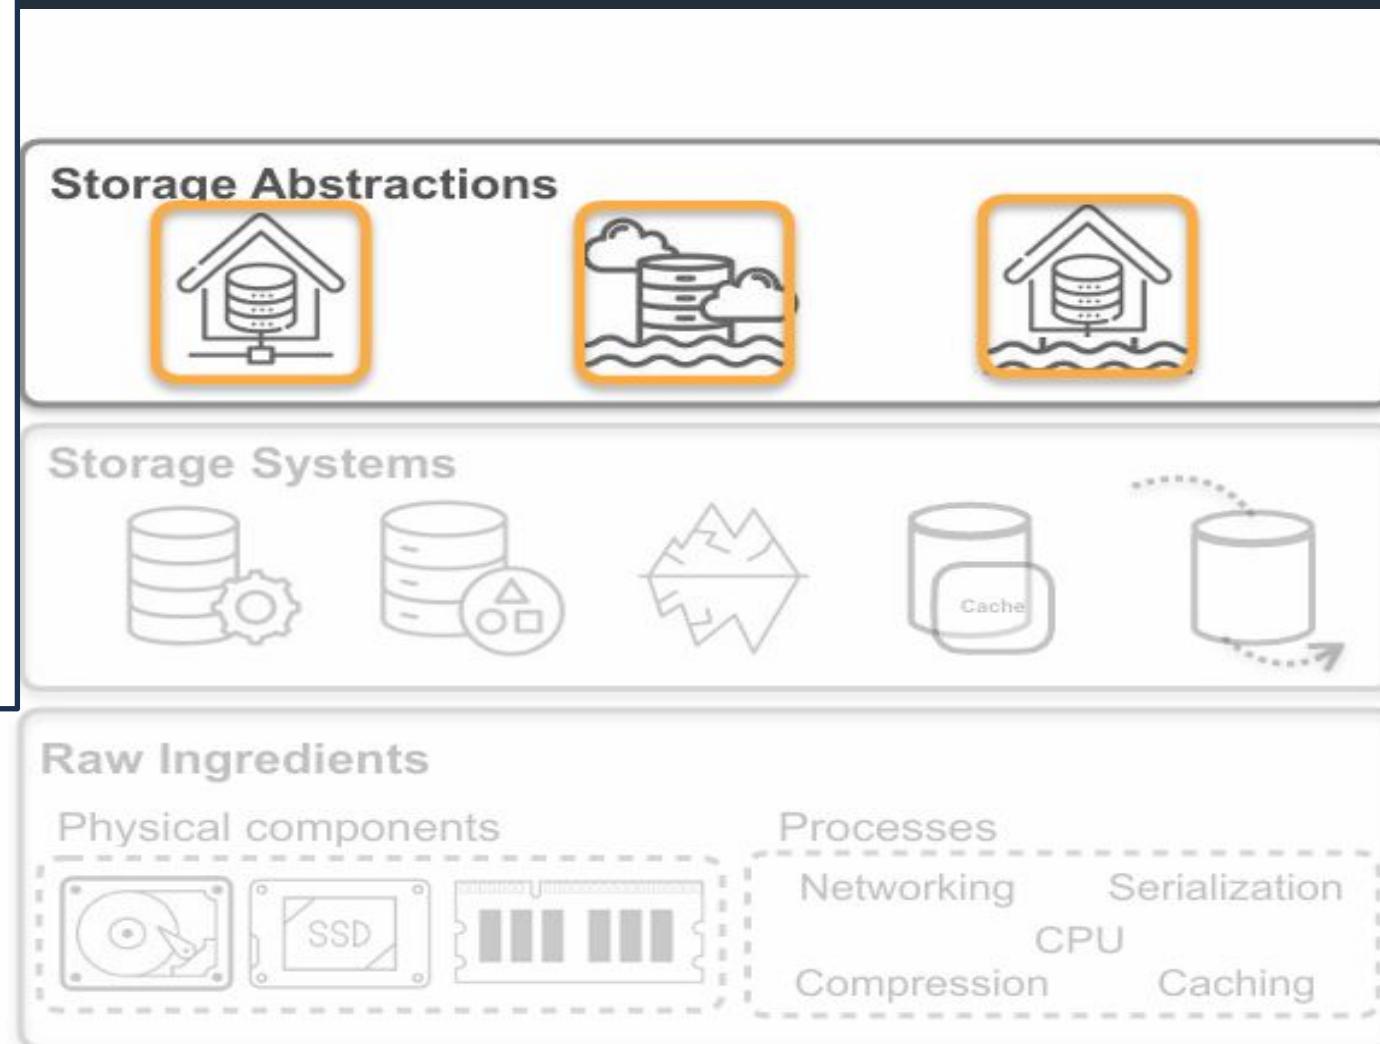
## Raw Ingredients (bottom layer)

Physical components: HDDs, SSDs, RAM, networking, serialization, compression. Data engineers rarely interact directly with them.

**Storage Systems** (middle layer) Built on raw ingredients. Managed systems: databases, object storage, OLTP vs. OLAP systems. Also includes specialized systems: graph databases, vector databases (for ML).

Storage

**Abstractions (top layer)** Cloud-level concepts: data warehouses, data lakes, data lakehouses. Organize systems for large-scale ingestion, transformation, and serving.



# Raw Ingredients: Physical Components of Data Storage

# Raw Storage Ingredients

## Persistent Storage Medium

Magnetic disk

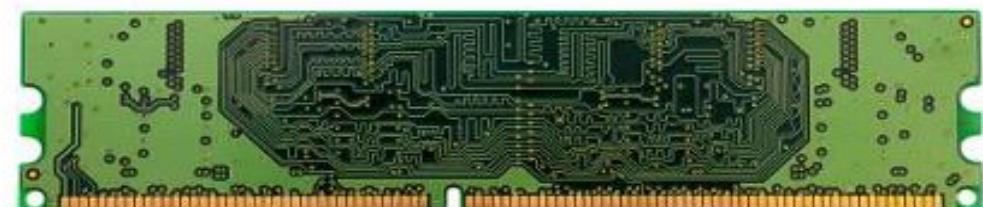


Solid-state storage



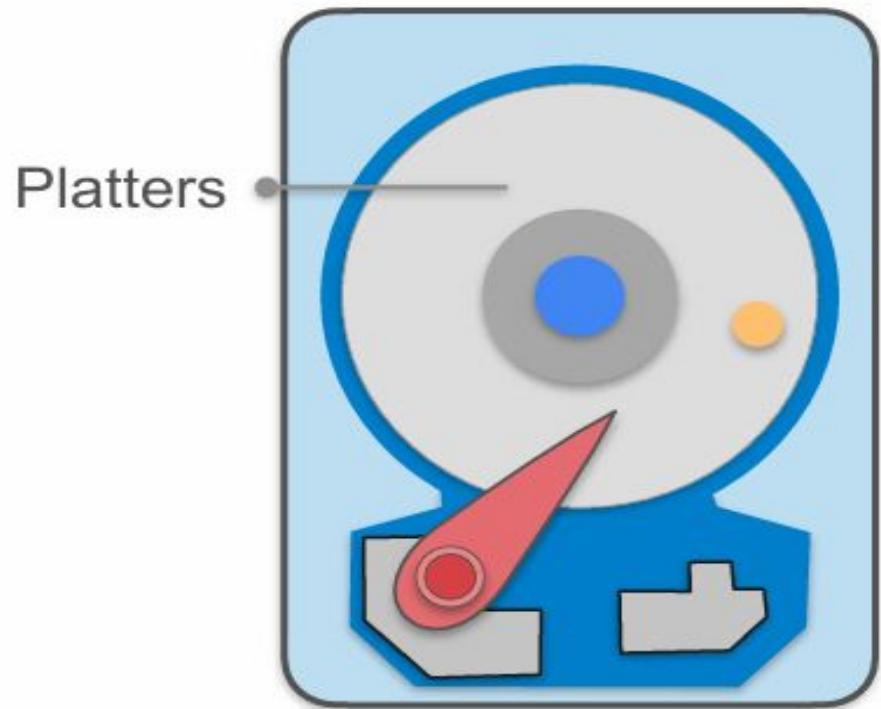
## Volatile Memory

RAM



CPU cache

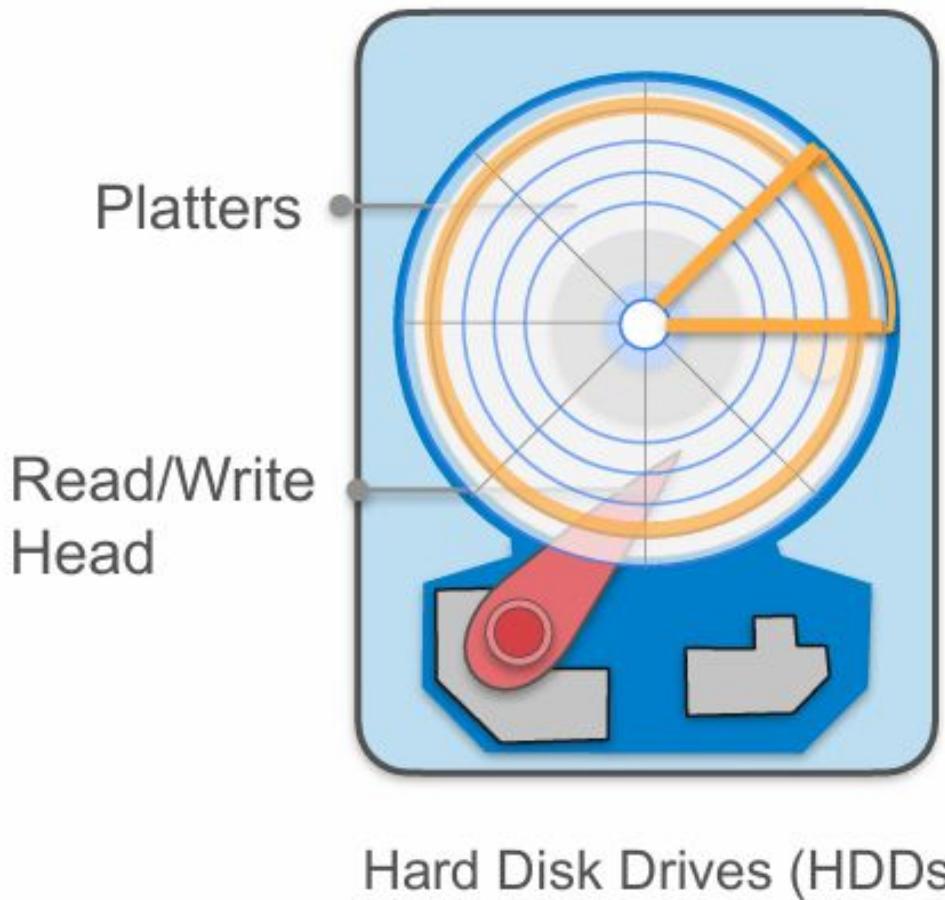
# Magnetic Disks



Hard Disk Drives (HDDs)



# Magnetic Disks

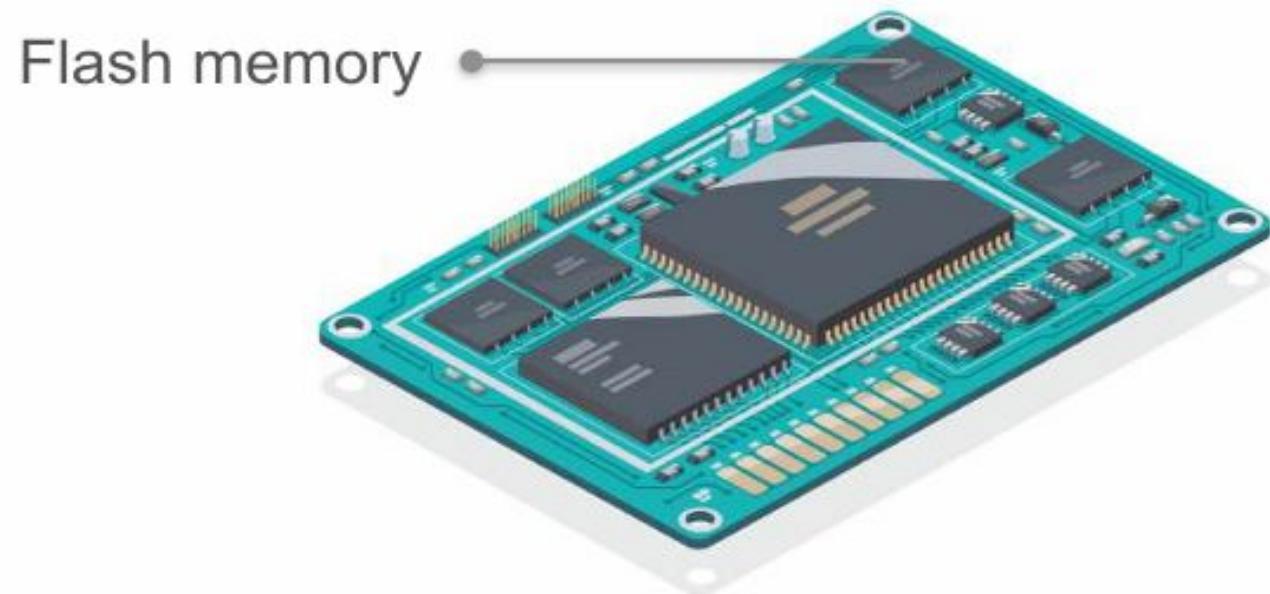


**Track + Sector**  
= **Address**

**Write:**  
Encode binary data by changing the magnetic field

**Read:**  
Converts magnetic field into binary data

# Solid State Drives



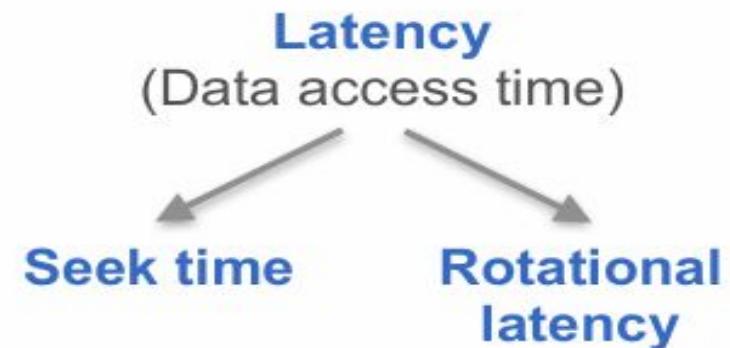
Solid-State Drives (SSDs)

SSDs read and write data  
much faster

# Performance Comparison

	Magnetic Disk	SSD
<b>Latency</b>	4 milliseconds	
<b>IOPS</b> (Input/output operations per second)	Hundreds	

**Commercial magnetic disk drive:**  
Rotates at 7200 revs/min



# Performance Comparison



	Magnetic Disk	SSD
<b>Latency</b>	4 milliseconds	0.1 milliseconds
<b>IOPS</b> (Input/output operations per second)	Hundreds	Tens of thousands



## Solid-State Drives (SSDs)



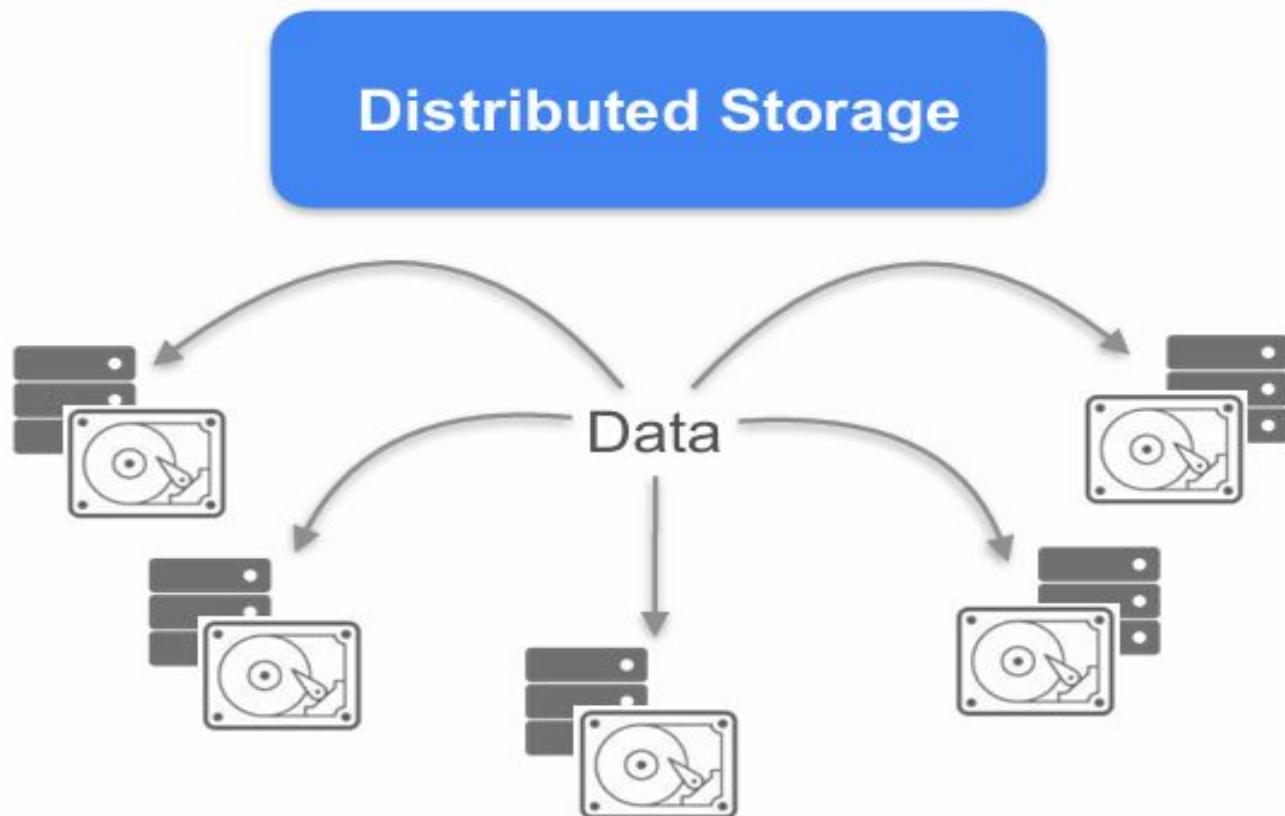
Electrical charges

# Performance Comparison



	Magnetic Disk	SSD
<b>Latency</b>	4 milliseconds	0.1 milliseconds
<b>IOPS</b> (Input/output operations per second)	Hundreds	Tens of thousands
<b>Data Transfer Speed</b> (number of bytes read/written from disk to memory in a second)	Up to 300 MB/s	4 GB/s

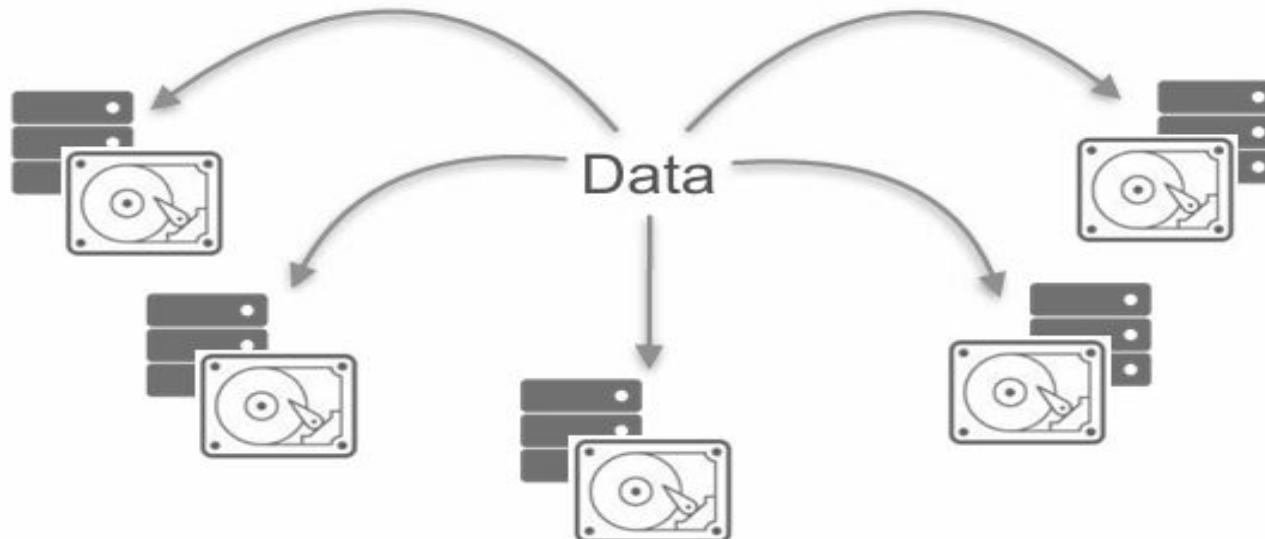
# Improving Performance



Data transfer speed limited by network performance

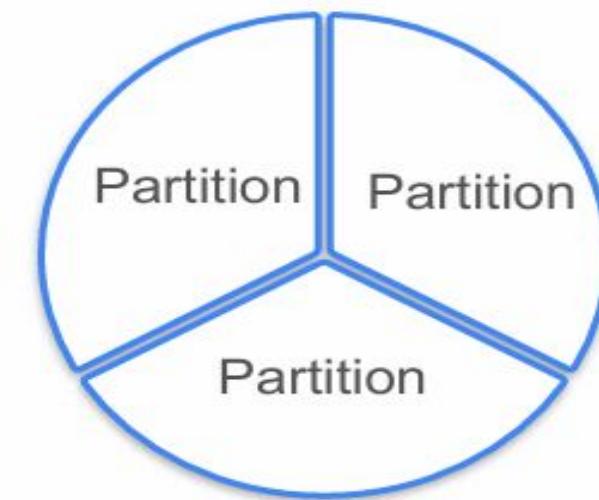
# Improving Performance

## Distributed Storage



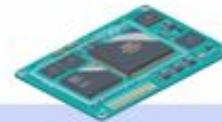
Data transfer speed limited by network performance

## Partitioning



Slicing SSDs into partitions

# Performance Comparison

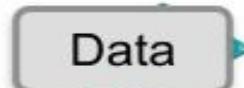


	Magnetic Disk	SSD
<b>Latency</b>	4 milliseconds	0.1 milliseconds
<b>IOPS</b> (Input/output operations per second)	Hundreds	Tens of thousands
<b>Data Transfer Speed</b> (number of bytes read/written from disk to memory in a second)	Up to 300 MB/s	4 GB/s
<b>Cost</b>	\$0.03–0.06/GB	\$0.08–0.10/GB

2-3 times cheaper

# Volatile Memory Ingredients

*\*\*Note: these metrics can vary*

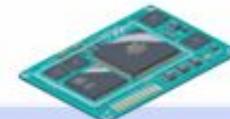


	Magnetic Disk	SSD	RAM (Random Access Memory)	CPU Cache
<b>Latency</b>	4 milliseconds	0.1 milliseconds	0.1 microseconds	
<b>IOPS</b> (Input/output operations per second)	Hundreds	Tens of thousands	Millions	
<b>Data Transfer Speed</b> (number of bytes read/written from disk to memory in a second)	Up to 300 MB/s	4 GB/s	100 GB/s	
<b>Cost</b>	\$0.03–0.06/GB	\$0.08–0.10/GB	> \$3/GB	

30-50 times more expensive

# Volatile Memory Ingredients

*\*\*Note: these metrics can vary*

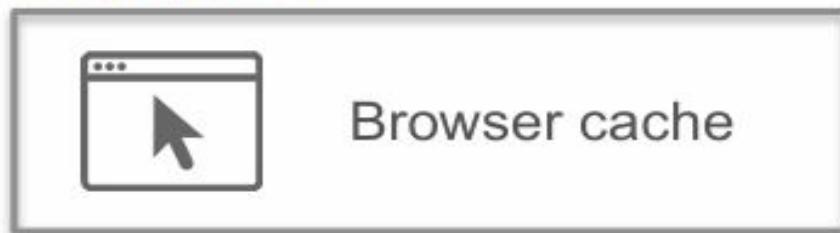


	Magnetic Disk	SSD	RAM (Random Access Memory)	CPU Cache
<b>Latency</b>	4 milliseconds	0.1 milliseconds	0.1 microseconds	1 nanosecond
<b>IOPS</b> (Input/output operations per second)	Hundreds	Tens of thousands	Millions	/
<b>Data Transfer Speed</b> (number of bytes read/written from disk to memory in a second)	Up to 300 MB/s	4 GB/s	100 GB/s	1 TB/s
<b>Cost</b>	\$0.03–0.06/GB	\$0.08–0.10/GB	> \$3/GB	/

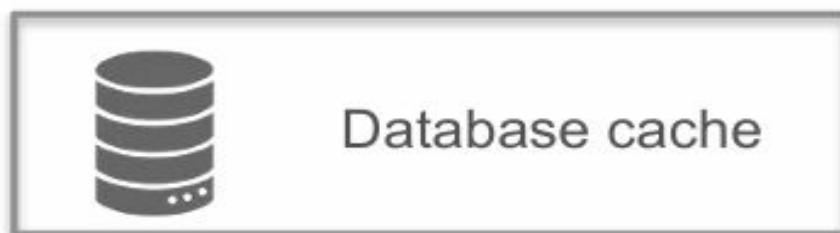
# CPU Cache Use Cases

- CPU caching
- Store frequently and recently accessed data in a fast access layer

## Examples



Store downloaded web resources

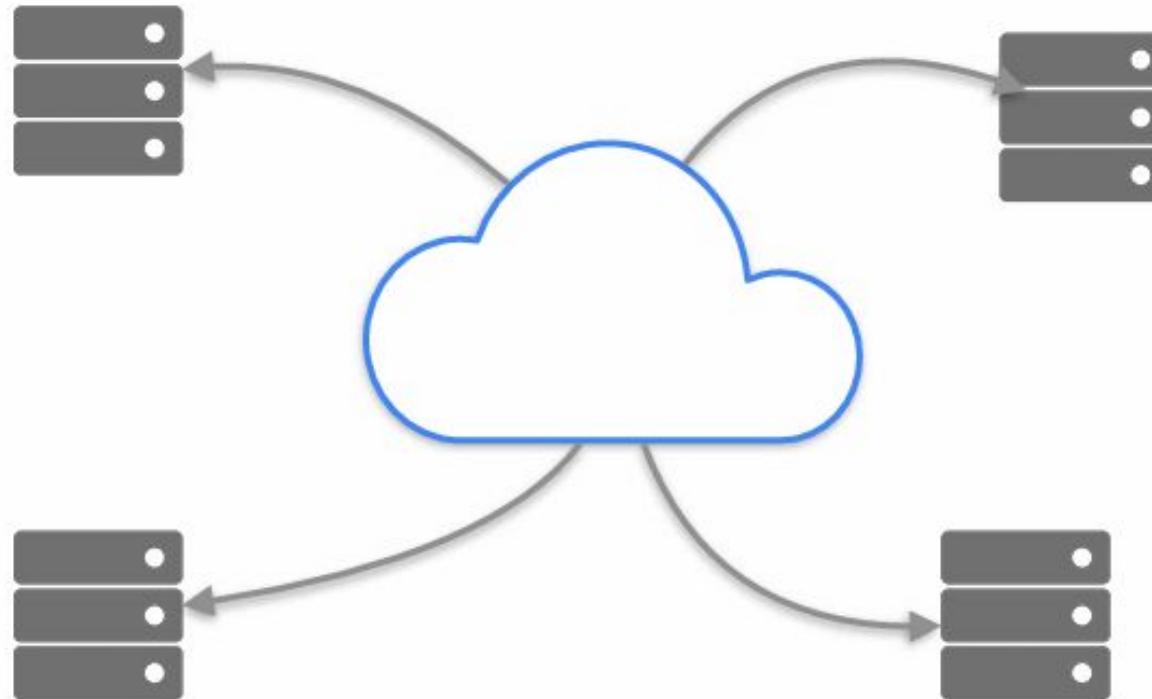


Store frequently used queries



# Raw Ingredients: Processes Required for Data Storage

# Networking and CPU — “Raw Ingredients” of Storage Systems

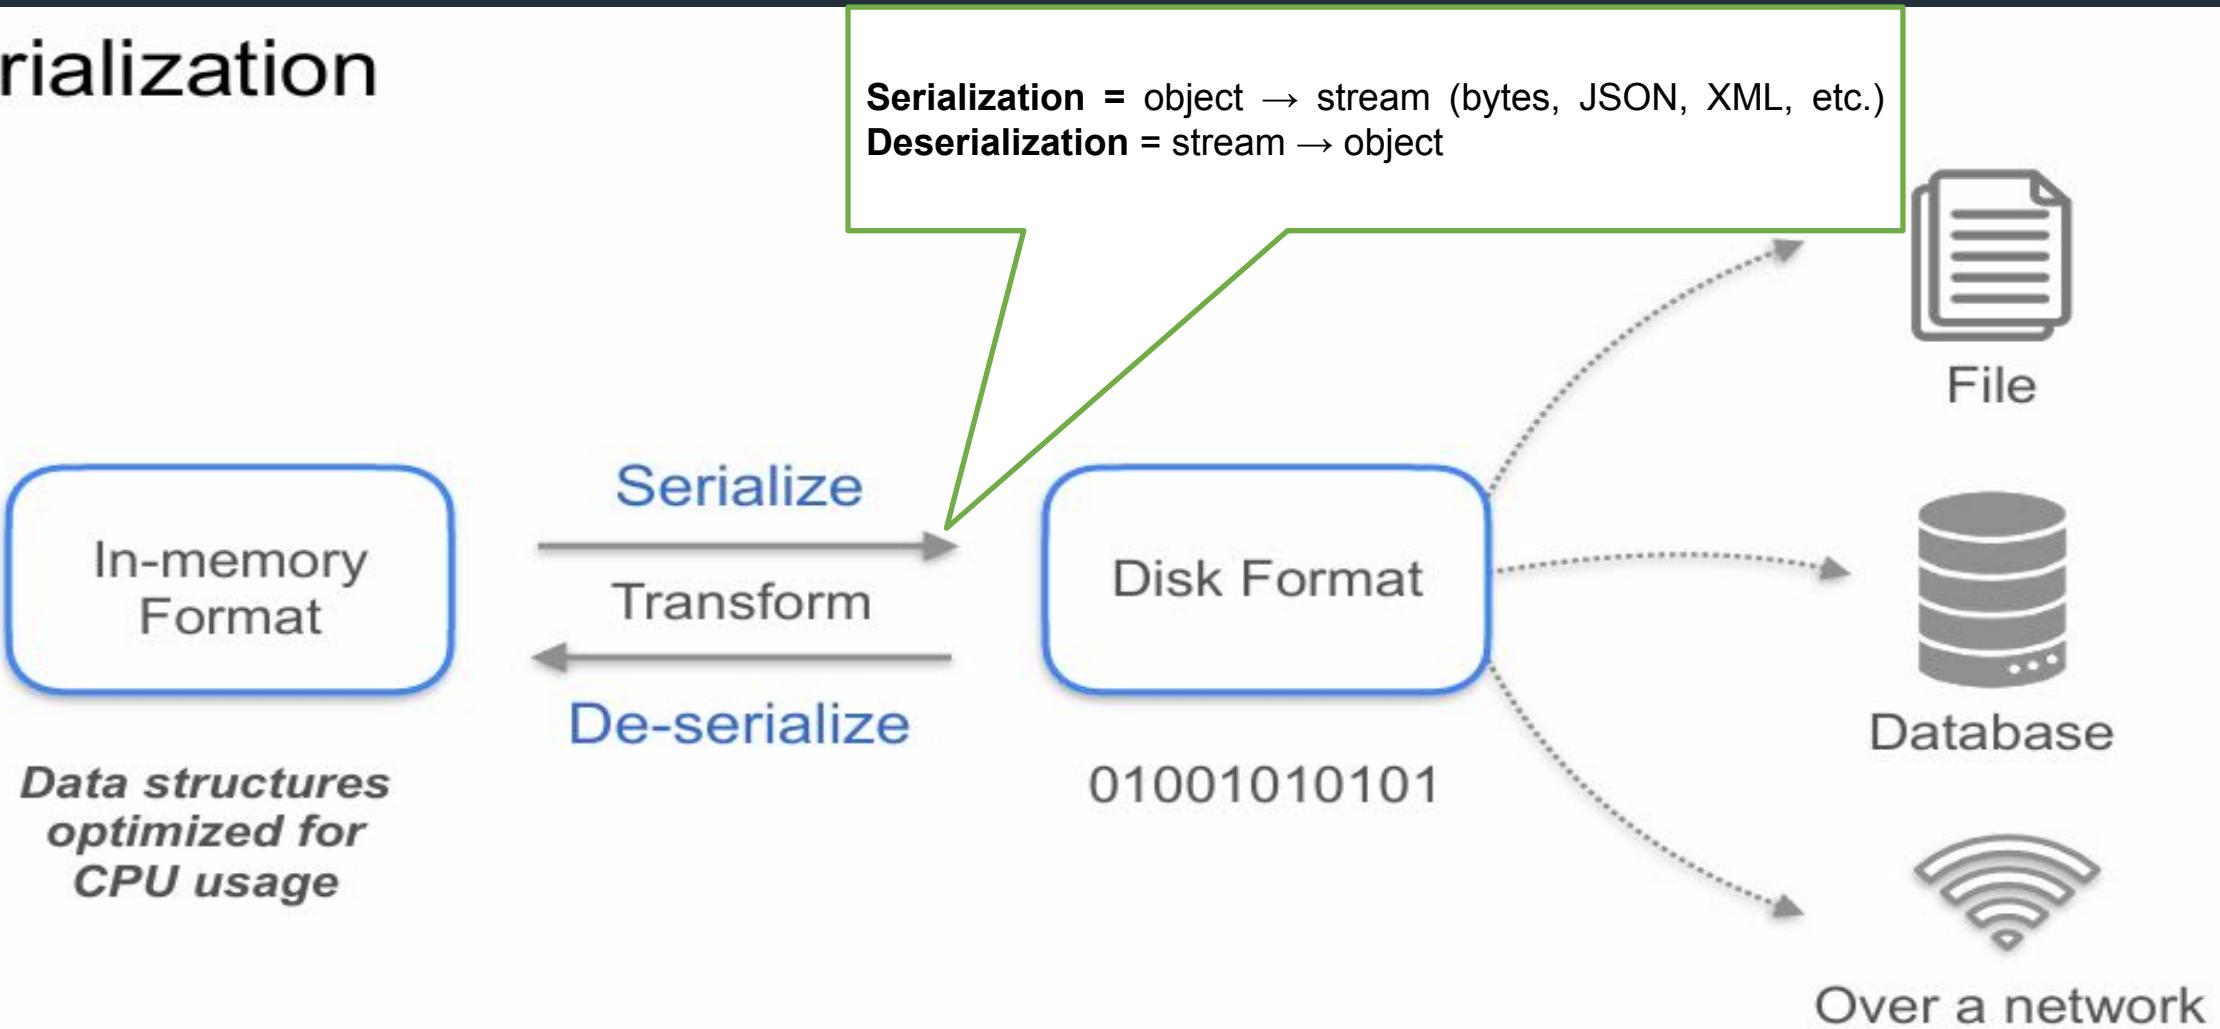


Data in memory (RAM) is stored in formats optimized for CPU access, not for long-term storage or network transfer.

## Enhance:

- read and write performance
- data durability
- data availability

# Serialization





# Serialization

Transactional operations

Order ID	Price	Product SKU	Quantity	Customer ID
1	40	45865	10	67t
2	23	90234	14	56t
3	45	12558	12	87q
4	50	45682	13	98q

Physical Storage



## Row-Based Serialization

bytes representing the 1st object | bytes representing the 2nd object | ... | bytes representing the last object

## Column-Based Serialization

--	--	--	--

Data is stored **row by row**, making it efficient for **transactional workloads (OLTP)** where entire records are frequently read or written

# Serialization

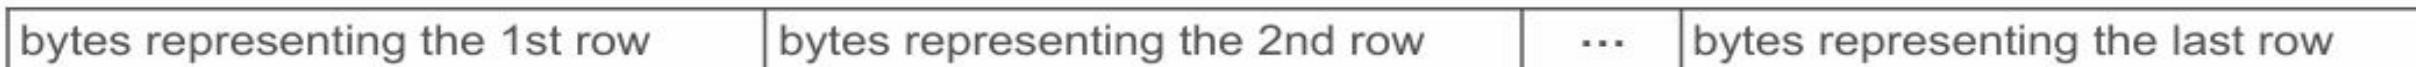
Order ID	Price	Product SKU	Quantity	Customer ID
1	40	45865	10	67t
2	23	90234	14	56t
3	45	12558	12	87q
4	50	45682	13	98q

Analytical queries

Physical Storage



## Row-Based Serialization

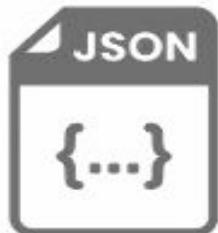


## Column-Based Serialization



# Serialization Formats

## Human-Readable Textual Formats



## Binary Formats



# Serialization Formats

## Human-Readable Textual Formats



- Row-based format
- Prone to error (no defined schema)
- Adding new rows or columns requires manual handling



## Binary Formats



# Serialization Formats

## Human-Readable Textual Formats



- Row-based format
- Prone to error (no defined schema)
- Adding new rows or columns requires manual handling



- Extensible markup language
- Viewed as a legacy format
- Slow to serialize and deserialize



## Binary Formats



# Serialization Formats

## Human-Readable Textual Formats



- Row-based format
- Prone to error (no defined schema)
- Adding new rows or columns requires manual handling



- Extensible markup language
- Viewed as a legacy format
- Slow to serialize and deserialize



- Used for plain-text object serialization
- Viewed as new standard for data exchange over APIs

## Binary Formats



# Serialization Formats

## Human-Readable Textual Formats



- Row-based format
- Prone to error (no defined schema)
- Adding new rows or columns requires manual handling



- Extensible markup language
- Viewed as a legacy format
- Slow to serialize and deserialize



- Used for plain-text object serialization
- Viewed as new standard for data exchange over APIs

## Binary Formats



- Column-based format
- For efficient storage and big data processing



# Serialization Formats

## Human-Readable Textual Formats



- Row-based format
- Prone to error (no defined schema)
- Adding new rows or columns requires manual handling



- Extensible markup language
- Viewed as a legacy format
- Slow to serialize and deserialize



- Used for plain-text object serialization
- Viewed as new standard for data exchange over APIs

## Binary Formats



- Column-based format
- For efficient storage and big data processing



- Row-based format
- Uses a schema to define its data structure
- Supports schema evolution

# Serialization Formats

## Human-Readable Textual Formats



- Row-based format
- Prone to error (no defined schema)
- Adding new rows or columns requires manual handling



- Extensible markup language
- Viewed as a legacy format
- Slow to serialize and deserialize



- Used for plain-text object serialization
- Viewed as new standard for data exchange over APIs

## Binary Formats

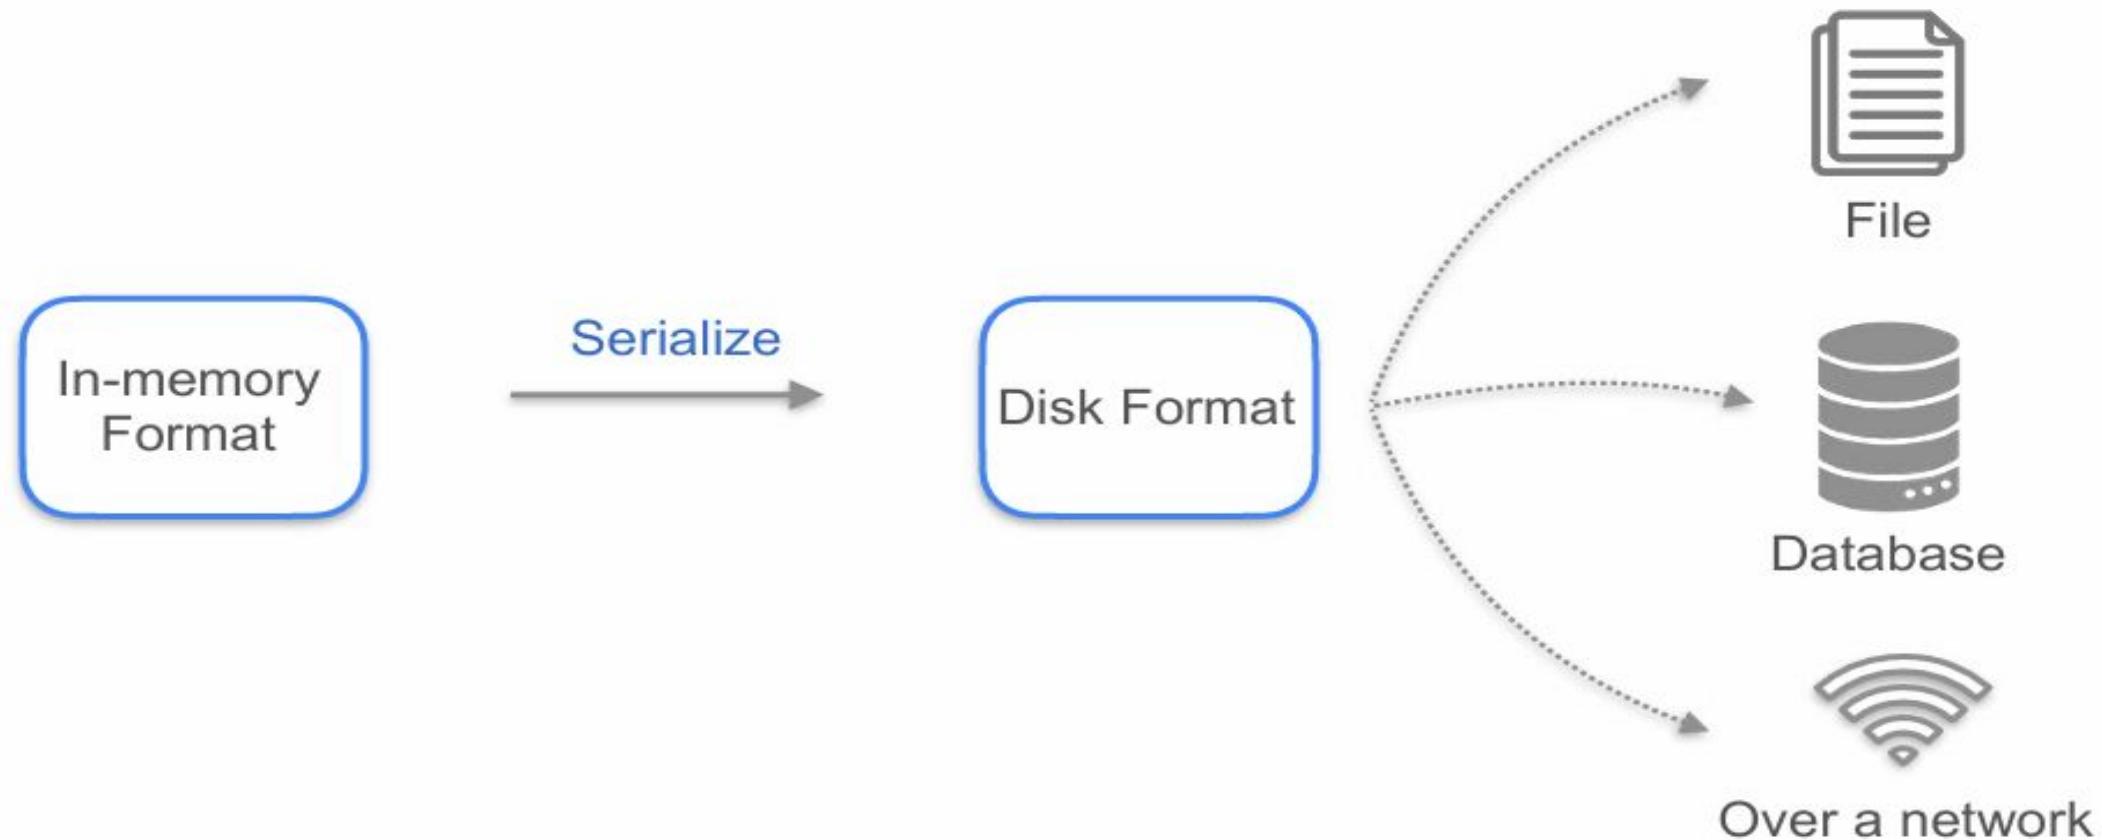


- Column-based format
- For efficient storage and big data processing

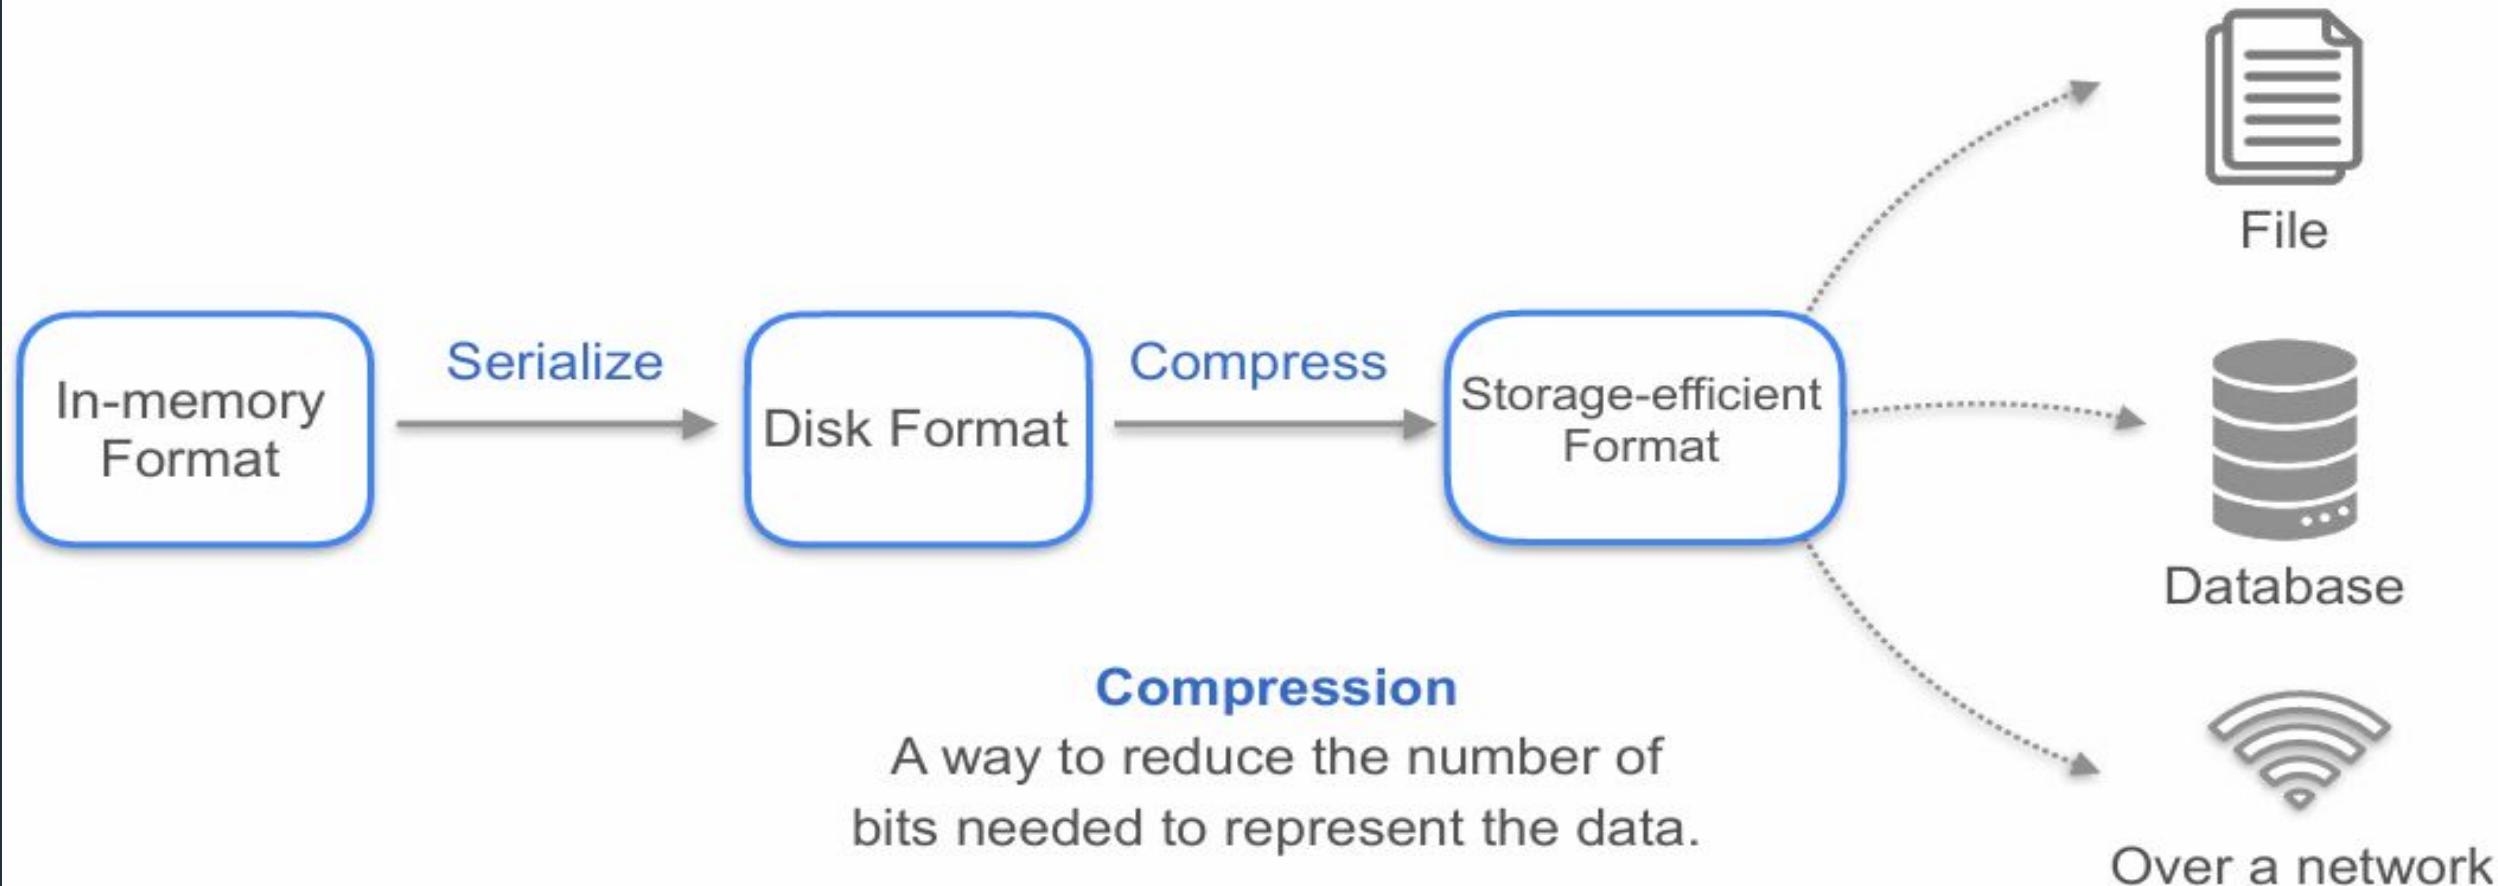


- Row-based format
- Uses a schema to define its data structure
- Supports schema evolution

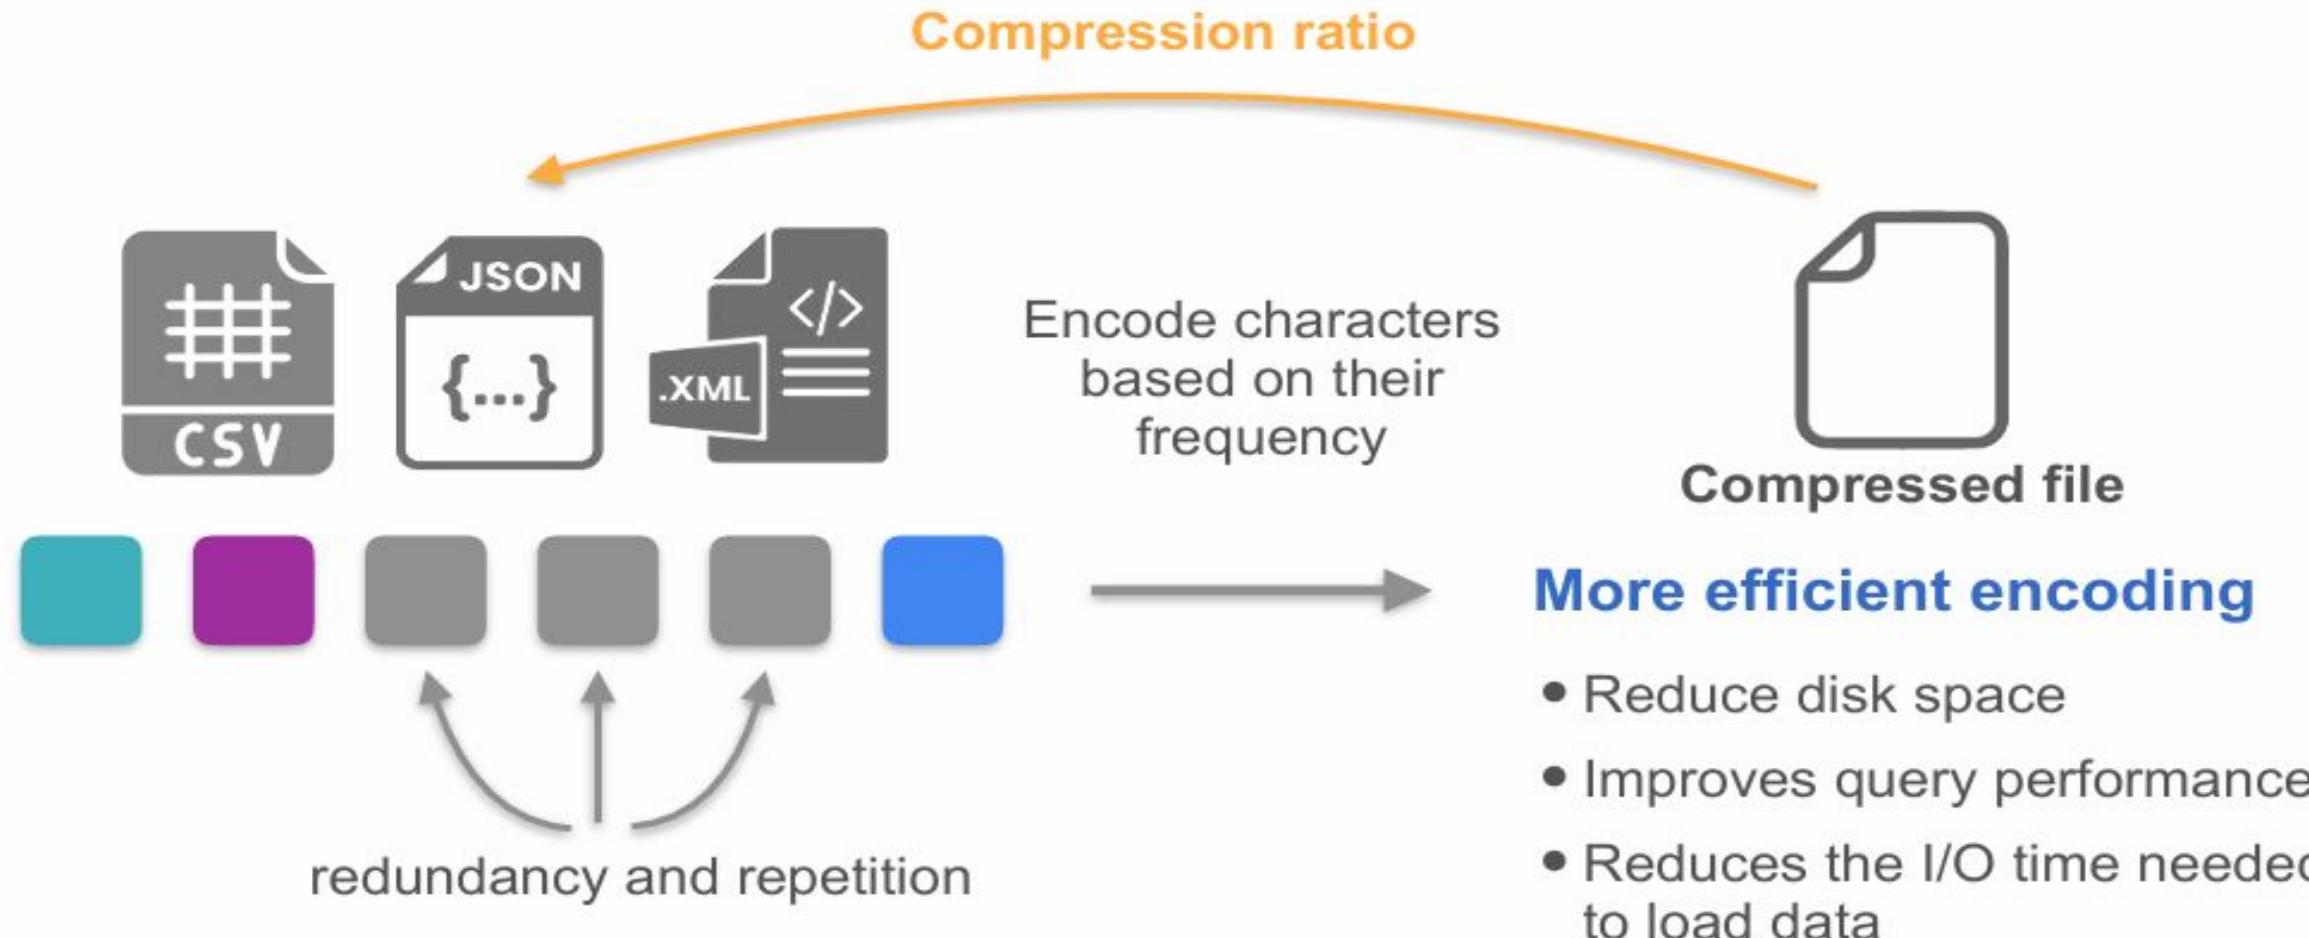
# Serialization



# Serialization

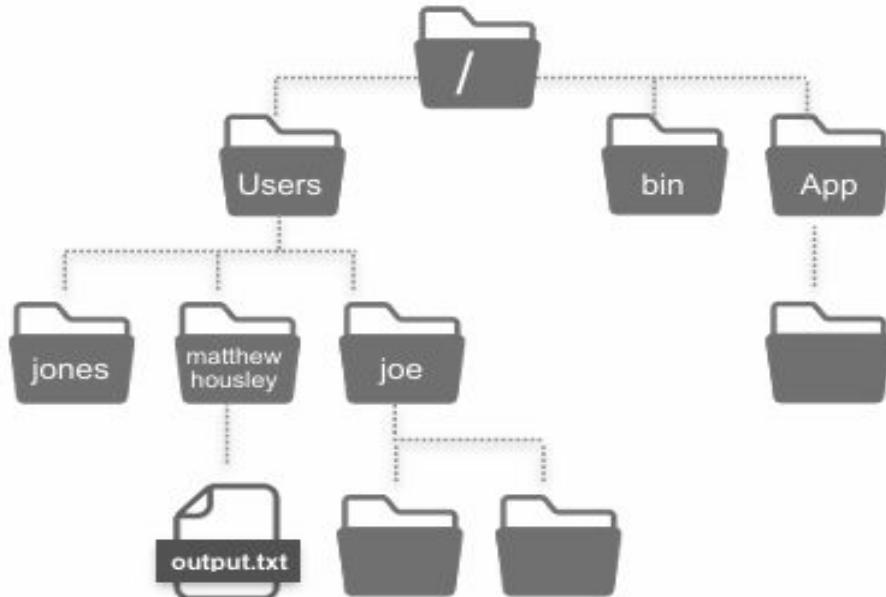


# Compression



# Cloud Storage Options: Block, Object and File storage

# File Storage



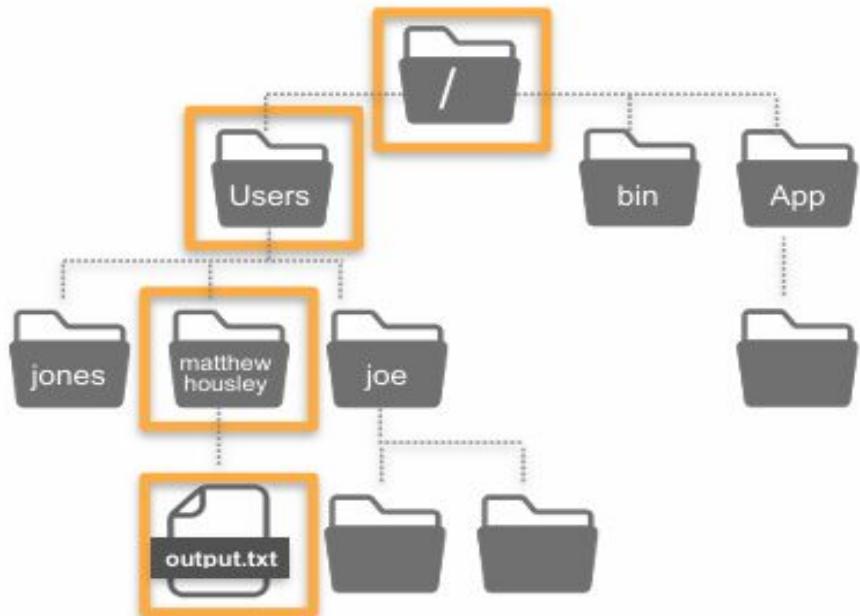
## File Storage

Organizes files into a directory tree

Each directory contains metadata about its files and subfolders :

- Name
- Owner
- Last modified date
- Permissions
- Pointer to the actual entity

# File Storage



/Users/matthewhousley/output.txt

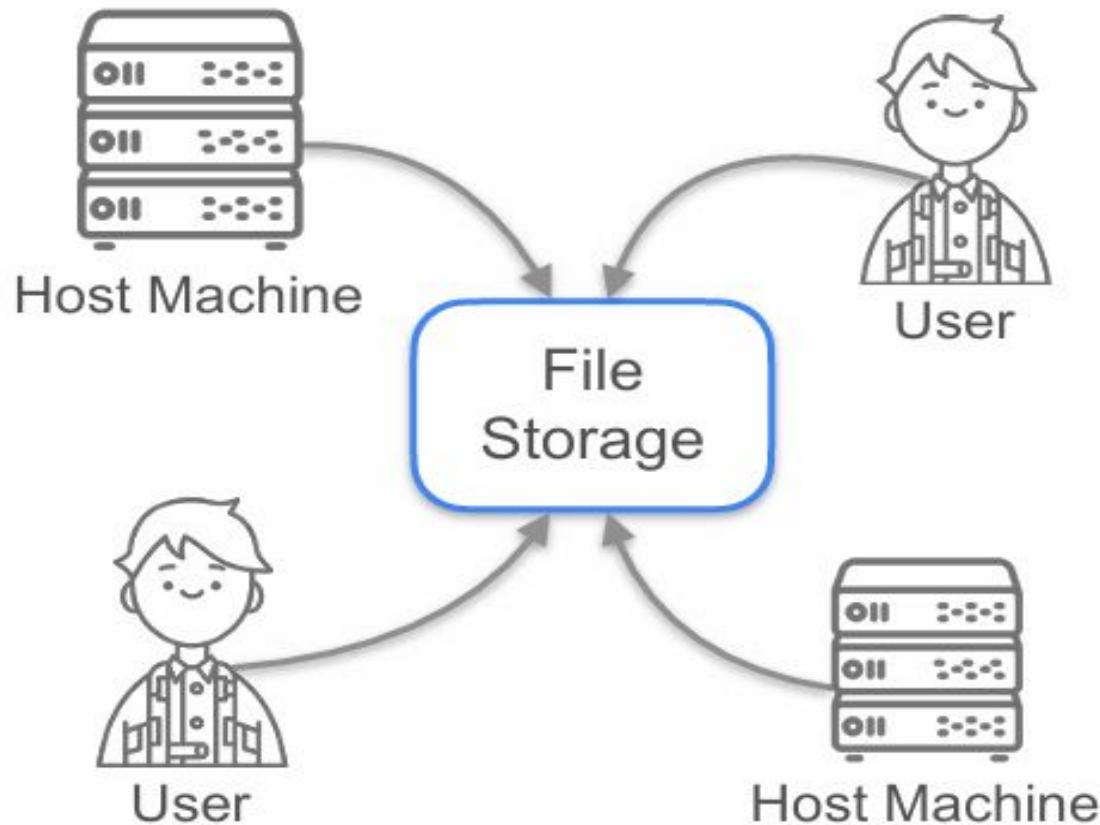
## File Storage

Organizes files into a directory tree

Each directory contains metadata about its files and subfolders :

- Name
- Owner
- Last modified date
- Permissions
- Pointer to the actual entity

# File Storage Use Cases



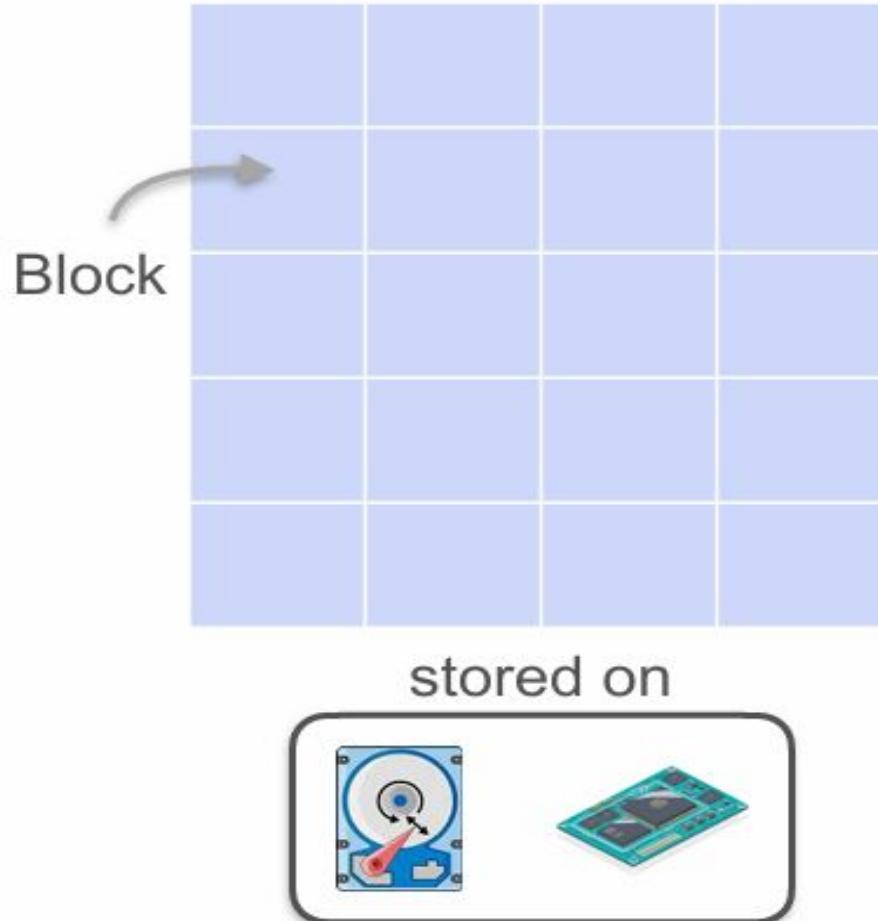
## Cloud File Storage Service



Amazon Elastic File System (EFS)

- Provides you access to shared files over a network
- Networking, scaling, and configuration are handled by the cloud vendor

# Block Storage

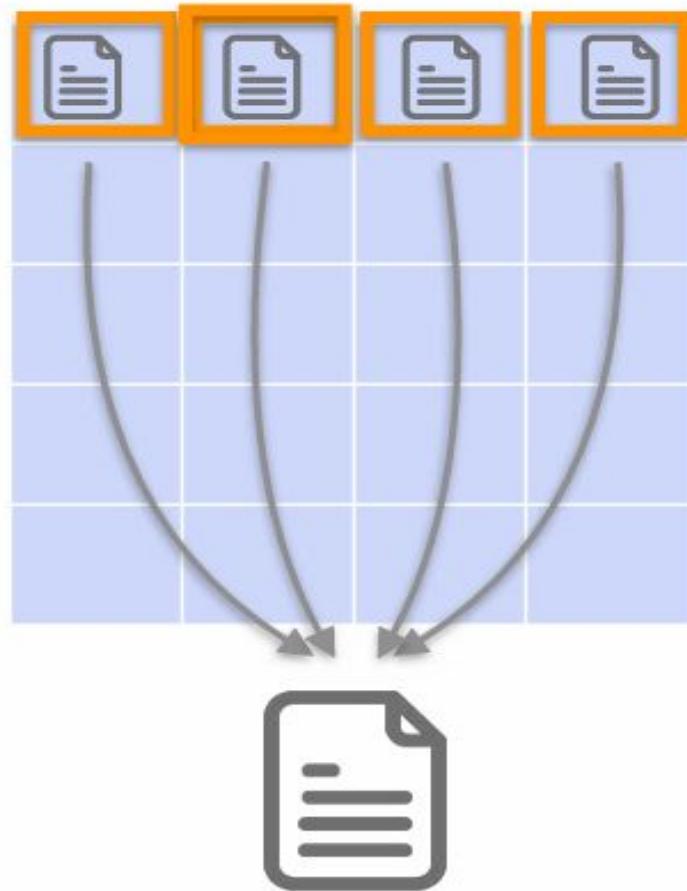


## Block Storage

Divides files into small, fixed-size blocks of data and stores them on disk

- Each block has a unique identifier
- You can efficiently retrieve and modify data in individual blocks
- You can distribute blocks of data across multiple storage disks
  - Higher scalability
  - Stronger data durability

# Block Storage

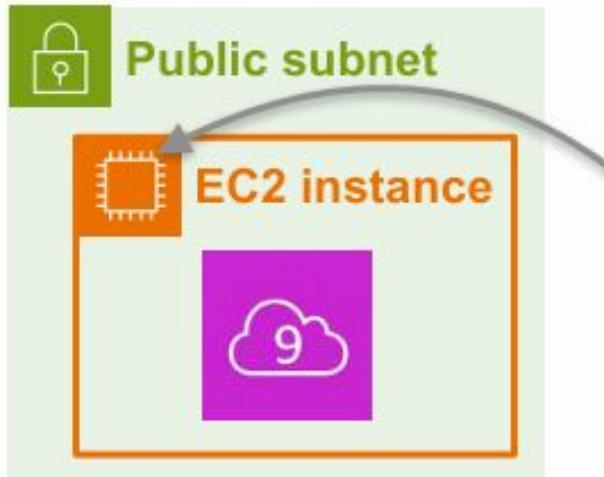


**Lookup Table**

File Piece	Block Identifier
First piece	1232
Second piece	1234
Third piece	1236
Fourth piece	1238

# Block Storage Use Cases

- Ideal for frequent access and modification
- Enables OLTP systems to perform small and frequent read and write operations with low latency
- Provides persistent storage for virtual machines



Attach a root storage device backed by a block storage volume

## Default storage for EC2



Amazon Elastic Block Store (EBS)

1. SSD for latency-sensitive workloads
2. Magnetic disks to store infrequently-accessed data

# Object Storage

## Object Storage

Stores immutable files as data objects in a flat structure



# Object Storage

## Object Storage

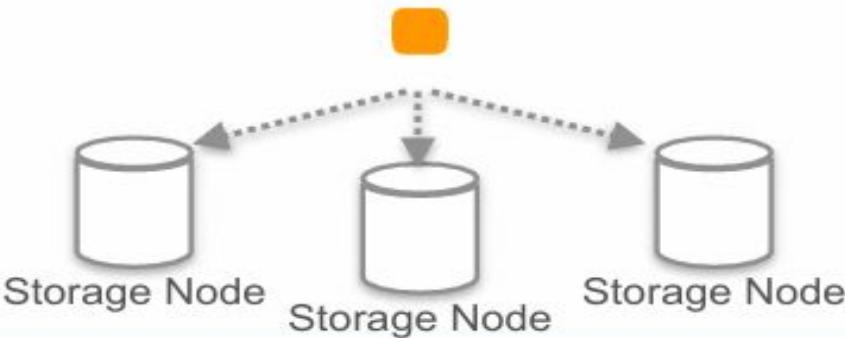
Stores immutable files as data objects in a flat structure



[s3://o'reilly-data-engineering-book/data-example.json](https://s3://o'reilly-data-engineering-book/data-example.json)

**The Bucket**

**Object key**

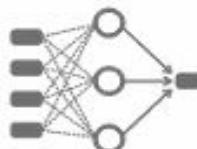


# Object Storage Use Cases



## Ideal for...

- Storage layer of cloud data warehouses or data lakes
- Storing data needed in OLAP systems
- Machine learning pipelines
  - Raw text
  - Images
  - Videos
  - Audio



## Not ideal for...

- Not good at supporting transactional workloads

# Cloud Storage Options

**File storage** → easy file sharing.  
**Block storage** → OLTP, frequent read/write.  
**Object storage** → OLAP, big data, massive scalability.

File Storage	Block Storage	Object Storage
<ul style="list-style-type: none"><li>• Supports data sharing</li><li>• Easy to manage with low performance and scalability requirements</li></ul>	<ul style="list-style-type: none"><li>• Supports transactional workloads</li><li>• Allows frequent read and write operations with low latency</li></ul>	<ul style="list-style-type: none"><li>• Supports analytical queries on massive datasets</li><li>• Offers high scalability and parallel data processing</li></ul>

# Storage Tiers – Hot, Warm, & Cold Data

## Cloud Storage Tiers



### Hot

Frequently accessed



### Warm

Occasionally accessed



### Cold

Rarely accessed



High Cost



Moderate Cost



Low Cost



# Storage Tiers

	Hot Storage	Warm Storage	Cold Storage
<b>Access Frequency</b>			
<b>Example</b>			
<b>Storage Medium</b>			
<b>Storage Cost</b>			
<b>Retrieval Cost</b>			

# Storage Tiers

	Hot Storage	Warm Storage	Cold Storage
<b>Access Frequency</b>	Very frequent		
<b>Example</b>	Product recommendation application		
<b>Storage Medium</b>	SSD & Memory		
<b>Storage Cost</b>	High		
<b>Retrieval Cost</b>	Low		

# Storage Tiers

	Hot Storage	Warm Storage	Cold Storage
Access Frequency	Very frequent	Less frequent	
Example	Product recommendation application	Regular reports and analyses	
Storage Medium	SSD & Memory	Magnetic disks or hybrid storage systems	
Storage Cost	High	Medium	
Retrieval Cost	Low	Medium	

# Storage Tiers

	Hot Storage	Warm Storage	Cold Storage
Access Frequency	Very frequent	Less frequent	Infrequent
Example	Product recommendation application	Regular reports and analyses	Archive
Storage Medium	SSD & Memory	Magnetic disks or hybrid storage systems	Low-cost magnetic disks
Storage Cost	High	Medium	Low
Retrieval Cost	Low	Medium	High

# Storage Tiers

Storage cost decreases as retrieval speed decreases (hot → warm → cold).  
Optimal design often uses a combination of tiers:

Hot for real-time access  
Warm for periodic access  
Cold for archival purposes

**Hot Storage**

**Warm Storage**

**Cold Storage**

<b>Access Frequency</b>	Very frequent	Less frequent	Infrequent
<b>Example</b>	Product recommendation application	Regular reports and analyses	Archive
<b>Storage Medium</b>	SSD & Memory	Magnetic disks or hybrid storage systems	Low-cost magnetic disks
<b>Storage Cost</b>	High	Medium	Low
<b>Retrieval Cost</b>	Low	Medium	High

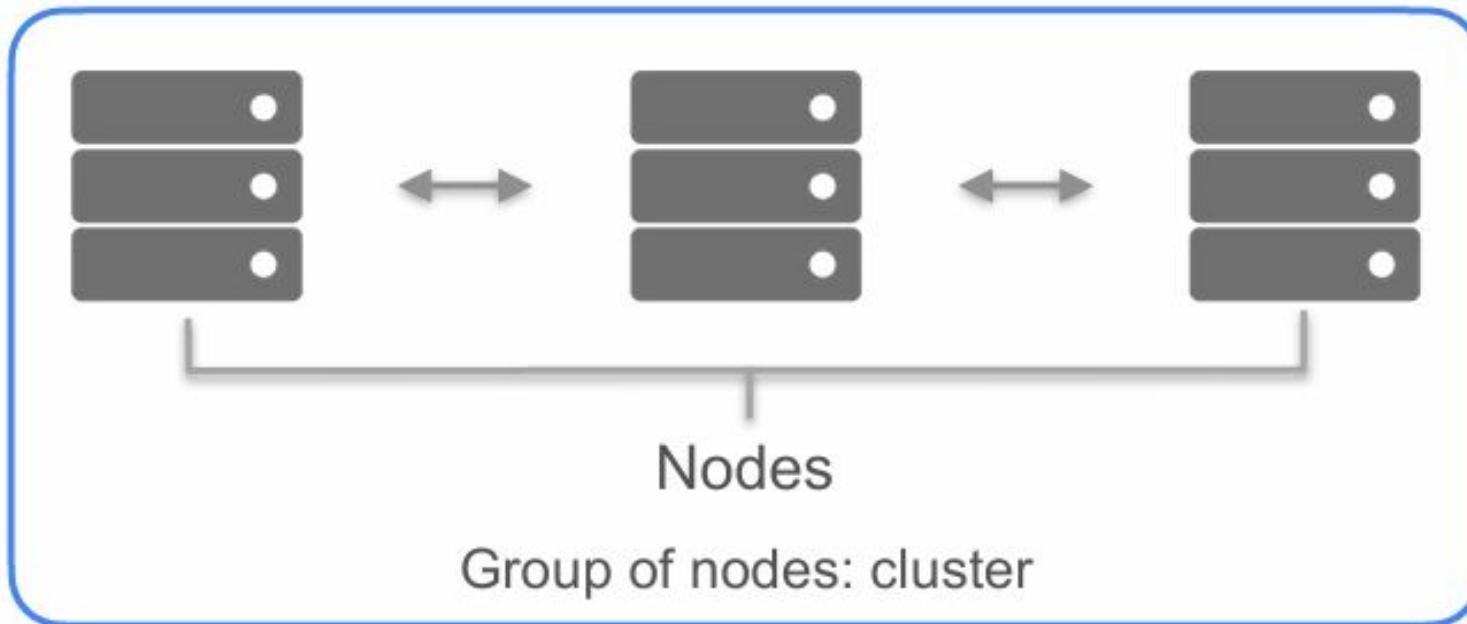


# AWS Storage Tiers

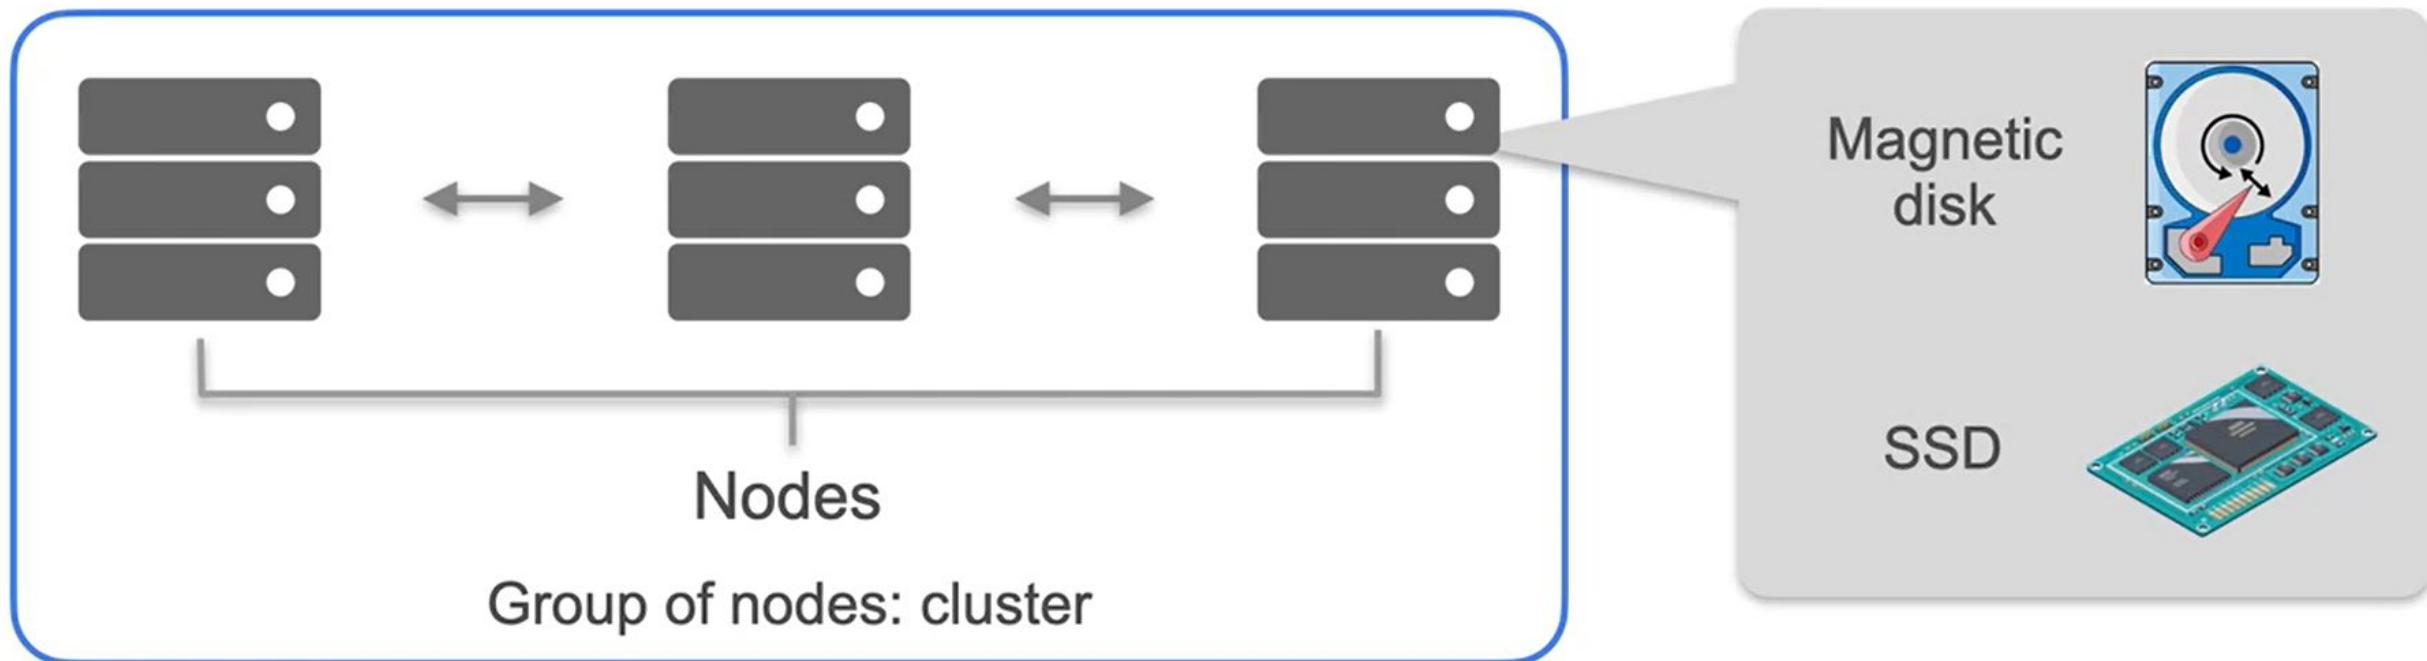


# Distributed Storage Systems

# How Distributed Storage Systems Work

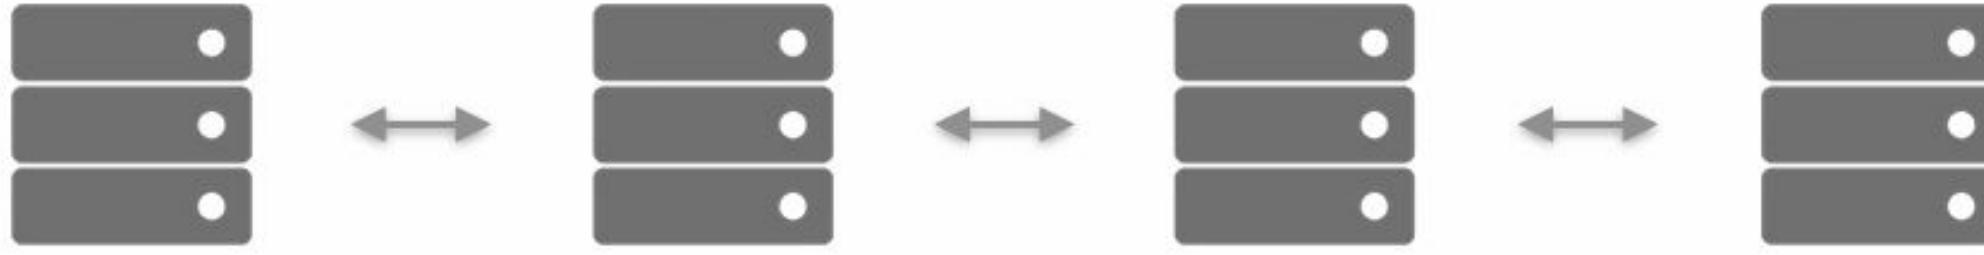


# How Distributed Storage Systems Work



$$\text{Total Capacity} = \text{Node 1} + \text{Node 2} + \text{Node 3}$$

# How Distributed Storage Systems Work



Horizontal Scaling

## Single Machine Storage Architecture

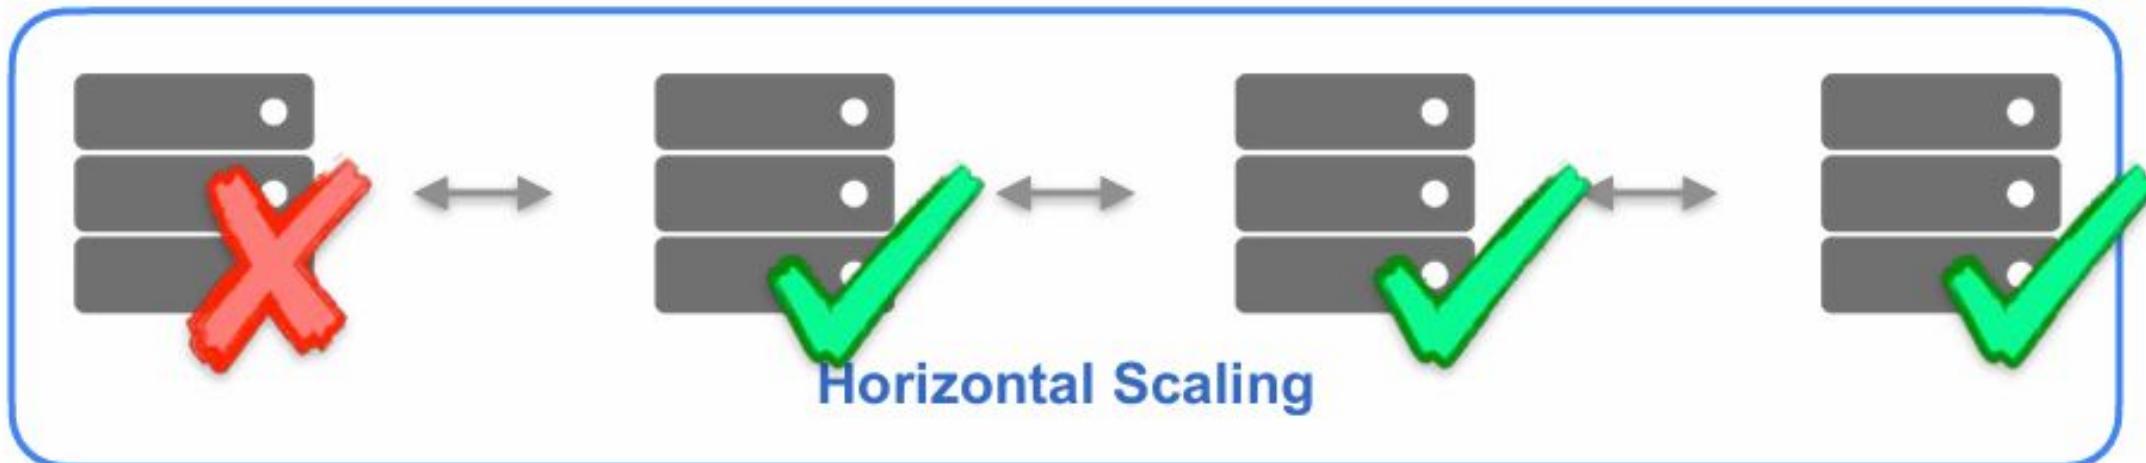


Vertical Scaling

**Vertical scaling:** Scale up a single server. Simple but limited.

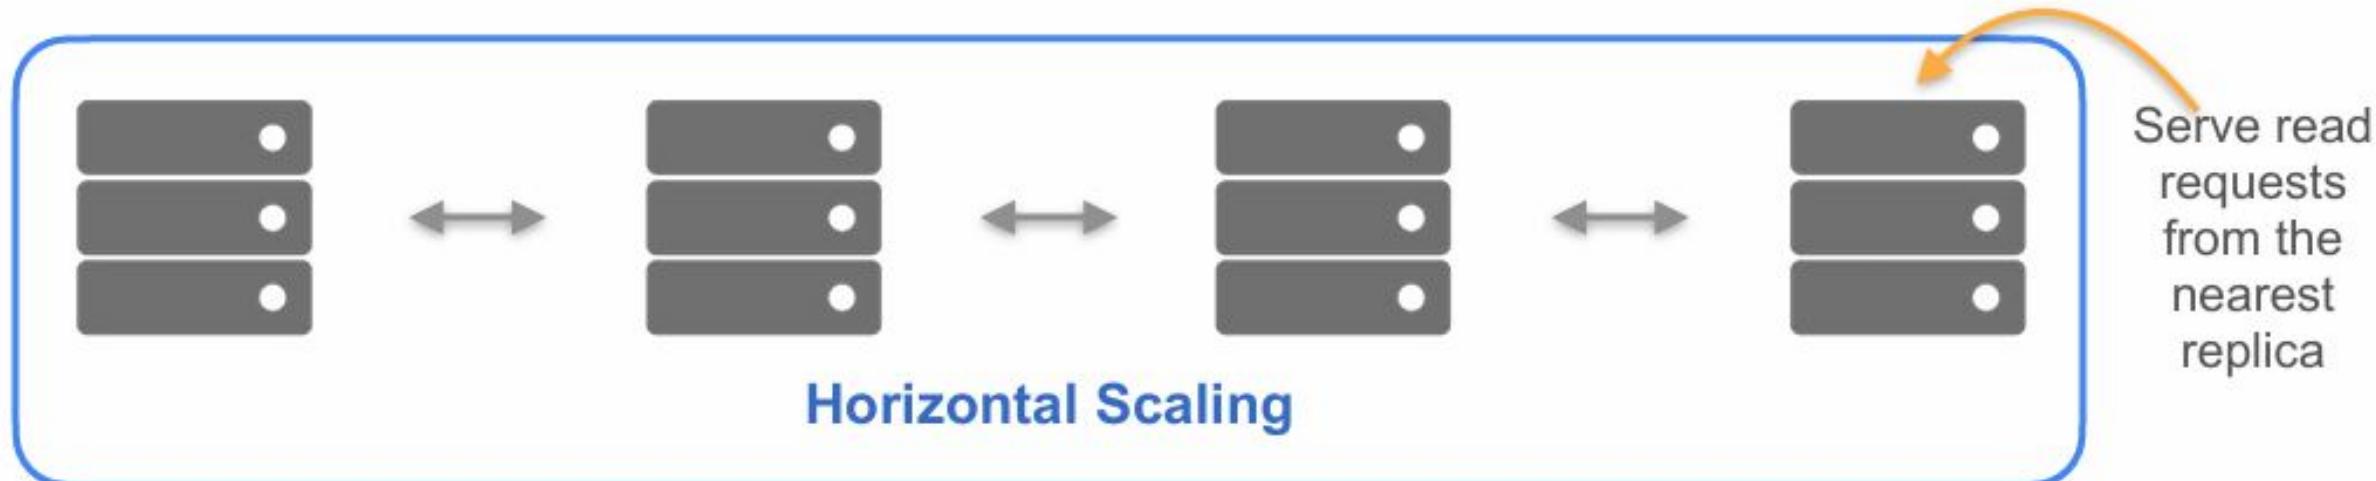
**Horizontal scaling:** Scale out across multiple servers. Complex but highly scalable.

# How Distributed Storage Systems Work

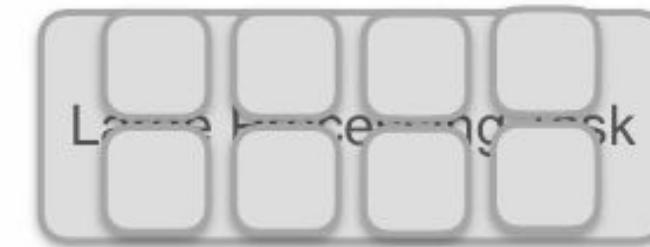


- Higher fault tolerance and data durability
- High availability

# How Distributed Storage Systems Work



- Higher fault tolerance and data durability
- High availability
- Process many read and write operations in parallel
- Fast data access



# Advantages of Distributed Storage Systems

## Distributed Storage Architecture



Object Storage



Cloud  
Data Warehouse



# Methods for Distributing Data

**Replication**

**Partitioning**

# Methods for Distributing Data

**Replication**

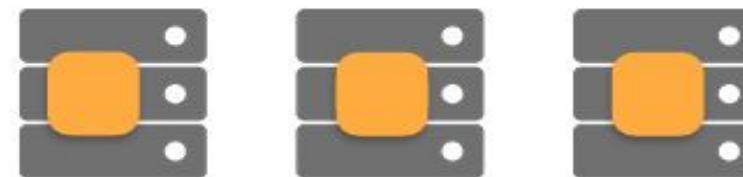


High availability and performance

**Partitioning**

# Methods for Distributing Data

**Replication**



High availability and performance

**Partitioning**

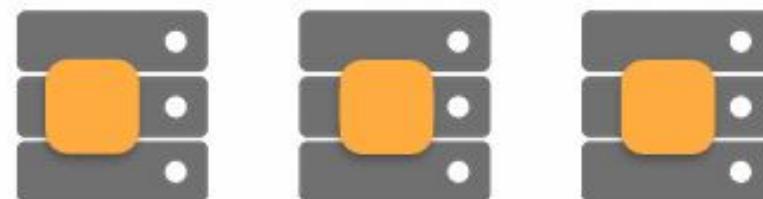


**Sharding**



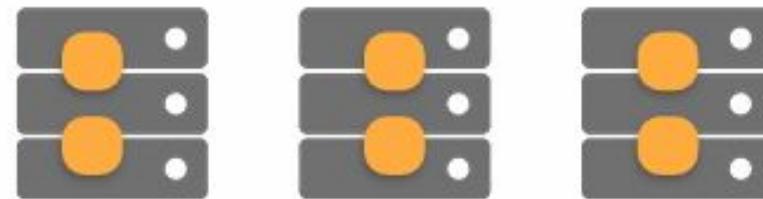
# Methods for Distributing Data

**Replication**



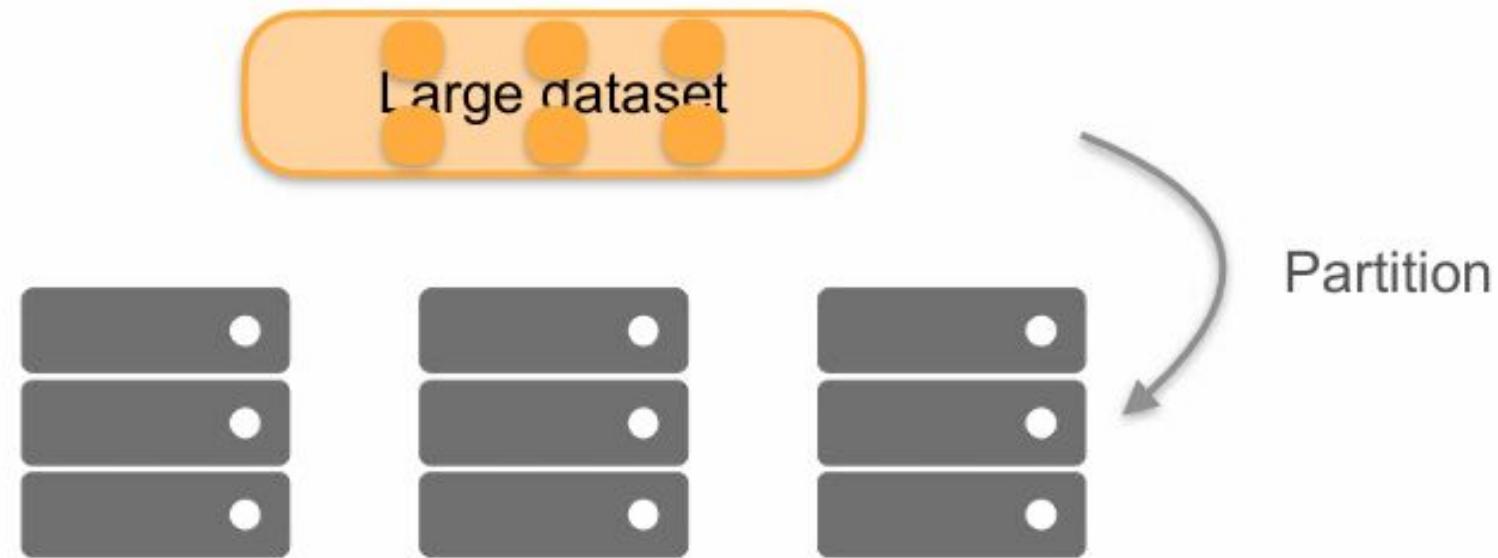
High availability and performance

**Partitioning**



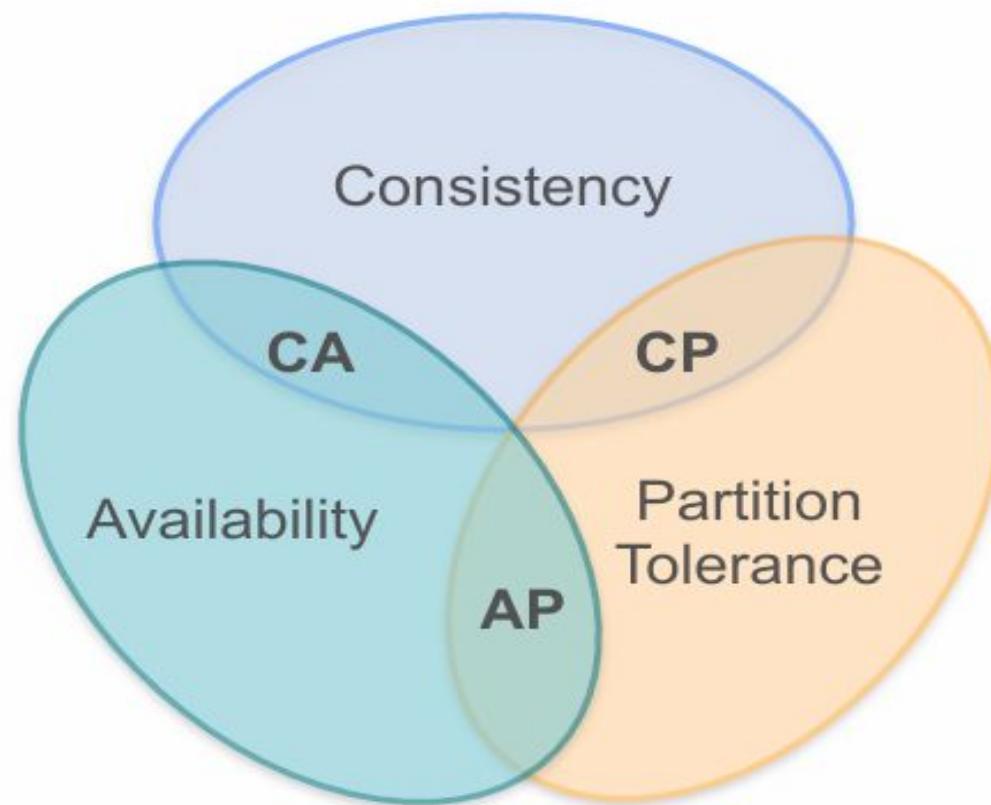
**Sharding**

# Methods for Distributing Data



# Distributed Storage Considerations – CAP Theorem

## The CAP theorem



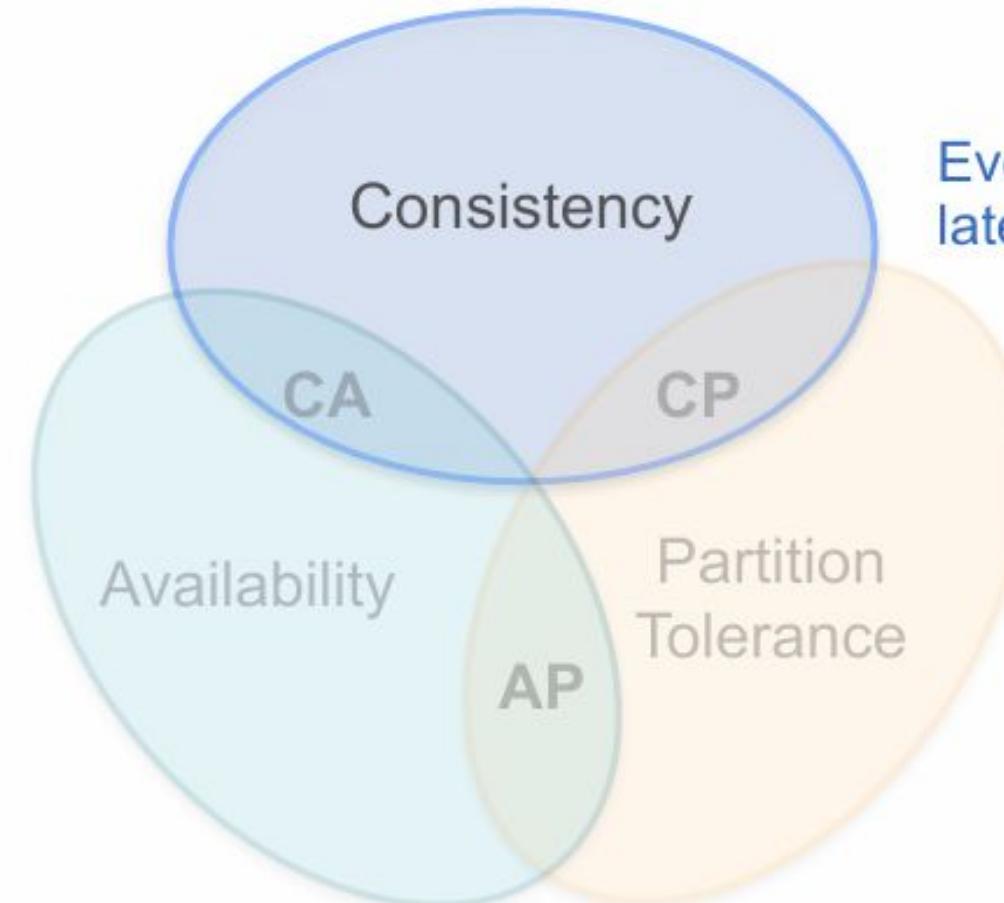
# Distributed Storage Considerations – CAP Theorem

## ACID



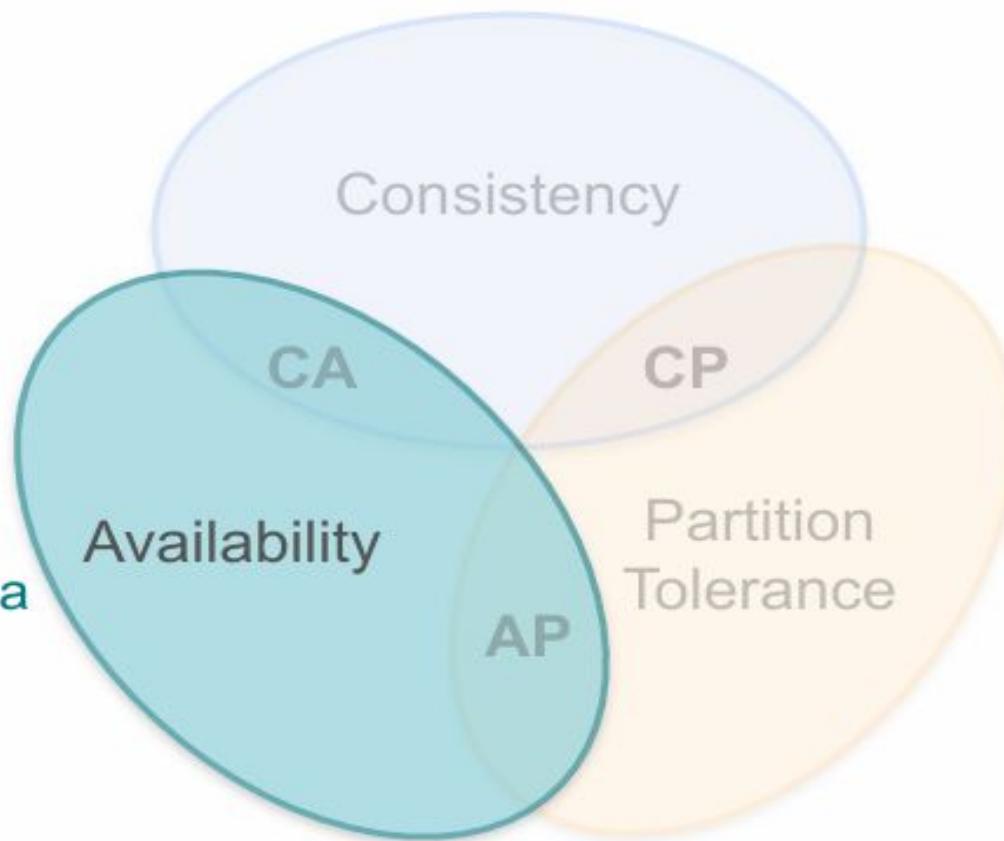
### Consistency

Any change to data must follow the set of rules defined by database schema



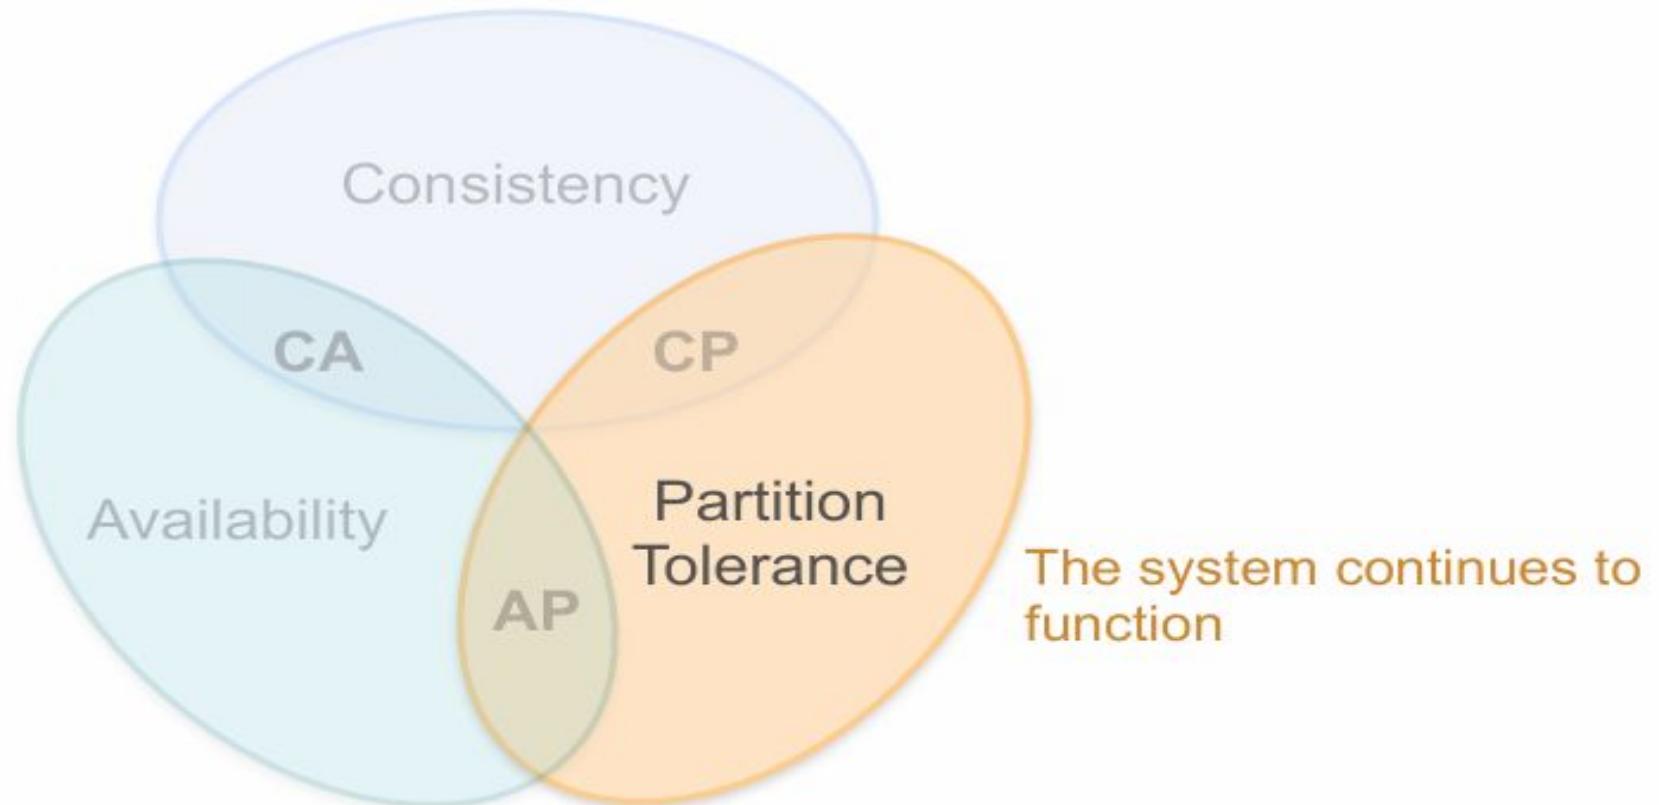
Every read reflects the latest write operation

# Distributed Storage Considerations – CAP Theorem



Every request will receive a response

# Distributed Storage Considerations – CAP Theorem



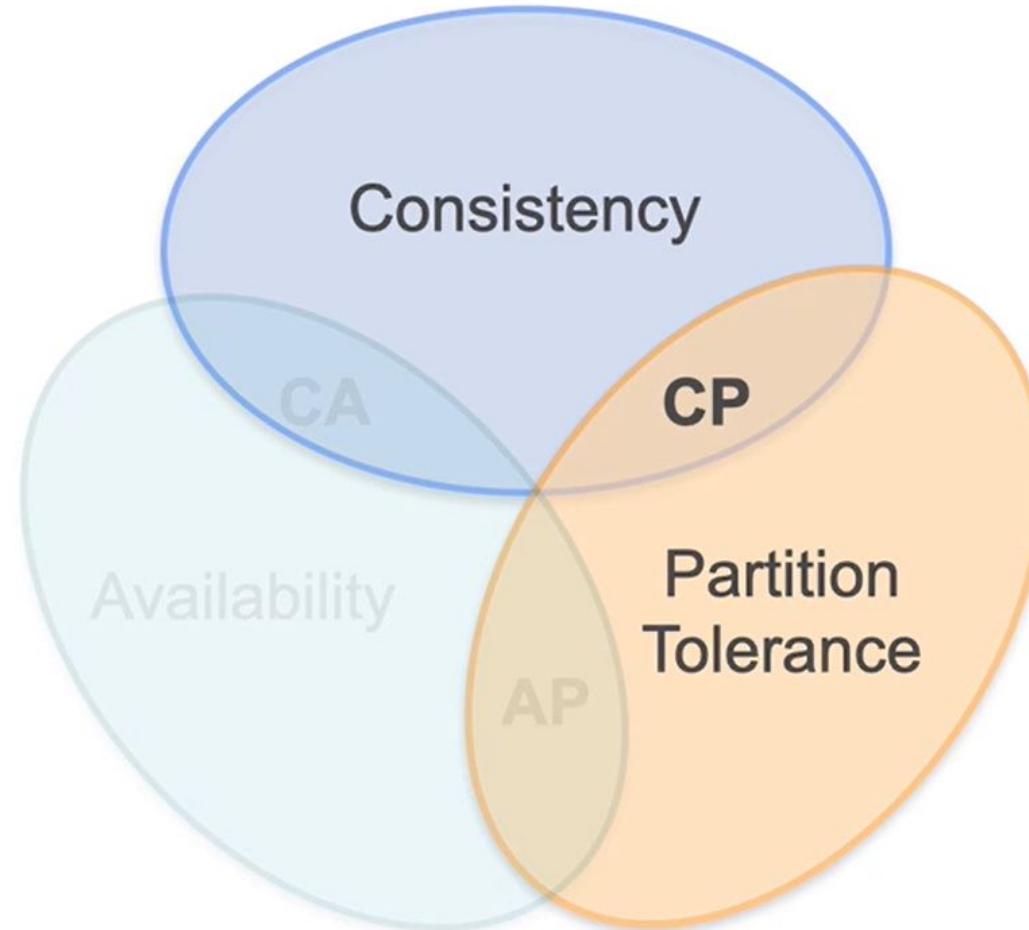
# Distributed Storage Considerations – CAP Theorem

## Scenario:

Accessing a node that's still being updated

## Option 1:

Cancel the request



# Distributed Storage Considerations – CAP Theorem

## Scenario:

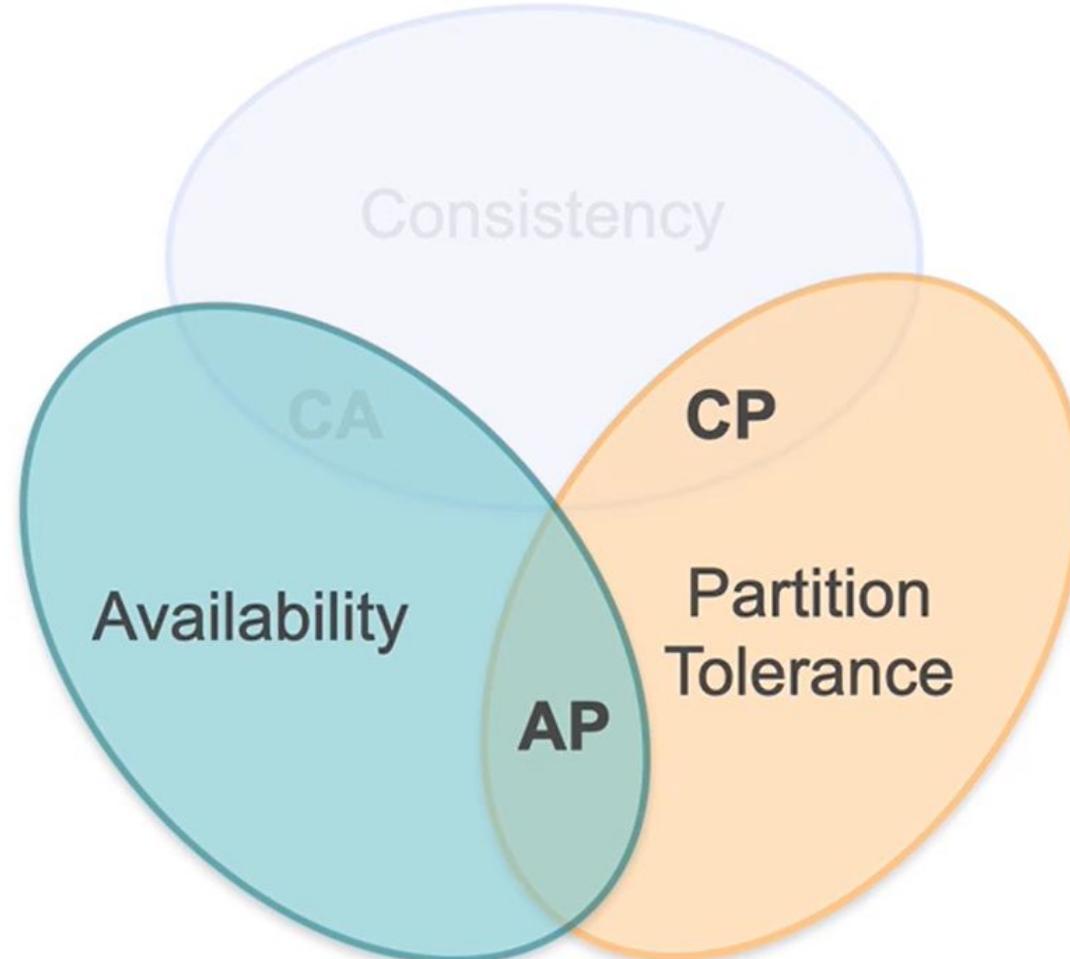
Accessing a node that's still being updated

## Option 1:

Cancel the request

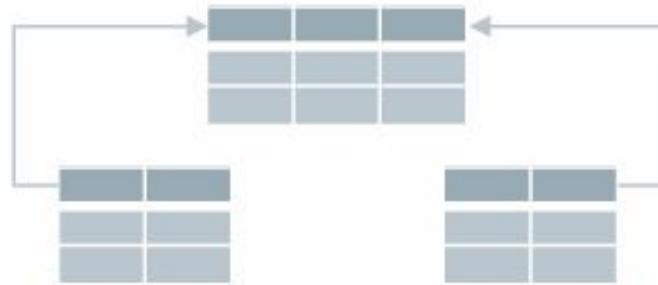
## Option 2:

Proceed with the read operation

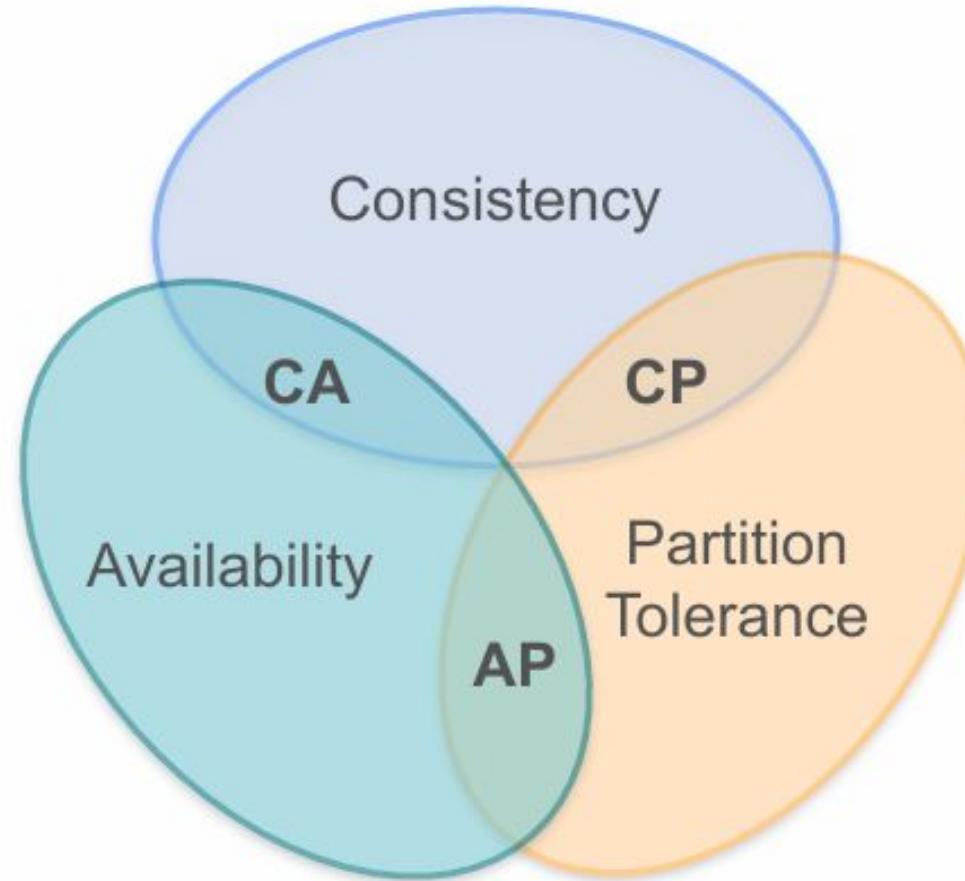


# Distributed storage considerations – ACID vs BASE

RDBMS



**ACID**  
compliant



NoSQL Databases



**BASE**  
principles

# Distributed storage considerations – ACID vs BASE

## ACID

compliant

**A**tomicity

**C**onsistency

**I**solation

**D**urability

## BASE

principles

**B**asically **A**vailable

**S**oft-state

**E**ventual Consistency

# Scenario

## Course 1



Data Scientist



**Main database (strong consistency):**  
Think of it as the **master copy** of your sales data.  
**Read replicas (eventual consistency):**  
These are **copies of the main database**.  
They update from the main database **after a short delay**.

### Main database instance (Strong consistency)



Read-replica of the prod database

- Ingest
- Transform
- Store
- Serve

### Read Replicas in RDS (Eventual consistency)



- Track changes in main database
- Update their own data



Amazon RDS

# Comparing Cloud Storage Options

**Object Storage**

**File Storage**

**Block Storage**

**Memory**

## Object Storage

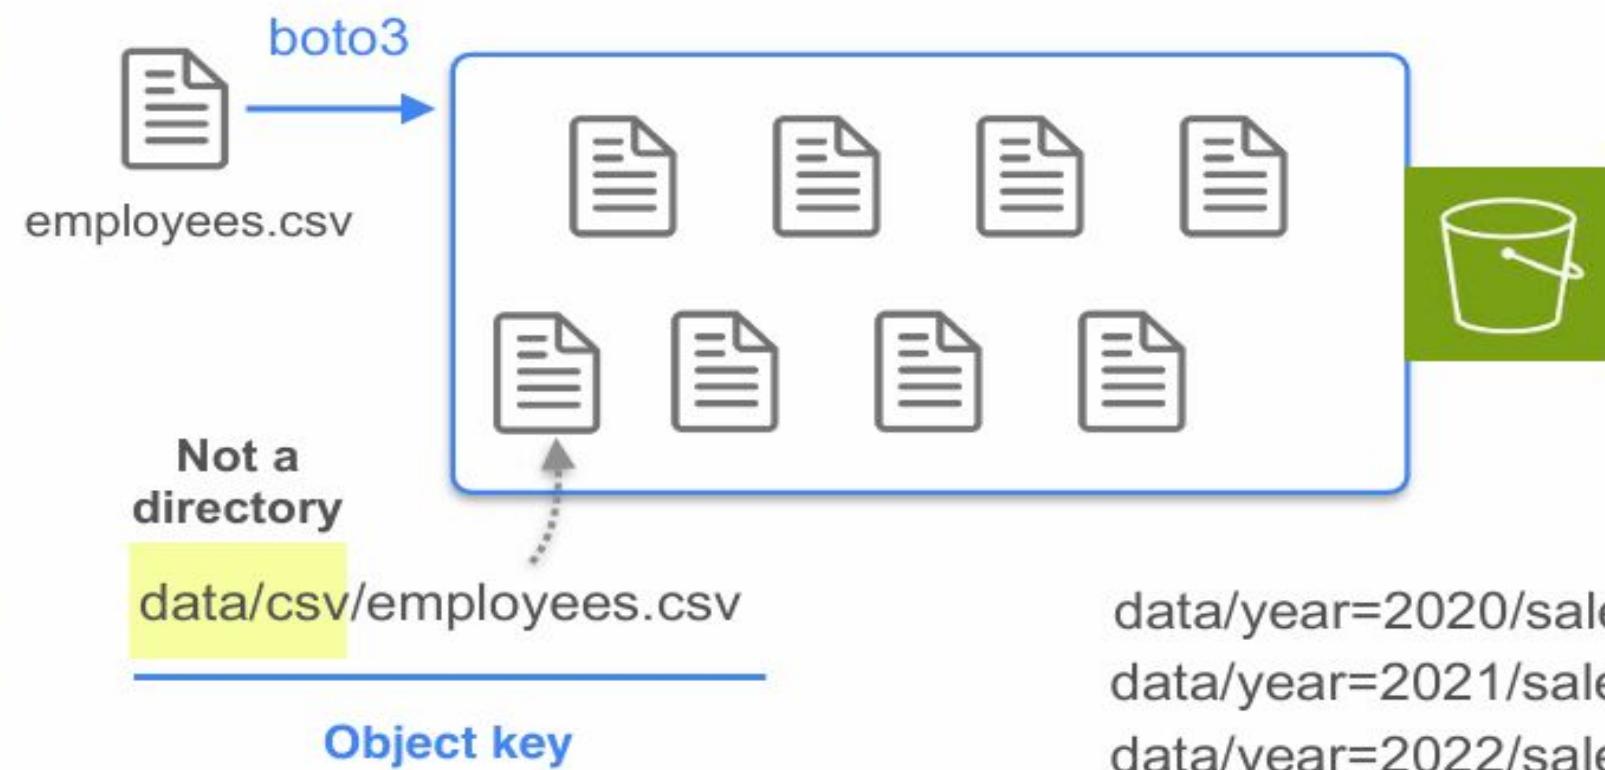
## File Storage

## Block Storage

## Memory

## Flat Structure

## Immutability



## Object Storage

## File Storage

## Block Storage

## Memory

## Flat Structure

## Immutability

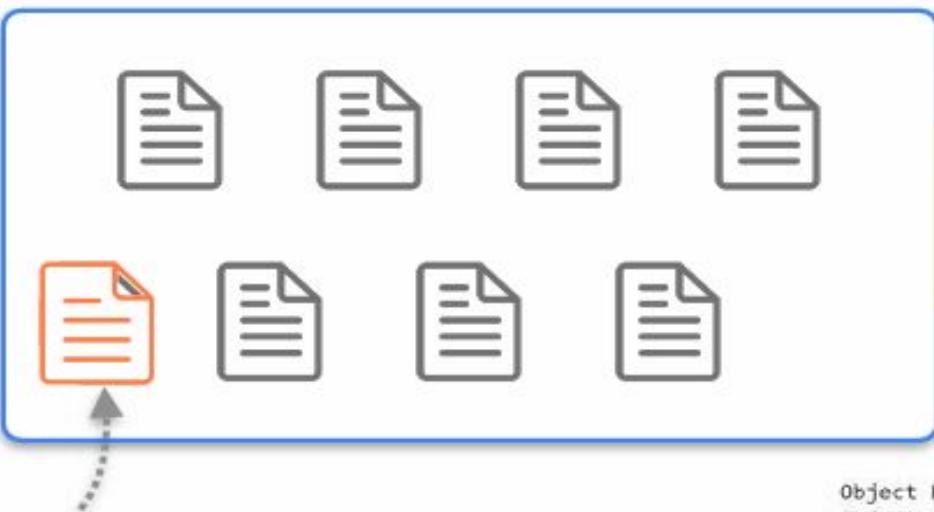
2. Modify employees' data



Not a  
directory

**data/csv/employees.csv**

**Object key**



1. Enable versioning



3. Use  
`list_object_version`

```
Object Key: data/csv/employees.csv
Object Version Id: q4A5B9CQZ10.u7.YKA7LabC_5GcBA.uZ
Is Latest: True
Last Modified: 2024-08-12 19:57:00+00:00
```

```
Object Key: data/csv/employees.csv
Object Version Id: WOOPNaVQTFHBl3CkIiRATRDCzt30Wq9k
Is Latest: False
Last Modified: 2024-08-12 19:39:56+00:00
```

## Object Storage

1. Navigate to the “data” directory
2. Explore the directory content and metadata
3. Explore how the data is modified in place

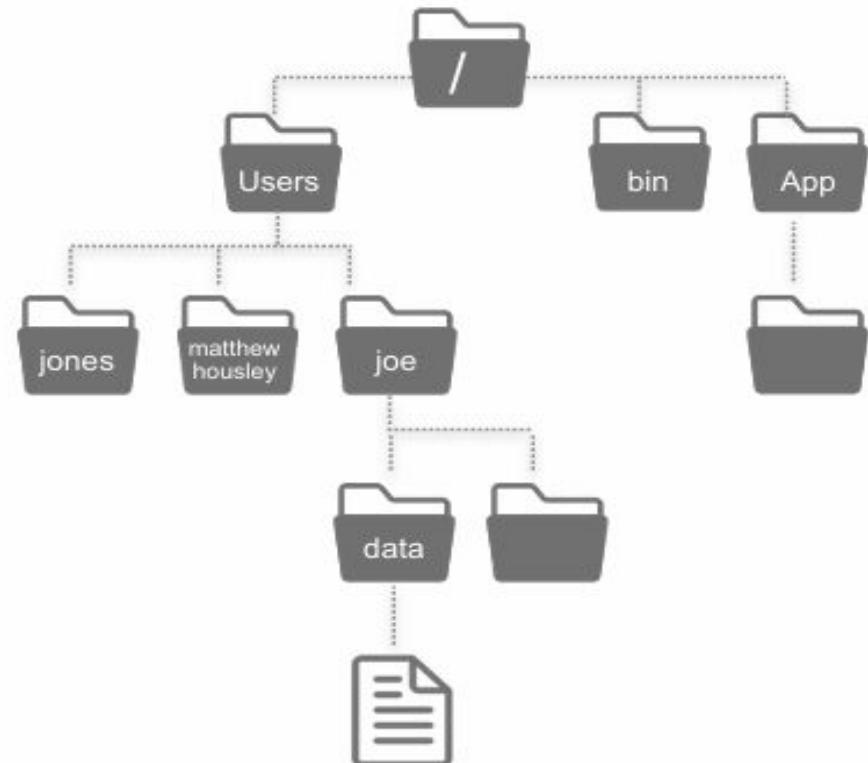
## File Storage

## Block Storage

## Memory

A directory

data/employees.csv



## Object Storage

## File Storage

## Block Storage

## Memory

1. Connect to the server
2. Send a file to the server

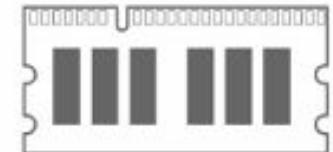


Server that emulates the behavior of block storage



## Object Storage

Transferring data from memory is faster than transferring data from disk.



## File Storage

Use the cache-pandas package:

- provides the “timed\_LRU\_cache” decorator to easily cache in memory pandas DataFrames generated by functions.

```
@timed_lru_cache(seconds=100, maxsize=None)
def read_csv_to_memory(path: str) -> pd.DataFrame:
    """Read CSV function with a cache decorator."""
    return pd.read_csv(path)
```

## Block Storage

## Memory

Compare the time it takes to read the file for the first time with the time it takes to read the same data stored in memory.

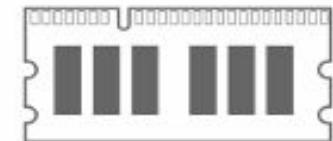
## Object Storage

## File Storage

## Block Storage

## Memory

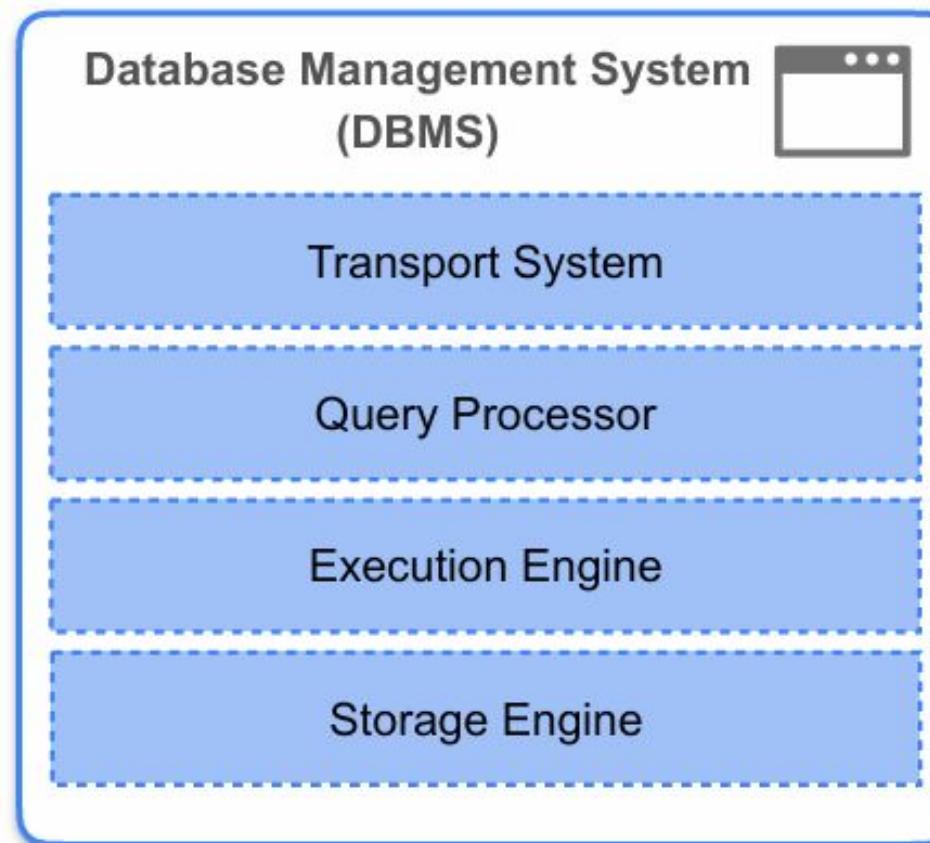
Monitor your memory storage capacity using htop command:



CPU Usage Report								
PID USER		PRI NI	VIRT	RES	S% CPU	%MEM	TIME+ COMMAND	
2880	root	20	0	218M	4385	3284.8	0.7 0.1	0:00.00 httpd
788	root	20	0	52944	17620	16364.5	0.0 0.9	0:00.72 /usr/lib/systemd/systemd-journald
1187	root	20	0	30880	18384	7572.5	0.0 0.5	0:00.07 /usr/lib/systemd/systemd-udevd
1387	systemd-re	20	0	21212	13876	18688.5	0.0 0.7	0:00.11 /usr/lib/systemd/systemd-resolved
1392	root	16 -4	28224	2300	1632.5	0.0 0.1	0:00.03 /sbin/auditd	
1393	root	16 -4	28224	2300	1632.5	0.0 0.1	0:00.00 /sbin/auditd	
1424	dbus	20	0	8388	3932	3344.5	0.0 0.2	0:00.02 /usr/bin/dbus-broker-launch --scope system --audit
1425	dbus	20	0	5264	2904	2404.5	0.0 0.1	0:00.19 dbus-broker --log 4 --controller 9 --machine-id ec20b71fb5d5e6c4dae675b053431827 --max
1426	root	20	0	15296	6444	5612.5	0.0 0.3	0:00.01 /usr/bin/systemd-inhibit --what=handle-suspend-key:handle-hibernate-key --who=noah --w
1428	root	20	0	81344	3824	2796.5	0.0 0.2	0:00.15 /usr/sbin/irqbalance --foreground
1429	libstorage	20	0	2756	2872	1988.5	0.0 0.1	0:00.03 /usr/bin/lsmd -d
1431	root	20	0	161M	6688	5396.5	0.0 0.3	0:37.39 /usr/sbin/rngd -f -x pkcs11 -x nist
1433	root	20	0	15888	7648	5672.5	0.0 0.4	0:00.02 /usr/lib/systemd/systemd-homed
1434	root	20	0	17744	9684	7520.5	0.0 0.5	0:00.16 /usr/lib/systemd/systemd-logind
1436	systemd-ne	20	0	2388	9228	7056.5	0.0 0.5	0:00.00 /usr/lib/systemd/systemd-networkd
1440	root	20	0	81344	3824	2796.5	0.0 0.2	0:00.00 /usr/sbin/irqbalance --foreground
1459	root	20	0	2668	1188	1828.5	0.0 0.1	0:00.00 /usr/sbin/acpid -f
1460	root	20	0	161M	6688	5396.5	0.0 0.3	0:18.47 /usr/sbin/rngd -f -x pkcs11 -x nist
1461	root	20	0	161M	6688	5396.5	0.0 0.3	0:18.87 /usr/sbin/rngd -f -x pkcs11 -x nist
1468	root	20	0	274M	3396	2632.5	0.0 0.2	0:00.00 /usr/sbin/gssproxy -D
1470	root	20	0	274M	3396	2632.5	0.0 0.2	0:00.00 /usr/sbin/gssproxy -D
1471	root	20	0	274M	3396	2632.5	0.0 0.2	0:00.00 /usr/sbin/gssproxy -D
1472	root	20	0	274M	3396	2632.5	0.0 0.2	0:00.00 /usr/sbin/gssproxy -D
1473	root	20	0	274M	3396	2632.5	0.0 0.2	0:00.00 /usr/sbin/gssproxy -D
1474	root	20	0	274M	3396	2632.5	0.0 0.2	0:00.00 /usr/sbin/gssproxy -D
1475	root	20	0	1776M	28824	8204.5	0.0 1.1	0:05.19 /usr/bin/containerd
1516	root	20	0	1776M	28824	8204.5	0.0 1.1	0:01.45 /usr/bin/containerd
1523	root	20	0	1776M	28824	8204.5	0.0 1.1	0:00.00 /usr/bin/containerd
1524	root	20	0	1776M	28824	8204.5	0.0 1.1	0:00.46 /usr/bin/containerd
1534	root	20	0	1776M	28824	8204.5	0.0 1.1	0:00.00 /usr/bin/containerd

# How Databases Store Data

# Database Management System



# Database Management System

## Storage Engine

- Serialization
- Arrangement of data on disk
- Indexing

## Modern Storage Engines

- Support the performance characteristics of SSDs
- Handle modern data types and structures
- Offer robust columnar storage support

## Average price of products purchased in the USA

```
SELECT AVG(price)
FROM my_table
WHERE country = "USA"
```

Order ID	Price	Product SKU	Quantity	Customer ID	Store ID	Country
1	40	458650	10	67t	3	Canada
2	23	902348	14	56t	3	Canada
3	45	1255893	12	87q	4	Canada
4	50	456829	13	98q	1	USA
5	34	568298	12	98q	1	USA
6	44	563783	4	67t	1	USA
7	22	234589	5	56u	2	Brazil
8	30	267895	12	78y	3	Canada
9	60	545659	14	13t	5	Mexico

.....

## Average price of products purchased in the USA

[Index](#)

```
SELECT AVG(price)
FROM my_table
WHERE country = "USA"
```

A data structure that helps you efficiently locate data

Order ID	Price	Product SKU	Quantity	Customer ID	Store ID	Country
1	40	458650	10	67t	3	Canada
2	23	902348	14	56t	3	Canada
3	45	1255893	12	87q	4	Canada
4	50	456829	13	98q	1	USA
5	34	568298	12	98q	1	USA
6	44	563783	4	67t	1	USA
7	22	234589	5	56u	2	Brazil
8	30	267895	12	78y	3	Canada
9	60	545659	14	13t	5	Mexico
.....						

scan all rows

## Average price of products purchased in the USA

```
SELECT AVG(price)
FROM my_table
WHERE country = "USA"
```

Order ID	Price	Product SKU	Quantity	Customer ID	Store ID	Country
1	40	458650	10	67t	3	Canada
2	23	902348	14	56t	3	Canada
3	45	1255893	12	87q	4	Canada
4	50	456829	13	98q	1	USA
5	34	568298	12	98q	1	USA
6	44	563783	4	67t	1	USA
7	22	234589	5	56u	2	Brazil
8	30	267895	12	78y	3	Canada
9	60	545659	14	13t	5	Mexico

...Scanning all rows:  $O(n)$

Binary search on rows:  $O(\log n)$

## Index

A data structure that helps you efficiently locate data

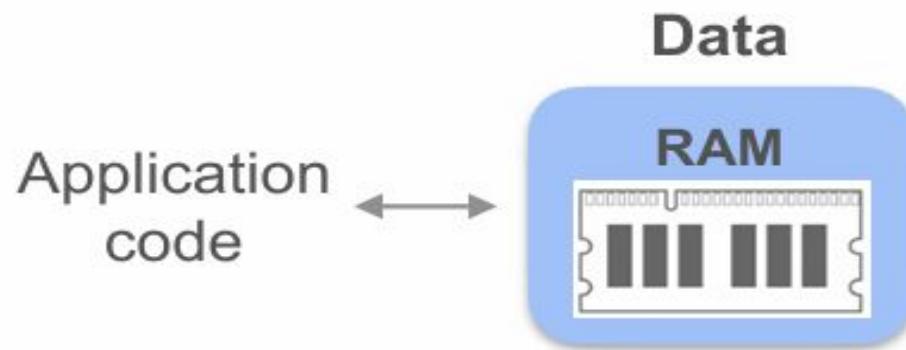
## Index table

Country	Row Address
Brazil	###
Canada	###
Mexico	###
USA	###
USA	###
USA	###



Use binary search to locate the USA rows

# In-Memory Storage Systems



- Excellent transfer speed and low latency
- Volatile
- Used to present data for ultra-fast retrieval:
  - Caching applications
  - Real-time bidding
  - Gaming leaderboards

## 1. Memcached

- Key-value store to cache database query results or API calls
- Used when it's acceptable for data to be lost

## 2. Redis

- Key-value store that supports more complex data types
- Supports high-performance applications that can tolerate minor data loss

# Row vs Column Storage

# Row-Oriented Storage

Order ID	Price	Product SKU	Quantity	Customer ID
1	40	45865	10	67t
2	23	90234	14	56t
3	45	12558	12	87q
4	50	45682	13	98q

↓  
Stores data  
row by row

## Physical Storage

1	40	45865	10	67t	2	23	90234	14	56t	...	4	50	45682	13	98q
bytes representing the 1st row					bytes representing the 2nd row					bytes representing the last row					

# Row-Oriented Storage

Order ID	Price	Product SKU	Quantity	Customer ID
1	40	45865	10	67t
2	23	90234	14	56t
3	45	12558	12	87q
4	50	45682	13	98q

**Row Storage is perfect for OLTP**

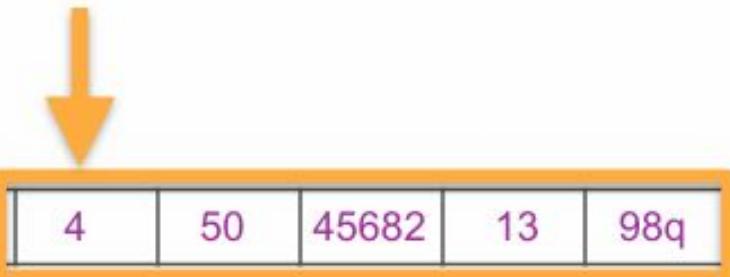
*Perform read and write operations with low latency*

← Locate this order

Stores data row by row

## Physical Storage

1	40	45865	10	67t	2	23	90234	14	56t	...	4	50	45682	13	98q
---	----	-------	----	-----	---	----	-------	----	-----	-----	---	----	-------	----	-----



# Row-Oriented Storage

Order ID	Price	Product SKU	Quantity	Customer ID
1	40	45865	10	67t
2	23	90234	14	56t
3	45	12558	12	87q
4	50	45682	13	98q

Stores data  
row by row

**Analytical queries** focus on summarizing or aggregating columns

- Total revenue?
- Most popular product?
- Average quantity?

## Physical Storage

1	40	45865	10	67t	2	23	90234	14	56t	...	4	50	45682	13	98q
---	----	-------	----	-----	---	----	-------	----	-----	-----	---	----	-------	----	-----

# Row-Oriented Storage

Order ID	Price	Product SKU	Quantity	Customer ID	...
1	40	45865	10	67t	...
2	23	90234	14	56t	...
3	45	12558	12	87q	...
4	50	45682	13	98q	...
...	...	...	...	...	...

1 million rows

30 columns

100 bytes per entry

```
SELECT SUM(price)
FROM my_table
```

## Physical Storage

1	40	45865	10	67t	2	23	90234	14	56t	...	4	50	45682	13	98q	...
---	----	-------	----	-----	---	----	-------	----	-----	-----	---	----	-------	----	-----	-----

# Row-Oriented Storage

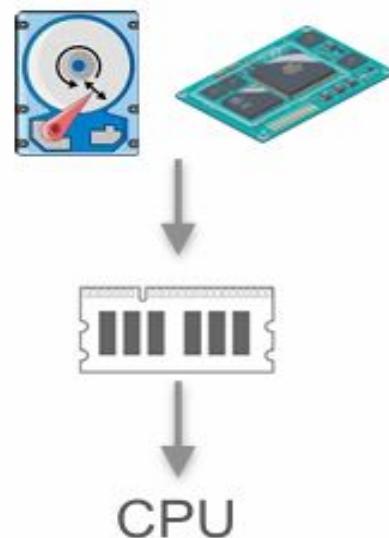
Order ID	Price	Product SKU	Quantity	Customer ID	...
1	40	45865	10	67t	...
2	23	90234	14	56t	...
3	45	12558	12	87q	...
4	50	45682	13	98q	...
...	...	...	...	...	...

1 million rows

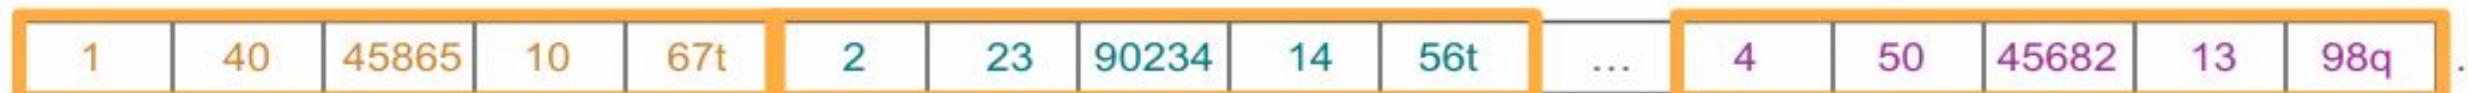
30 columns

100 bytes per entry

```
SELECT SUM(price)
FROM my_table
```



## Physical Storage



# Row-Oriented Storage

Order ID	Price	Product SKU	Quantity	Customer ID	...
1	40	45865	10	67t	...
2	23	90234	14	56t	...
3	45	12558	12	87q	...
4	50	45682	13	98q	...
				...	...

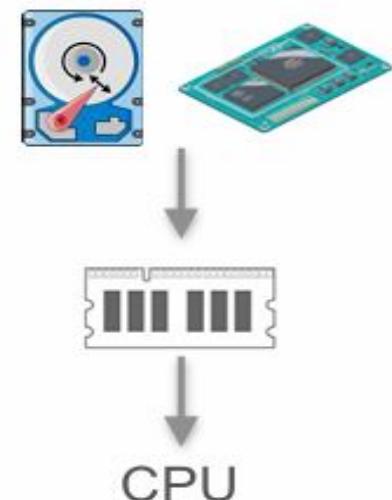
Total size = rows × columns × size per entry

1 million rows  $\times$  30 columns  $\times$  100 bytes per entry = 3 GB

Data transfer speed: 200 MB/s

Total transfer time?  $\frac{3\text{GB or } 3000\text{ MB}}{200\text{ MB/s}} = 15\text{ s}$

```
SELECT SUM(price)
FROM my_table
```



# Row-Oriented Storage

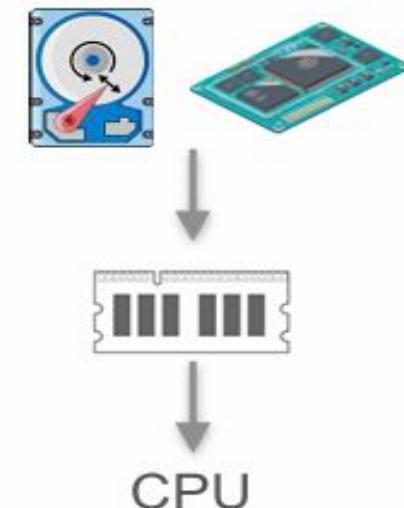
Order ID	Price	Product SKU	Quantity	Customer ID	...
1	40	45865	10	67t	...
2	23	90234	14	56t	...
3	45	12558	12	87q	...
4	50	45682	13	98q	...
...	...	...	...	...	...

$$1 \text{ billion rows} \times 30 \text{ columns} \times 100 \text{ bytes per entry} = 3000 \text{ GB}$$

Data transfer speed: 200 MB/s

$$\text{Total transfer time? } \frac{3000 \text{ GB}}{200 \text{ MB/s}} = 4 \text{ hours !}$$

```
SELECT SUM(price)
FROM my_table
```



# Column-Oriented Storage

Order ID	Price	Product SKU	Quantity	Customer ID	...
1	40	45865	10	67t	...
2	23	90234	14	56t	...
3	45	12558	12	87q	...
4	50	45682	13	98q	...
...	...	...	...	...	...

↓  
Stores data  
Column by column

## Physical Storage

1	2	3	4	40	23	45	50	45865	90234	12558	45682	...
---	---	---	---	----	----	----	----	-------	-------	-------	-------	-----

bytes representing 1st column

bytes representing 2nd column

bytes representing 3rd column

# Column-Oriented Storage — Suitable for OLAP systems!

Order ID	Price	Product SKU	Quantity	Customer ID	...
1	40	45865	10	67t	...
2	23	90234	14	56t	...
3	45	12558	12	87q	...
4	50	45682	13	98q	...
...	...	...	...	...	...

1 billion rows

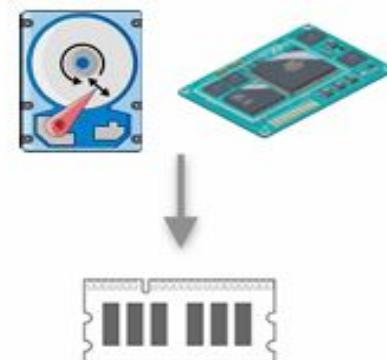
30 columns

100 bytes per entry = 100 GB

Data transfer speed: 200 MB/s

Total transfer time?  $\frac{100 \text{ GB or } 100,000 \text{ MB}}{200 \text{ MB/s}} = 8.33 \text{ minutes}$

```
SELECT SUM(price)
FROM my_table
```



## Row-oriented Storage

Transfer 1 billion rows  
from disk to memory

4 hours

# Column-Oriented Storage

Order ID	Price	Product SKU	Quantity	Customer ID	...
1	40	45865	10	67t	...
2	23	90234	14	56t	...
3	45	12558	12	87q	...
4	50	45682	13	98q	...
...	...	...	...	...	...

Terrible for  
transactional  
workloads!

Stores data  
Column by column

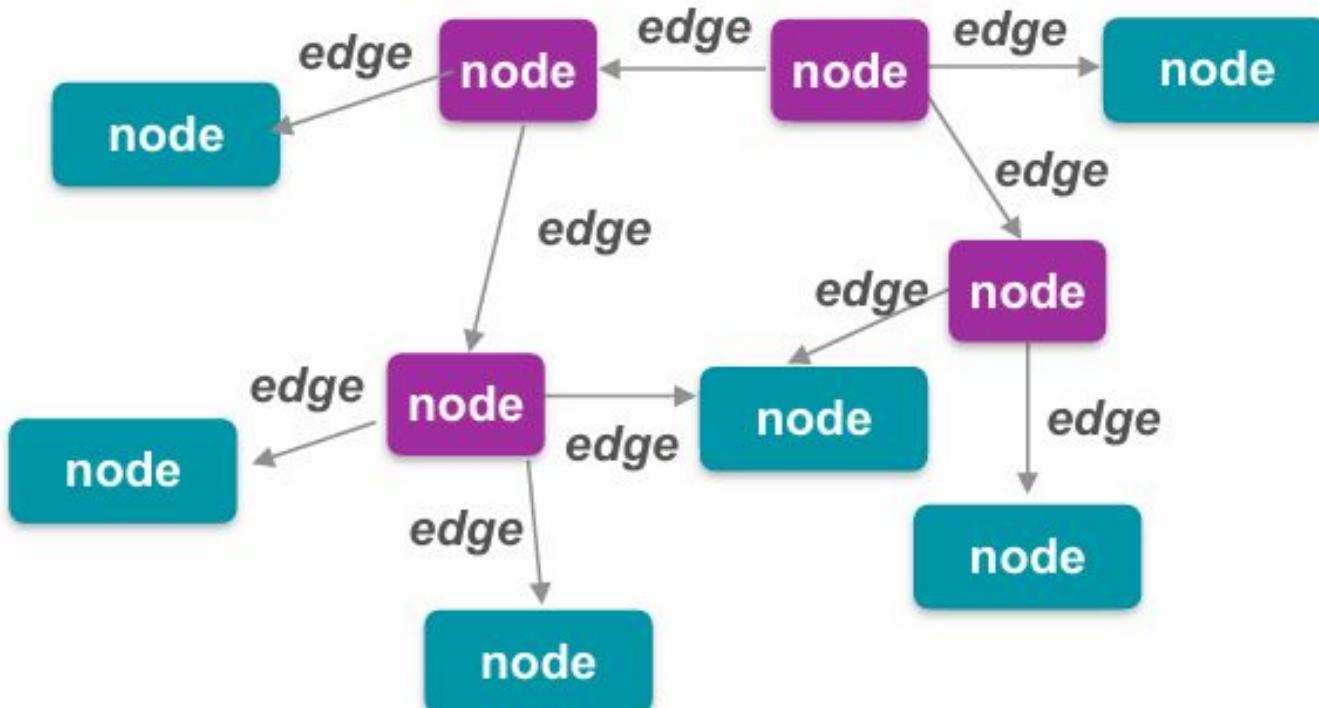
**Physical Storage**

Deserialize the column, modify it,  
then write it back to storage



# Graph Databases

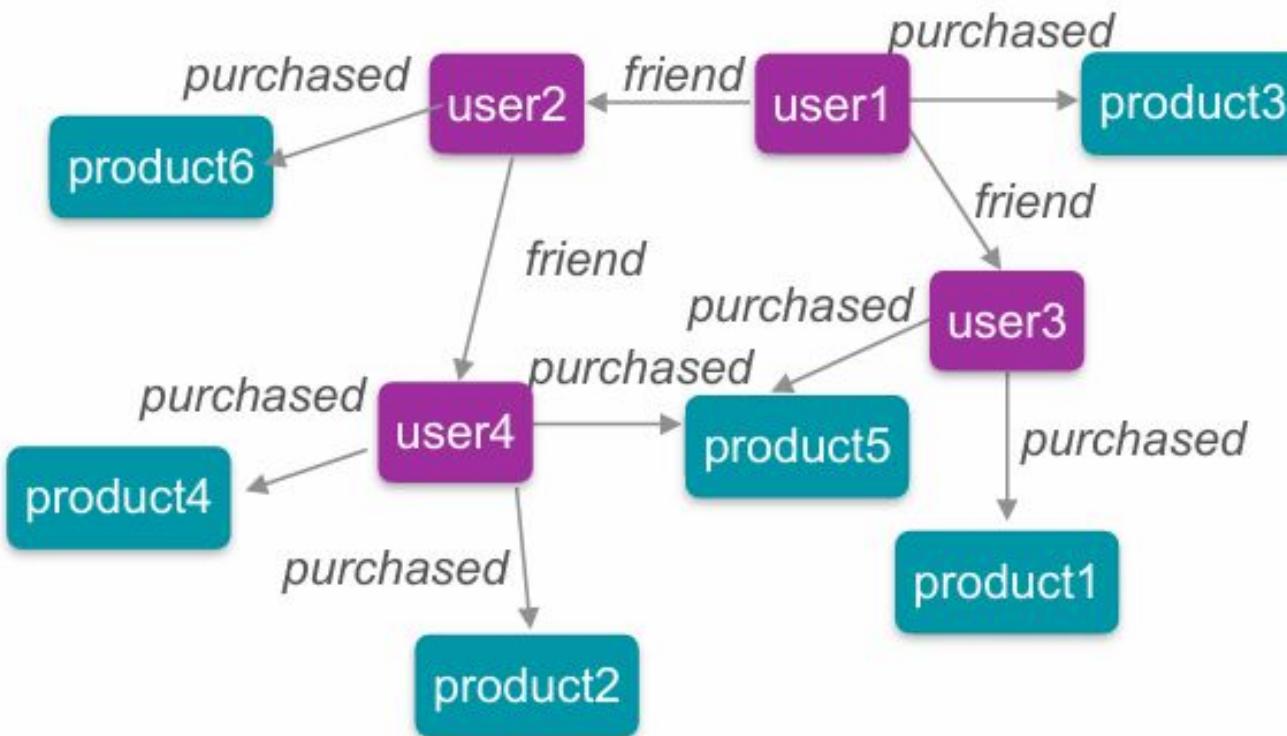
# Graph Database



- Nodes represent data items
  - Edges represent connection between the data items
  - Graph databases model complex connections between data entities

# Graph Database

**Relationships are first-class citizens**

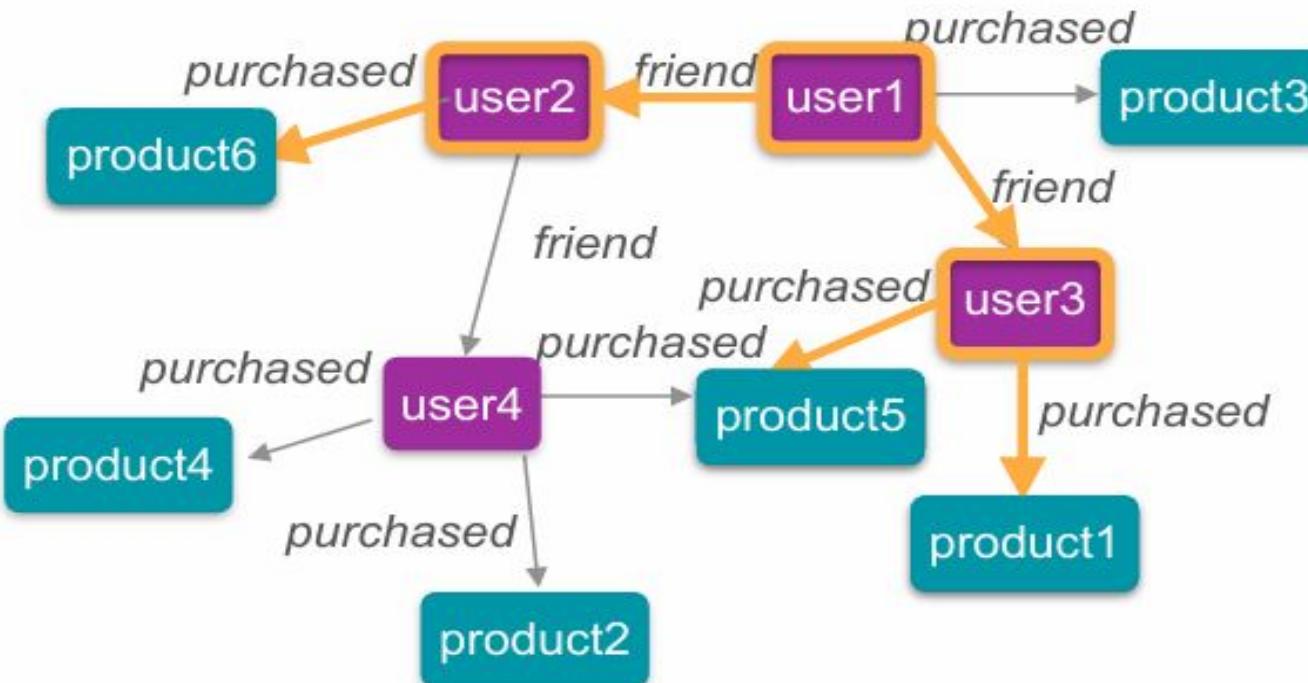


**Relational database**

purchase		friendship	
user	product	user	friend
user1	product3	user1	user3
user2	product6	user1	user2
user3	product1	user2	user4
user3	product5		
user4	product5		
user4	product4		
user4	product2		

# Querying Data

Traverse the graph to query relationships



Recommendations  
for user1

Relational database

purchase

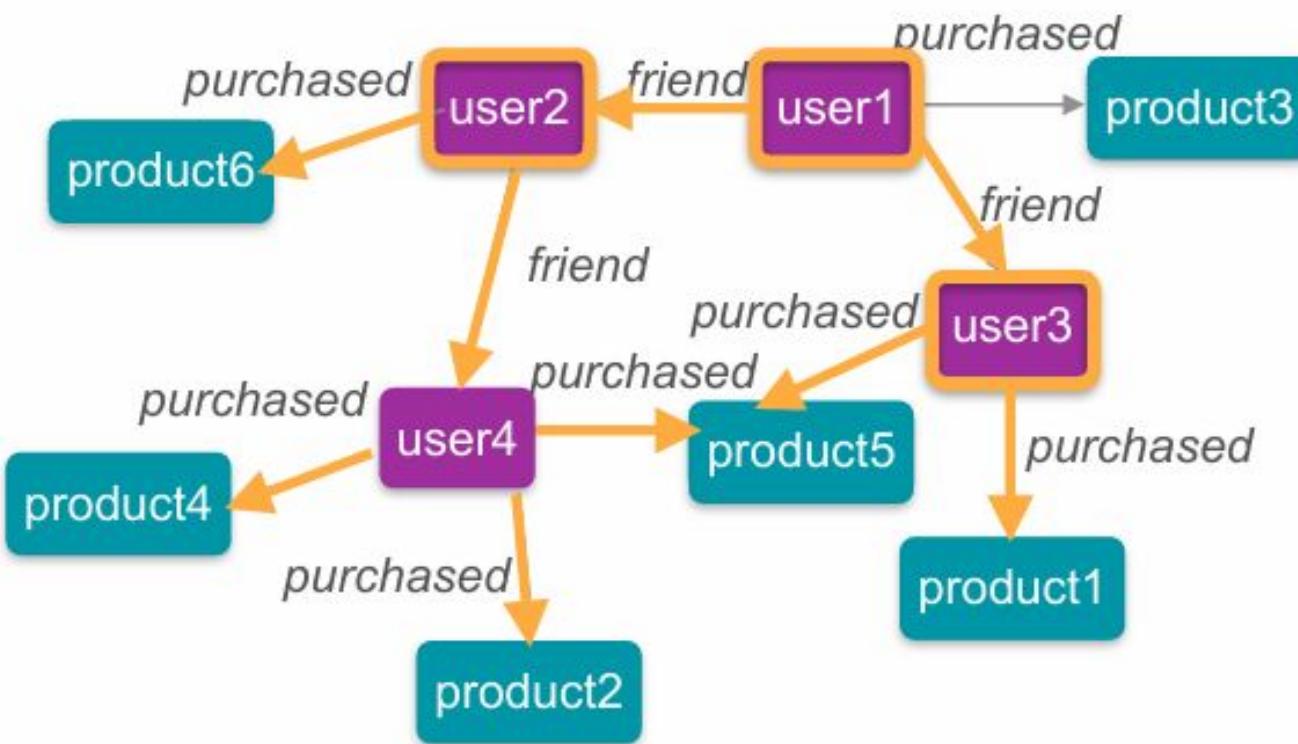
user	product
user1	product3
user2	product6
user3	product1
user3	product5
user4	product5
user4	product4
user4	product2

friendship

user	friend
user1	user3
user1	user2
user2	user4

# Querying Data

Traverse the graph to query relationships



Recommendations  
for user1

product1

product5

product6

Relational database

purchase

user	product
user1	product3
user2	product6
user3	product1
user3	product5
user4	product5
user4	product4
user4	product2

friendship

user	friend
user1	user3
user1	user2
user2	user4

Less efficient in  
querying complex  
relationships!

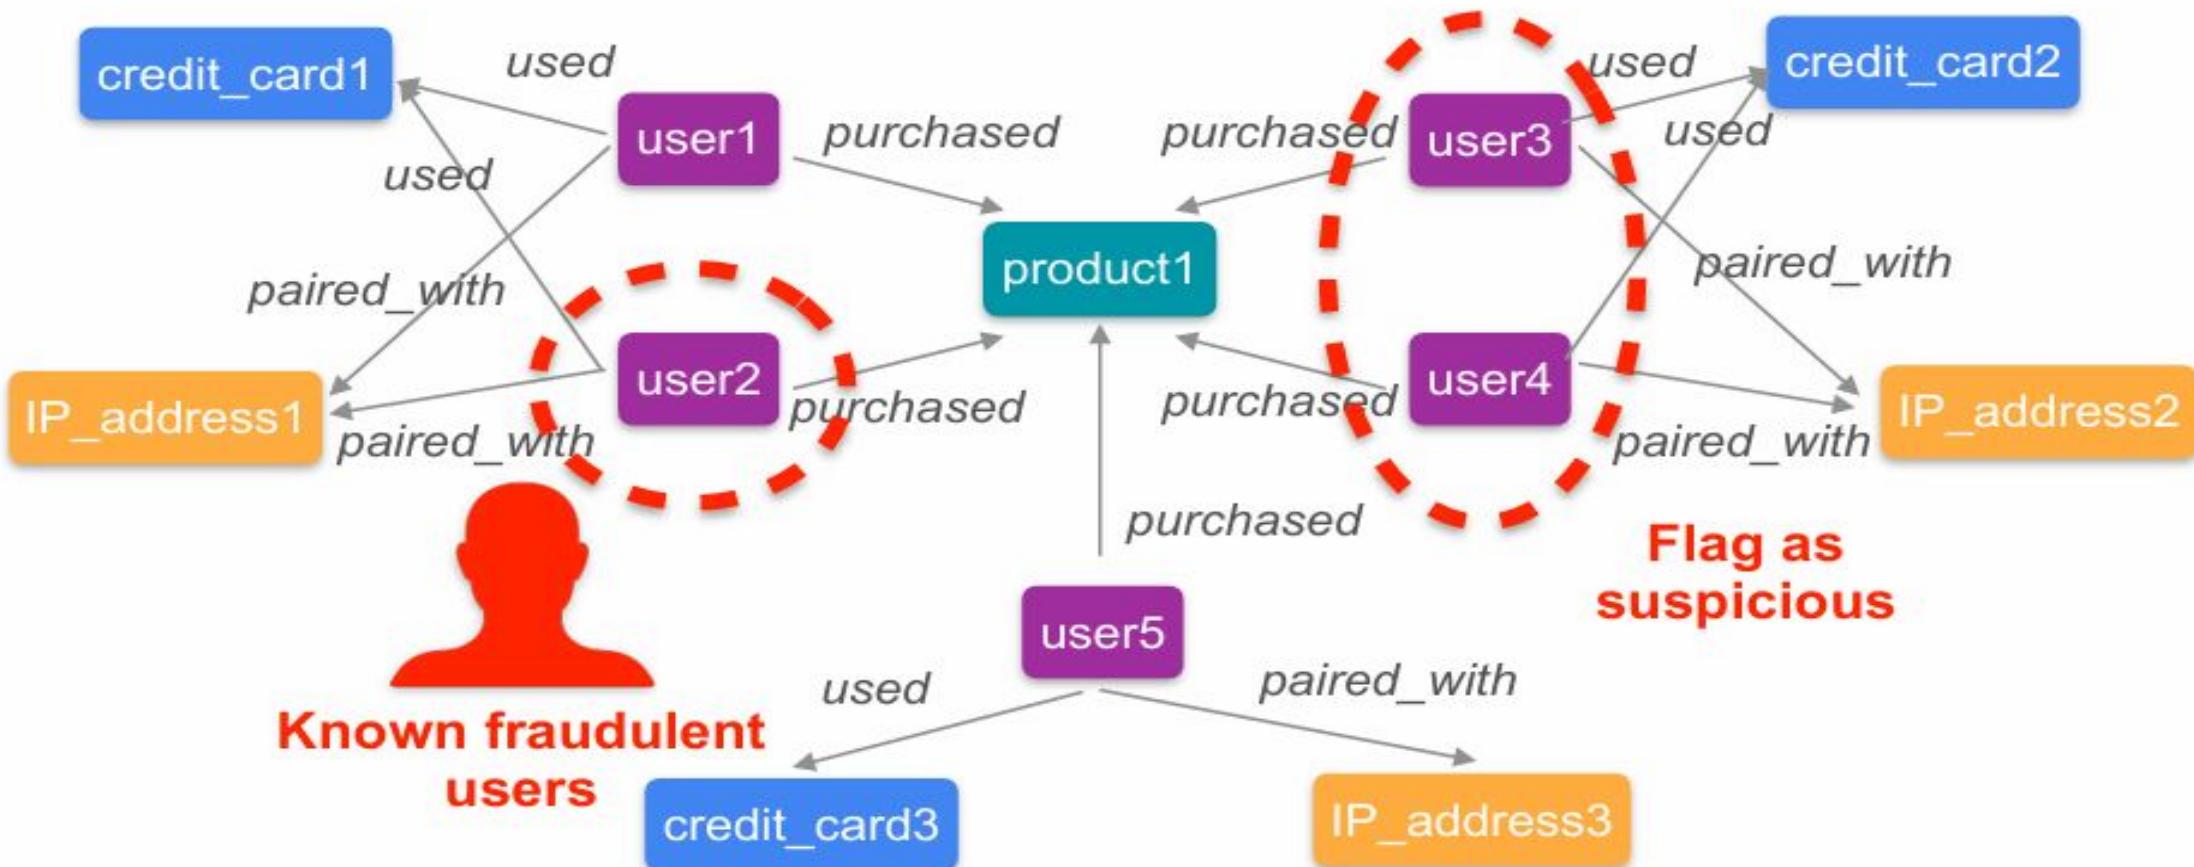
SELECT DISTINCT purchase.product  
FROM friendship

JOIN purchase ON friendship.friend = purchase.user  
WHERE friendship.user = 'user1'

# Graph Database - Use Cases

- Recommending products
- Modeling social networks
- Representing network and IT operations
- Simulating supply chains logistics
- Tracing data lineage

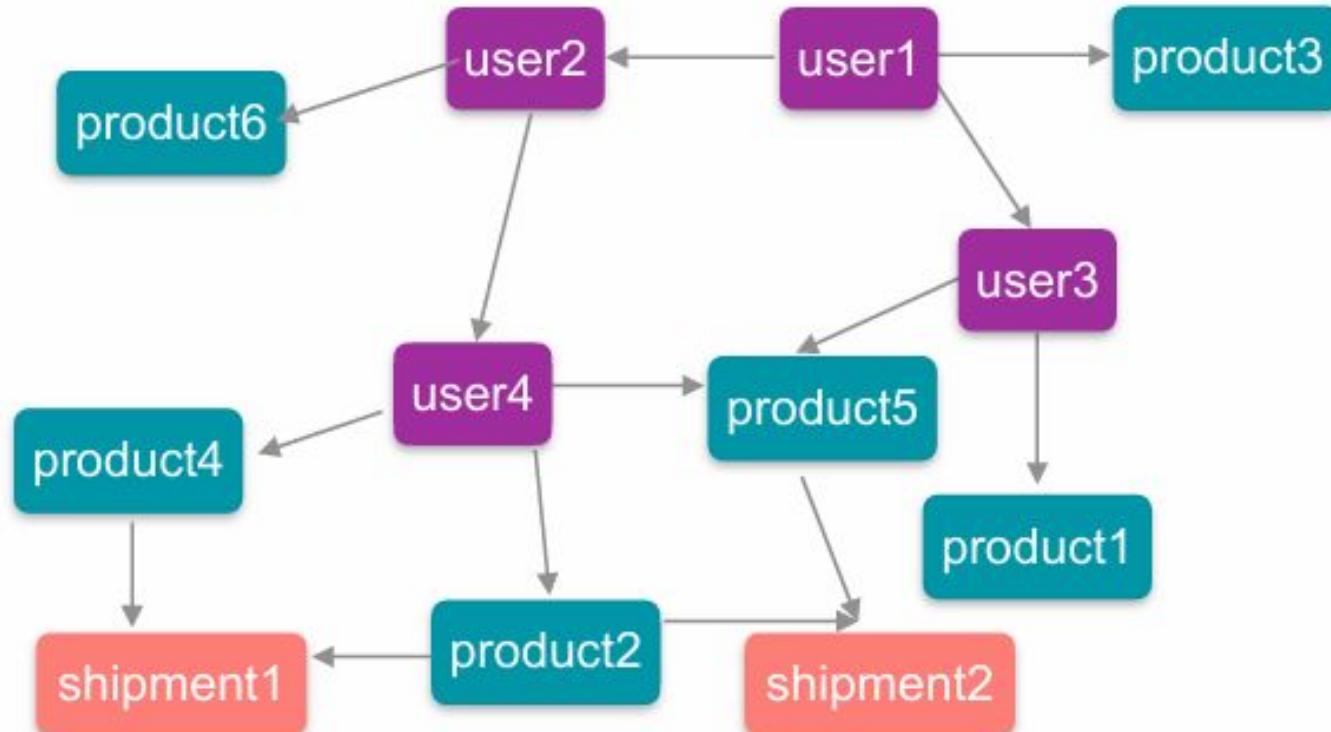
# Use Case - Fraud Detection



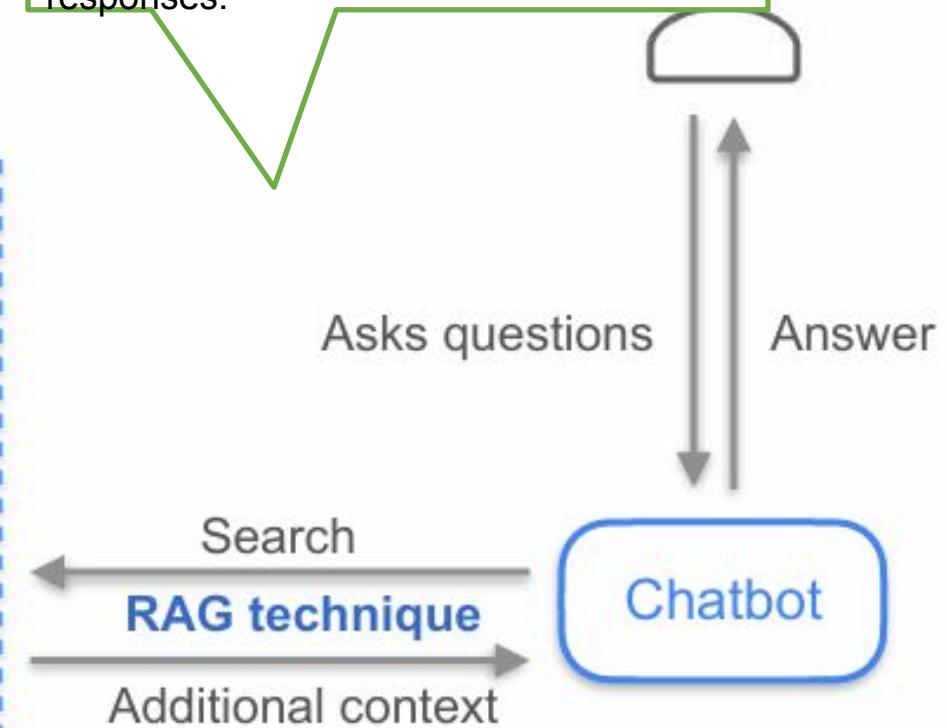
# Use Case - Knowledge Graph

## Knowledge Graph

Connect diverse data from disparate sources



A knowledge graph connects e-commerce entities like customers, products, and orders, enabling chatbots to find relevant relationships quickly. Using RAG, the LLM accesses this fresh, company-specific data to provide accurate and up-to-date responses.



# Graph Databases

## Examples of Graph Databases



Amazon Neptune

## Examples of Graph Query Language

Cypher

Gremlin

SparQL

# Vector Databases

## Vector data

Consists of numerical values arranged in an array



Image



0.5	0.5	0.7	0.5	0.3	0.1	0.7	0.2
0.4	0.8	0.9	0.1	0.3	0.1	0.4	0.2
0.3	0.5	0.7	0.8	0.3	0.1	0.6	0.2

## Vector embeddings

Capture semantic meaning of an item, like a text document or image



- Can convert an entire database of docs or text into embeddings
- Embeddings help you more efficiently find and retrieve similar items
- Example: Finding similar text
  - Compute embeddings for the query item
  - Database returns similar vectors (based on closeness)

# Distance Metric

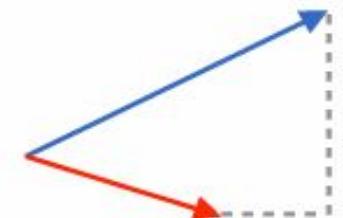
**Vector database uses a distance metric to find similar vectors**



Euclidean distance



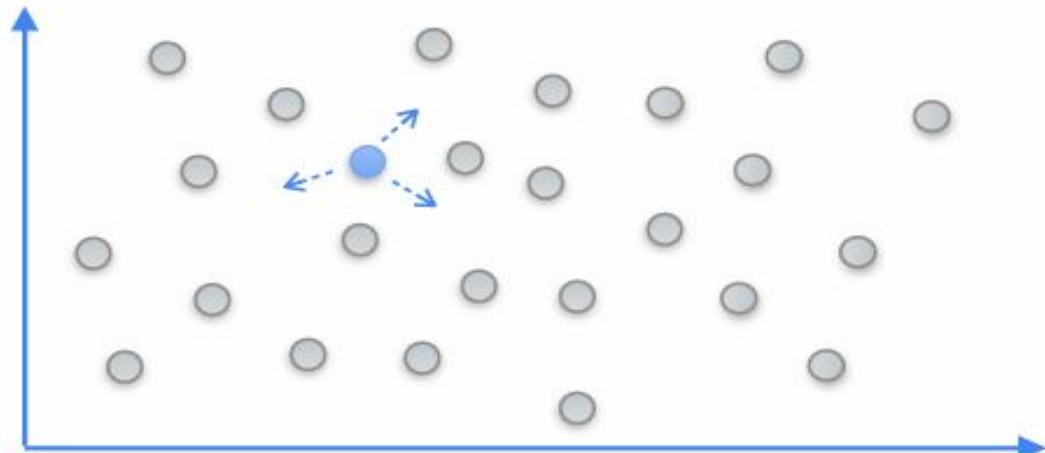
Cosine distance



Manhattan distance

# Similarity Search - Popular Algorithm

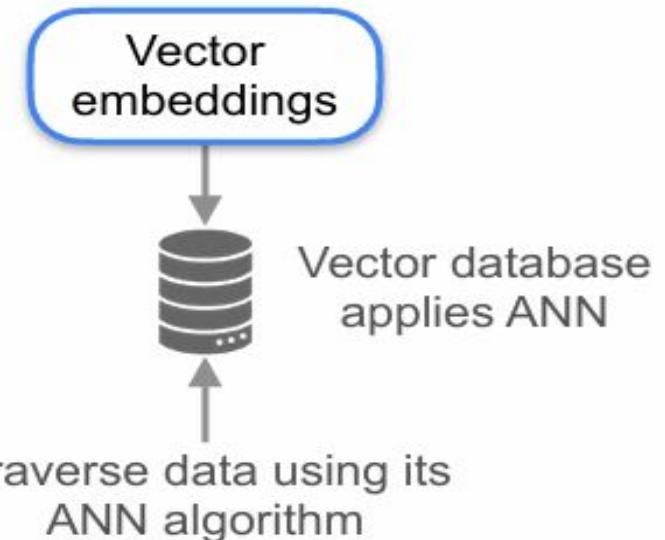
## K-nearest neighbors (KNN)



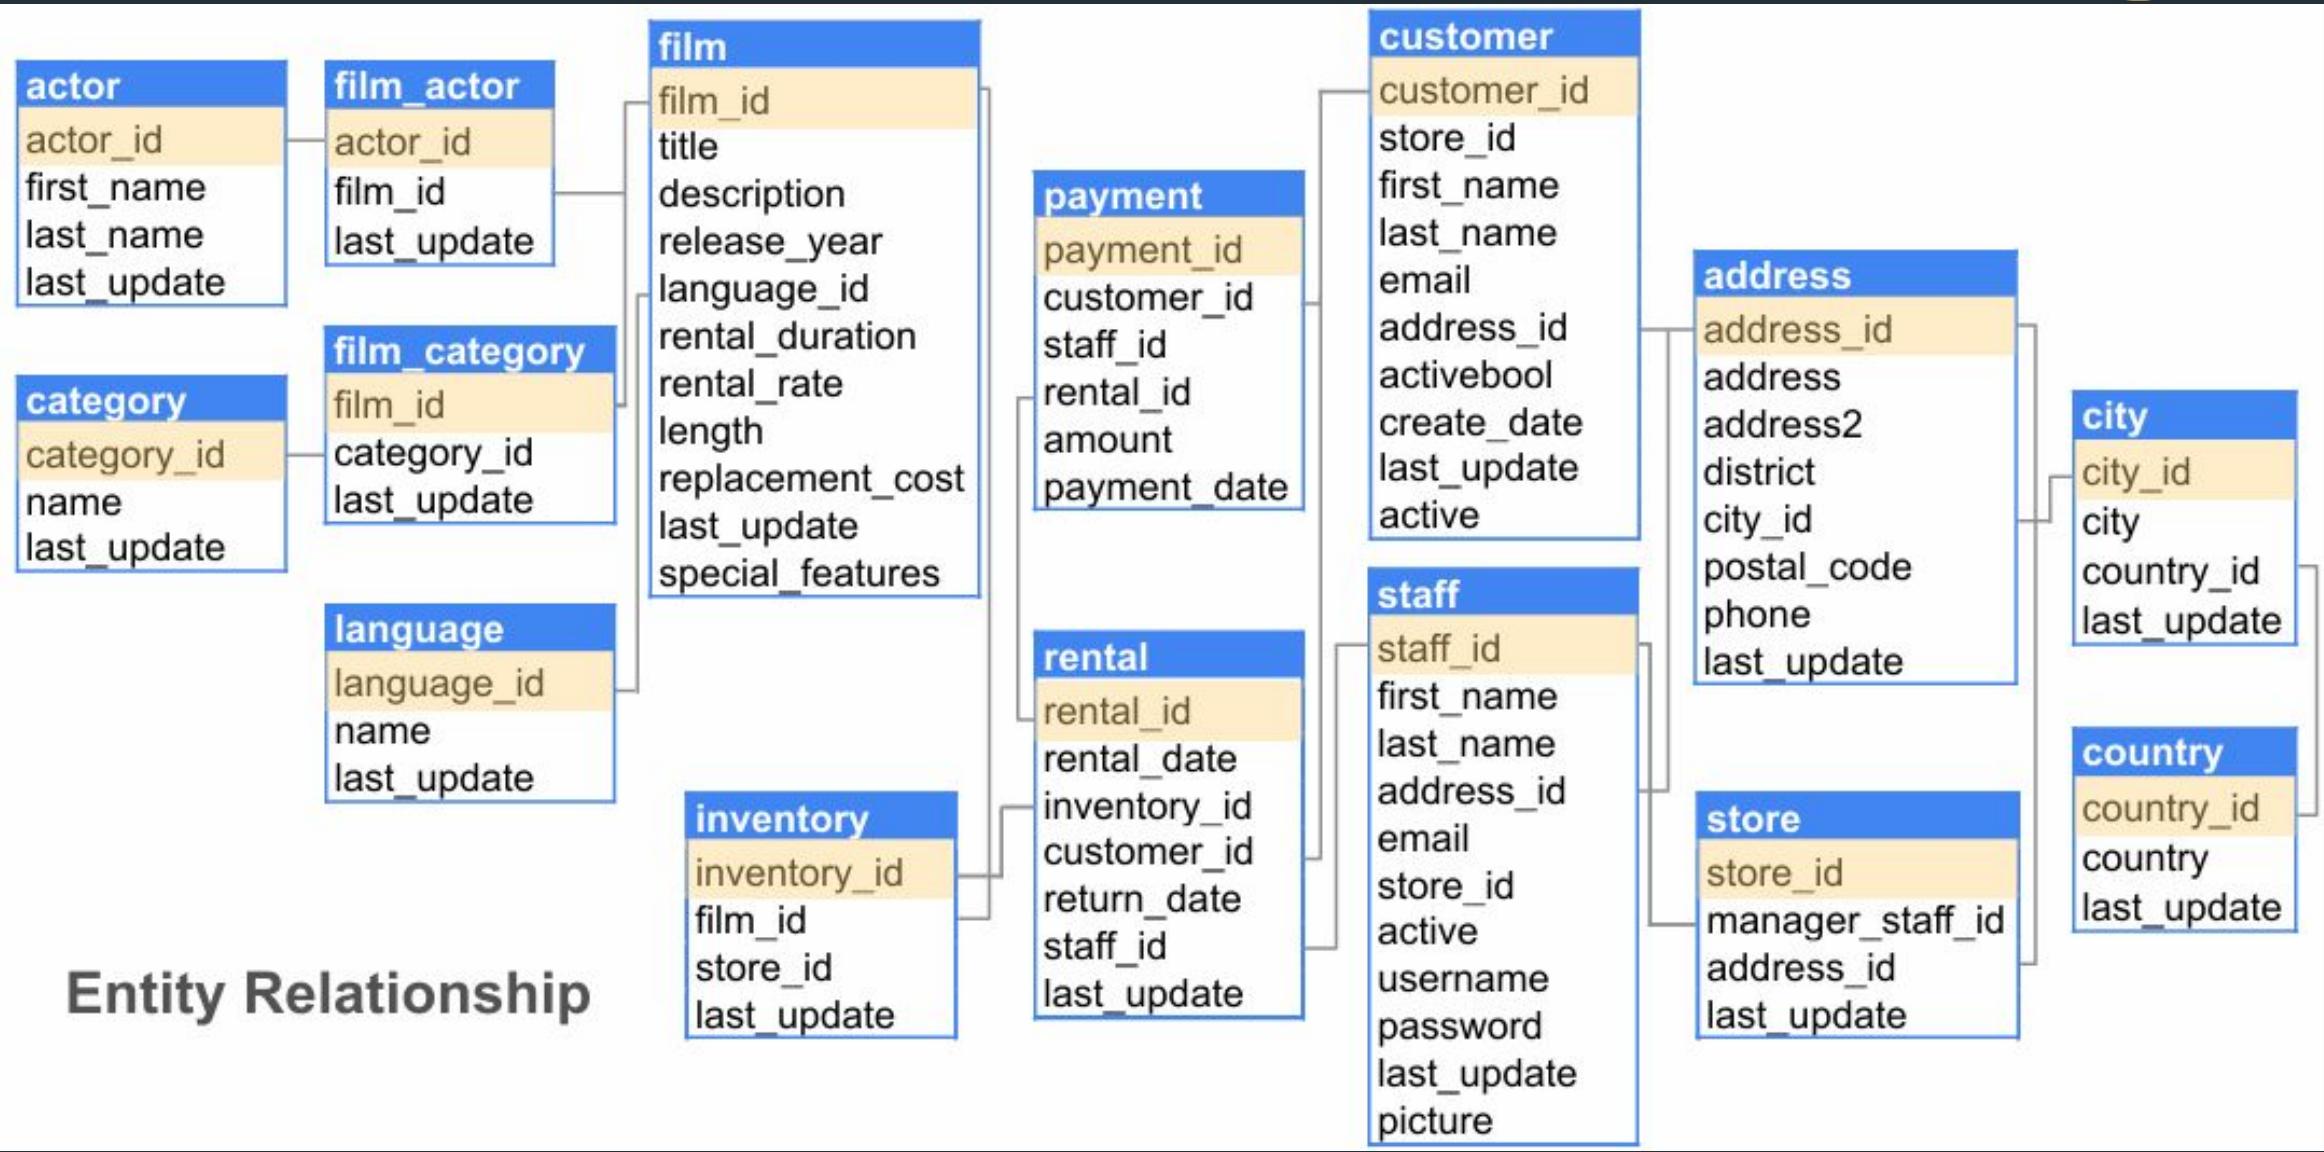
- Calculates distance to all vector embeddings
- Becomes inefficient when the data size increases
- Suffers from the curse of dimensionality

## ANN (Approximate Nearest Neighbors)

- Find a good guess for the nearest neighbors
- More efficient than K-NN
- Vector databases are built to support ANN algorithms

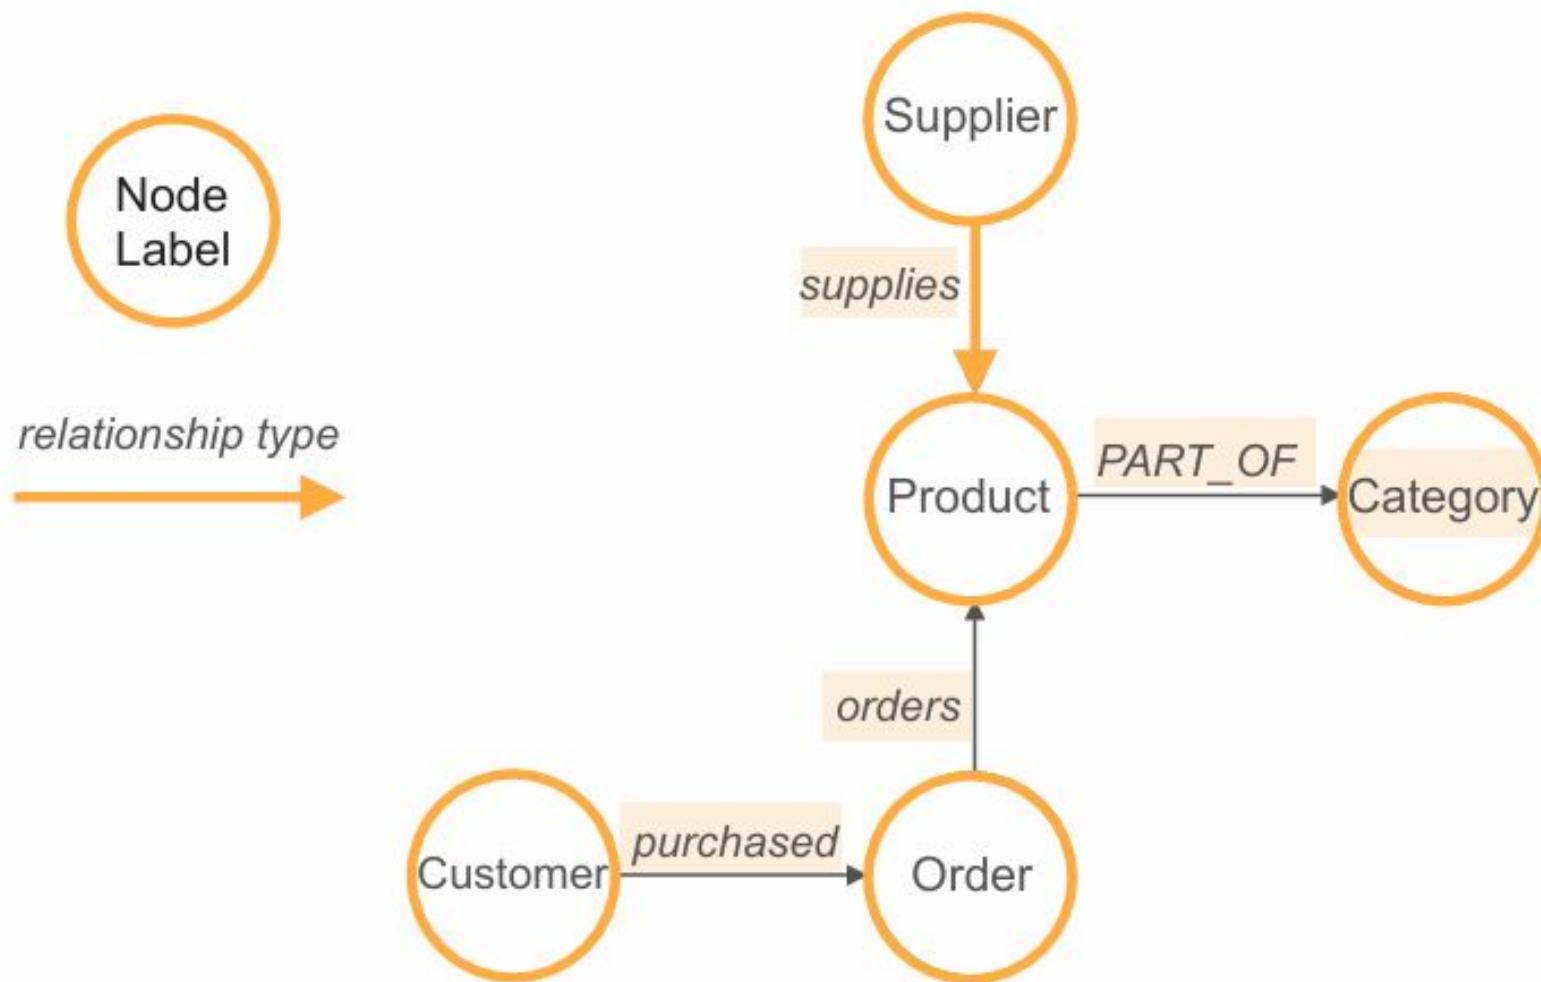


# Neo4j Graph Database & Cypher Query Language

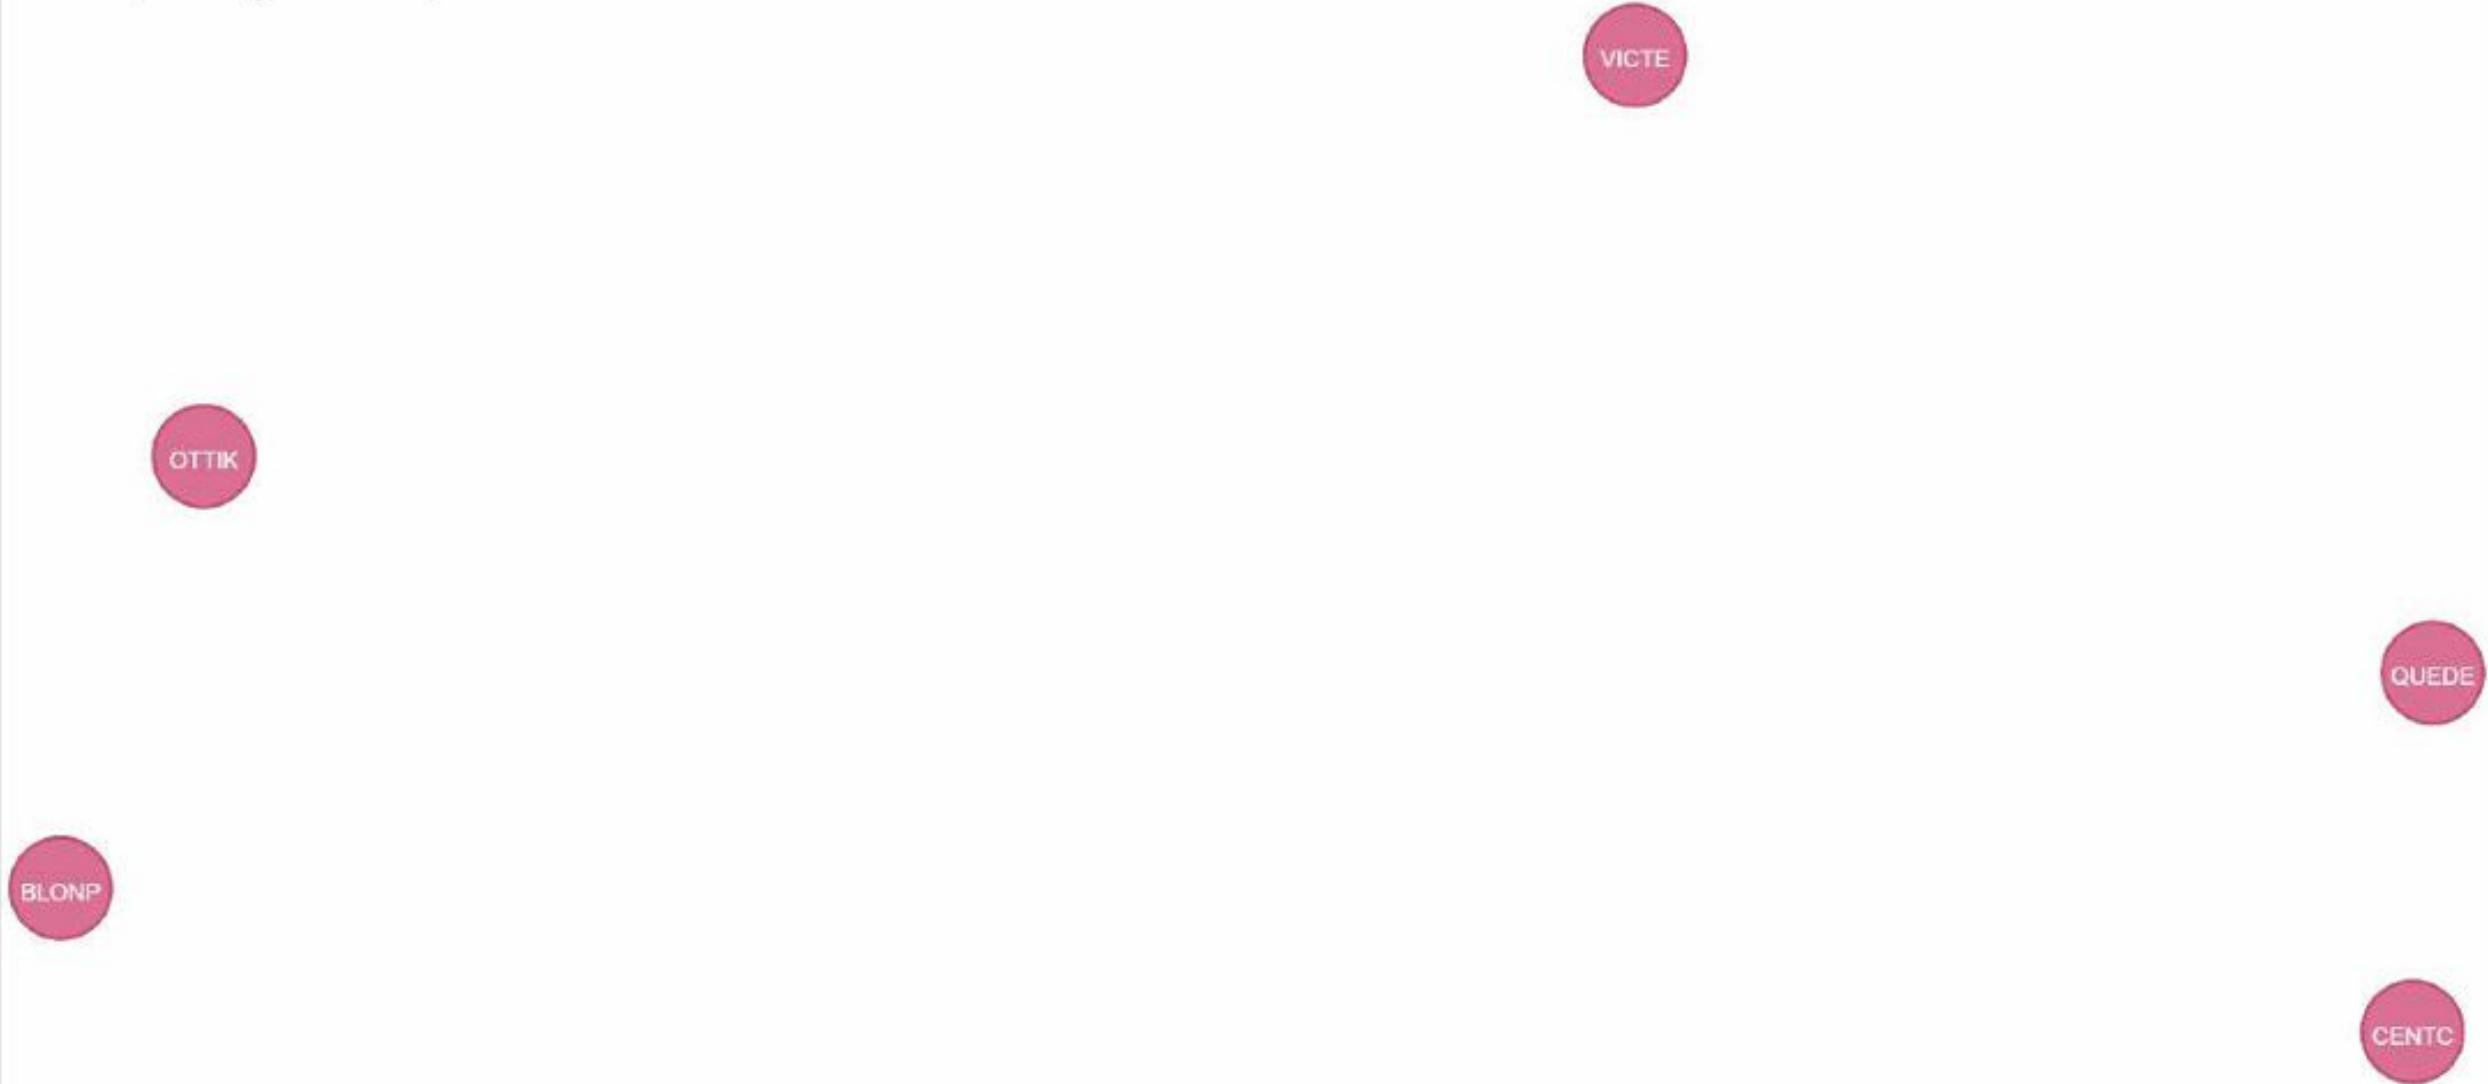


## Entity Relationship

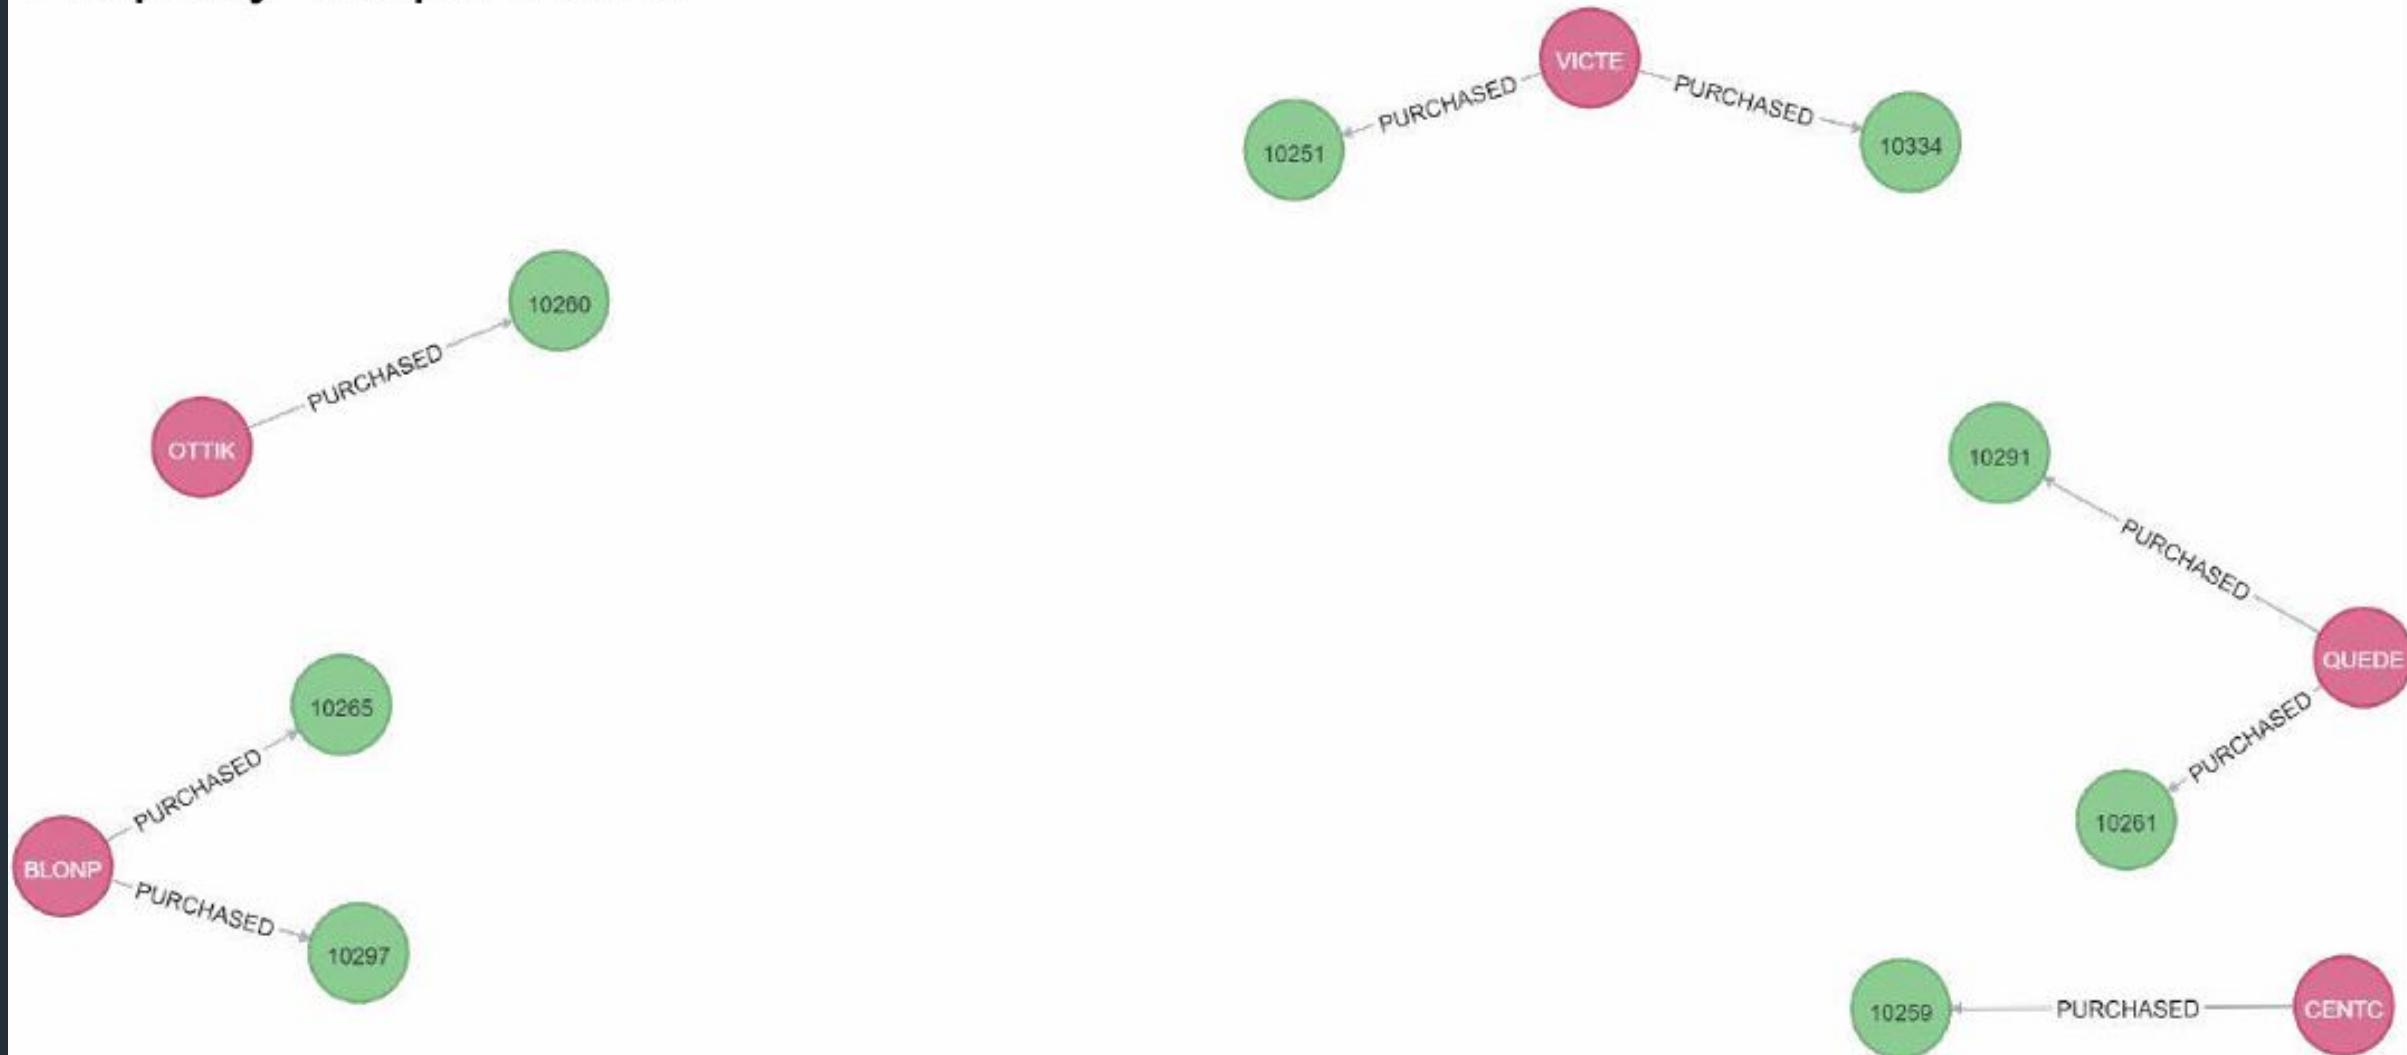
# Property Graph Model



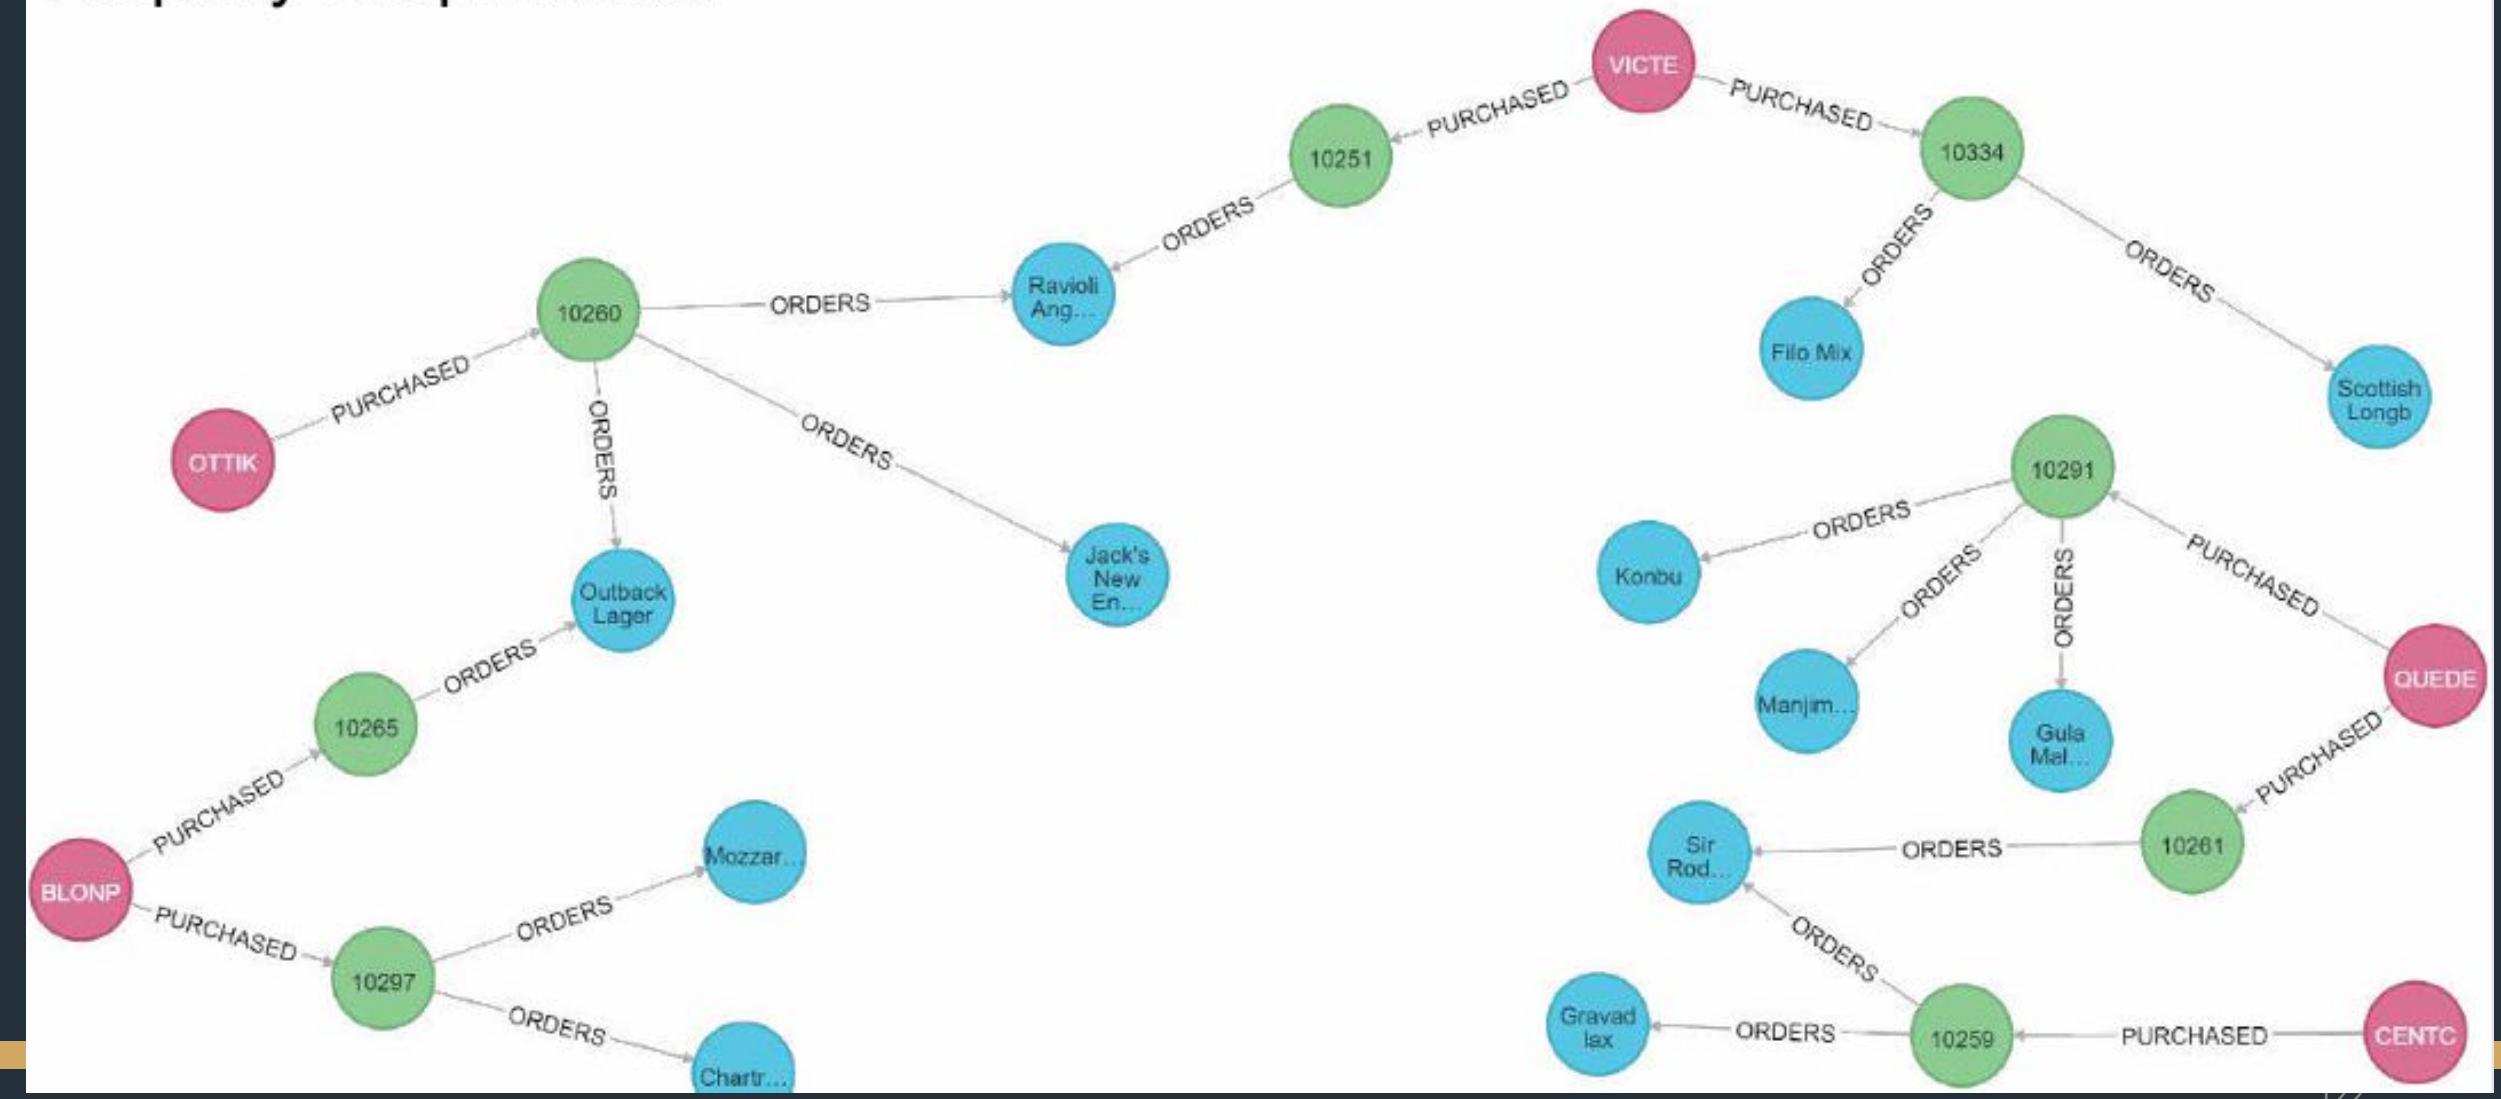
# Property Graph Model



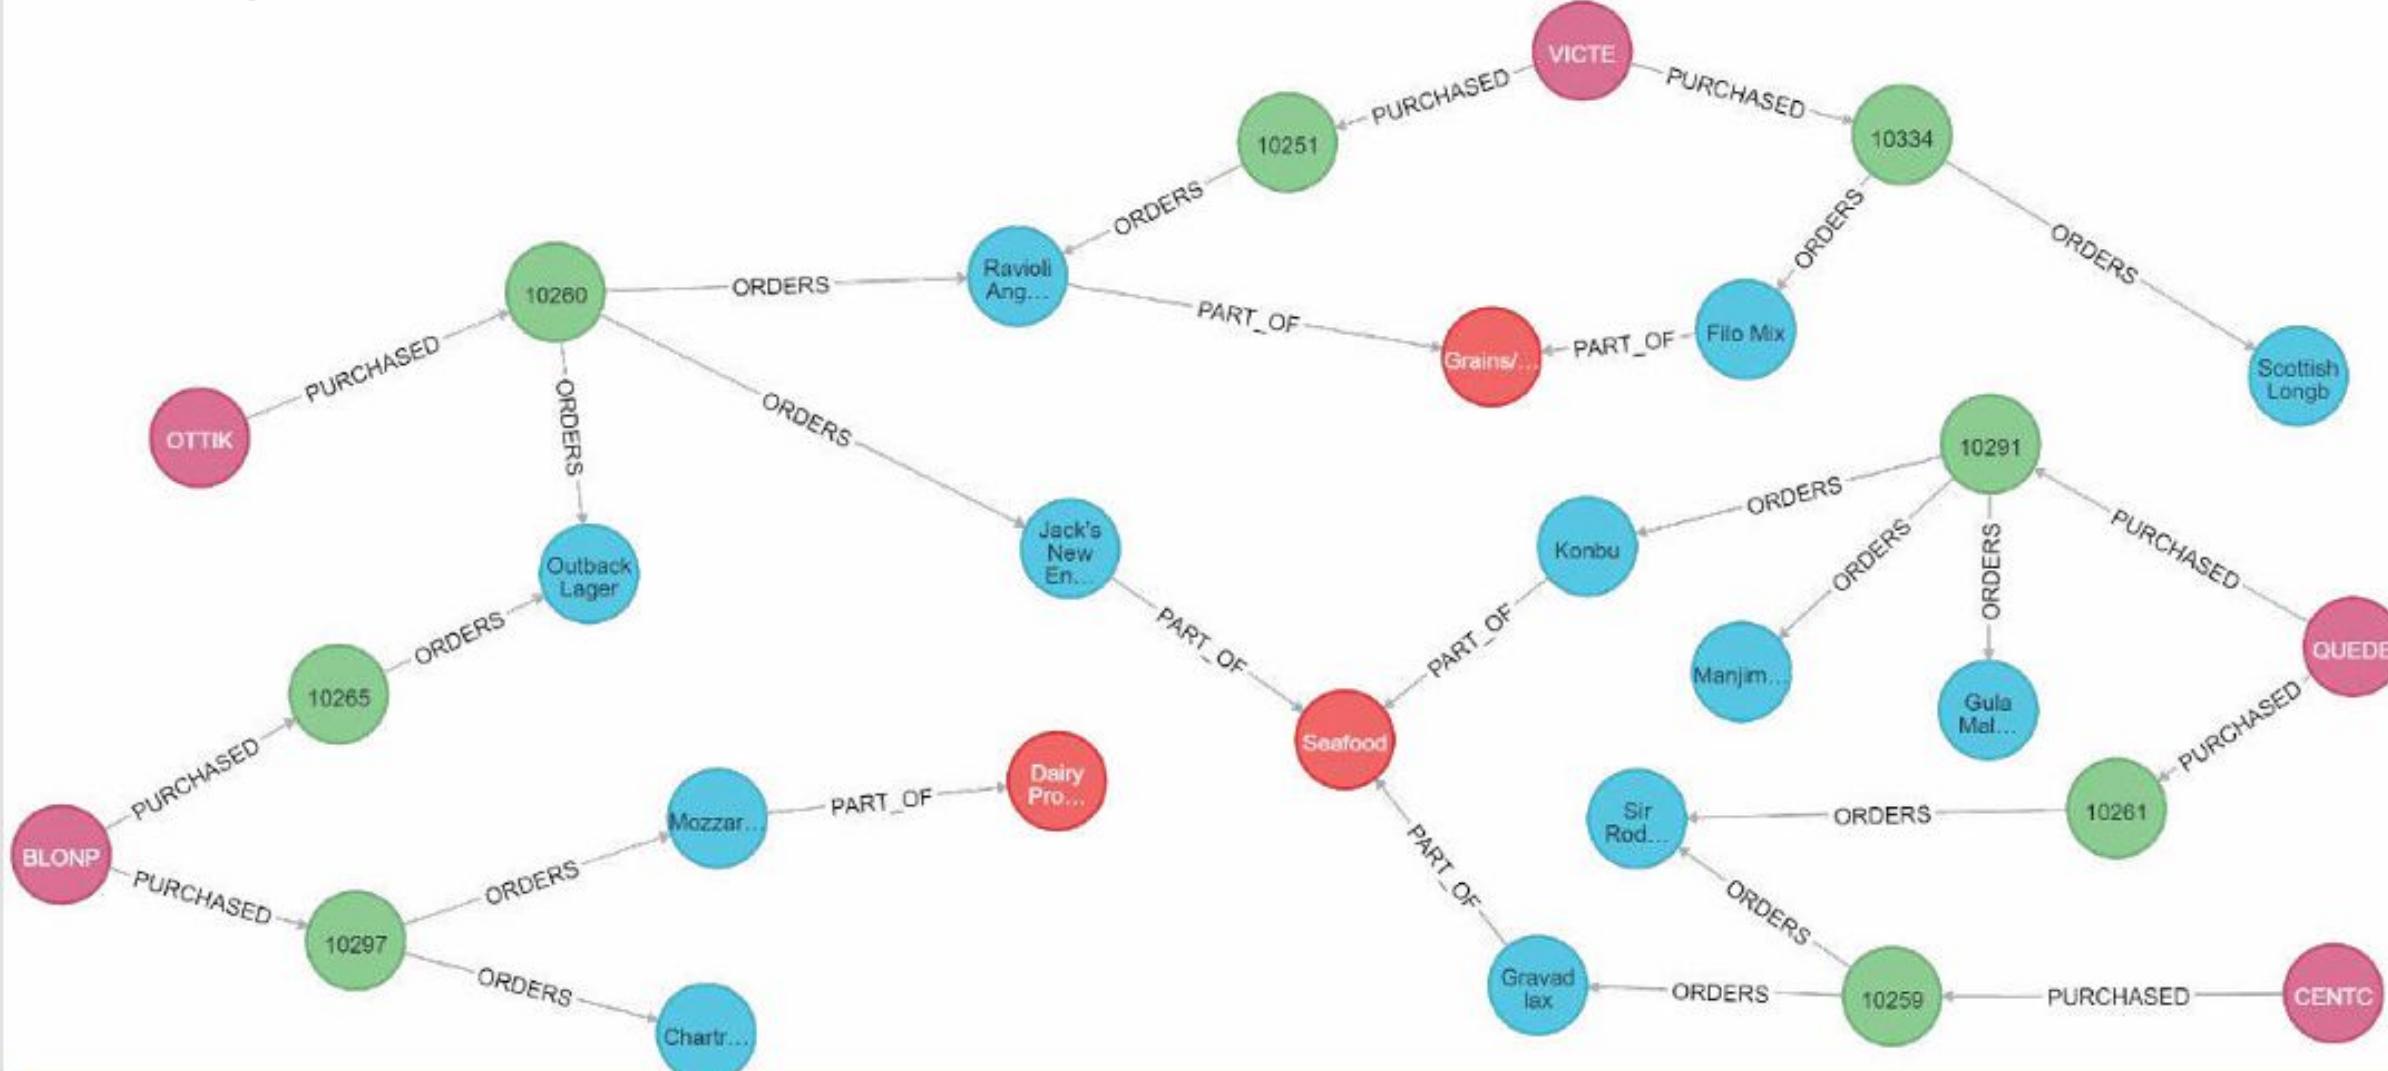
# Property Graph Model



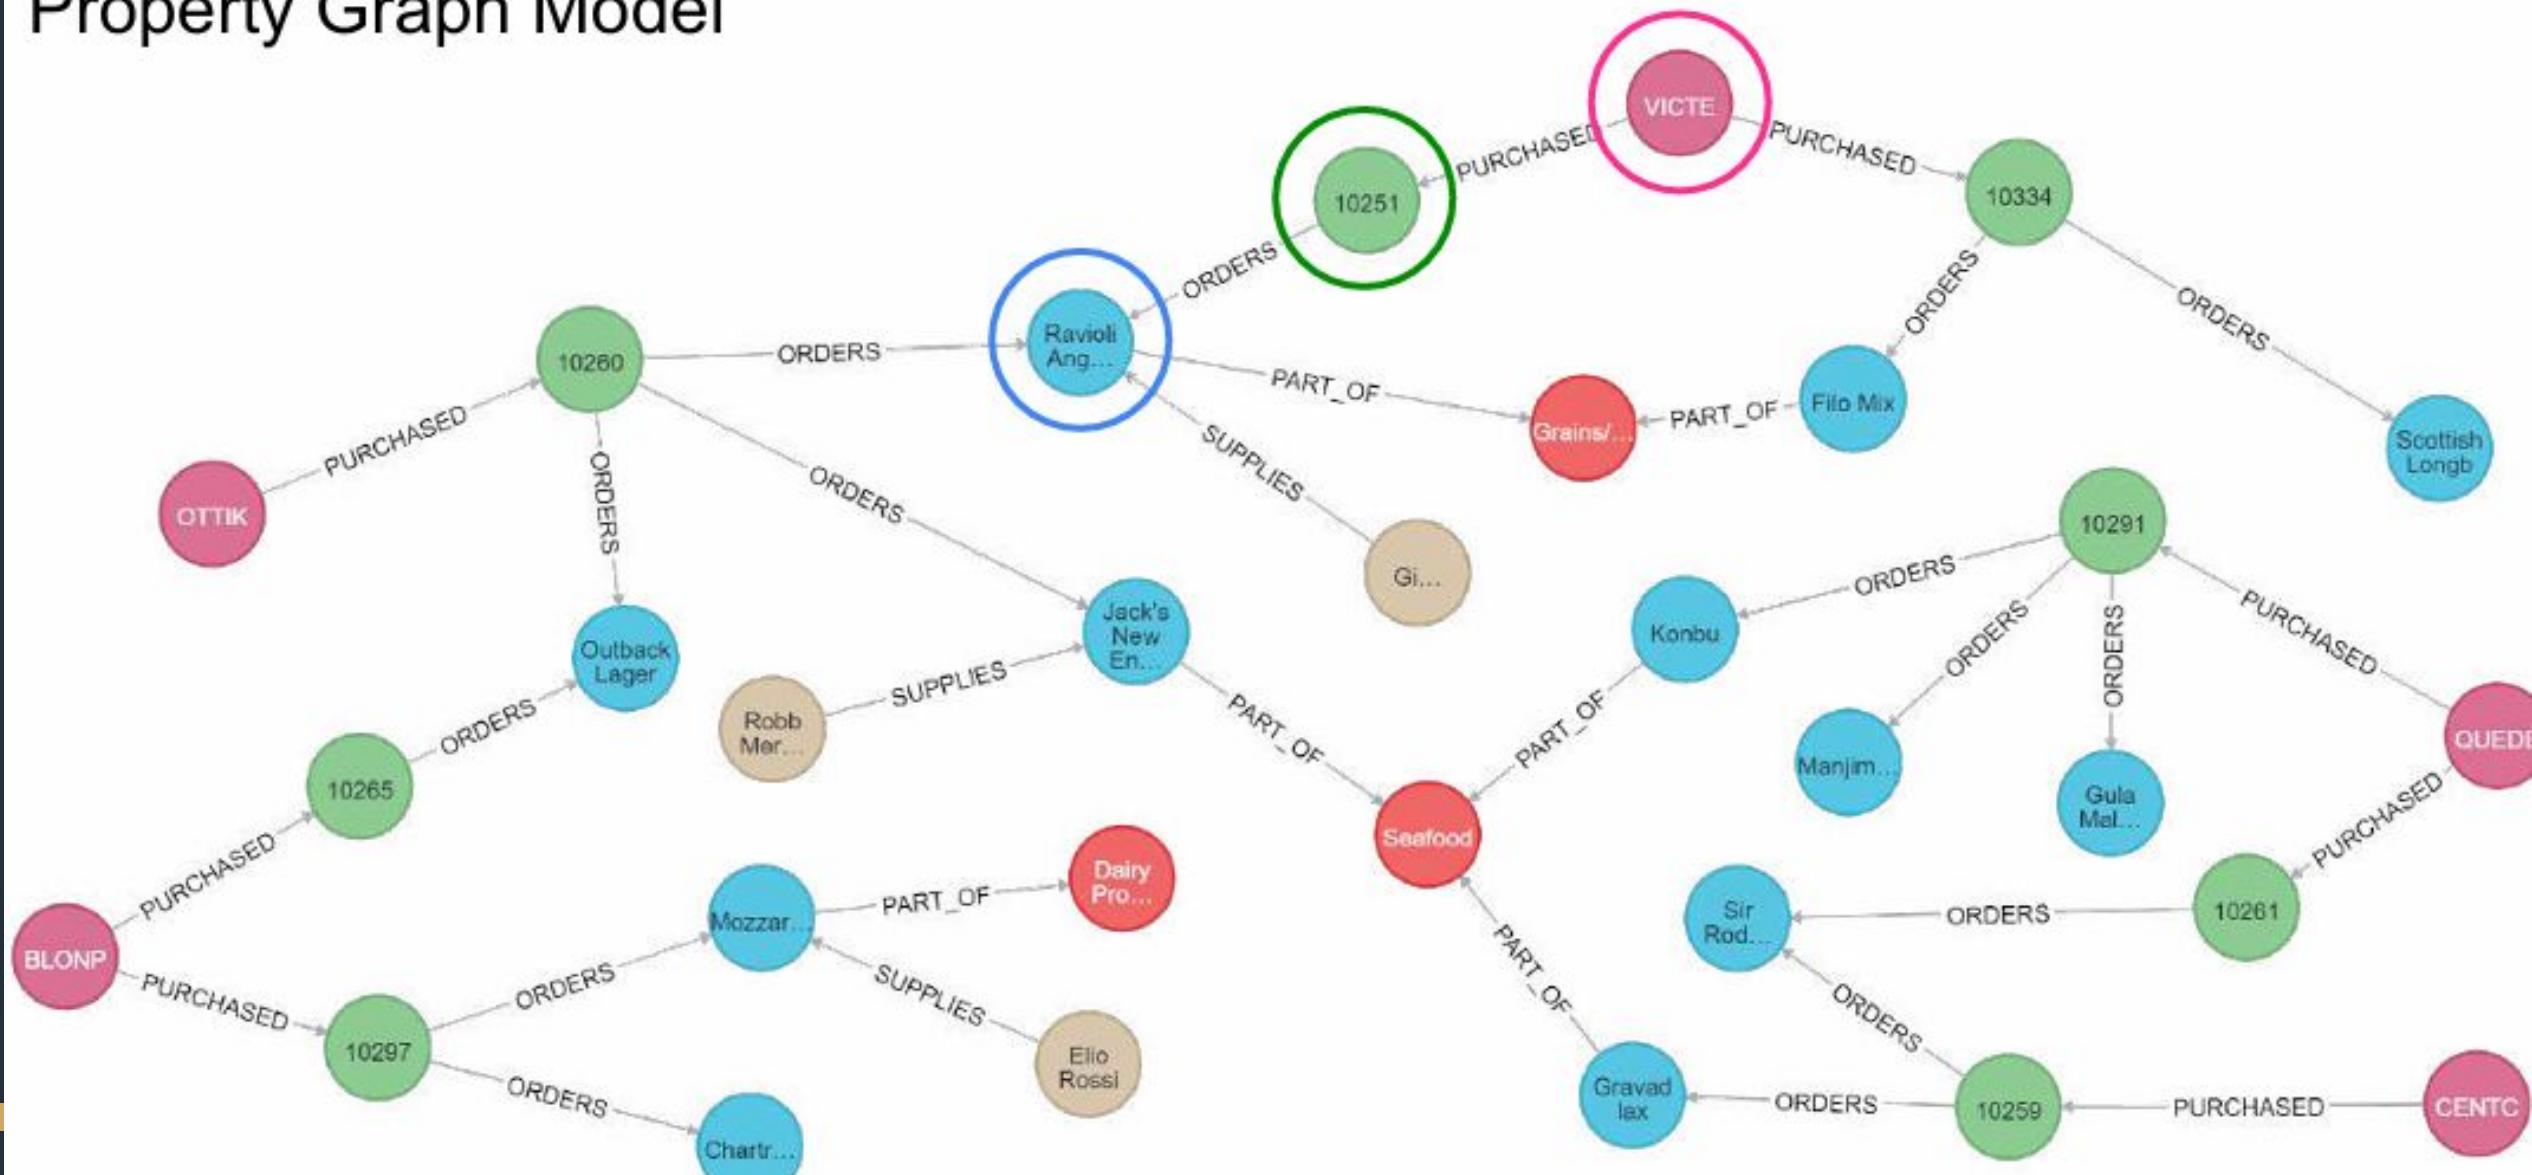
# Property Graph Model



## Property Graph Model



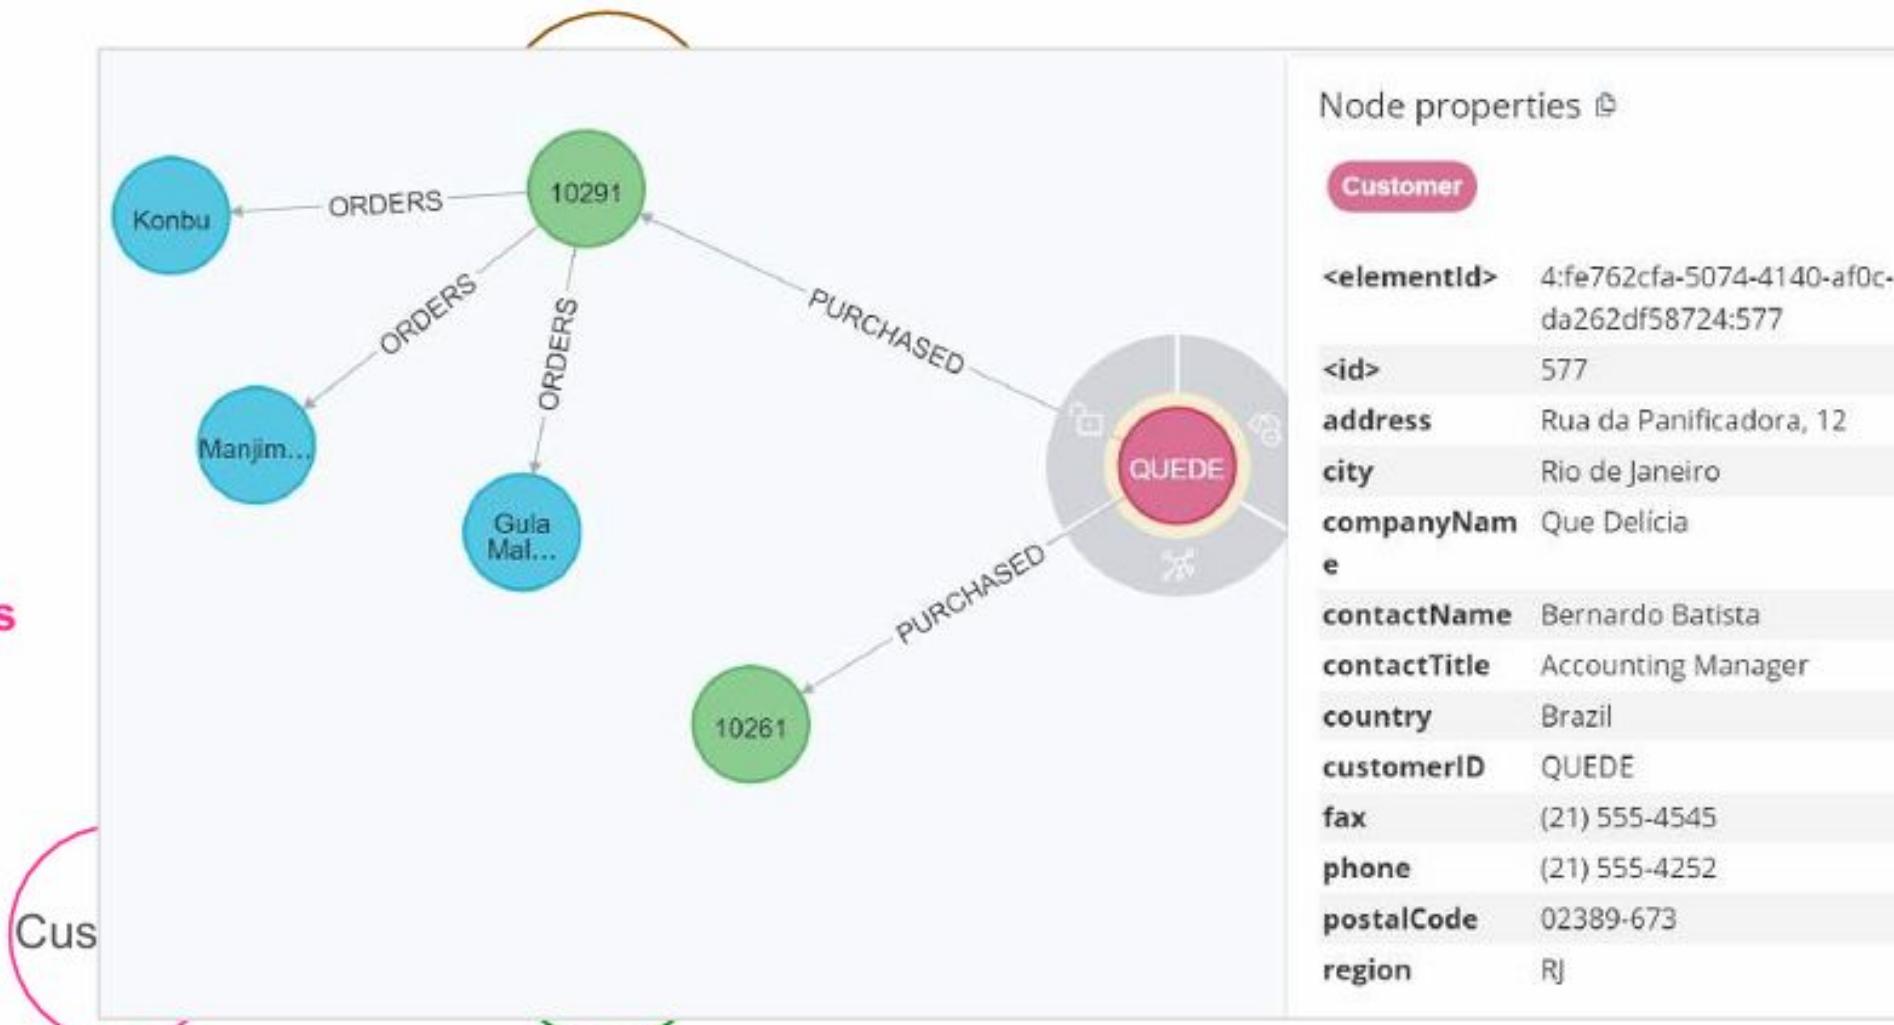
# Property Graph Model



# Property Graph Model

## Customer Properties

- address
- city
- companyName
- contactName
- contactTitle
- country
- customerID



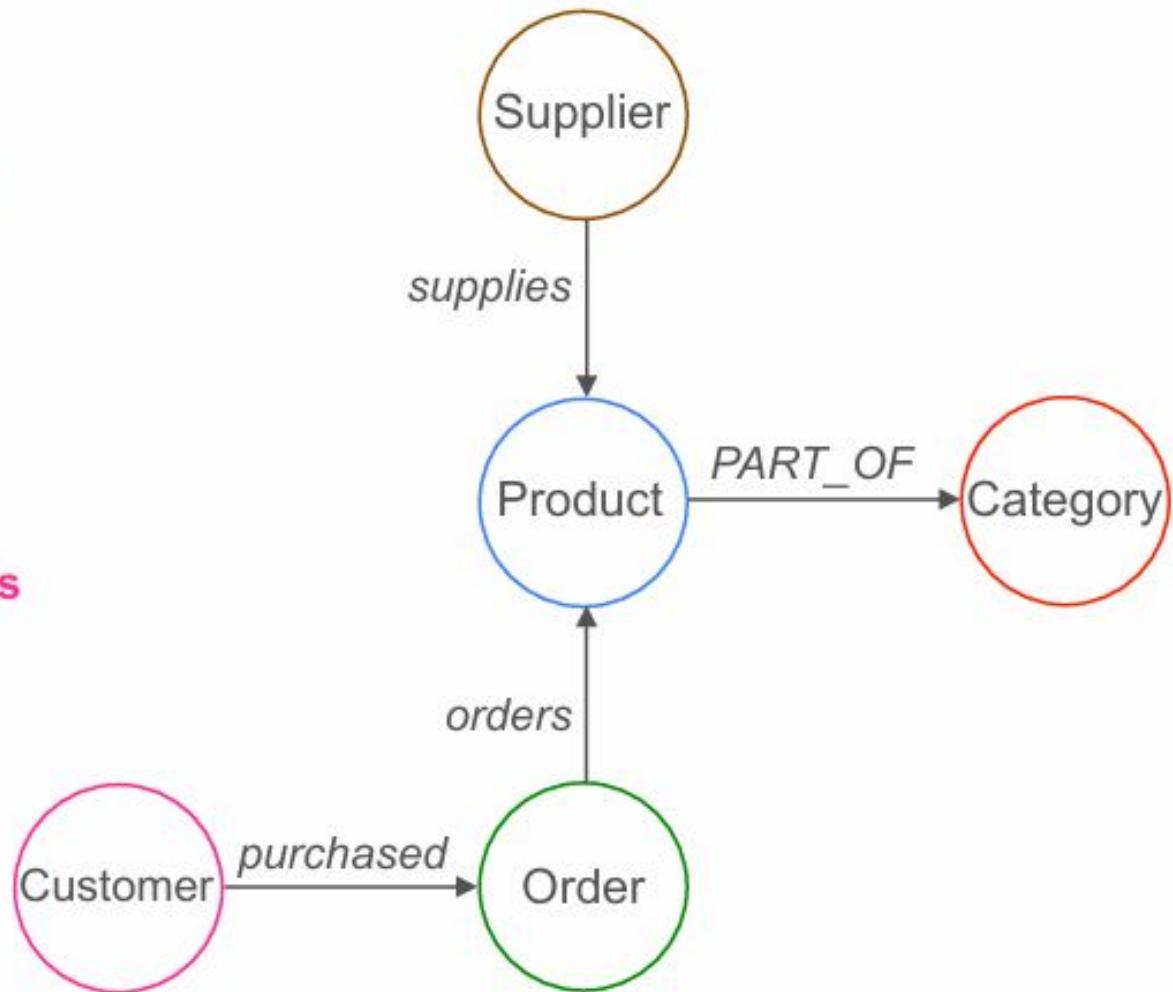
# Property Graph Model

## Product Properties

- productID
- productName
- unitPrice
- unitsInStock
- unitsOnOrder

## Customer Properties

- address
- city
- companyName
- contactName
- contactTitle
- country
- customerID
- .....



## Supplier Properties

- address
- city
- contactName
- fax
- region
- supplierID
- postalCode

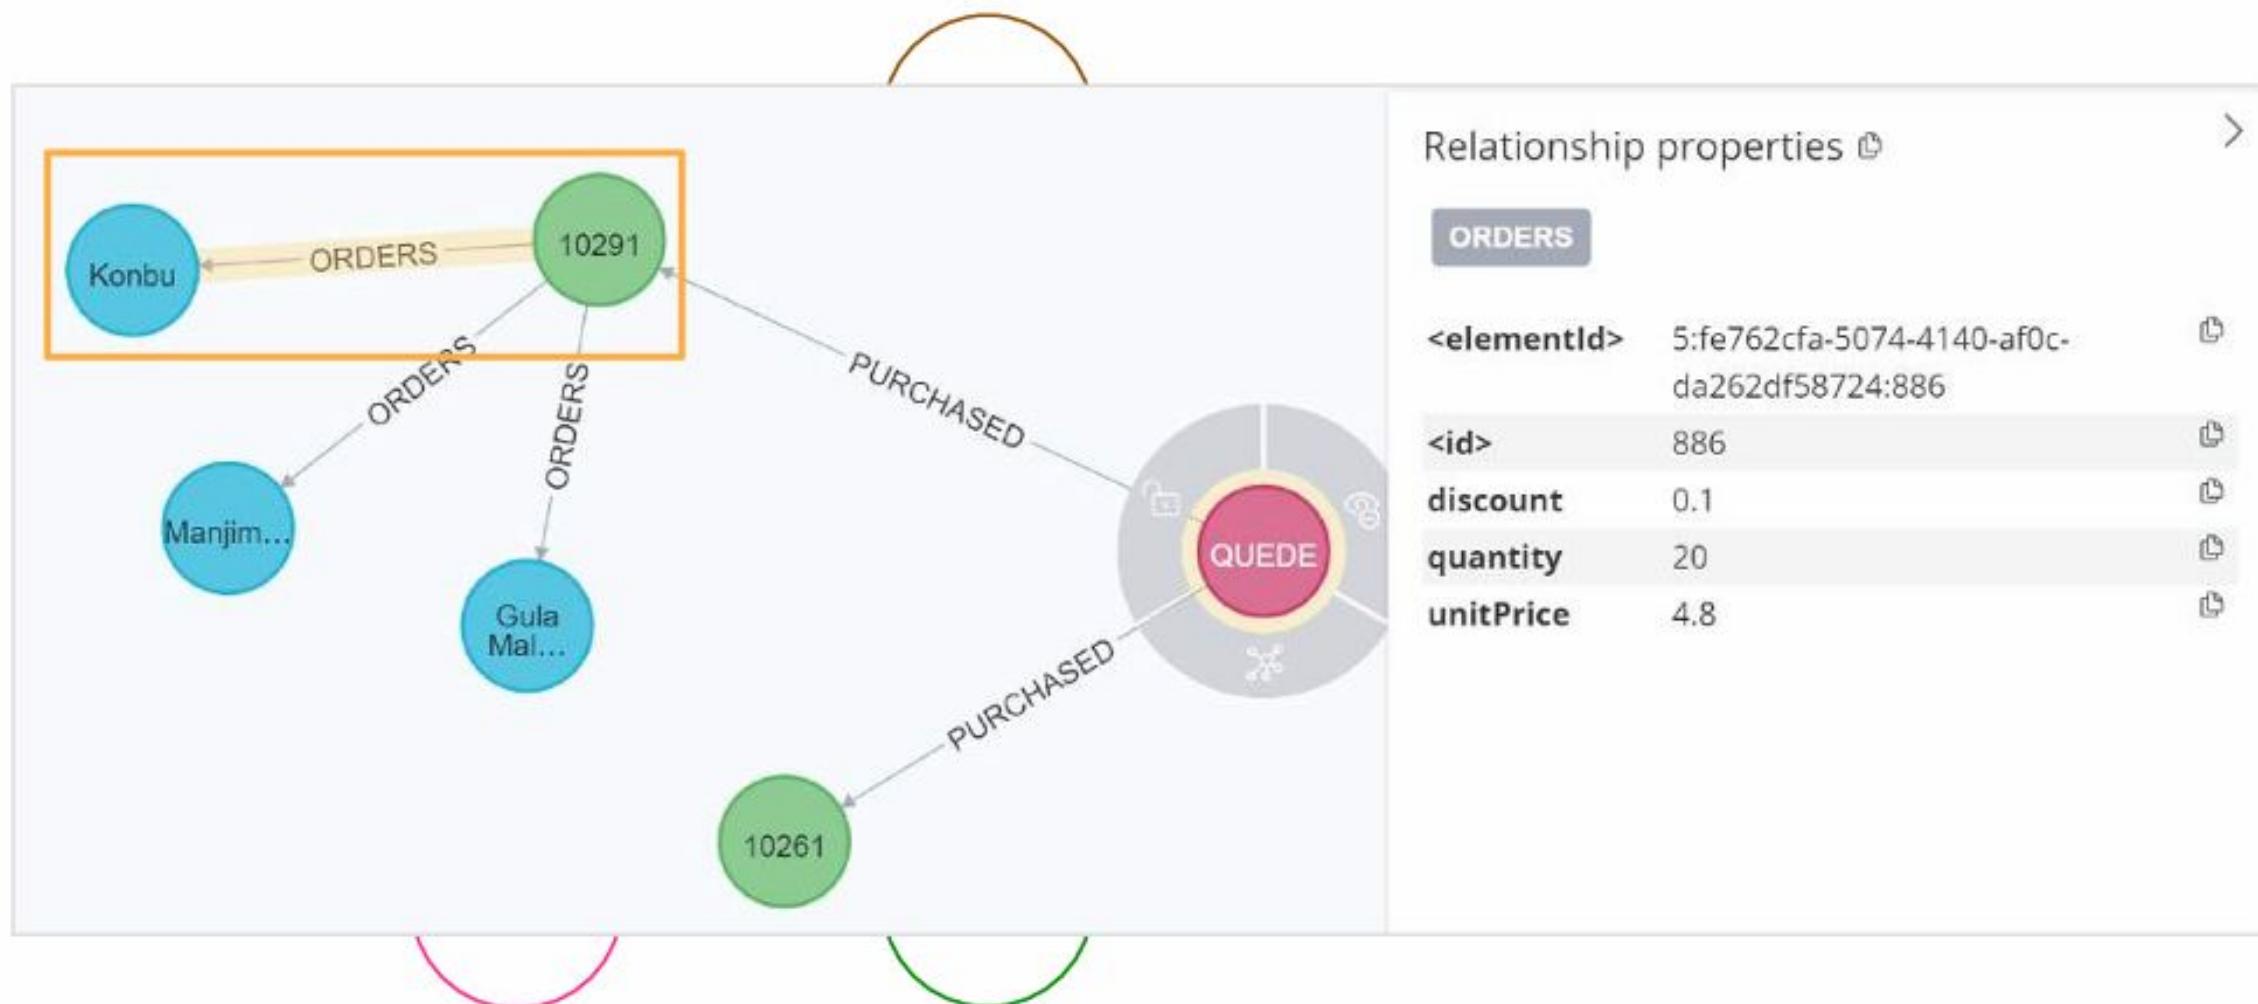
## Category Properties

- categoryName

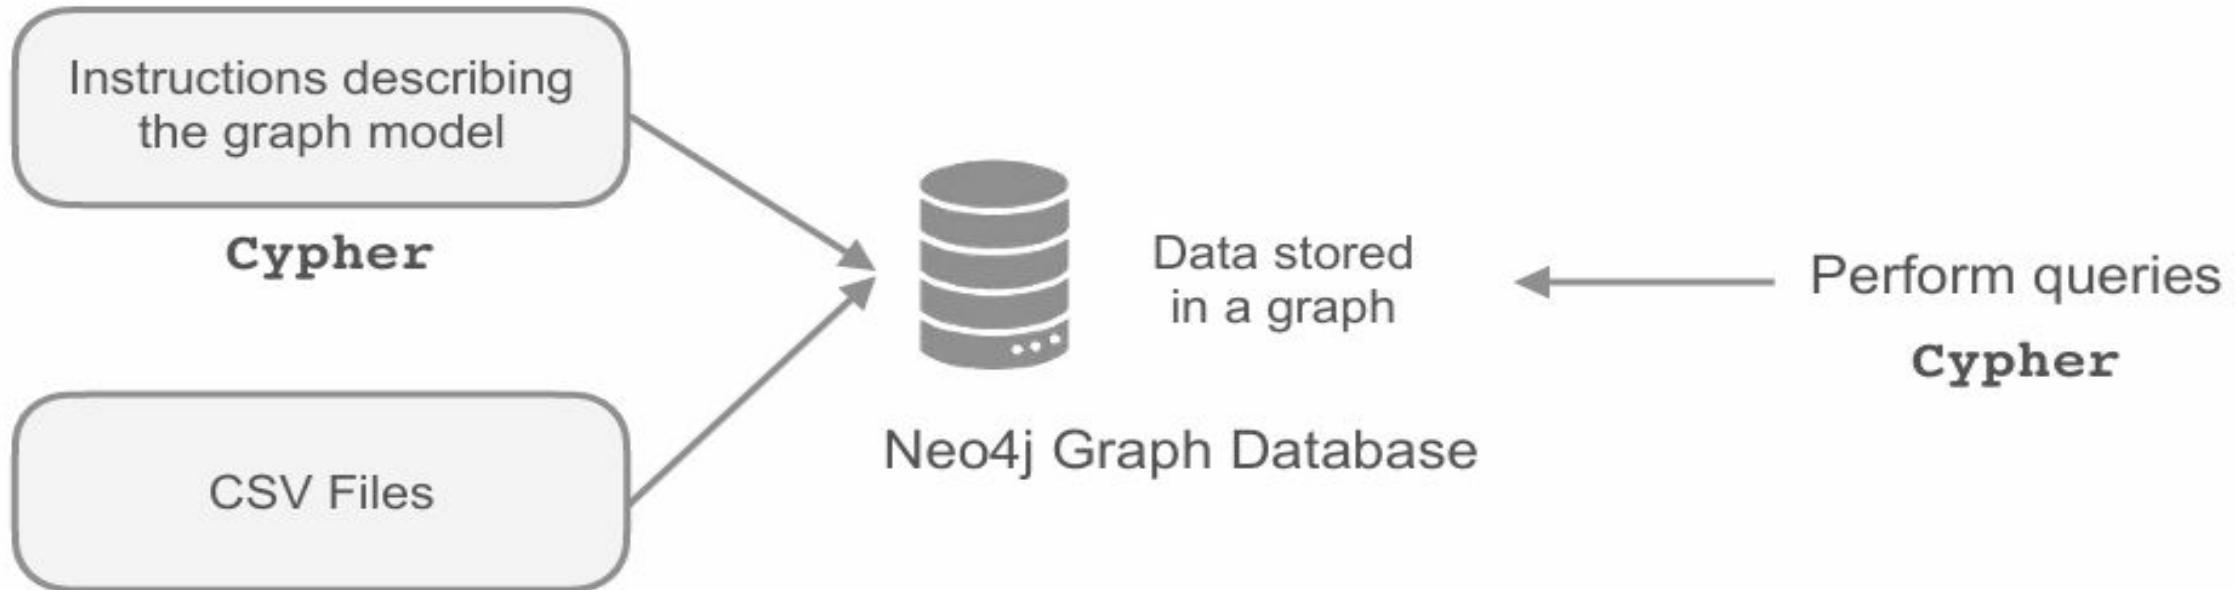
## Order Properties

- freight
- orderDate
- orderID
- requiredDate
- shipAddress
- .....

# Property Graph Model



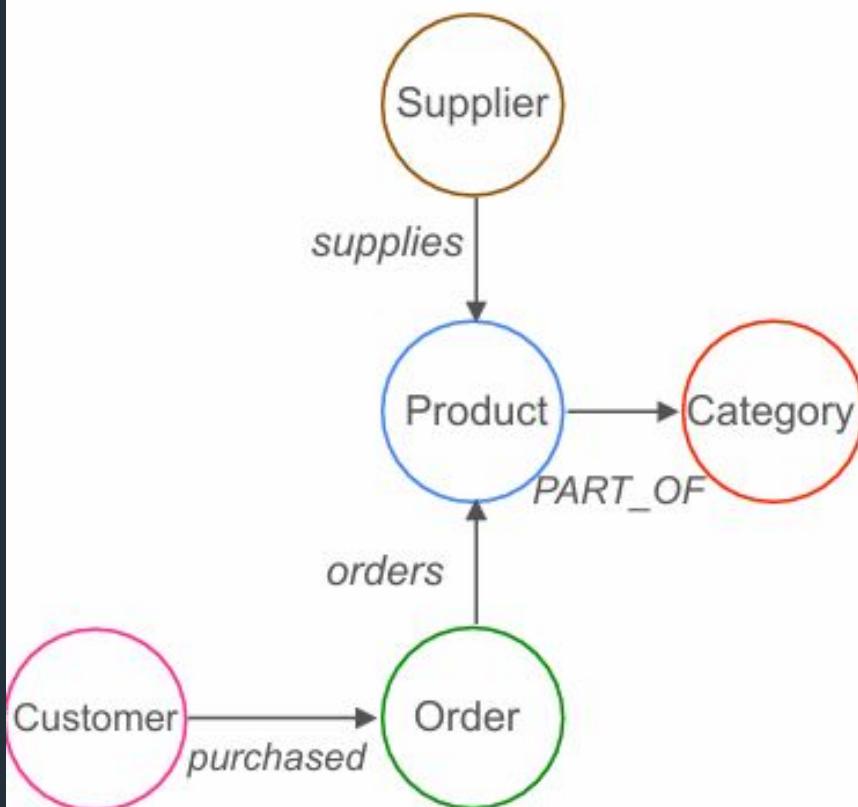
# Creating a Graph Database



- In the next video, we'll go through some queries examples.
- In the lab, you'll also practice CRUD operations.

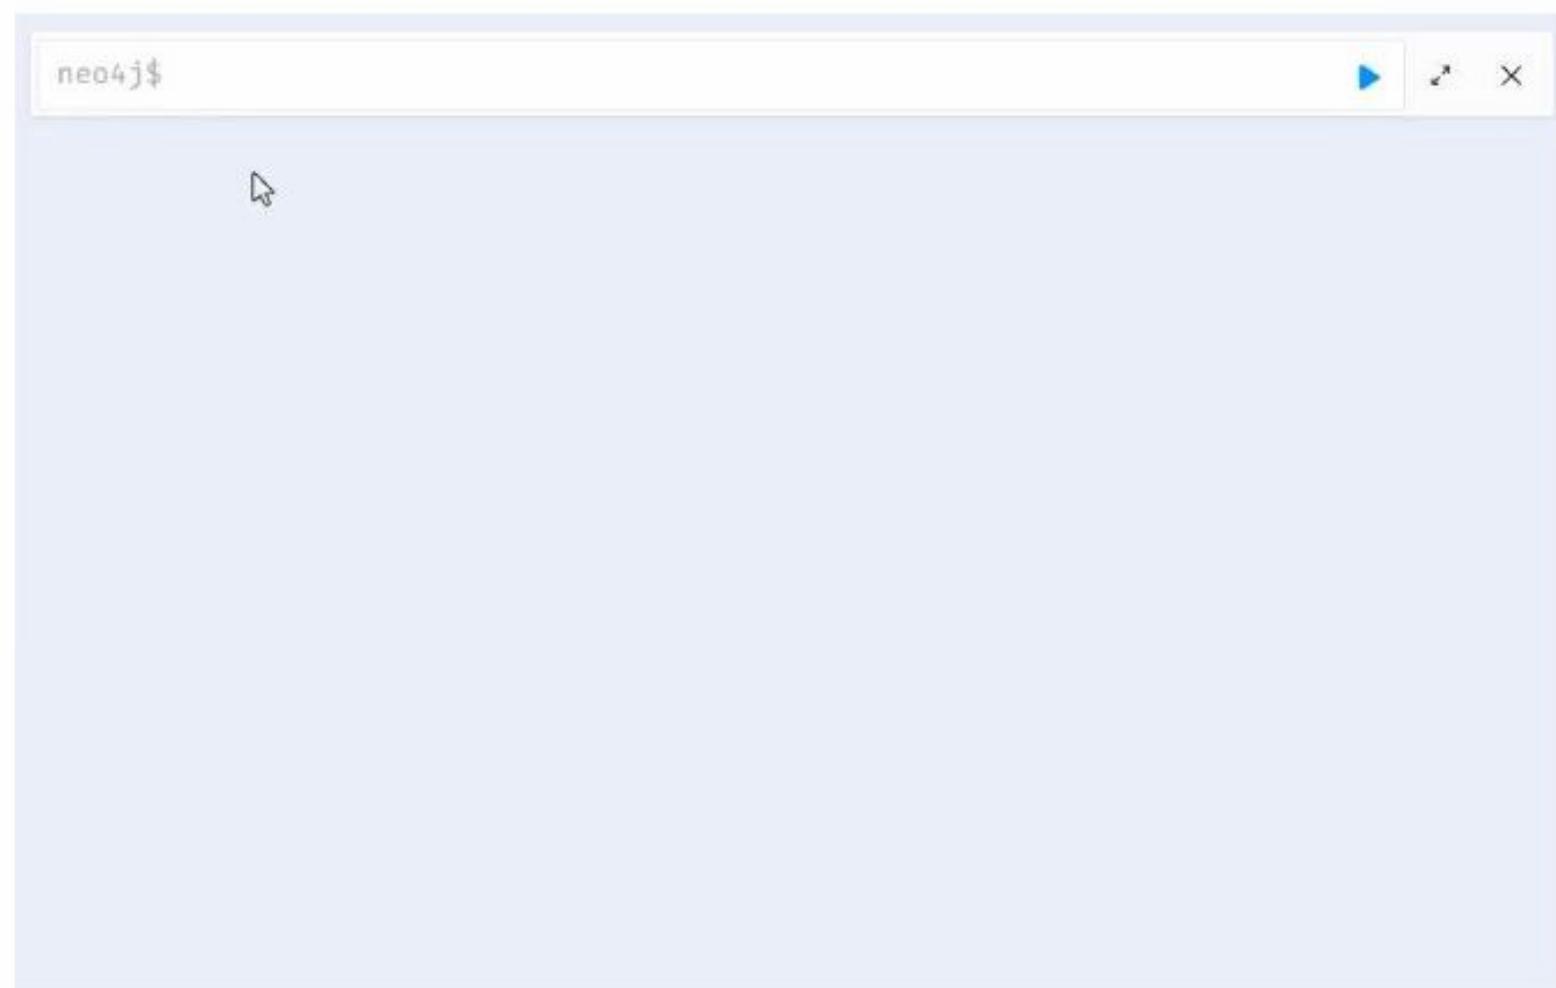
# Neo4j Graph Database & Cypher Query Language

```
MATCH pattern RETURN result  
node ()
```



## MATCH Statement

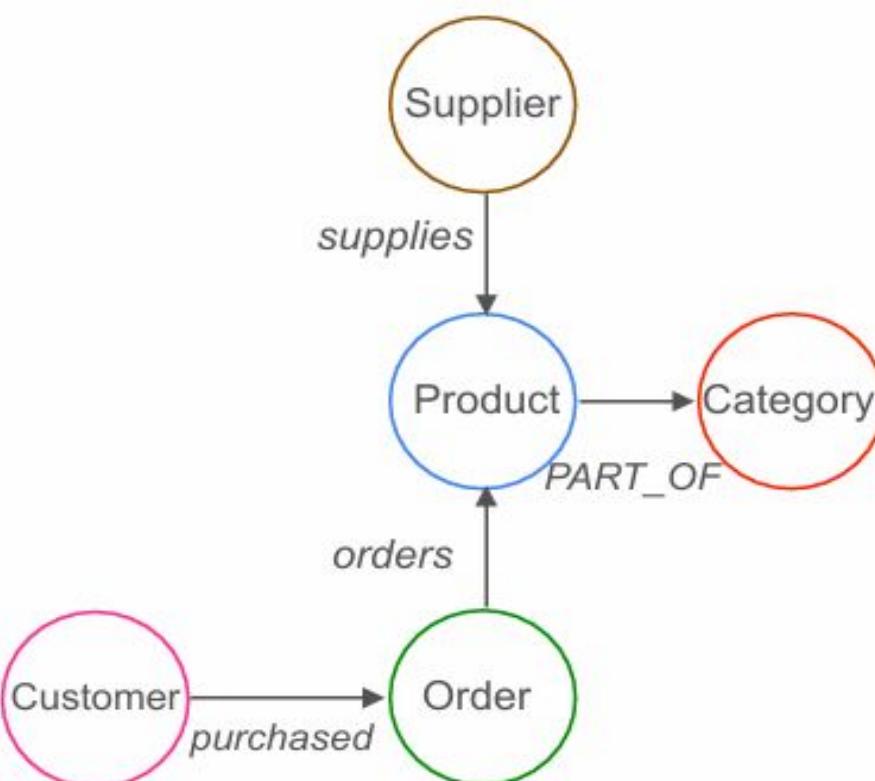
Retrieve all nodes



neo4j\$

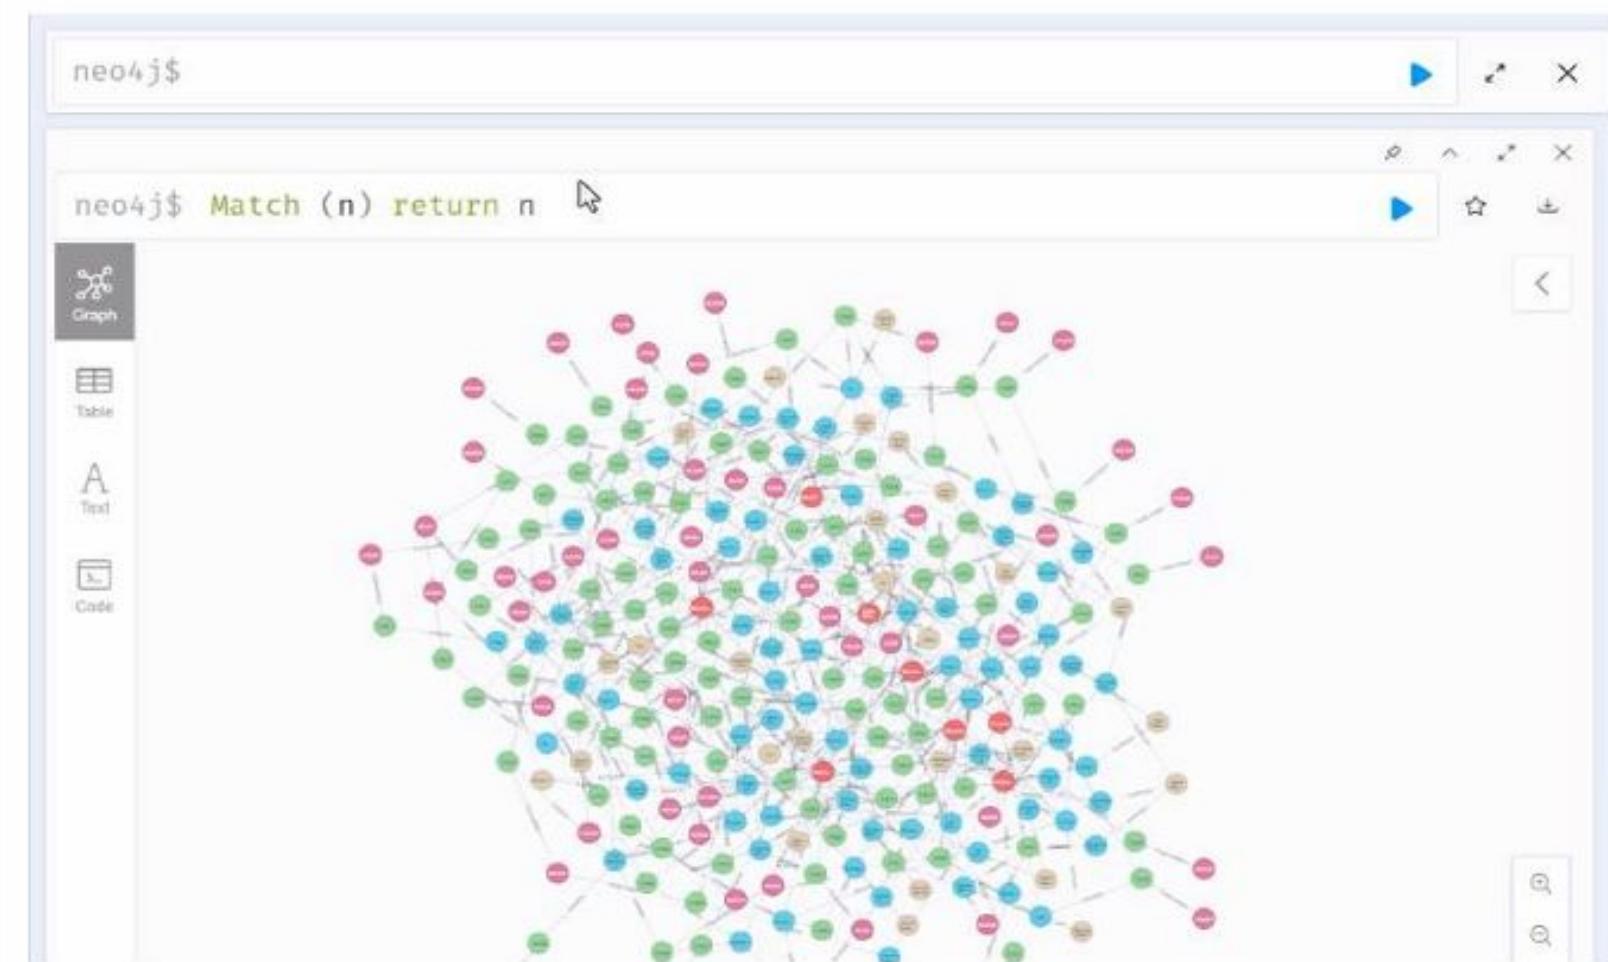
The screenshot shows the Neo4j browser interface with an empty results window. The title bar says "neo4j\$". There are three buttons in the top right corner: a blue play button, a green refresh button, and a red close button. A cursor arrow is visible in the center of the screen.

```
MATCH pattern RETURN result  
node ()
```



## MATCH Statement

Get the total number of nodes



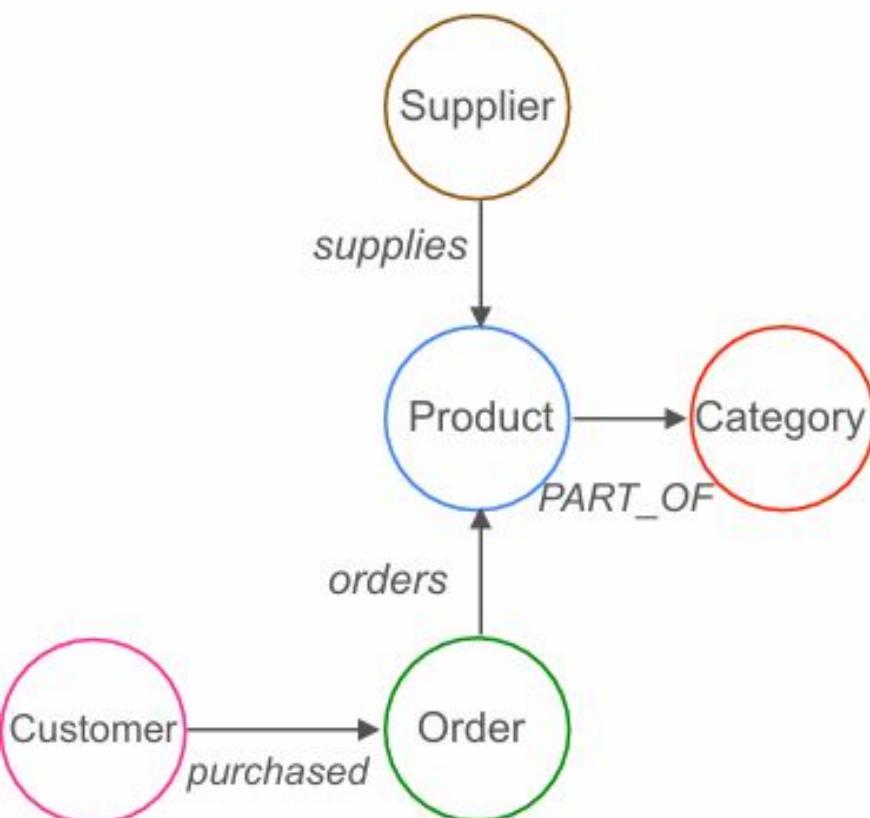
```
neo4j$  
neo4j$ MATCH (n) RETURN n
```

## MATCH Statement

Explore the node labels using the `labels` function

```
MATCH pattern RETURN result
```

```
node ()
```



```
neo4j$
```

```
neo4j$ Match (n) return count(n)
```

count(n)
1 265

```
A Text
```

```
Started streaming 1 records after 1 ms and completed after 1 ms.
```

```
neo4j$ Match (n) return n
```

```
Customer
```

MATCH pattern RETURN result

node ()



## MATCH Statement

Specify the label of the node



```
neo4j$ Match (n) return distinct labels(n)
```

Table	Text
1	["Supplier"]
2	["Category"]
3	["Product"]
4	["Order"]
5	["Customer"]

Started streaming 5 records in less than 1 ms and completed after 23 ms.

MATCH *pattern* RETURN *result*

node

( )



## MATCH Statement

Explore the properties of each order node using the Properties function



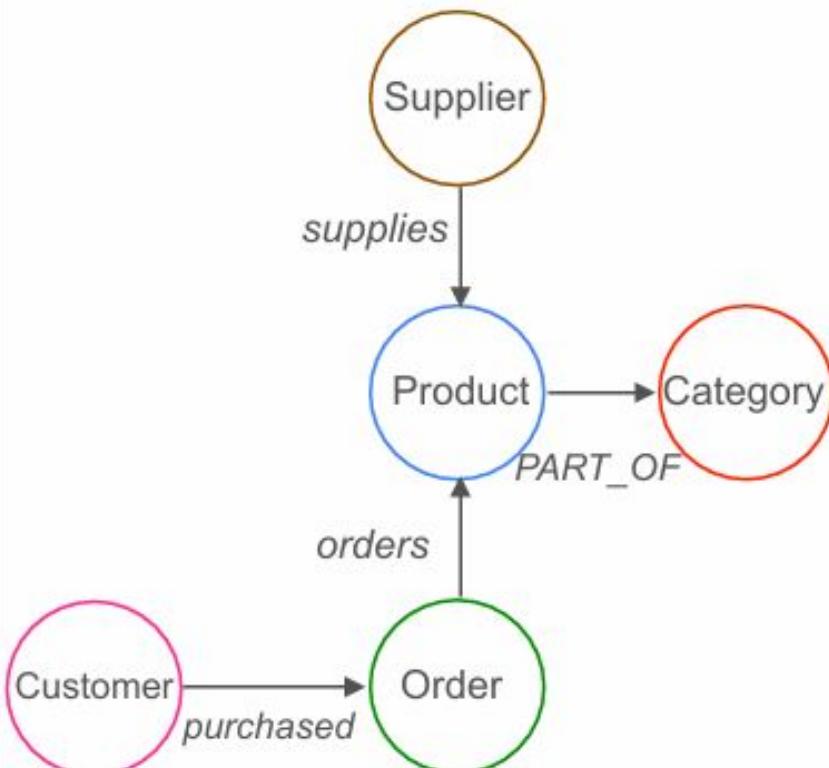
```

neo4j$ Match (n:Order) return count(n)
+-----+
| count(n) |
+-----+
| 99       |
+-----+

Started streaming 1 records after 2 ms and completed after 10 ms.

neo4j$ Match (n) return distinct labels(n)
+-----+
| labels(n) |
+-----+
| [Order]   |
+-----+
  
```

```
MATCH pattern RETURN result
node      ()
```



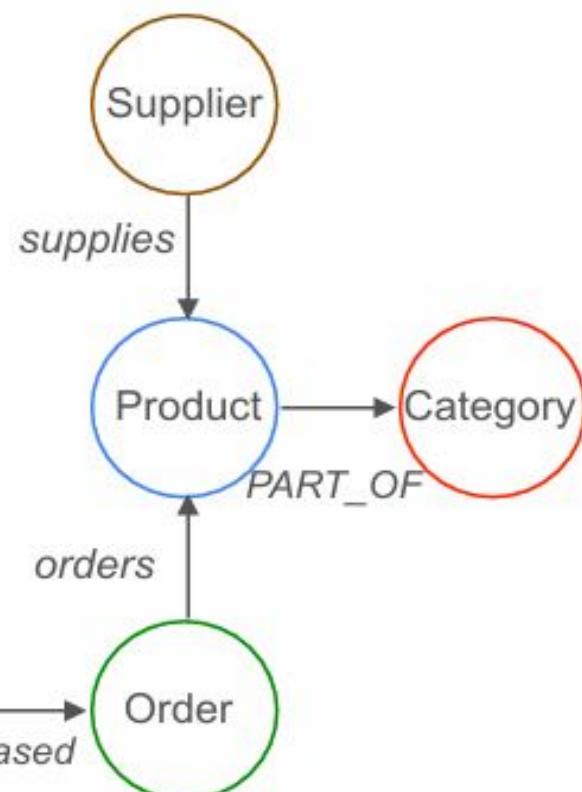
## MATCH Statement

Explore the properties of each order node using the Properties function

Properties(n)
<pre>{     "orderID": "10248",     "orderDate": "00:00:00",     "requiredDate": "00:00:00",     "shippedDate": "00:00:00",     "shipCity": "Reims",     "shipCountry": "France",     "shipName": "Vins et alcools Chevalier",     "shipPostalCode": "51100",     "shipRegion": "NULL",     "shipAddress": "59 rue de l'Abbaye",     "freight": "32.38" }</pre>

MATCH pattern RETURN result

node	( )
relationship	[ ]
path	(source node)-[r]->(target node)



## MATCH Statement

Count all the directed paths

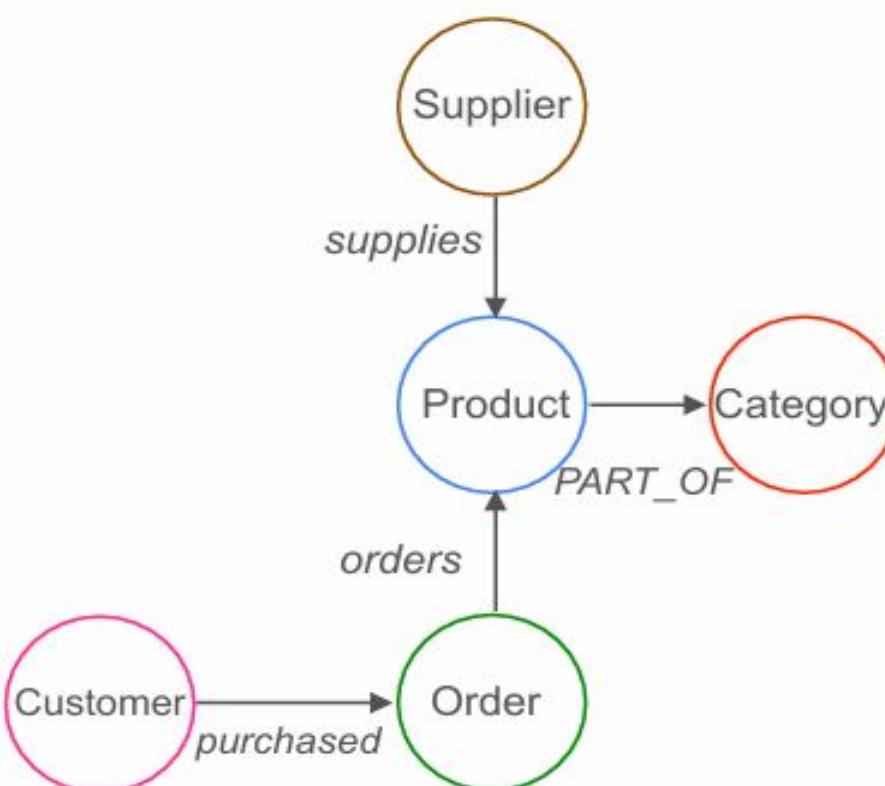
```

neo4j$ Match (n:Order) return Properties(n) limit 1
Properties(n)
{
    "shipCity": "Reims",
    "orderID": "10248",
    "shippedDate": "00:00.0",
    "orderDate": "00:00.0",
    "shipRegion": "NULL",
    "freight": "32.38",
    "shipName": "Vins et alcools Chevalier",
    "shipCountry": "France",
    "shipAddress": "59 rue de l'Abbaye",
    "requiredDate": "00:00.0",
    "shipPostalCode": "51100"
}
  
```

Started streaming 1 records after 2 ms and completed after 8 ms.

MATCH pattern RETURN result

node	( )
relationship	[ ]
path	(source node)-[r]->(target node)



## MATCH Statement

Return the types of relationships

```

neo4j$ Match ()-[r]->() return count(r)
  
```

count(r)
518

Started streaming 1 records in less than 1 ms and completed after 5 ms.

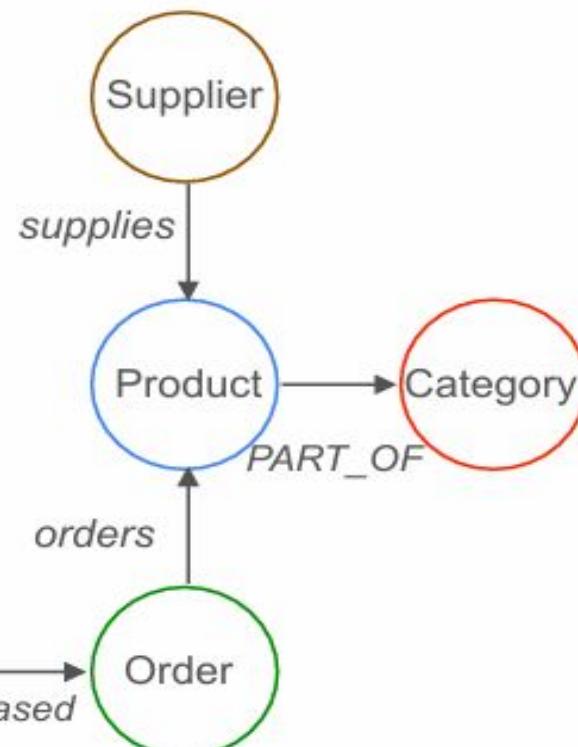
```

neo4j$ Match (n:Order) return Properties(n) limit 1
  
```

Properties(n)
---------------

```
MATCH pattern RETURN result
```

node	( )
relationship	[ ]
path	(source node)-[r]->(target node)



## MATCH Statement

Specify the type of the relationship

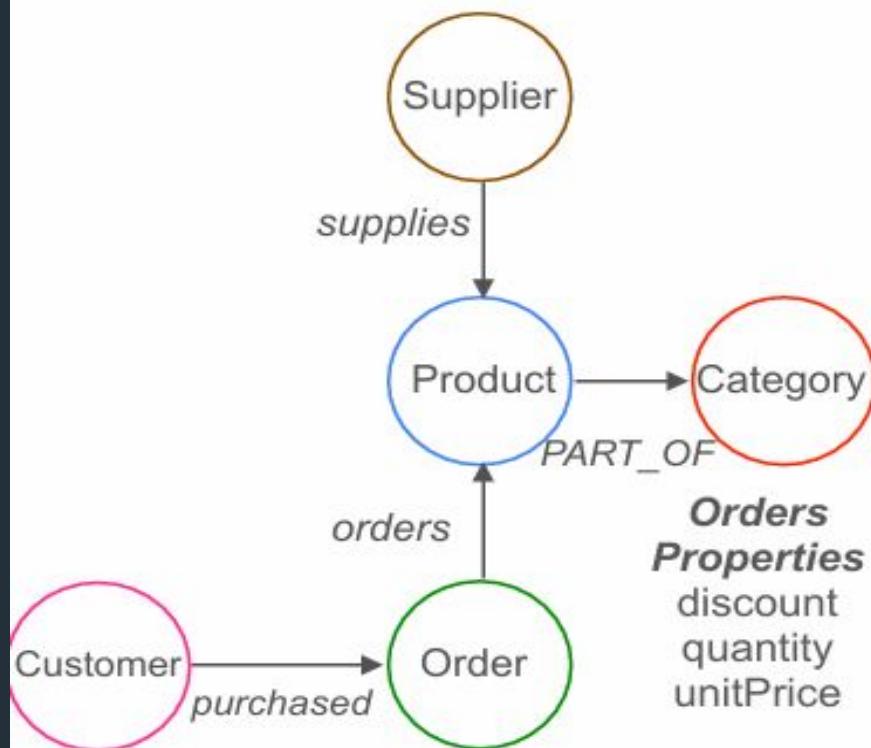
```

neo4j$ Match ()-[r]->() return distinct type(r)
+-----+
| type(r) |
+-----+
| "SUPPLIES" |
| "PART_OF" |
| "PURCHASED" |
| "ORDERS" |
+-----+
Started streaming 4 records after 1 ms and completed after 21 ms.

neo4j$ Match (n:Order) return Properties(n) limit 1
  
```

```
MATCH pattern RETURN result
```

node	( )
relationship	[ ]
path	(source node)-[r]->(target node)



## MATCH Statement

Return the properties of a relationship

```
neo4j$ Match ()-[r:ORDERS]->()
```

```
neo4j$ Match ()-[r]->() return distinct type(r)
```

Table

type(r)

"SUPPLIES"

"PART\_OF"

"PURCHASED"

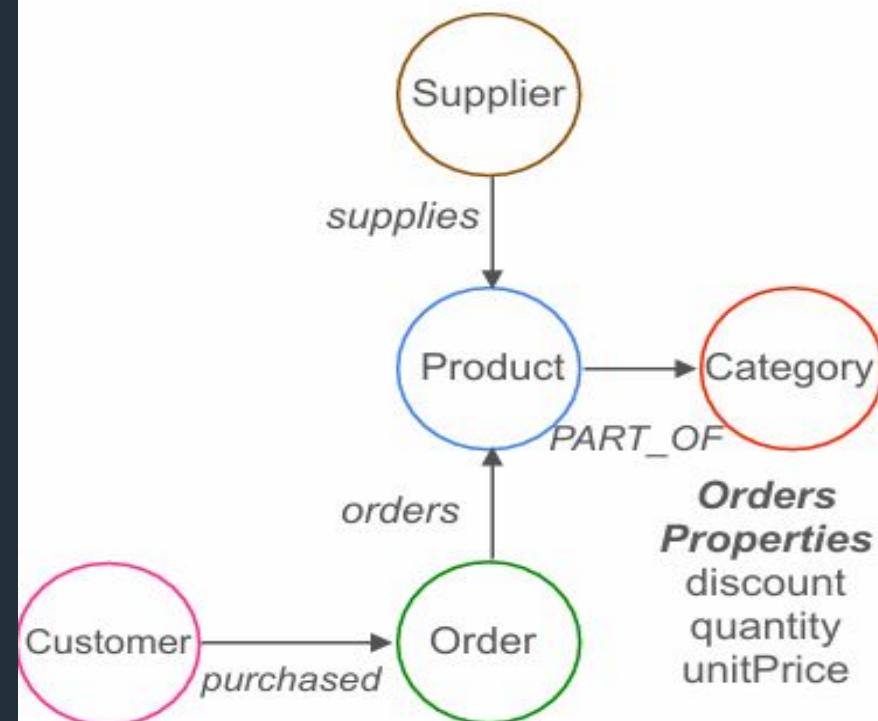
"ORDERS"

Started streaming 4 records after 1 ms and completed after 21 ms.

```
neo4j$ Match (n:Order) return Properties(n) limit 1
```

MATCH *pattern* RETURN *result*

node	( )
relationship	[ ]
path	(source node)-[ ]->(target node)



## MATCH Statement

Return the properties of a relationship

```

neo4j$ Match ()-[r:ORDERS]->() return AVG(r.quantity*r.unitPrice)
AVG(r.quantity*r.unitPrice)
1
502.4645283018867

Started streaming 1 records after 1 ms and completed after 25 ms.

neo4j$ Match ()-[r]->() return distinct type(r)
type(r)

```

## MATCH Statement

`MATCH pattern RETURN result`

<code>node</code>	<code>()</code>
<code>relationship</code>	<code>[]</code>
<code>path</code>	<code>(source node)-[ ]-&gt;(target node)</code>



Get the average price for all orders grouped by product category

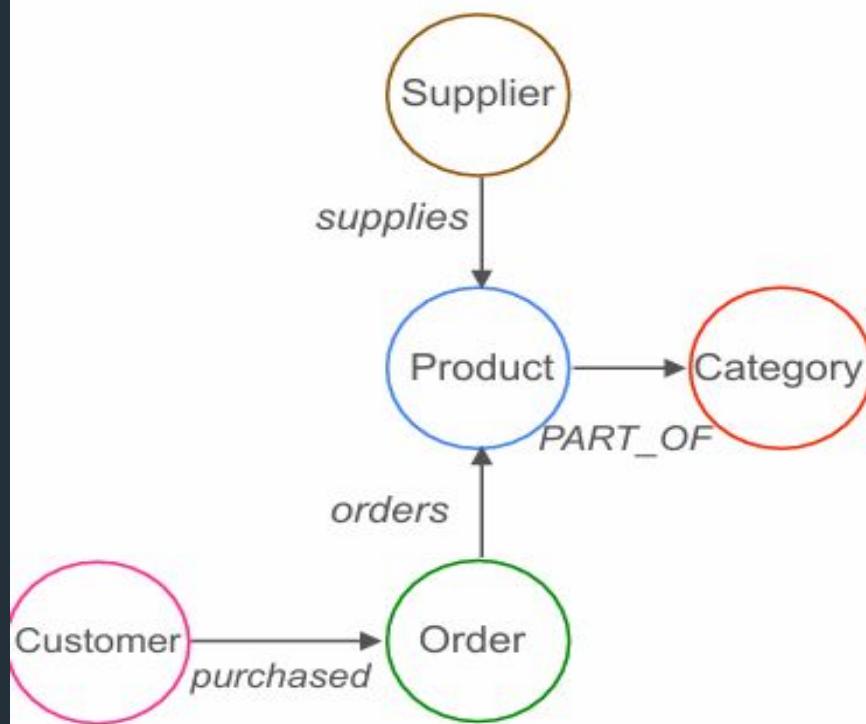
```

neo4j$ Match ()-[r:ORDERS]→() return AVG(r.quantity*r.unitPrice) as a...
average_price
502.4645283018867
Started streaming 1 records in less than 1 ms and completed after 25 ms.

neo4j$ Match ()-[r]→() return distinct type(r)
type(r)
  
```

MATCH *pattern* RETURN *result*

node	( )
relationship	[ ]
path	(source node)-[ ]->(target node)



## MATCH Statement

Get the average price for all orders grouped by product category

```
neo4j$ Match ()-[r:ORDERS]→()-[part:PART_OF]→(c:Category) return
      AVG(r.quantity*r.unitPrice) as average_price
```

```
neo4j$ Match ()-[r:ORDERS]→() return AVG(r.quantity*r.unitPrice) as a...
```

average_price
502,464,528,301,8867

A

Text

Code

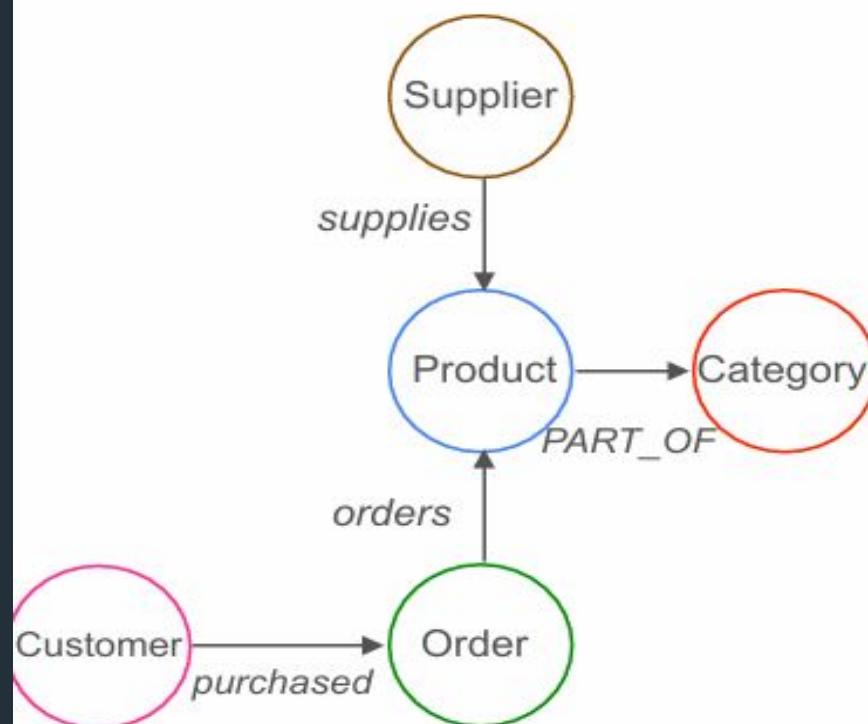
Started streaming 1 records in less than 1 ms and completed after 25 ms.

```
neo4j$ Match ()-[r]→() return distinct type(r)
```

type(r)
---------

MATCH *pattern* RETURN *result*

node	( )
relationship	[ ]
path	(source node)-[ ]->(target node)



## MATCH Statement

Retrieve the product name and product unit price of all products that belong to category “Meat/Poultry”

neo4j\$

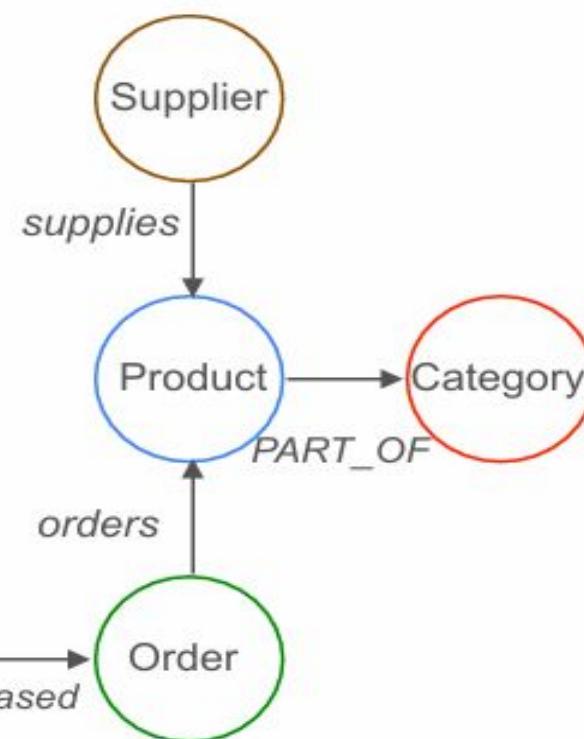
```

neo4j$ Match ()-[r:ORDERS]→()-[part:PART_OF]→(c:Category) return c.c...
  
```

c.categoryName	average_price
"Confections"	531.36666666666664
"Dairy Products"	504.31836734693877
"Grains/Cereals"	278.67
"Meat/Poultry"	766.9142857142859
"Produce"	645.6571428571427
"Seafood"	434.8282051282051

MATCH pattern RETURN result

node	( )
relationship	[ ]
path	(source node)-[ ]->(target node)



## MATCH Statement

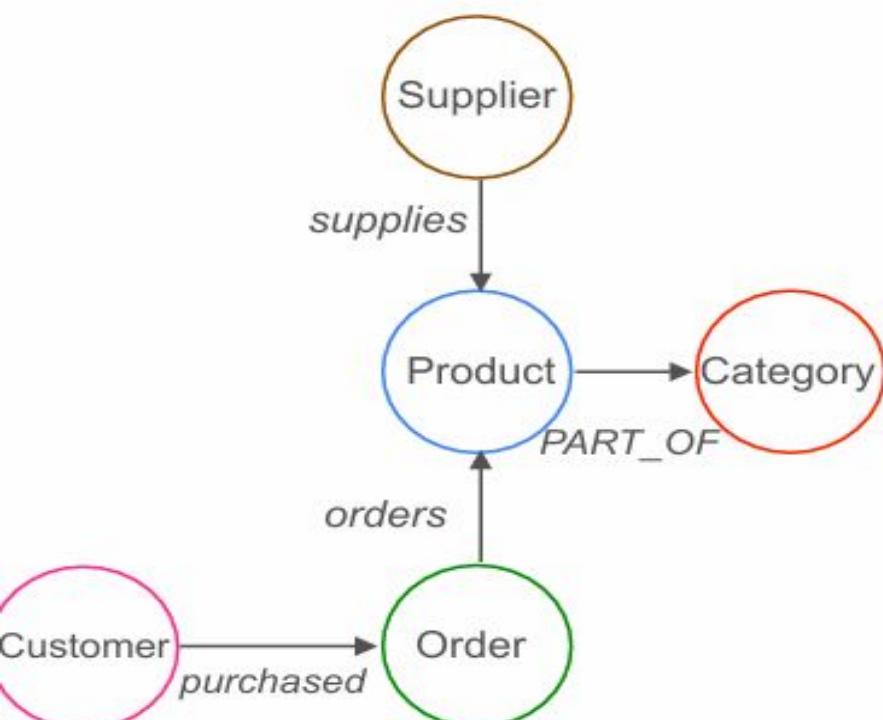
Retrieve the product name and product unit price of all products that belong to category "Meat/Poultry"

neo4j\$ Match (p:Product)-[:PART\_OF]→(c:Category)

neo4j\$ Match ()-[r:ORDERS]→()-[part:PART\_OF]→(c:Category) return c.c...

c.categoryName	average_price
"Confections"	531.3666666666666
"Dairy Products"	504.31836734693877
"Grains/Cereals"	278.67
"Meat/Poultry"	766.9142857142859
"Produce"	645.6571428571427
"Seafood"	434.8282051282051

```
MATCH pattern RETURN result
node      ()
relationship []
path      (source node)-[ ]->(target node)
```



## MATCH Statement

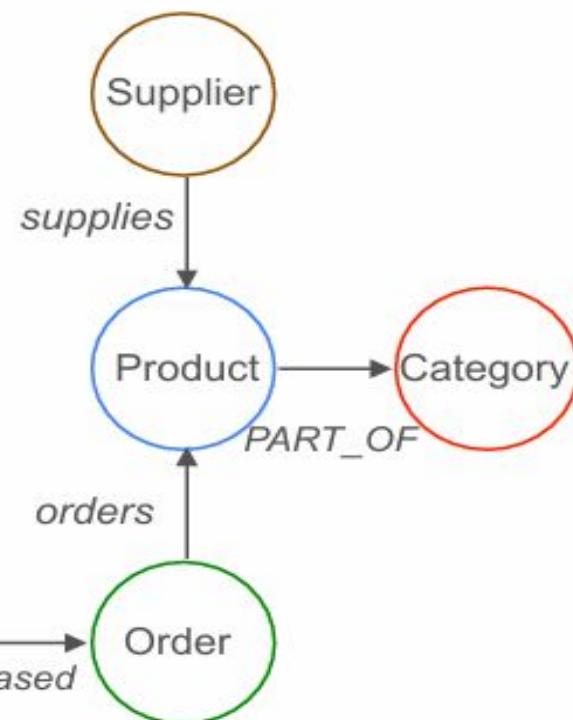
Retrieve the product name and product unit price of all products that belong to category “Meat/Poultry”

```
1 Match (p:Product)-[:PART_OF]→(c:Category)
2 where c.categoryName="Meat/Poultry"
```

c.categoryName	average_price
"Confections"	531.3666666666664
"Dairy Products"	504.31836734693877
"Grains/Cereals"	278.67
"Meat/Poultry"	766.9142857142859
"Produce"	645.6571428571427
"Seafood"	434.8282051282051

MATCH pattern RETURN result

node	( )
relationship	[ ]
path	(source node)-[ ]->(target node)



## MATCH Statement

Retrieve the product name and product unit price of all products that belong to category "Meat/Poultry"

neo4j\$

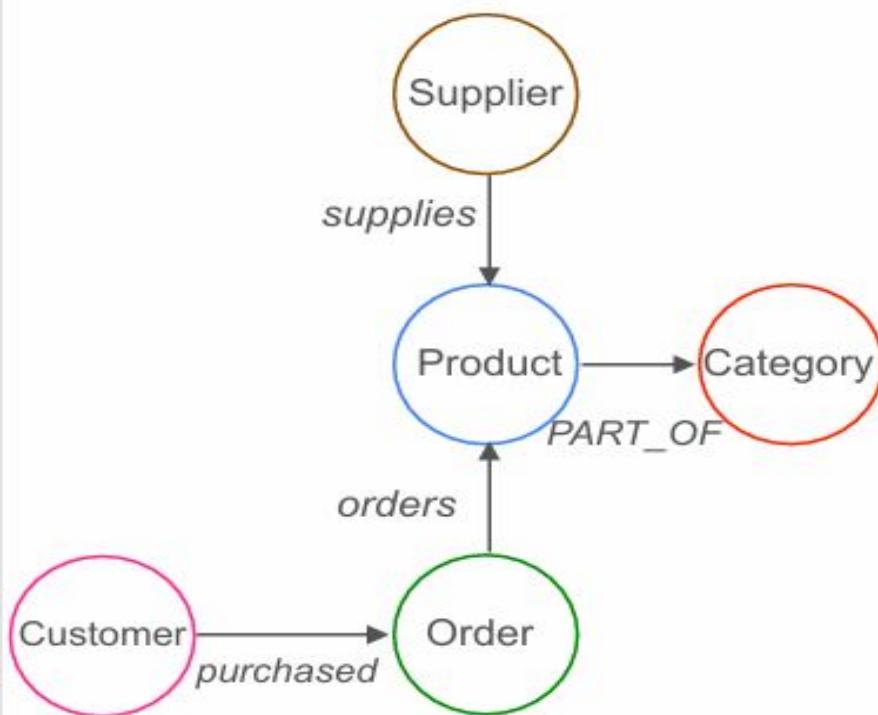
```

1 Match (p:Product)-[:PART_OF]→(c:Category)
2 where c.categoryName="Meat/Poultry"
3 return p.productName, p.unitPrice
  
```

	p.productName	p.unitPrice
1	"Perth Pasties"	32.8
2	"Tourtière"	7.45
3	"Alice Mutton"	39.0
4	"Pâté chinois"	24.0
5	"Thüringer Rostbratwurst"	123.79
6	"Mishi Kobe Niku"	97.0

MATCH *pattern* RETURN *result*

node	( )
relationship	[ ]
path	(source node)-[ ]->(target node)



## MATCH Statement

Retrieve the product name of all products ordered by the customer “QUEDE”

neo4j\$

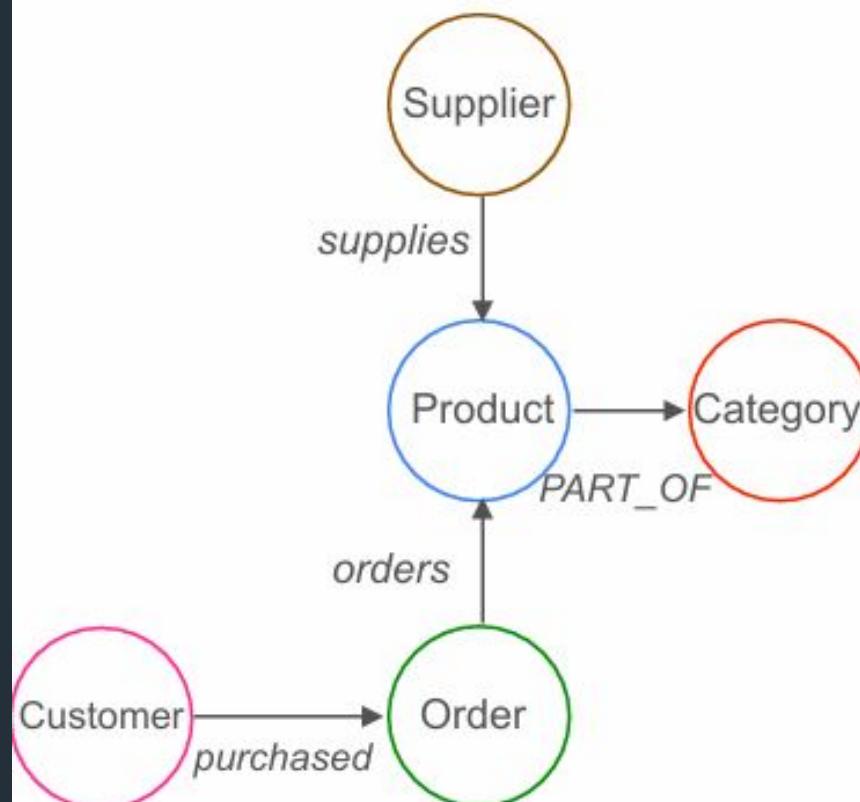
```

neo4j$ Match (p:Product)-[:PART_OF]→(c:Category {categoryName:"Meat/P...")
```

p.productName	p.unitPrice
"Perth Pasties"	32.8
"Tourtière"	7.45
"Alice Mutton"	39.0
"Pâté chinois"	24.0
"Thüringer Rostbratwurst"	123.79
"Mishi Kobe Niku"	97.0

MATCH pattern RETURN result

node	( )
relationship	[ ]
path	(source node)-[ ]->(target node)



## MATCH Statement

Get the ID of other customers who ordered the same products as “QUEDE”

```

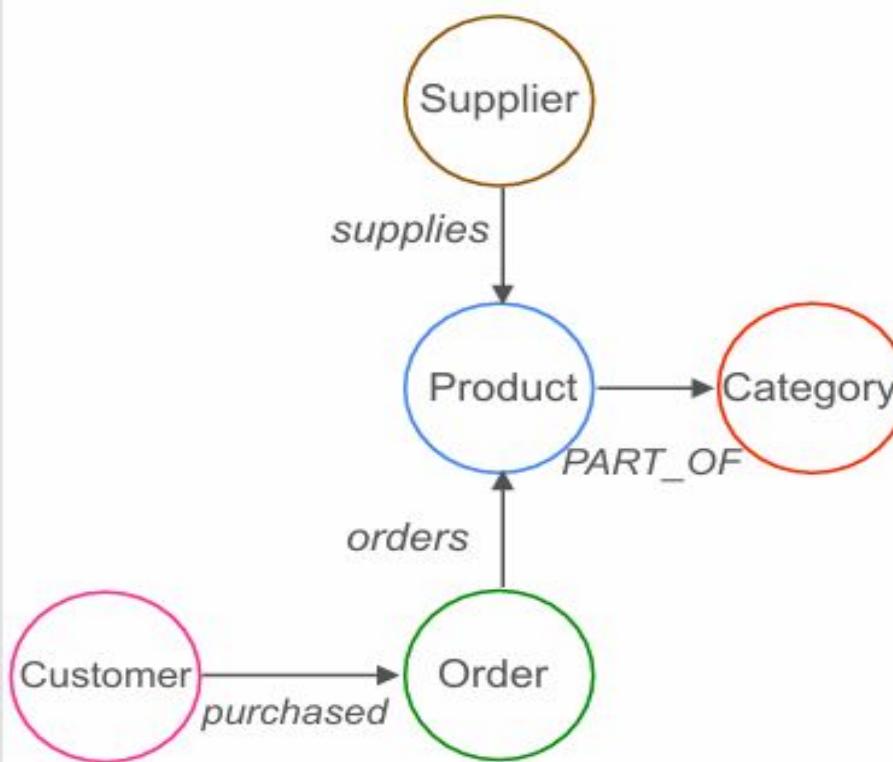
neo4j$ Match (c1:Customer {customerID:"QUEDE"}) -[:PURCHASED]→()-
[:ORDERS]→(p:Product)
  
```

Table	p.productName
Text	"Steeleye Stout"
Code	"Sir Rodney's Scones"
	"Gula Malacca"
	"Konbu"
	"Manjimup Dried Apples"

Started streaming 5 records after 1 ms and completed after 1 ms.

MATCH pattern RETURN result

node	( )
relationship	[ ]
path	(source node)-[ ]->(target node)



## MATCH Statement

Retrieve the orders that contain at most two products

```

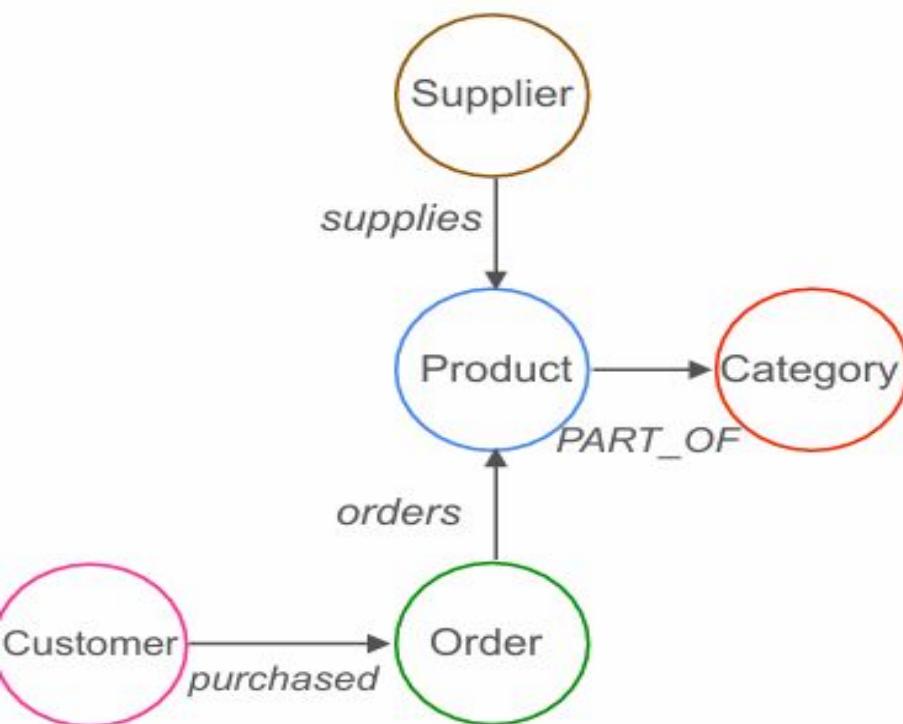
neo4j$ Match (c1:Customer {customerID:'QUEDE'}) -[:PURCHASED]→()-[ :OR...

```

c2.customerID
"QUICK"
"SAVEA"
"ROMEY"
"ISLAT"
"WARTH"
"CENTC"

MATCH pattern RETURN result

node	( )
relationship	[ ]
path	(source node)-[ ]->(target node)



## MATCH Statement

Retrieve the orders that contain at most two products

neo4j\$

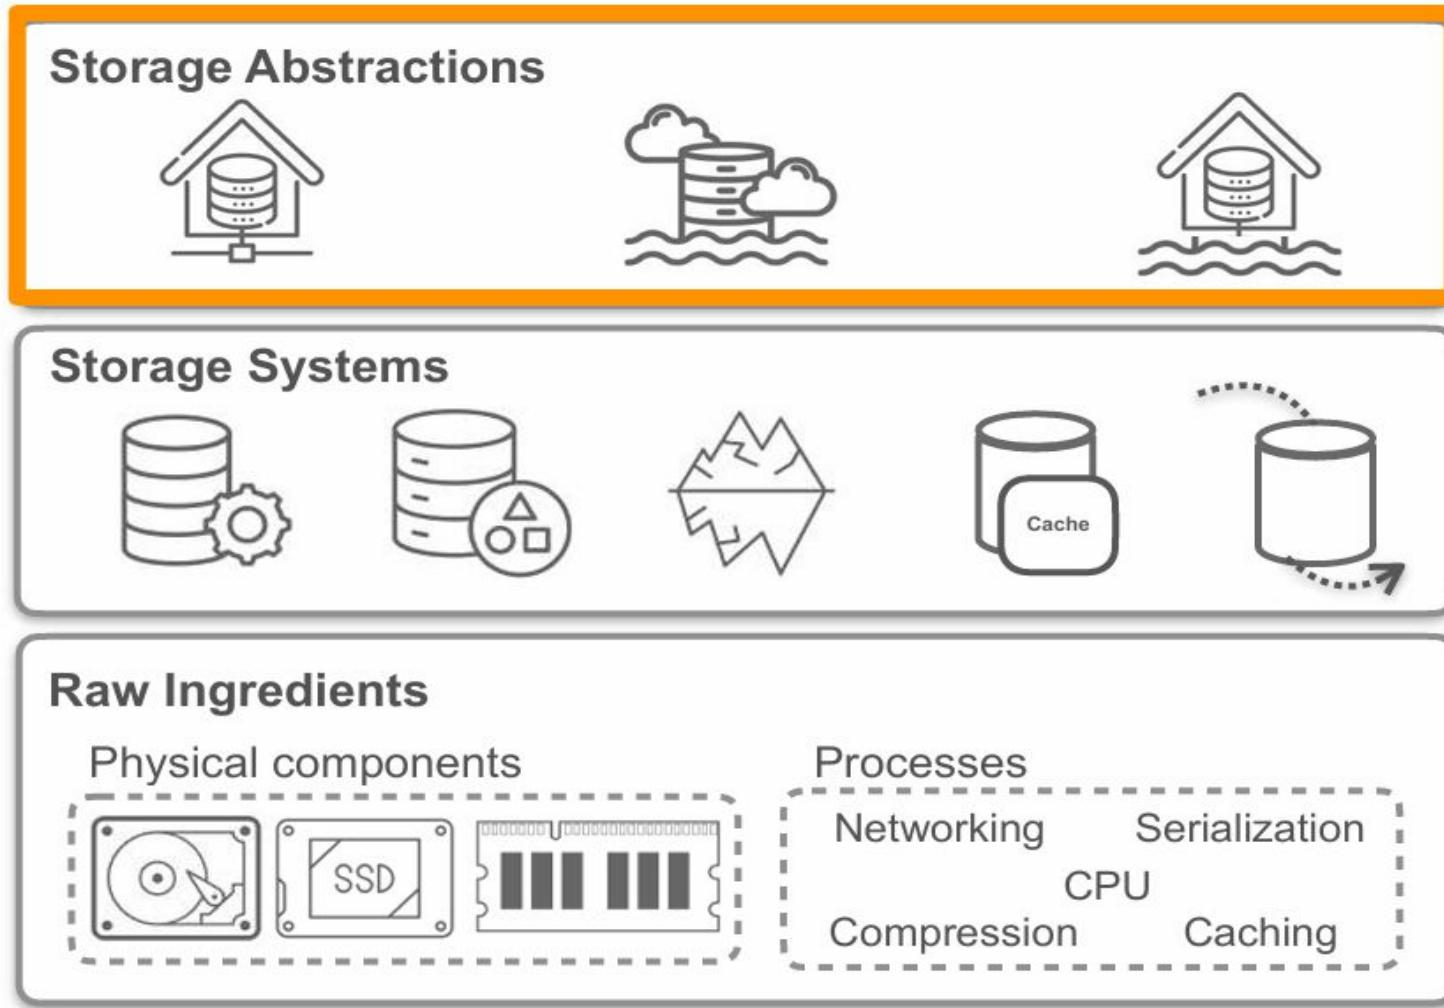
```

1 Match (o:Order)-[:ORDERS]→(p:Product)
2 return o.orderID as ID, count(p) as countProd
  
```

ID	countProd
"10294"	5
"10317"	1
"10285"	3
"10342"	4
"10255"	4
"10327"	4

# Storage Abstractions

# Storage Hierarchy



# Storage Hierarchy



## Data Warehouse

Cloud data  
warehouse

## Data Lake

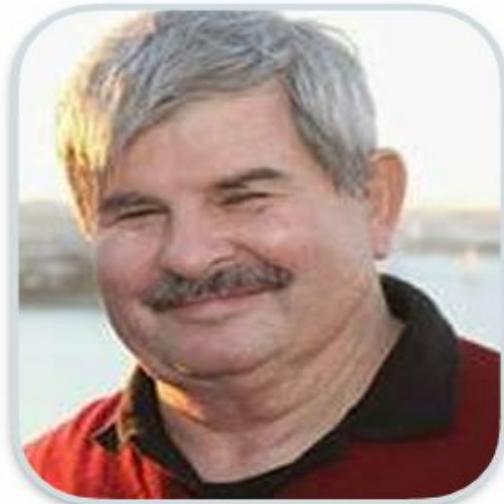
Supports  
growing storage  
needs

## Data Lakehouse

Combines the  
advantages of data  
warehouses and  
data lakes

# Data Warehouse - Key Architectural ideas

# Data Warehouse



**Bill Inmon**

“Father” of the  
Data Warehouse

## **Data Warehouse:**

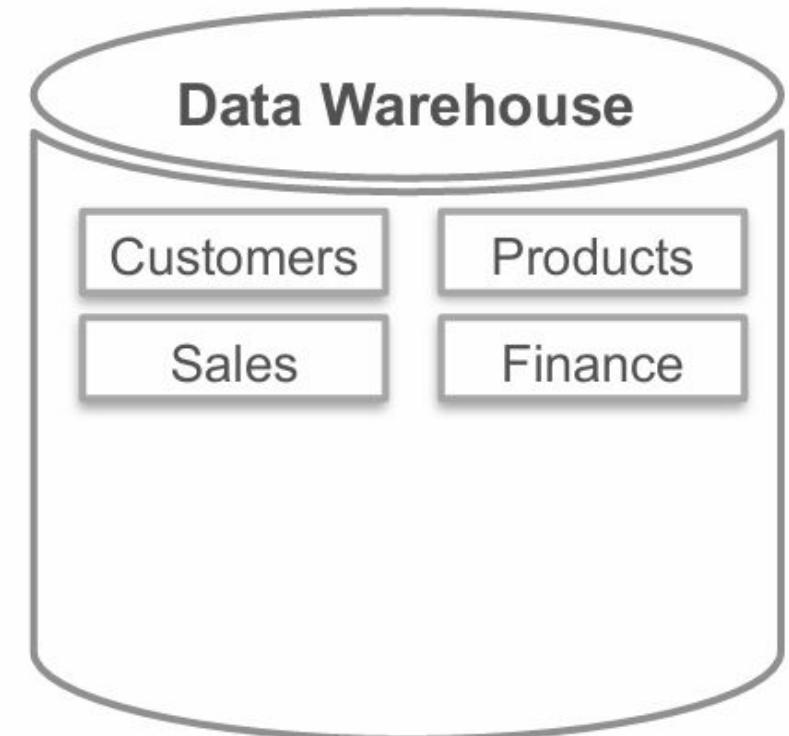
A subject-oriented, integrated, nonvolatile, and  
time-variant collection of data in support of  
management's decisions.

# Data Warehouse

## Subject-Oriented

Organizes and stores  
data around key  
business domains

(Models data to support  
decision making)



# Data Warehouse

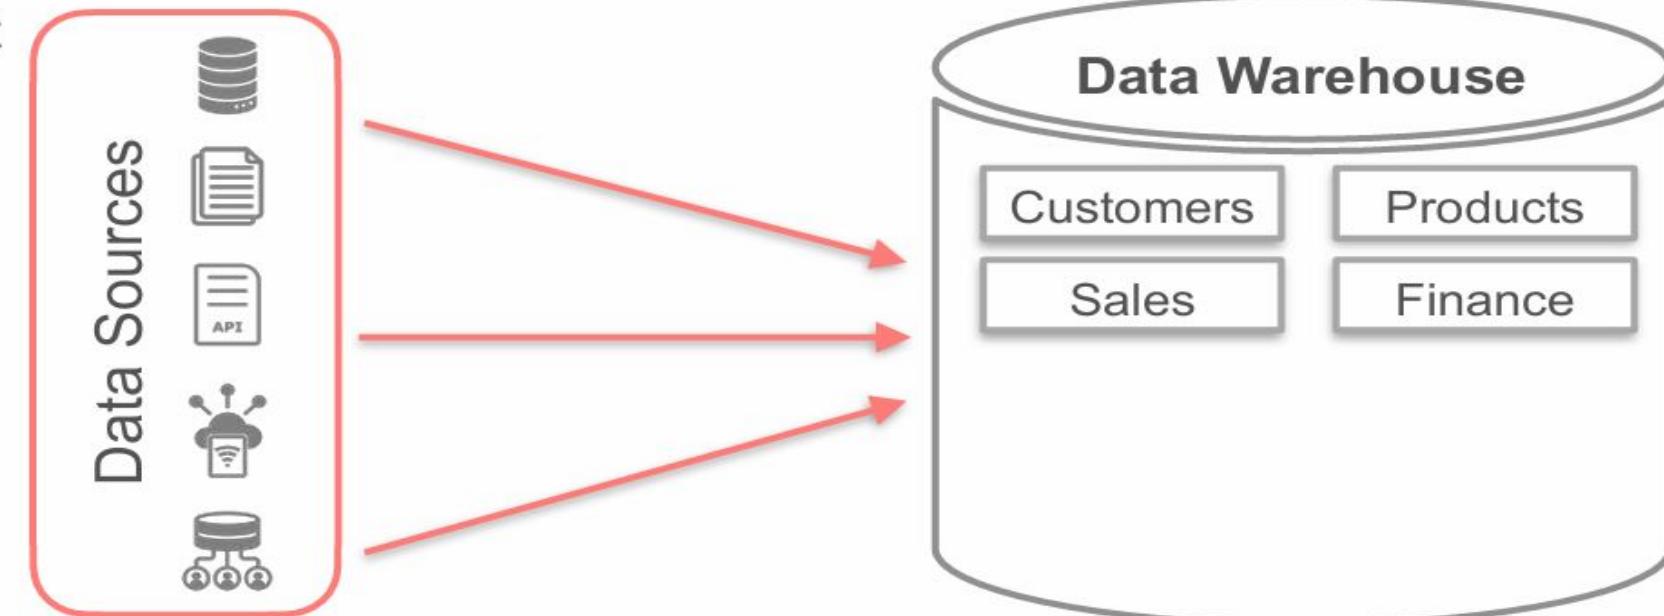
## Subject-Oriented

Organizes and stores data around key business domains

(Models data to support decision making)

## Integrated

Combines data from different sources into a consistent format



# Data Warehouse

## Subject-Oriented

Organizes and stores data around key business domains

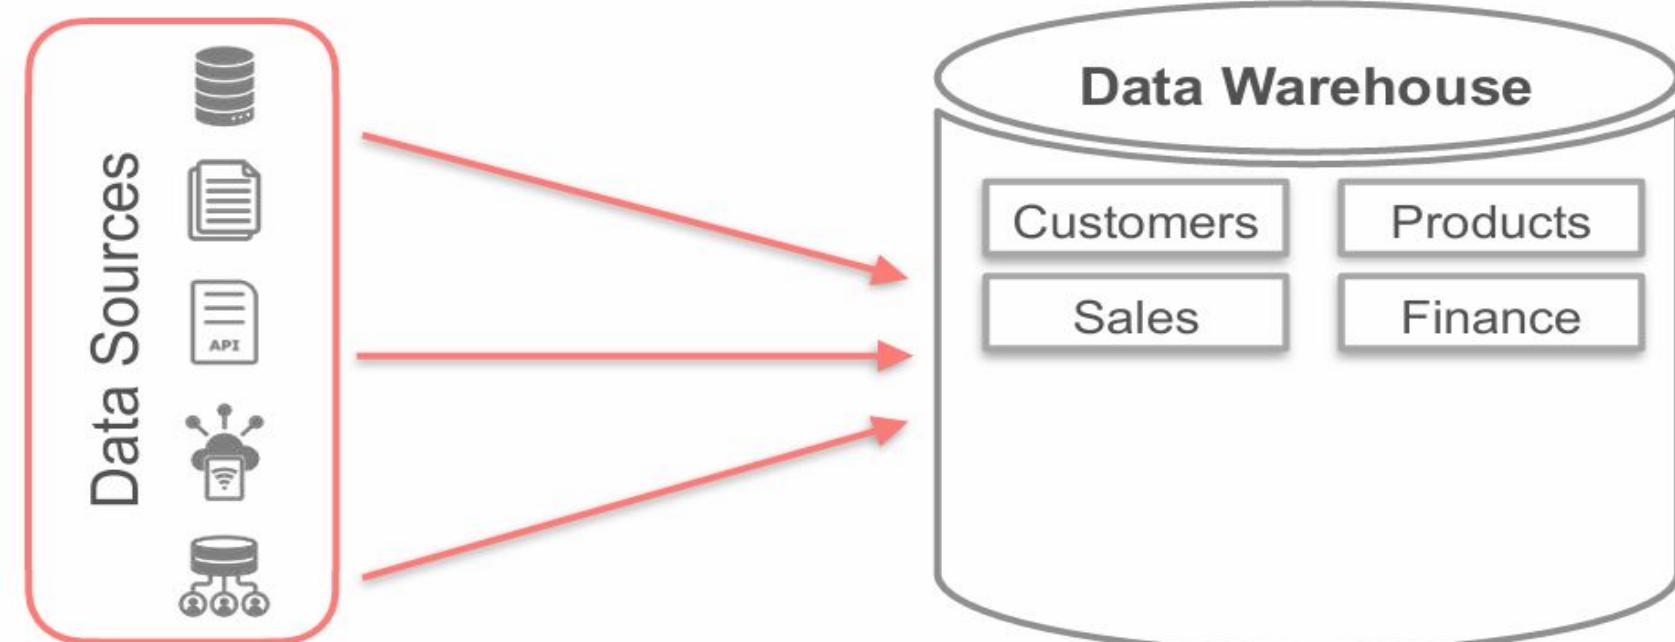
(Models data to support decision making)

## Integrated

Combines data from different sources into a consistent format

## Nonvolatile

Data is read-only and cannot be deleted or updated



# Data Warehouse

## Subject-Oriented

Organizes and stores data around key business domains

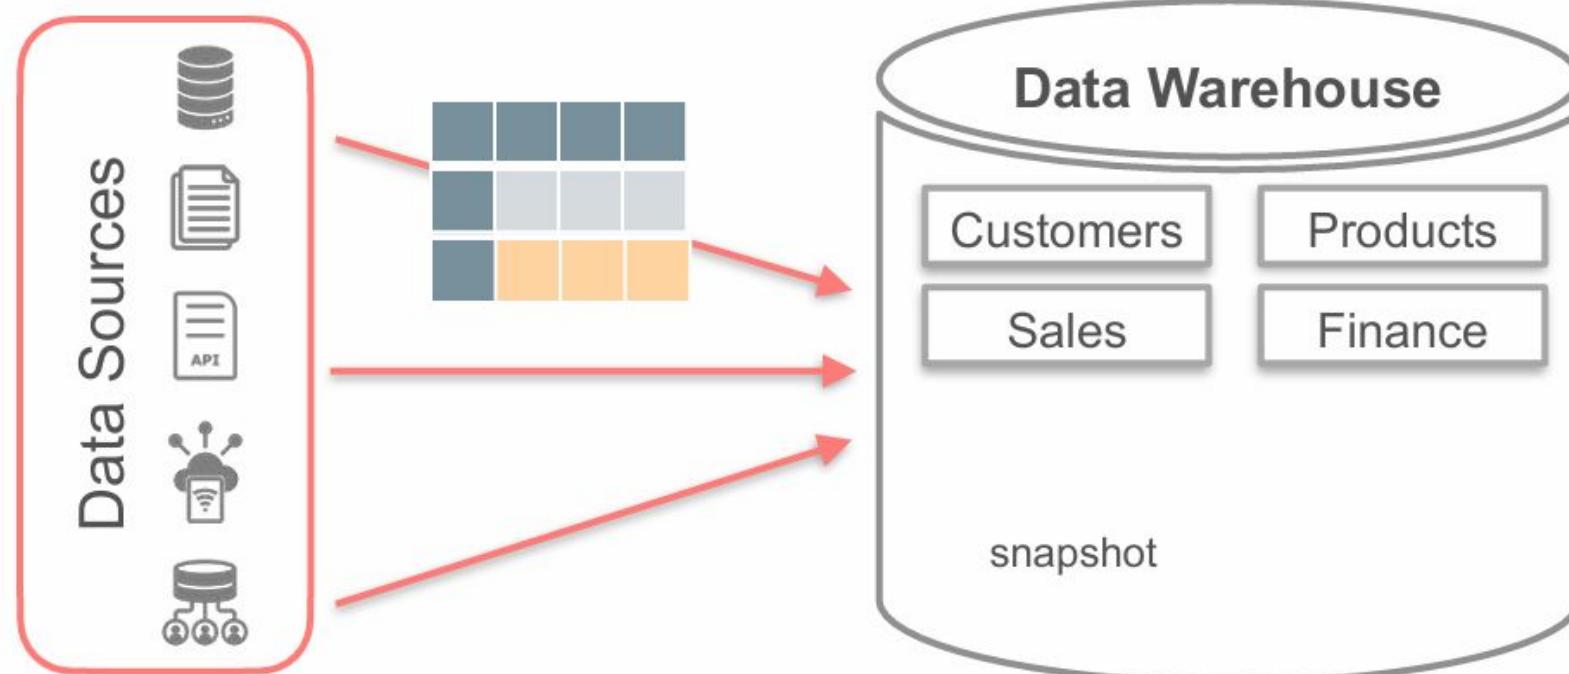
(Models data to support decision making)

## Integrated

Combines data from different sources into a consistent format

## Nonvolatile

Data is read-only and cannot be deleted or updated



# Data Warehouse

## Subject-Oriented

Organizes and stores data around key business domains

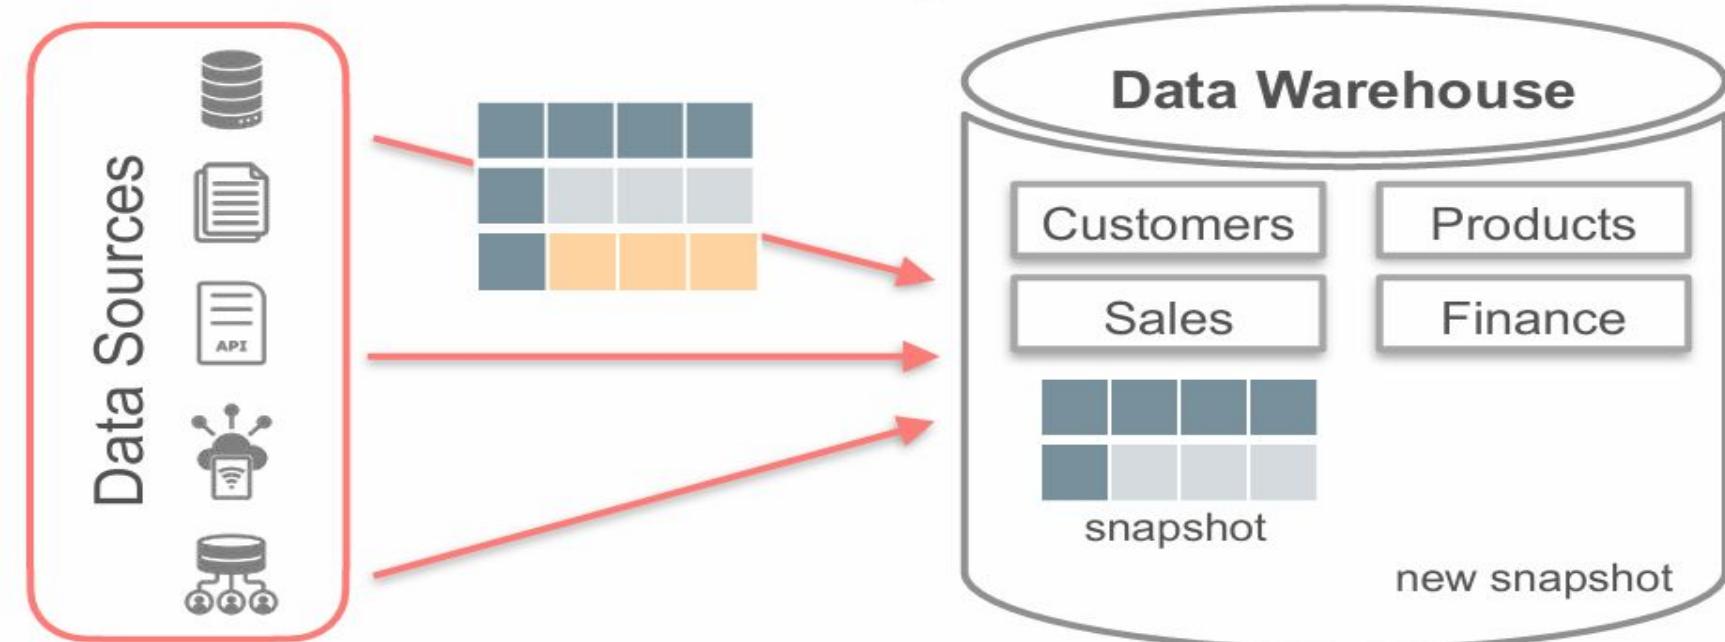
(Models data to support decision making)

## Integrated

Combines data from different sources into a consistent format

## Nonvolatile

Data is read-only and cannot be deleted or updated



# Data Warehouse

## Subject-Oriented

Organizes and stores data around key business domains

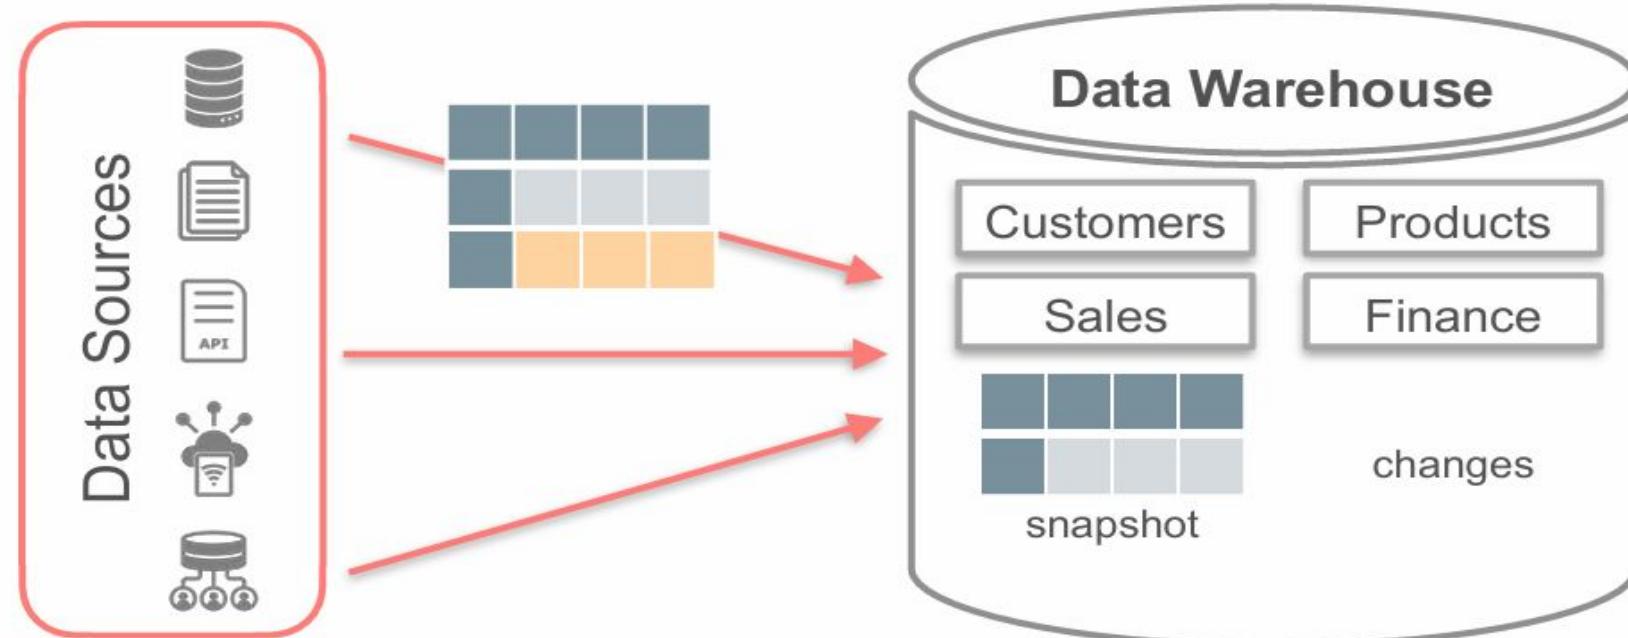
(Models data to support decision making)

## Integrated

Combines data from different sources into a consistent format

## Nonvolatile

Data is read-only and cannot be deleted or updated



# Data Warehouse

## Subject-Oriented

Organizes and stores data around key business domains

(Models data to support decision making)

## Integrated

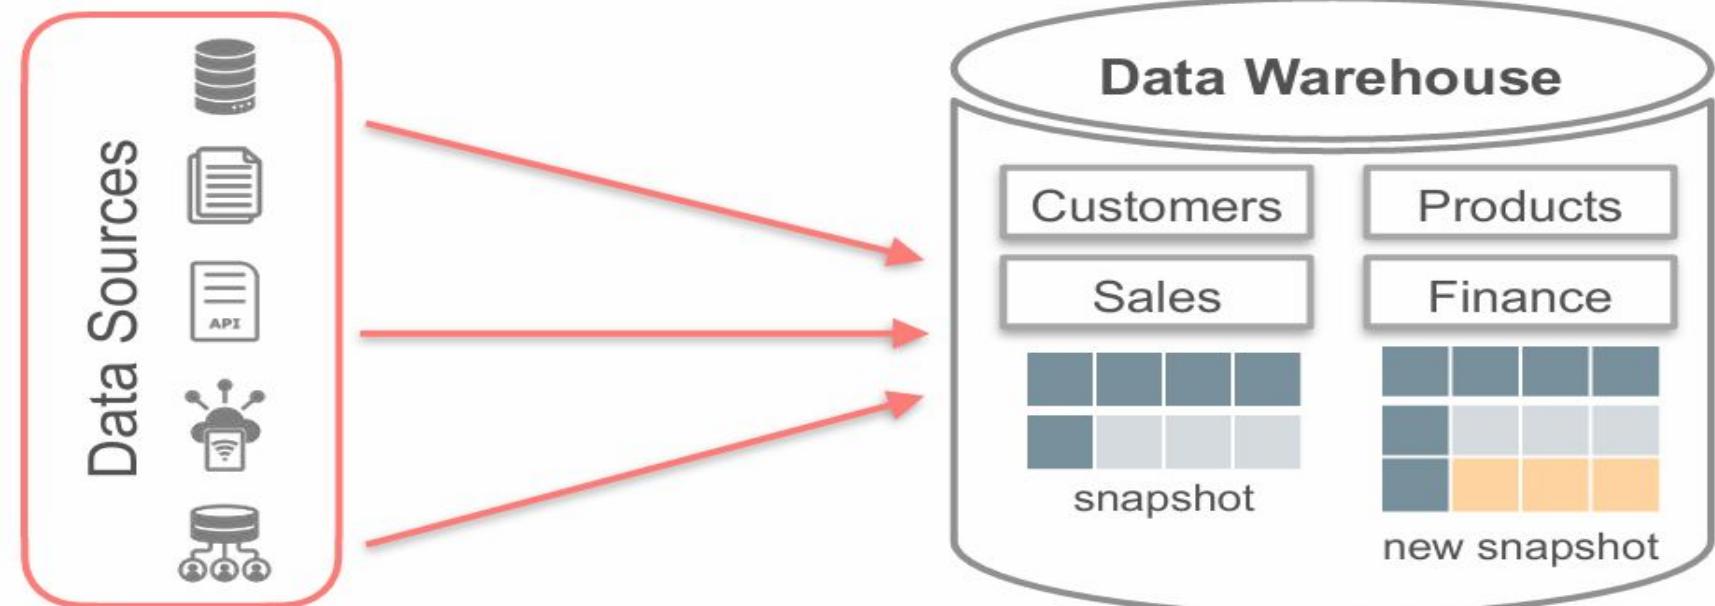
Combines data from different sources into a consistent format

## Nonvolatile

Data is read-only and cannot be deleted or updated

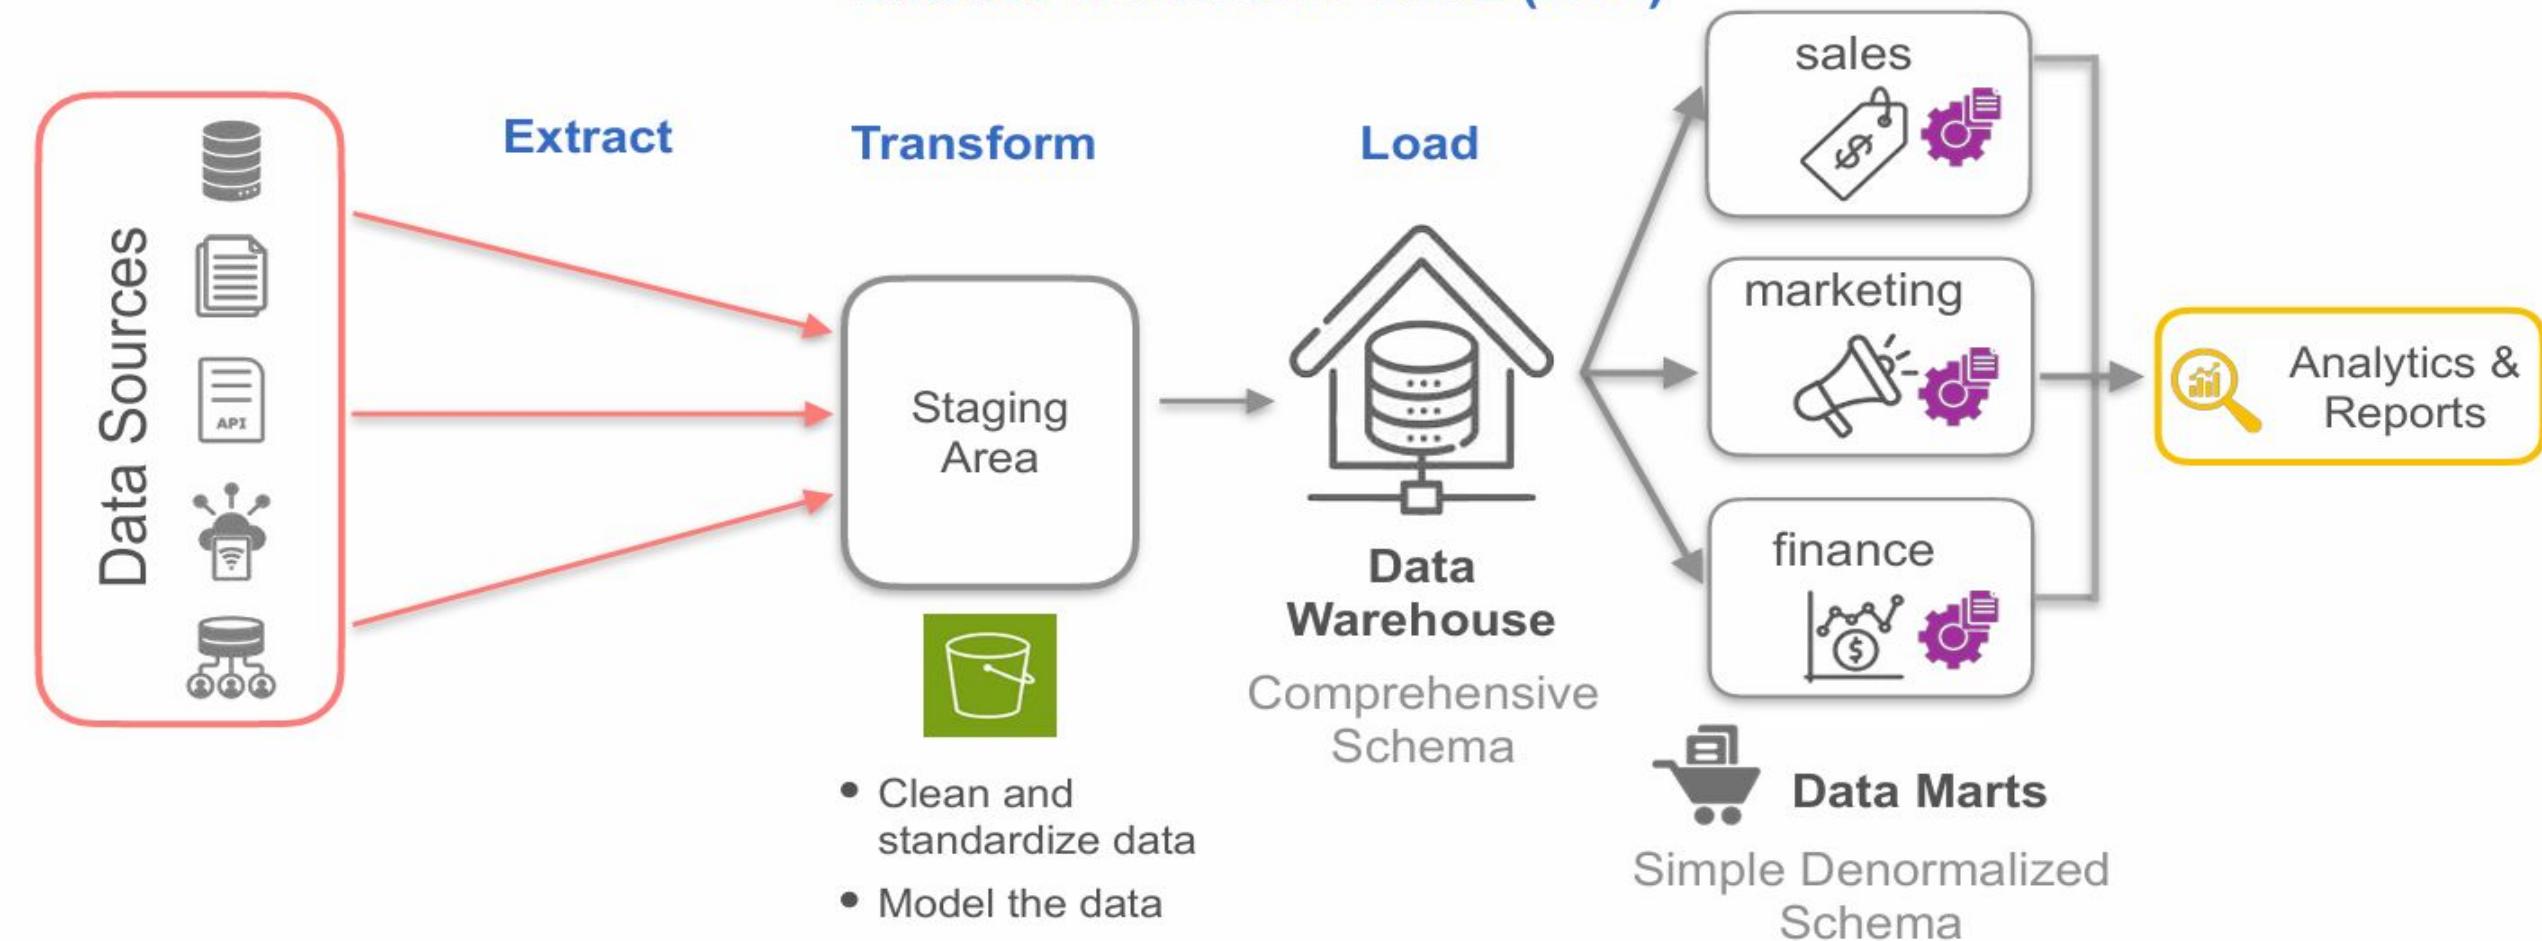
## Time-variant

Stores current and historical data (Unlike OLTP systems)

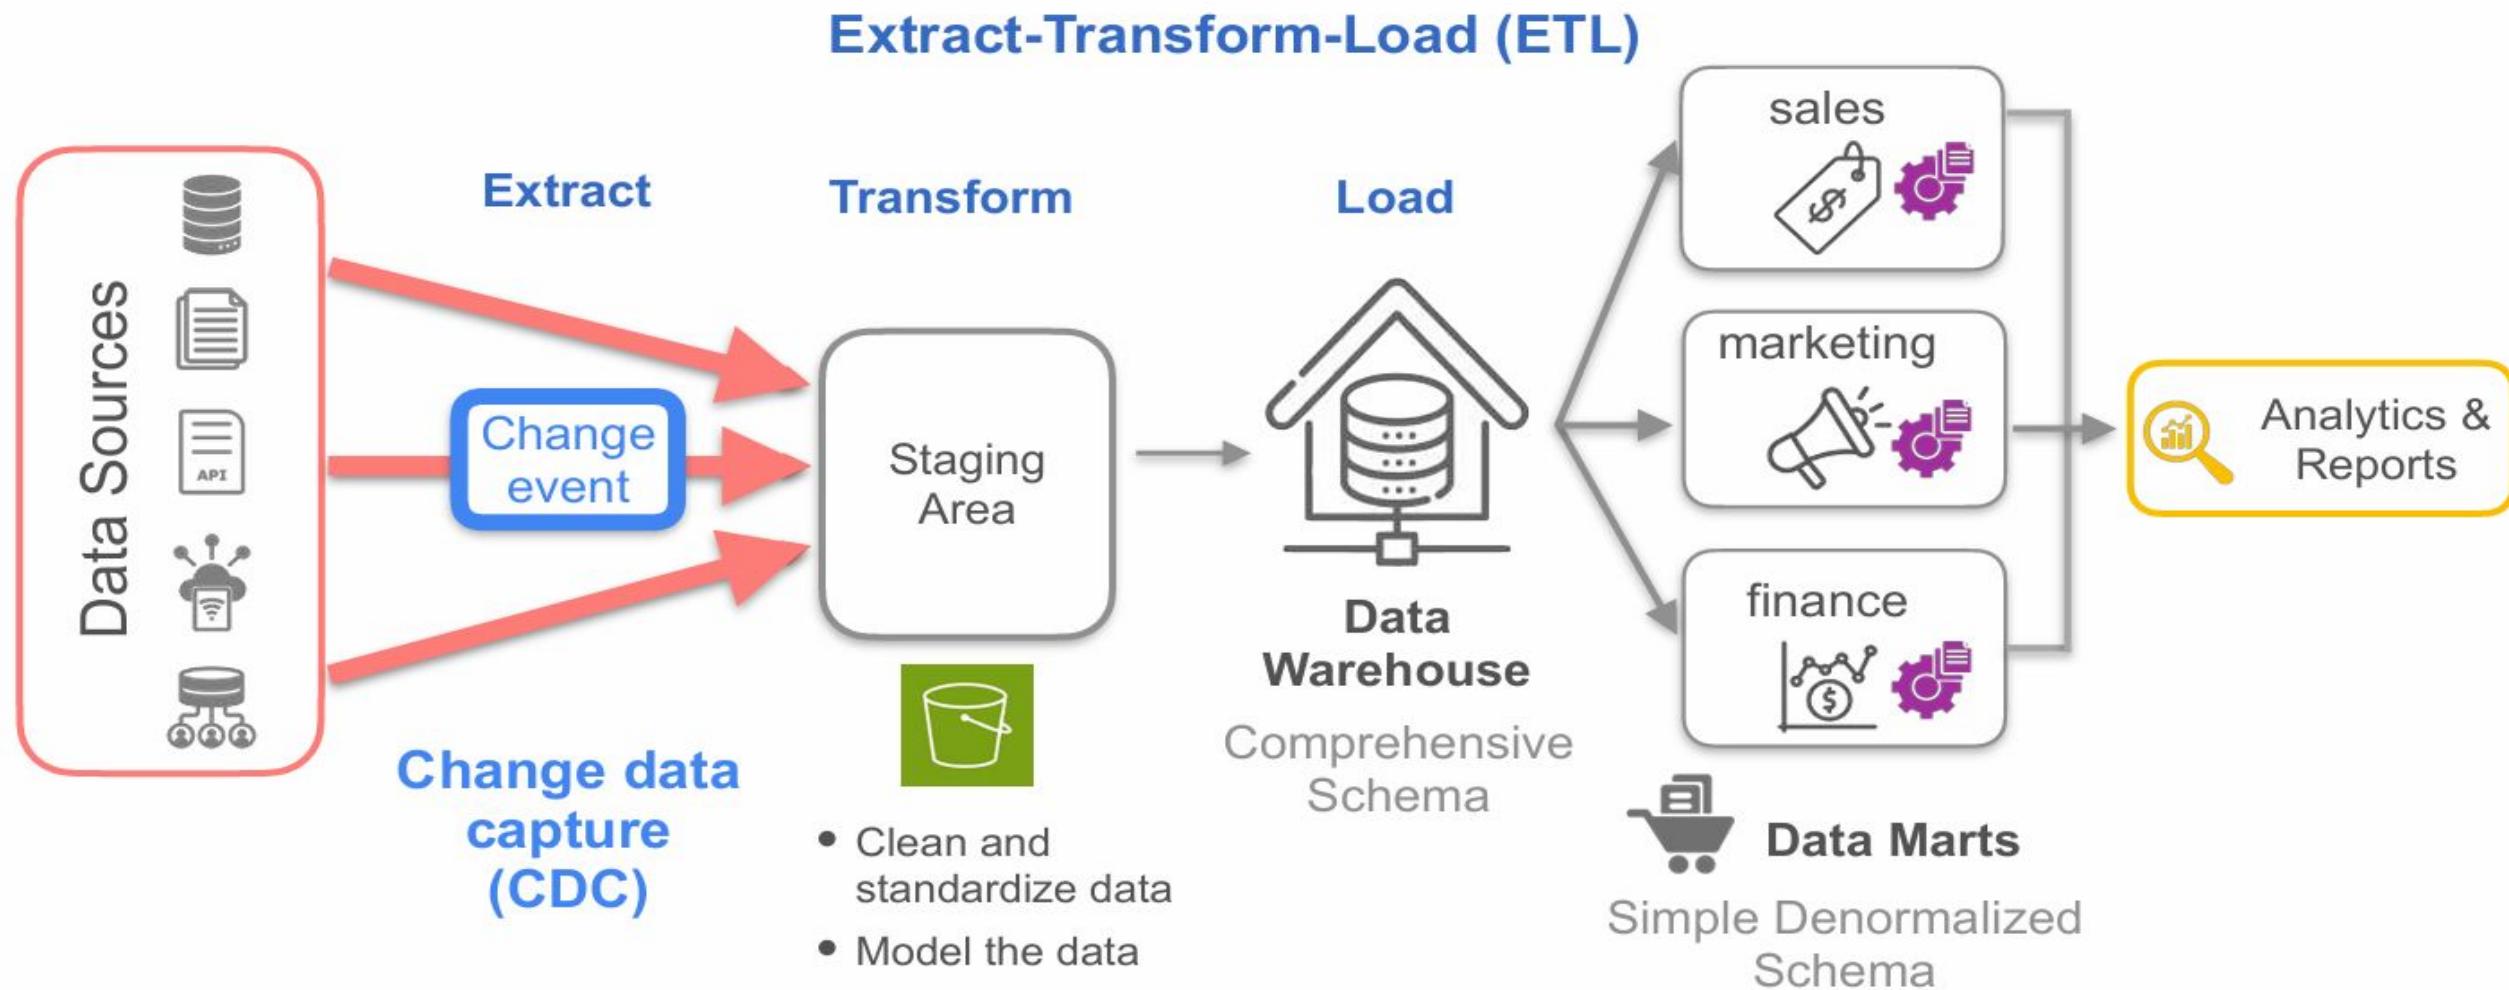


# Data Warehouse-Centric Architecture

## Extract-Transform-Load (ETL)



# Change Data Capture





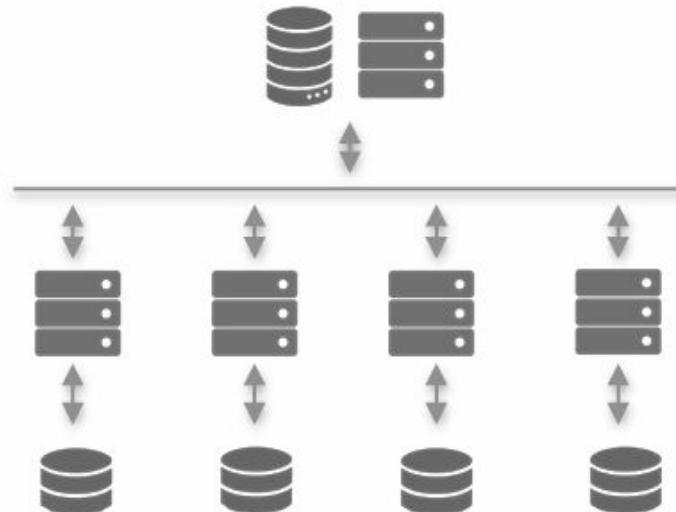
# Data Warehouse Implementation

## Early Data Warehouses



Big monolithic server

## Data Warehouses with Massively Parallel Processing (MPP)



- Scans large amounts of data in parallel
- Complex configurations and requires effort to maintain

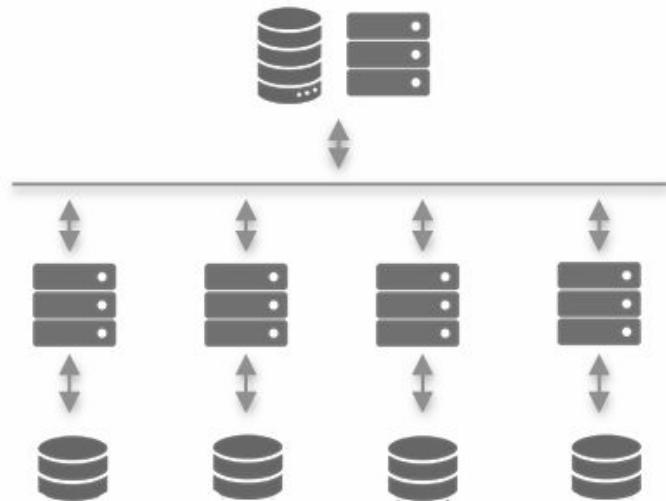
# Data Warehouse Implementation

## Early Data Warehouses



Big monolithic server

## Data Warehouses with Massively Parallel Processing (MPP)



- Scans large amounts of data in parallel
- Complex configurations and requires effort to maintain

## Modern Cloud Data Warehouses



Amazon Redshift



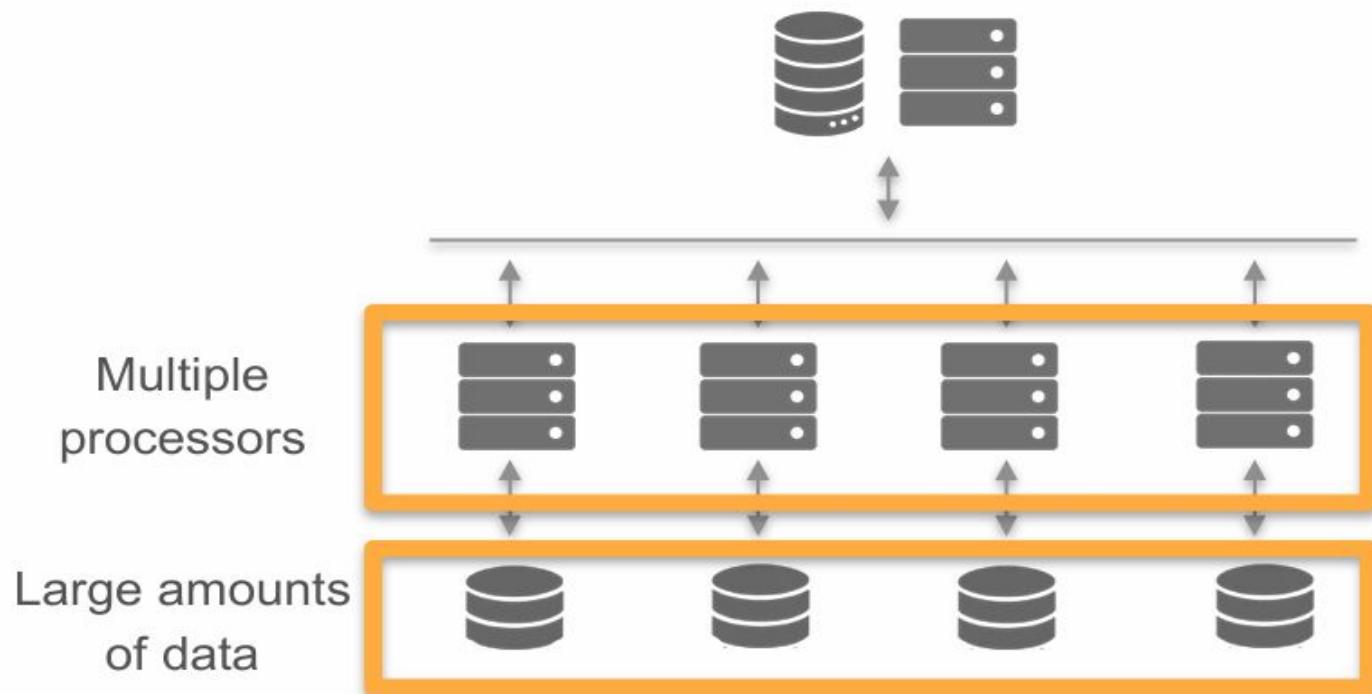
Google Big Query



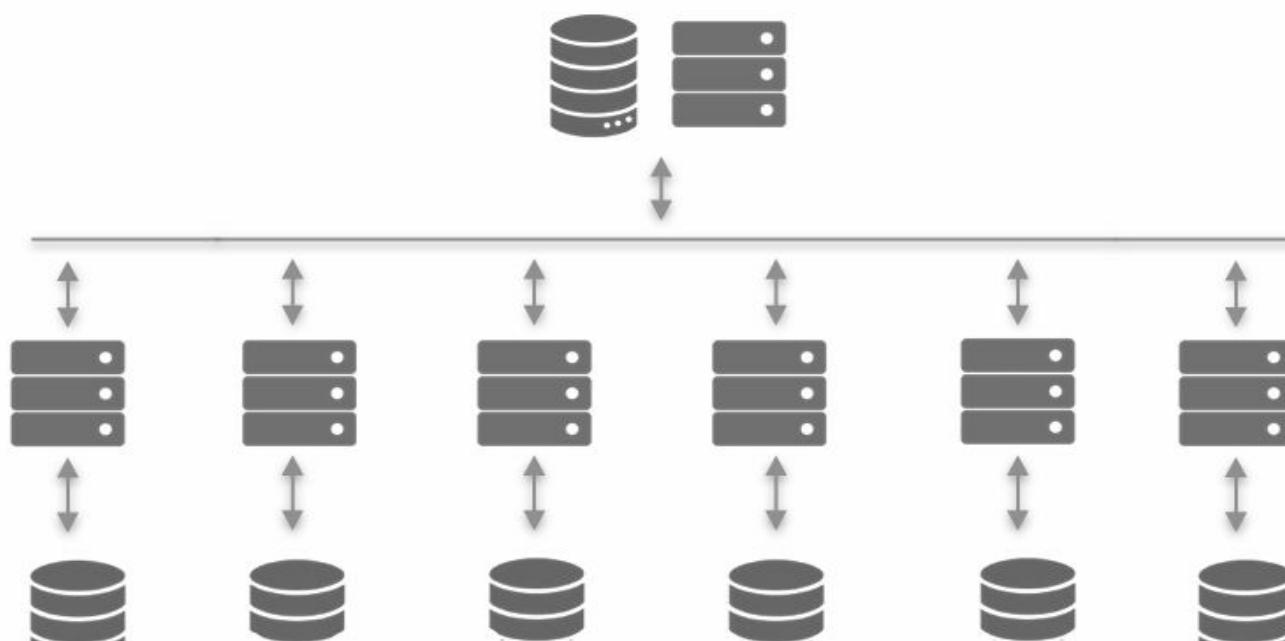
- Separates compute from storage
- Expands the capability of MPP systems

# Modern Cloud Data Warehouses

# Massively Parallel Processing



# Massively Parallel Processing - Cloud Data Warehouses



Massively Parallel Processing (MPP)



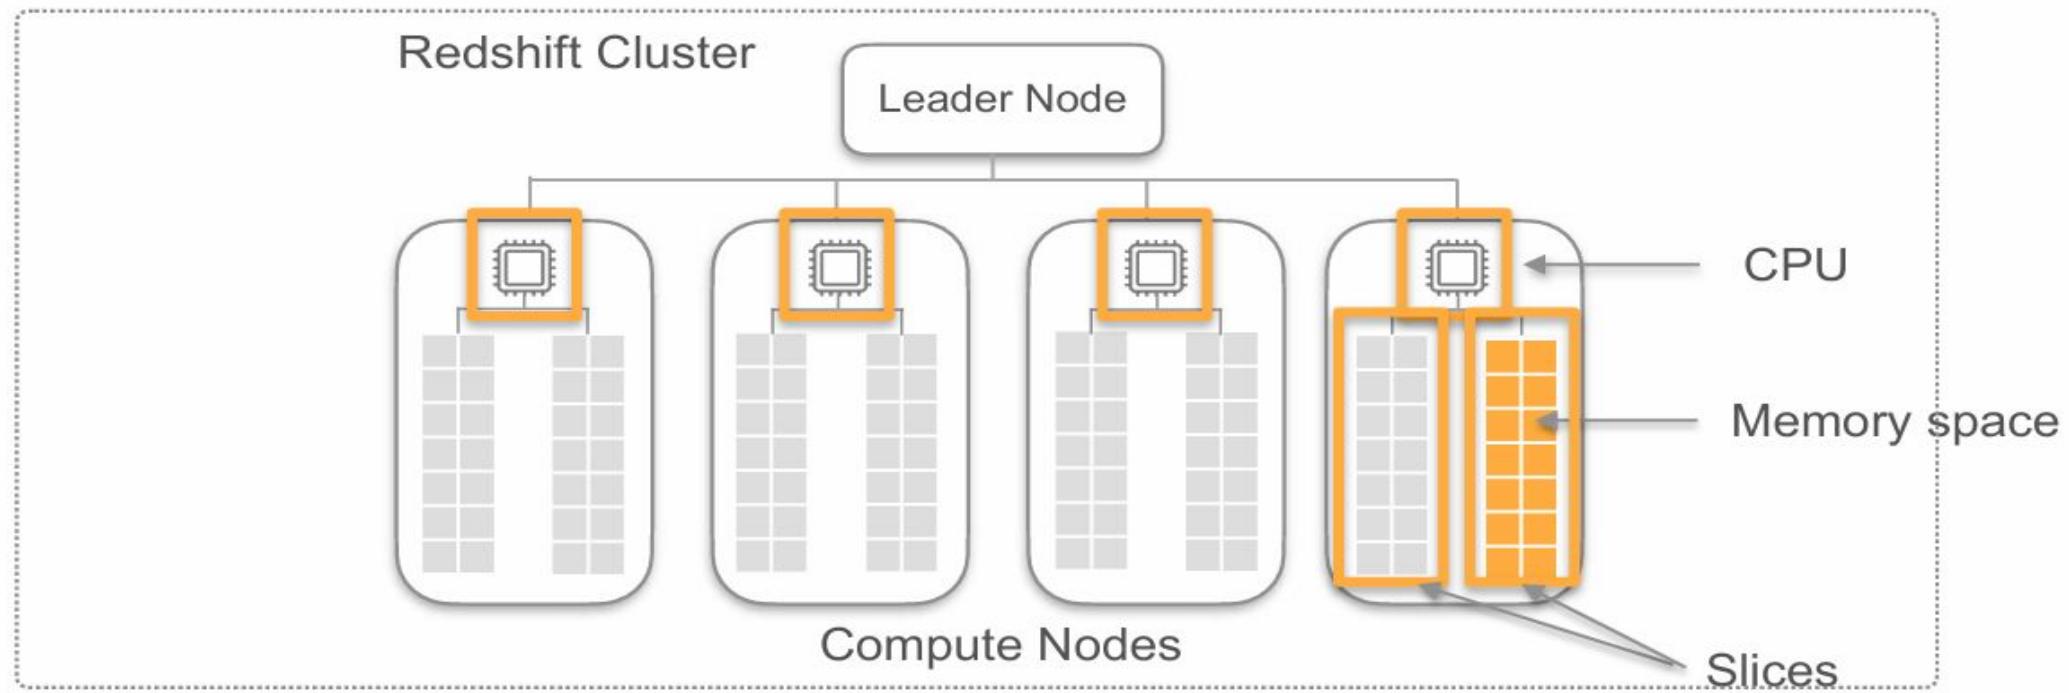
Amazon Redshift



Google  
Big Query

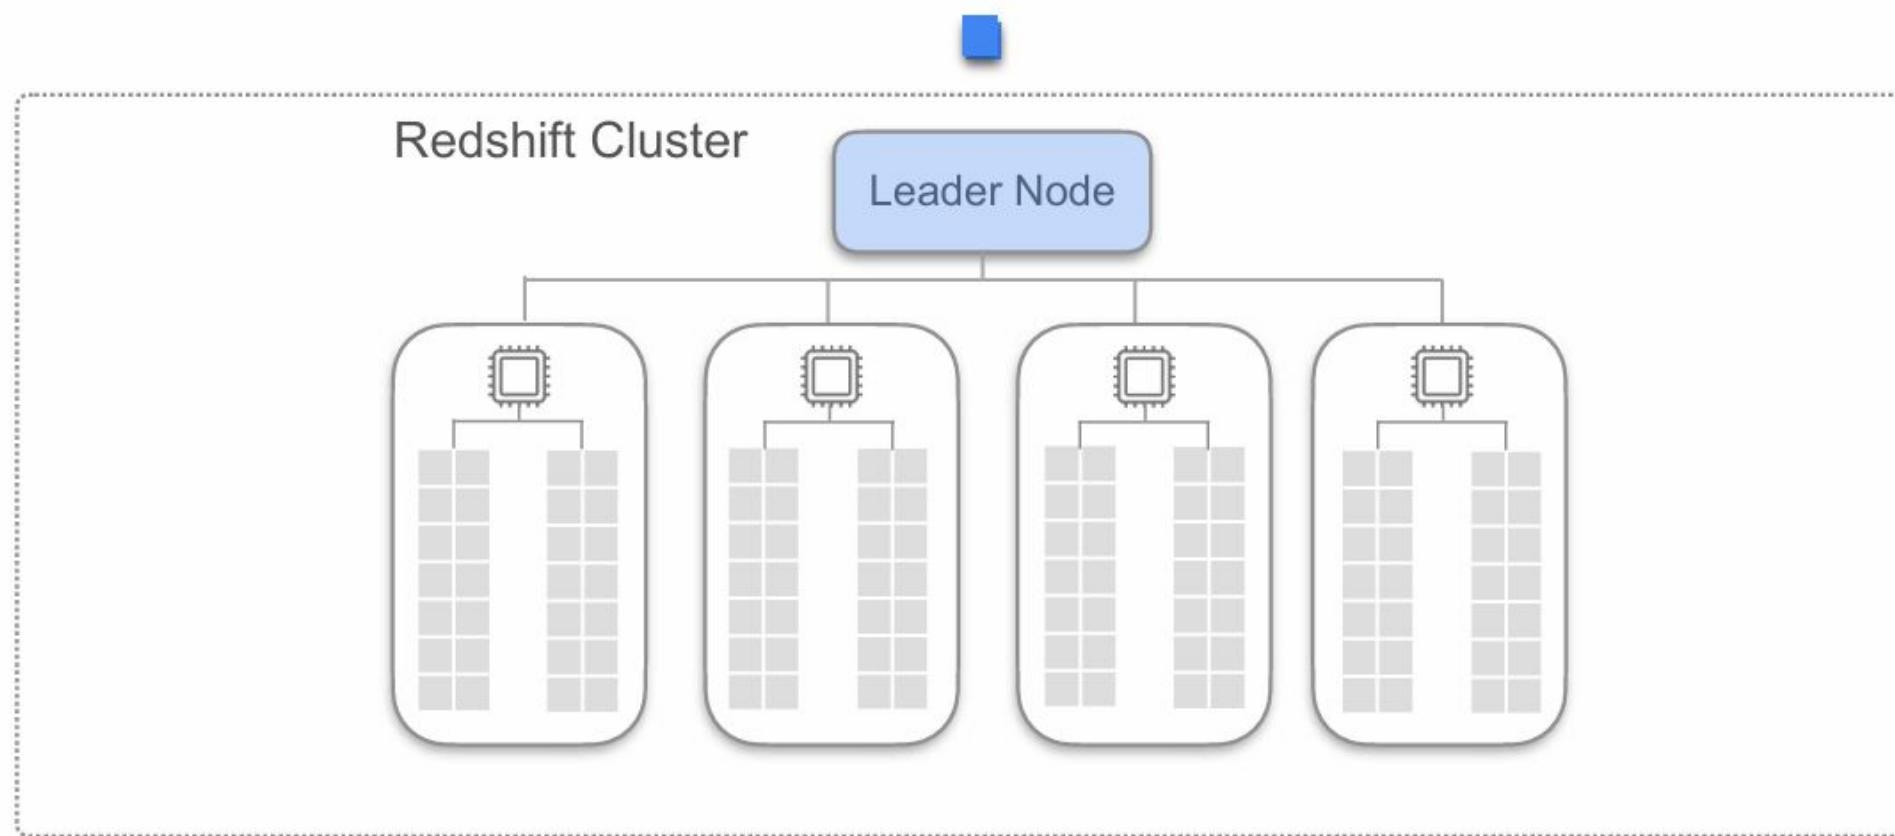


# MPP Architecture for Amazon Redshift



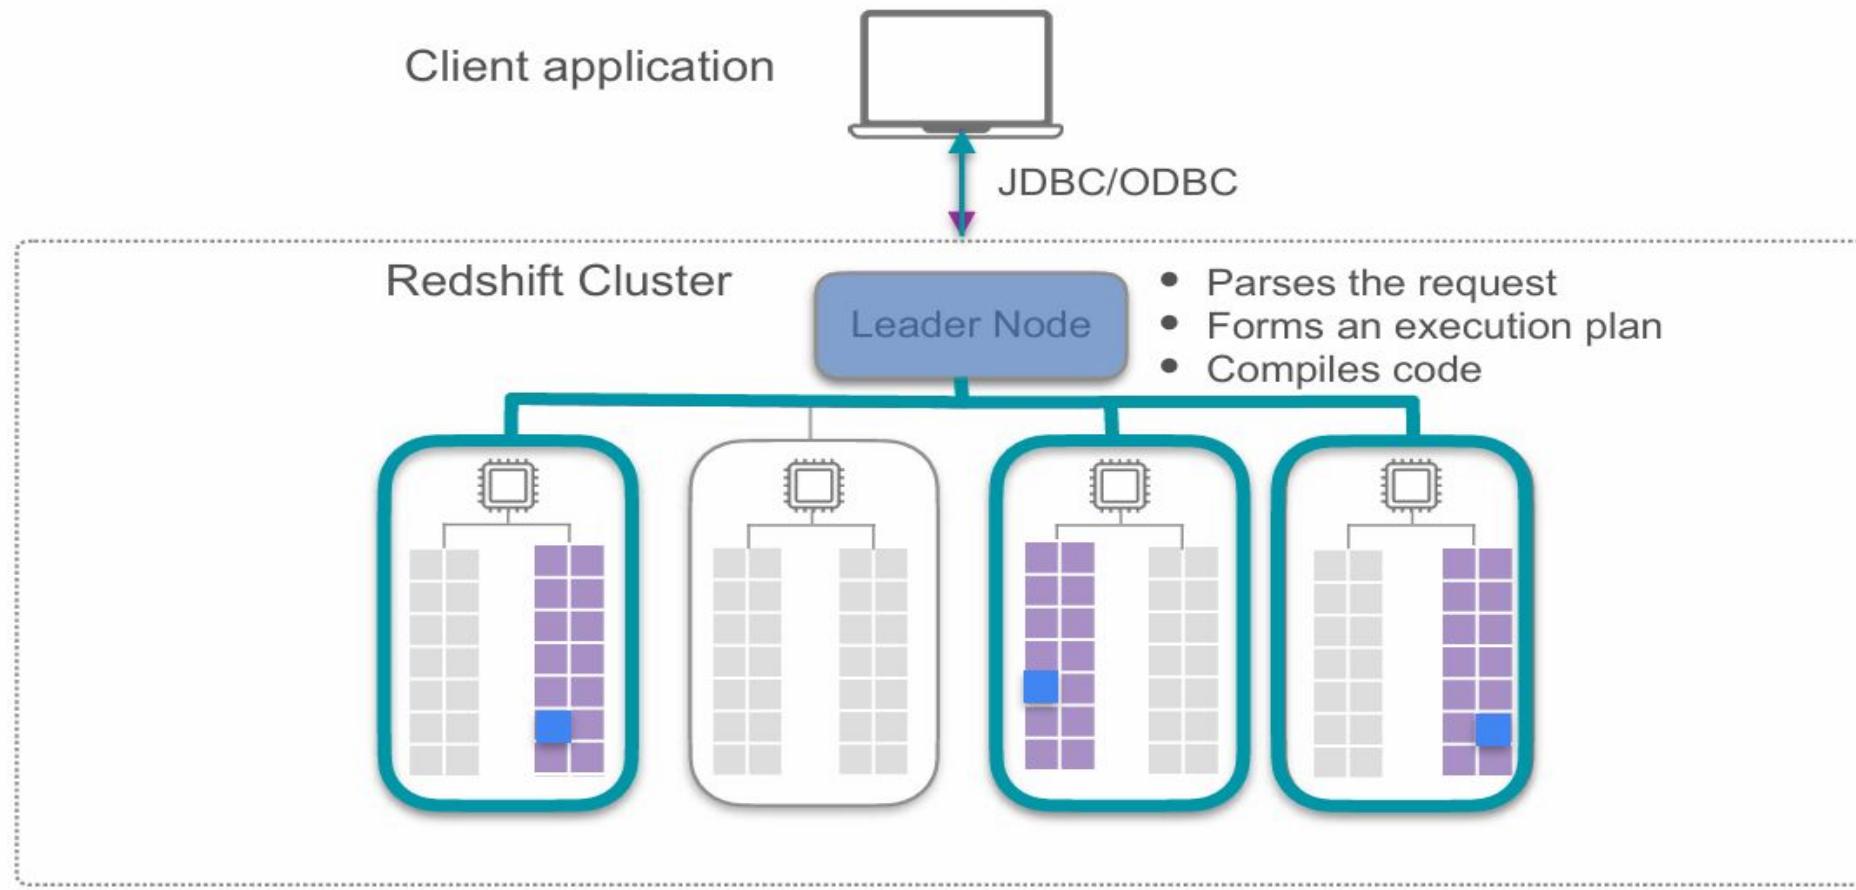
Amazon Redshift

# MPP Architecture for Amazon Redshift

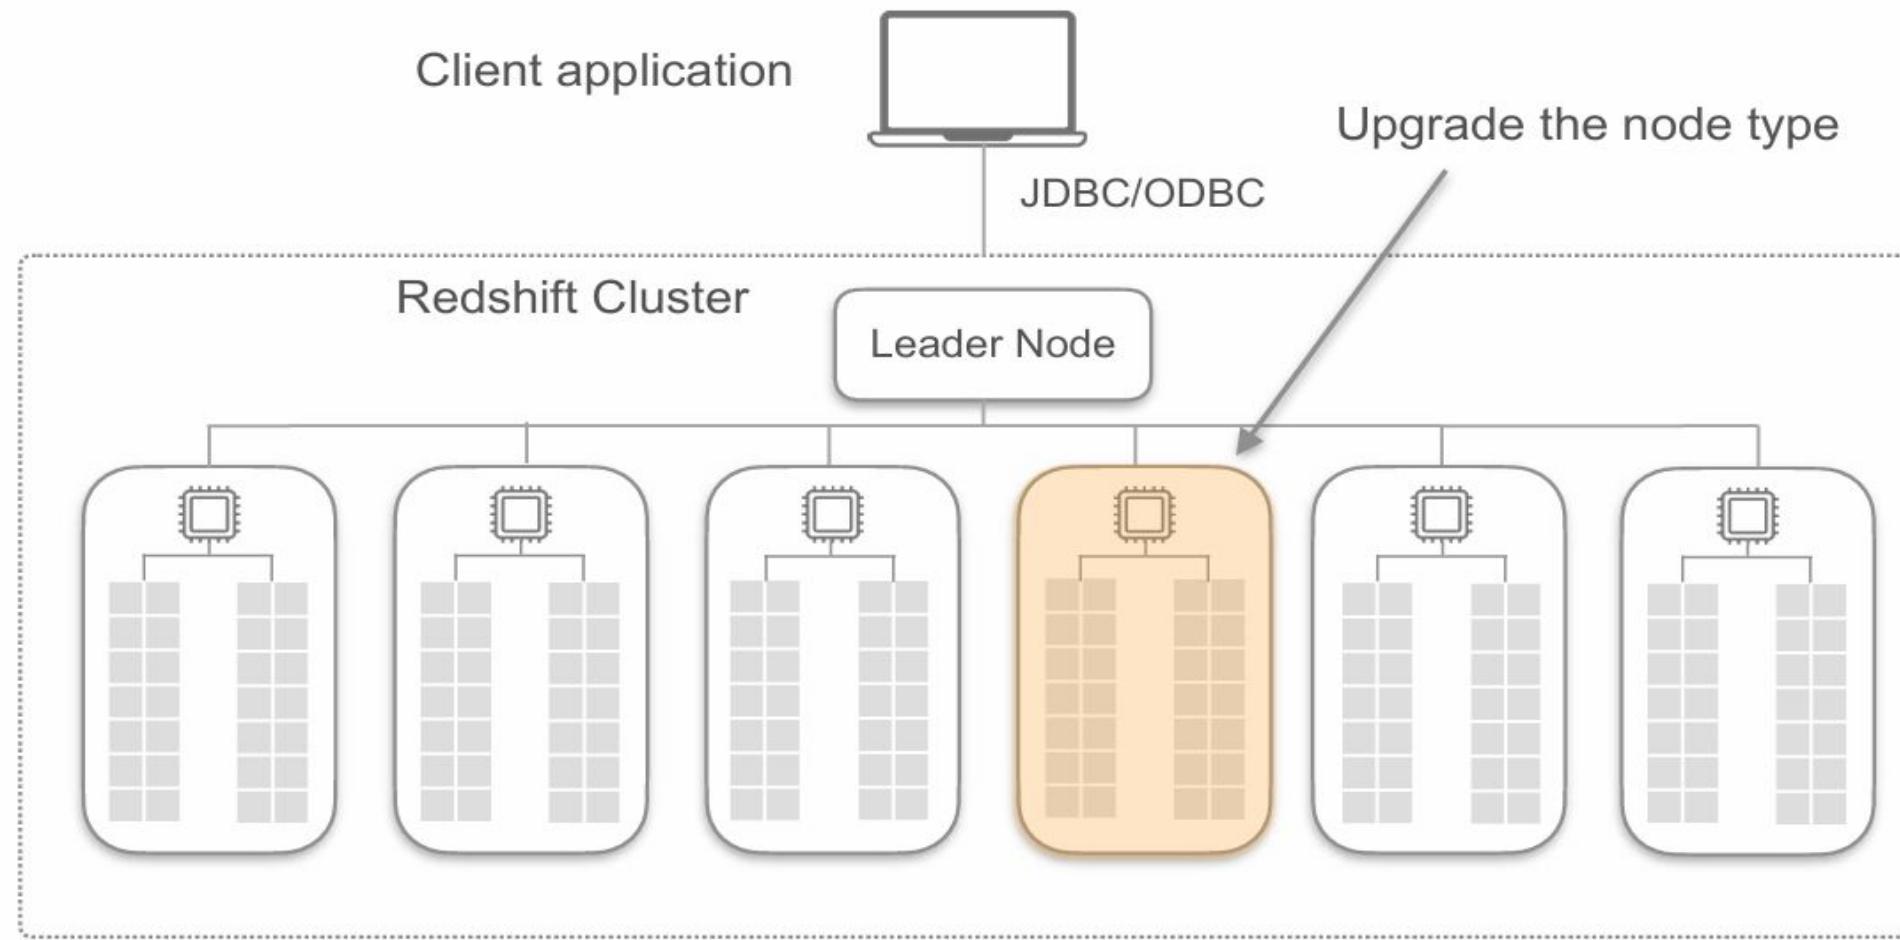


Amazon Redshift

# MPP Architecture for Amazon Redshift

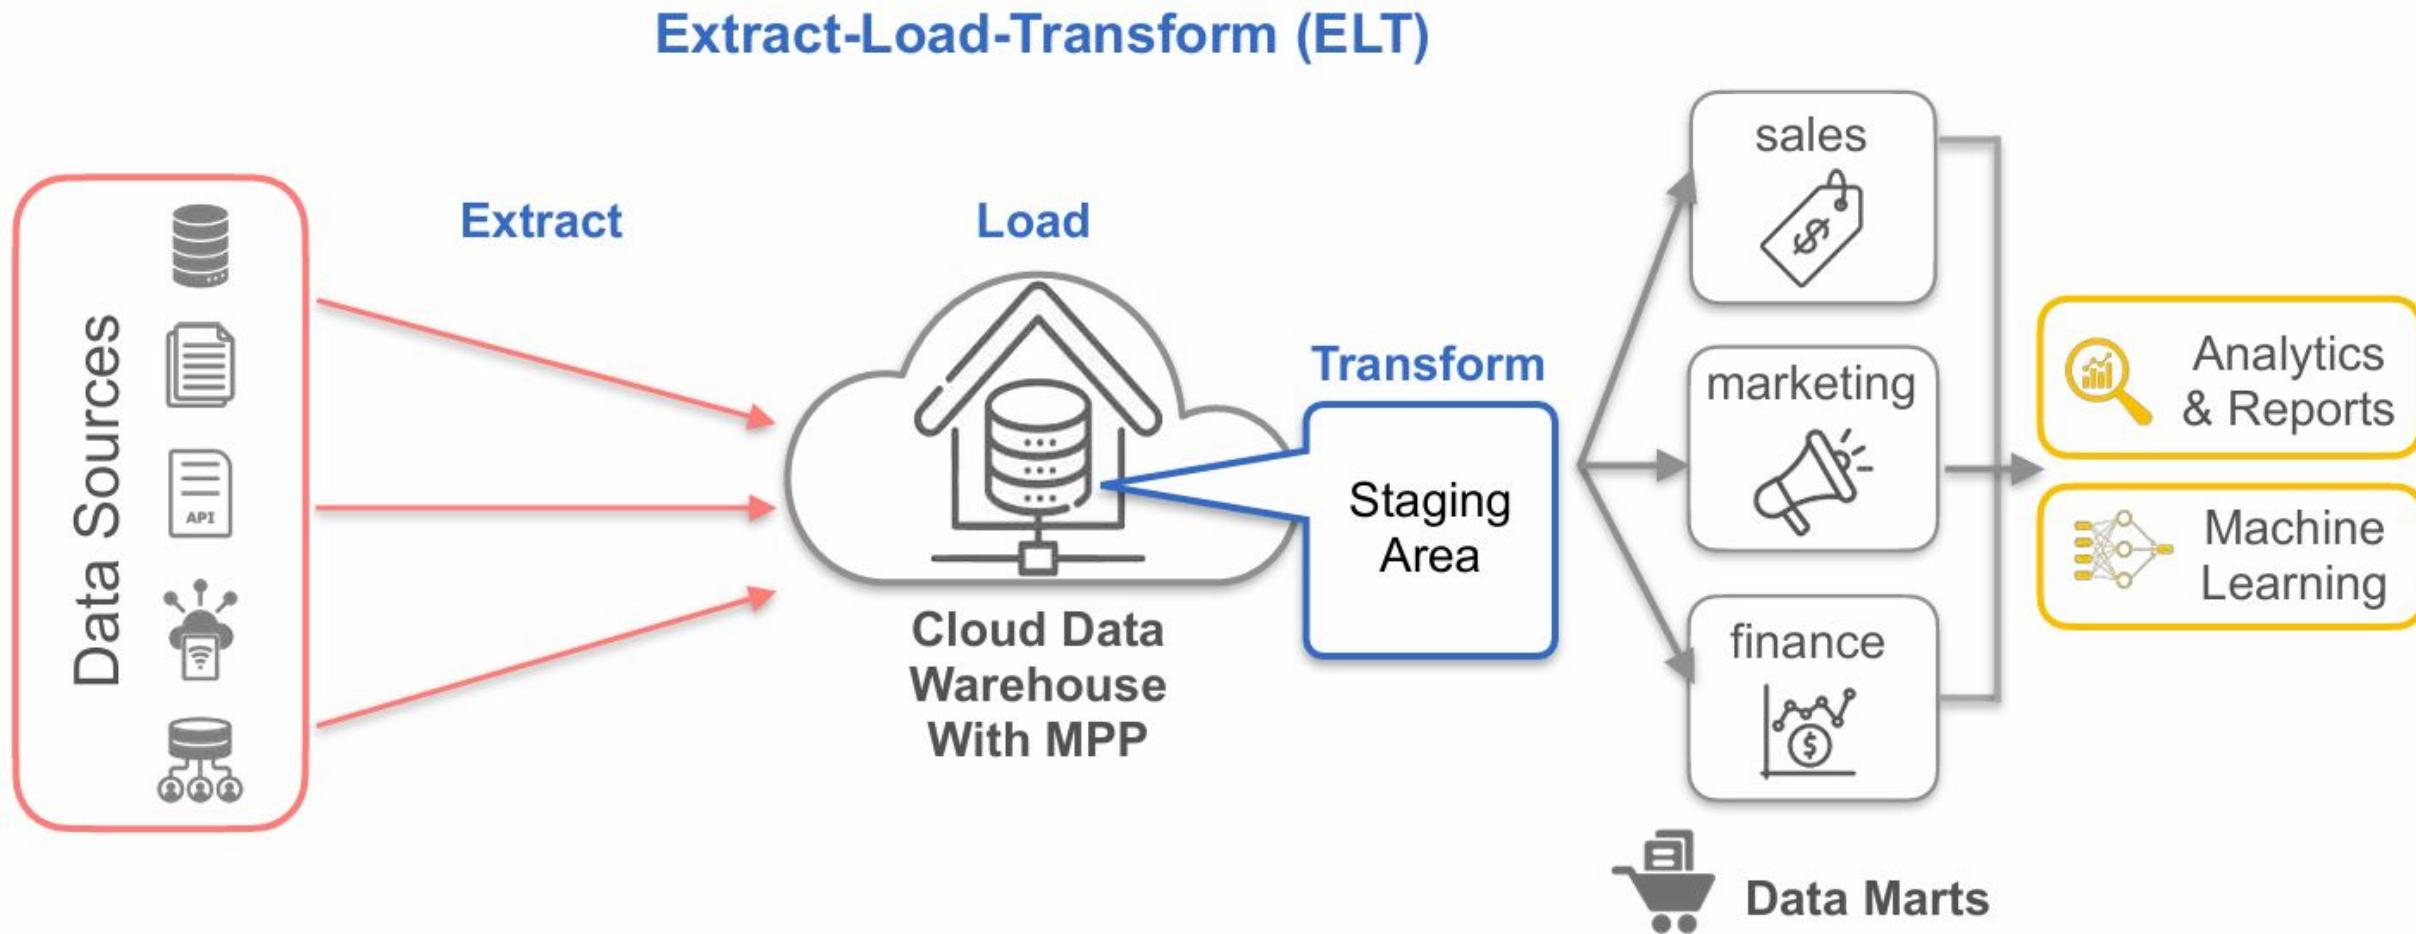


# MPP Architecture for Amazon Redshift



Amazon Redshift

# Data Warehouse-Centric Architecture



# Cloud Data Warehouse

## Columnar Architecture

Order ID	Price	Product SKU	Quantity	Customer ID
1	40	458650	10	67t
2	23	902348	14	56t
3	45	1255893	12	87q
4	50	456829	13	98q

Facilitates high performance analytical queries

## Separation of Compute and Storage



Compute



Storage

# Cloud Data Warehouse

## Traditional Data Warehouse

- Stored data is highly structured
- Data modeled to enable analytical queries

## Cloud Data Warehouse

- Stored data is highly structured
- Data modeled to enable analytical queries
- High processing from MPP
- Columnar Storage
- Separation of storage and compute

**Efficiently stores and processes data for high-volume analytical workloads**

# Data Lakes – Key Architectural ideas

# Data Lake



- Central repository for storing large volumes of data
- No fixed schema or predefined set of transformations
- **Schema-on-read pattern:**
  - Reader determines the schema when reading the data

## Data Lake 1.0

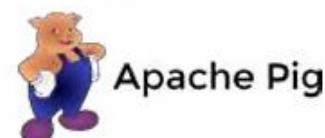
Combined different storage and processing technologies



Storage



Processing  
Tools



# Shortcomings of Data Lake 1.0



## Data Swamp

- No proper data management
- No data cataloging
- No data discovery tools
- No guarantee on the data integrity and quality

# Shortcomings of Data Lake 1.0

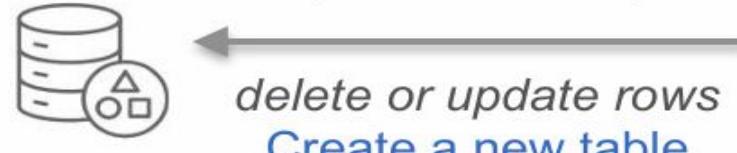


## Data Swamp

- No proper data management
- No data cataloging
- No data discovery tools
- No guarantee on the data integrity and quality

## Write-only storage

- Data Manipulation Language (DML) operations were painful to implement



- Difficult to comply with data regulations



# Shortcomings of Data Lake 1.0

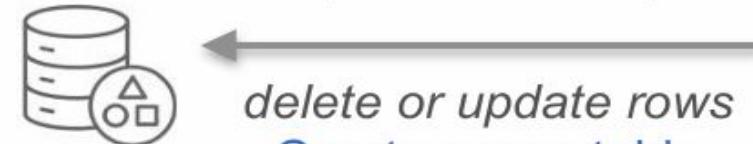


## Data Swamp

- No proper data management
- No data cataloging
- No data discovery tools
- No guarantee on the data integrity and quality

## Write-only storage

- Data Manipulation Language (DML) operations were painful to implement



- Difficult to comply with data regulations



## No schema management and data modeling

- Hard to process stored data
- Data not optimized for query operations such as joins



# Next-Generation Data Lakes

## Data Zones

Used to organize data in a data lake, where each zone houses data that has been processed to varying degrees



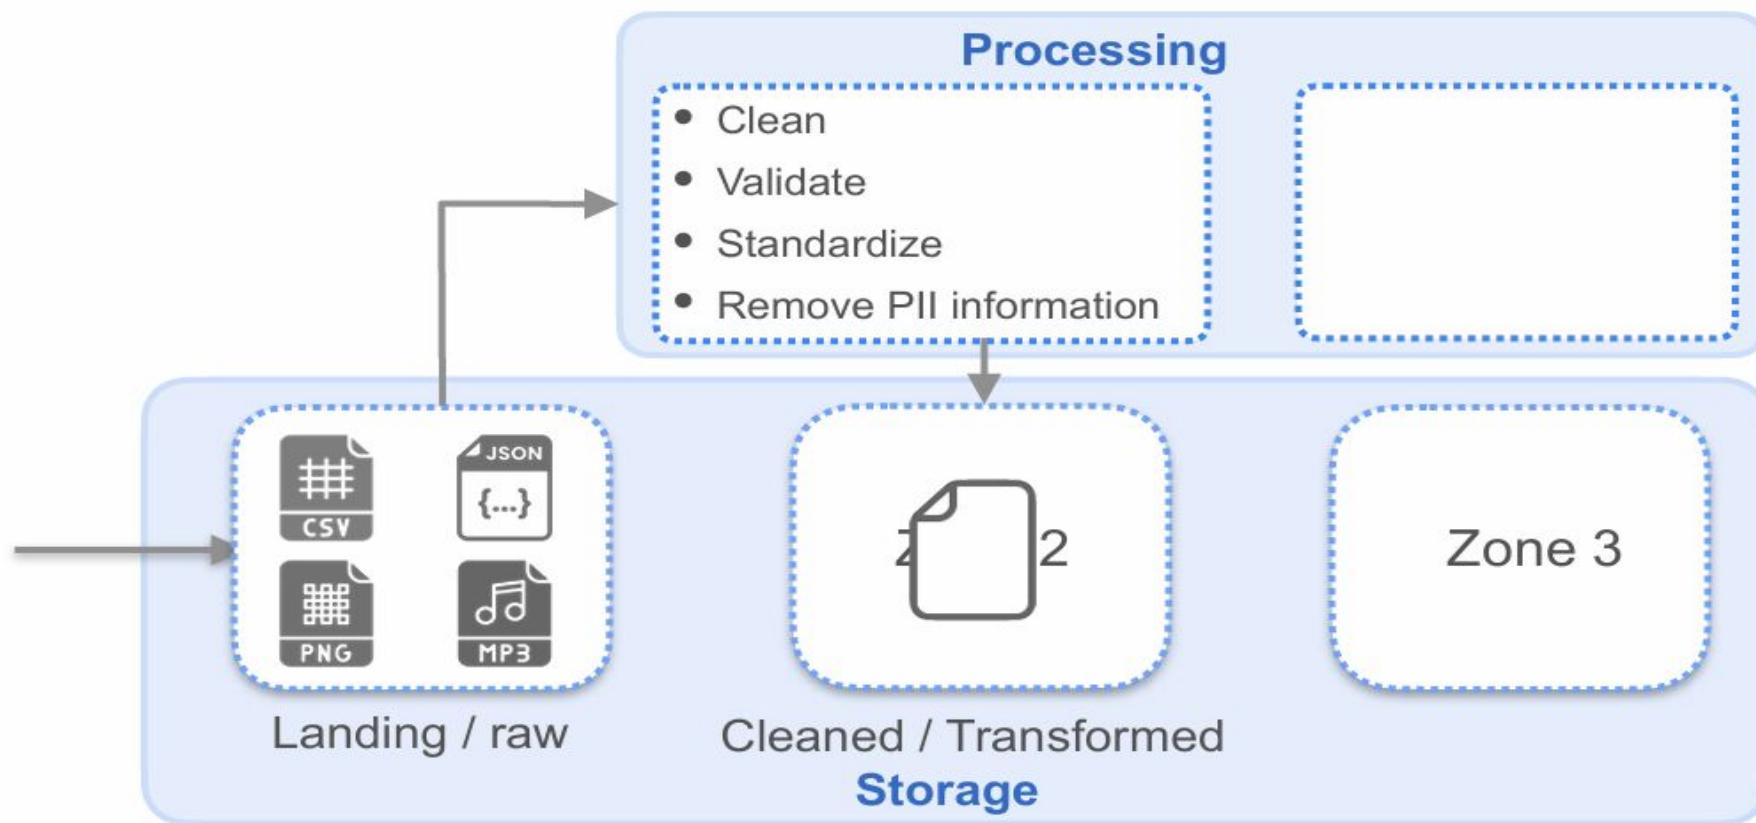
## Data Zones

Used to organize data in a data lake, where each zone houses data that has been processed to varying degrees



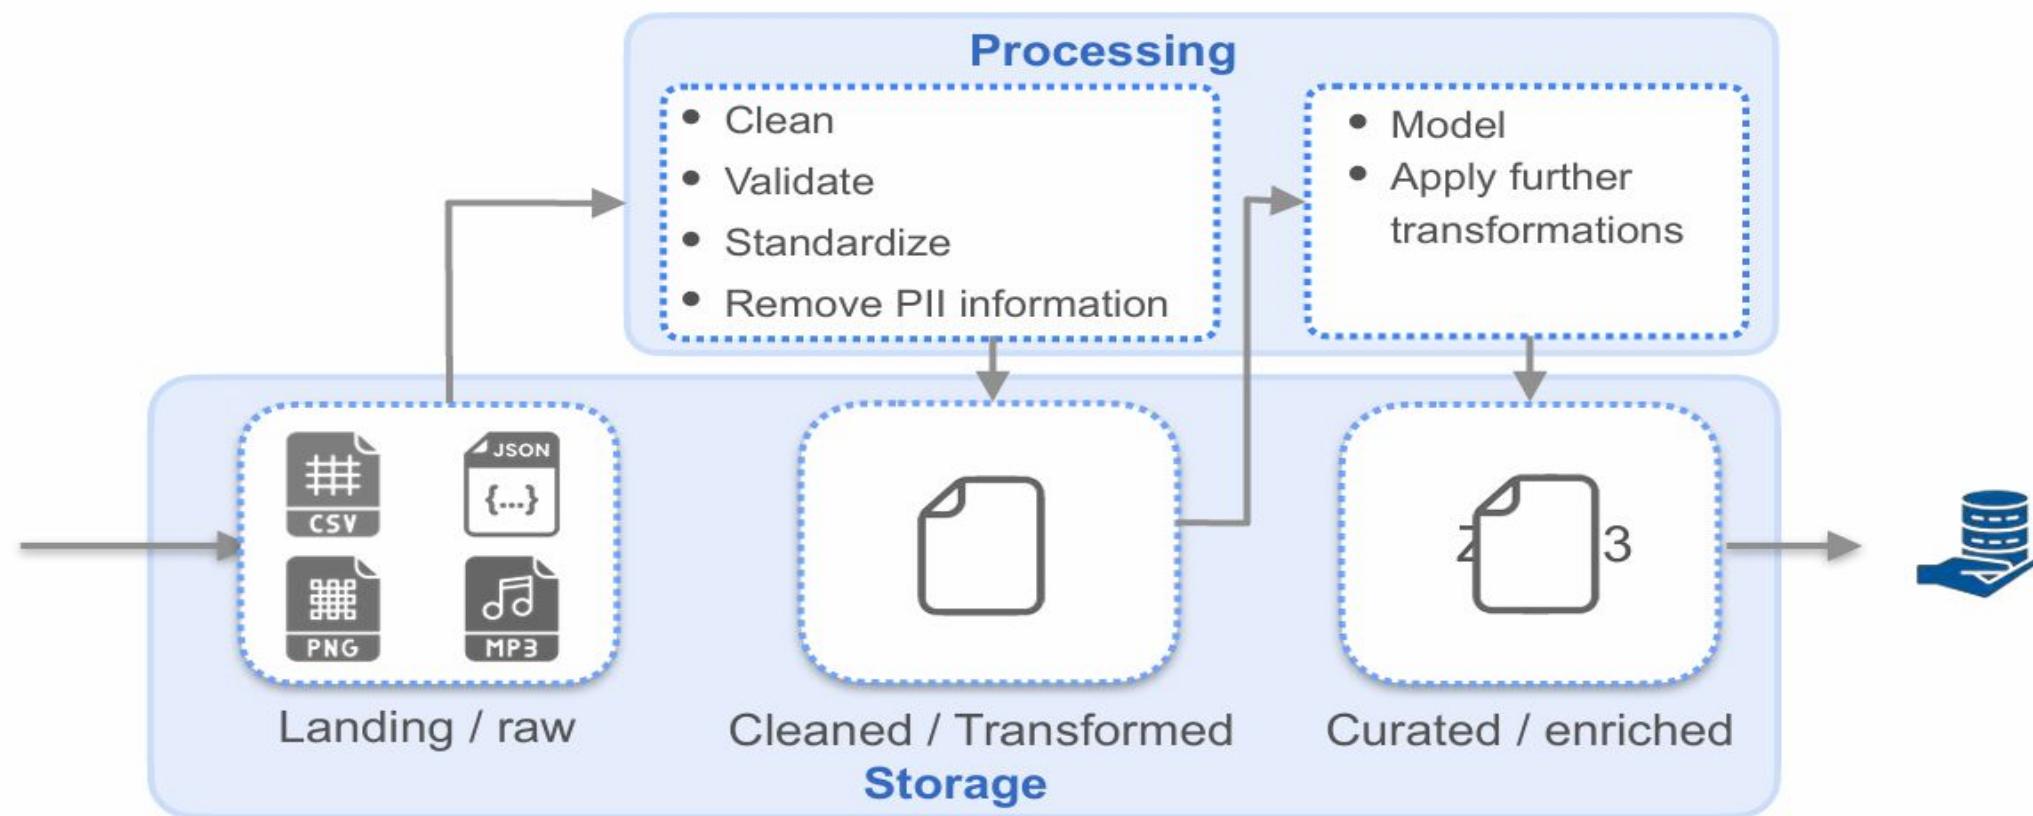
## Data Zones

Used to organize data in a data lake, where each zone houses data that has been processed to varying degrees



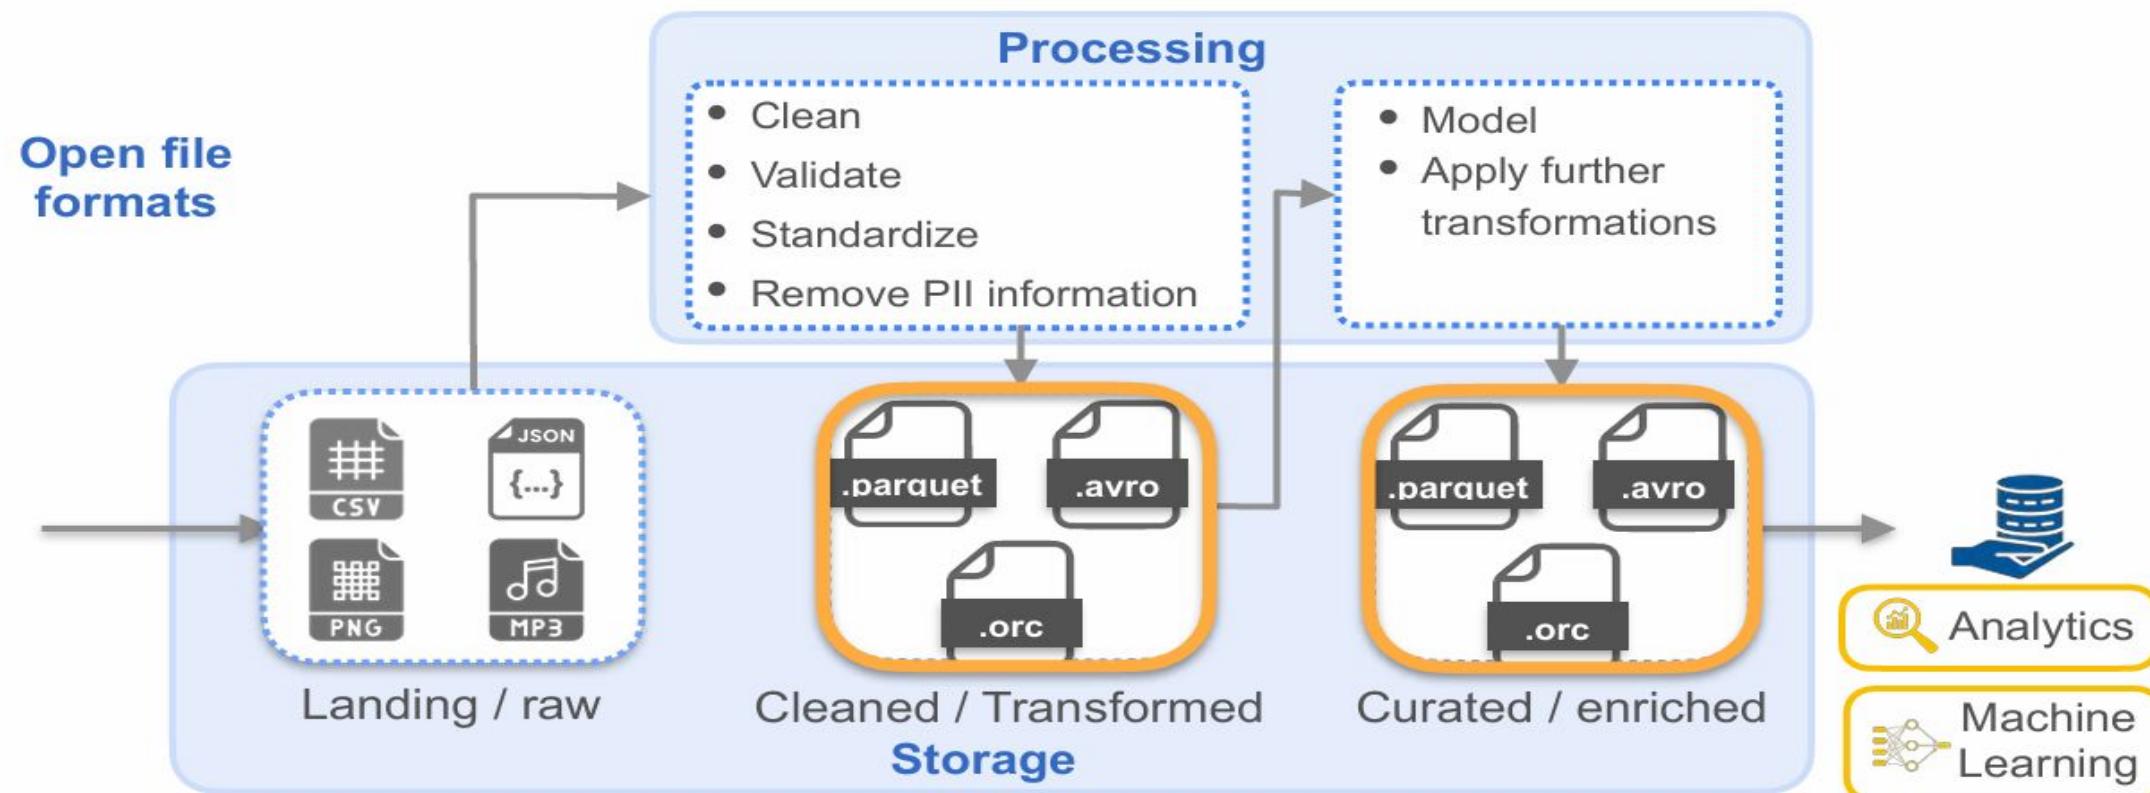
## Data Zones

Used to organize data in a data lake, where each zone houses data that has been processed to varying degrees



## Data Zones

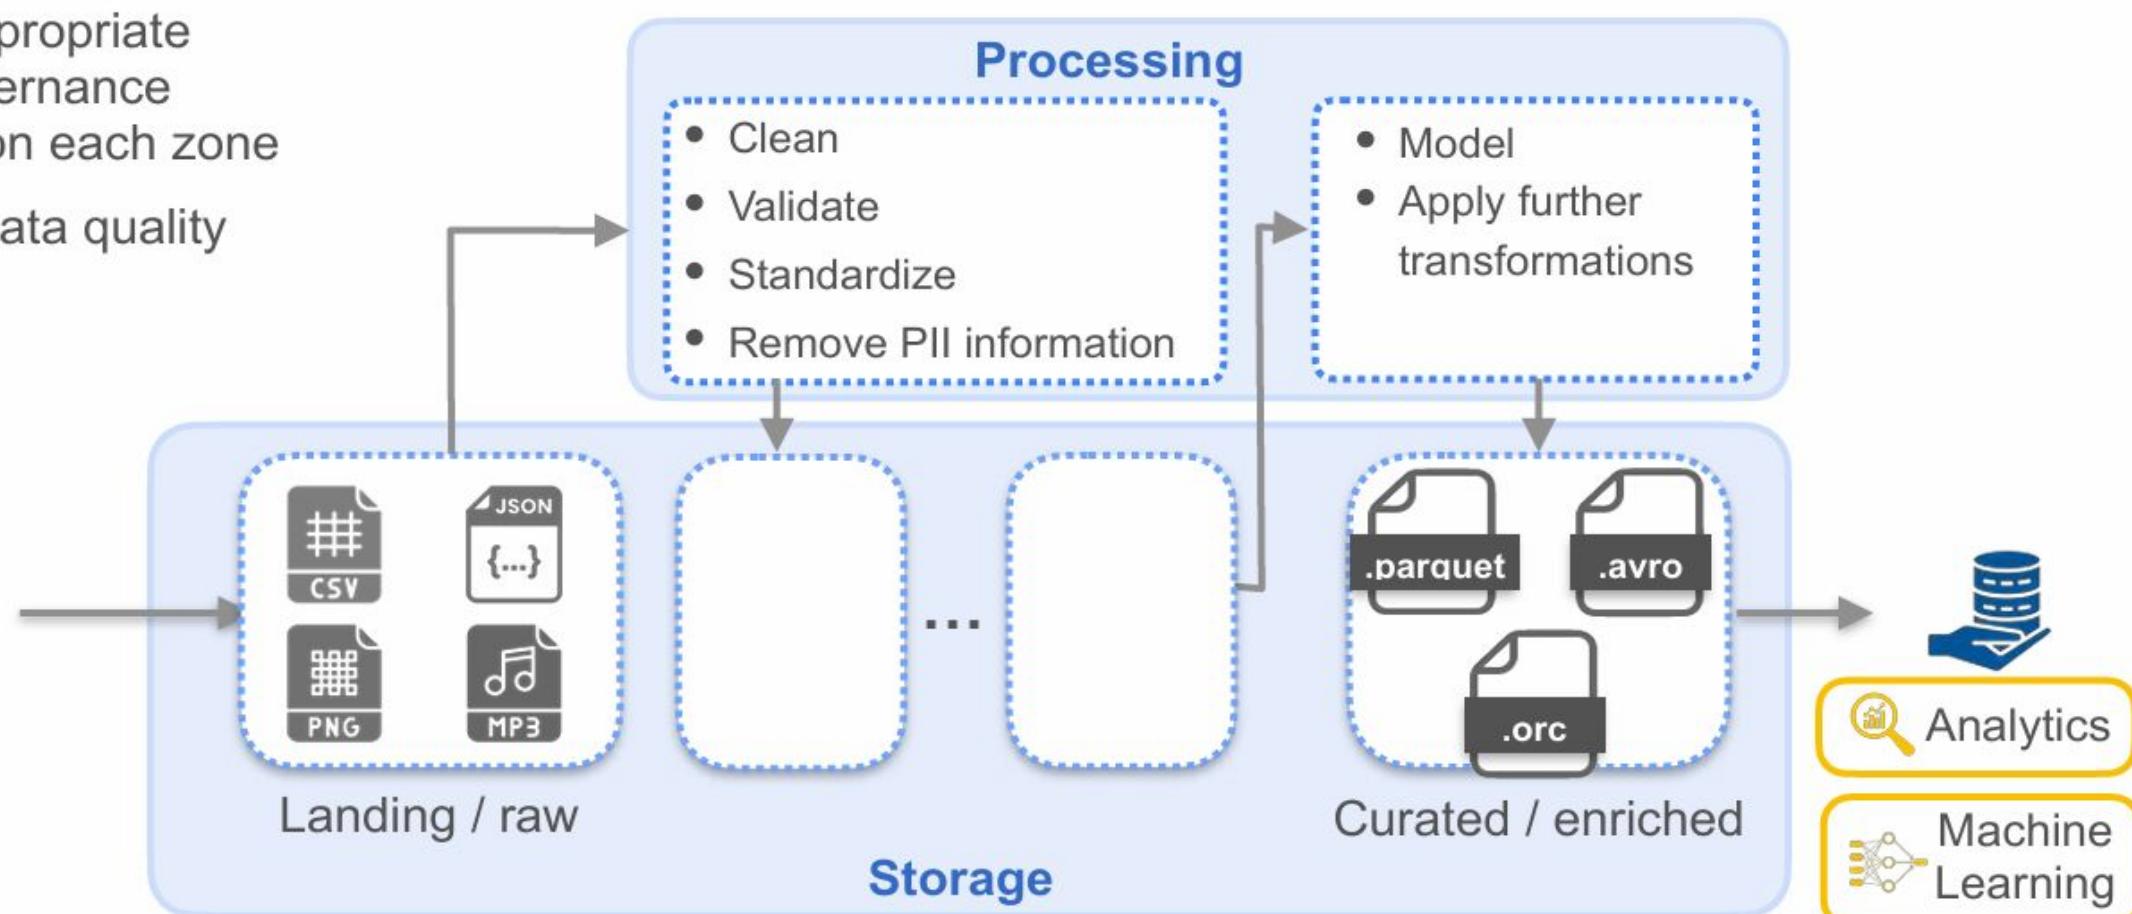
Used to organize data in a data lake, where each zone houses data that has been processed to varying degrees



## Data Zones

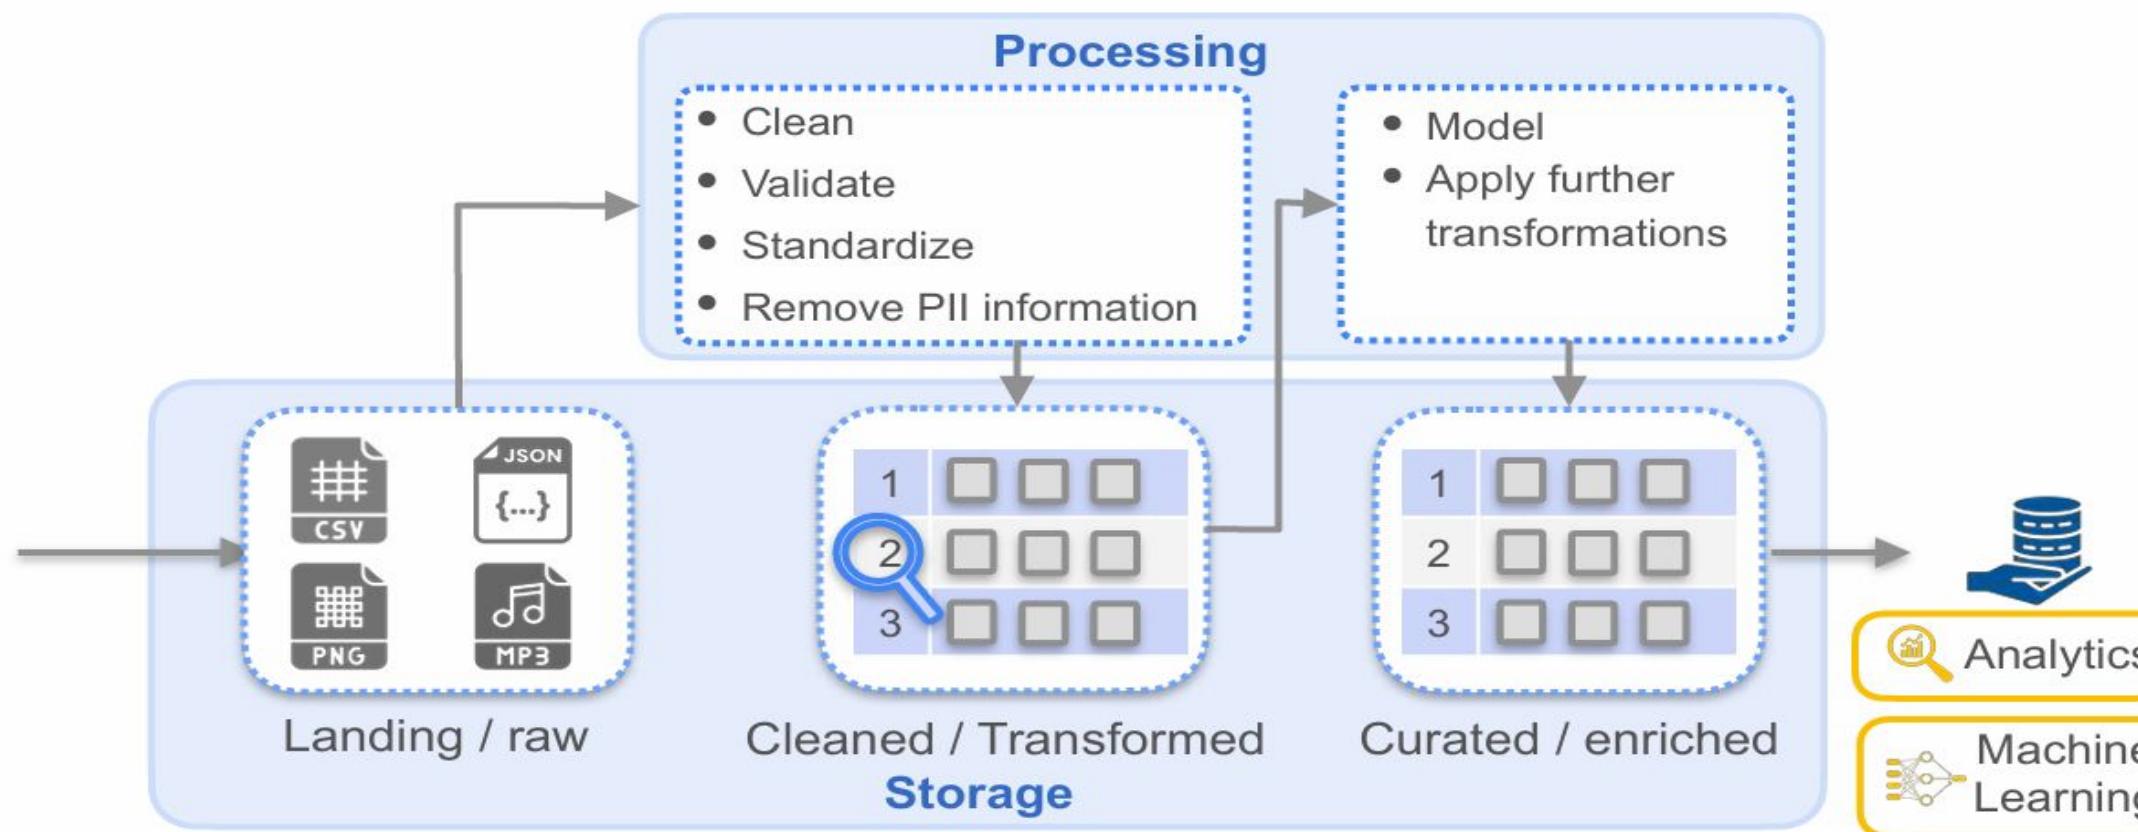
Used to organize data in a data lake, where each zone houses data that has been processed to varying degrees

- Apply appropriate data governance policies on each zone
- Ensure data quality



## Data Partitioning

Divide a dataset into smaller, more manageable parts based on a set of criteria (e.g. time, date, location recorded in the data)



## Data Catalog

Collection of metadata about the dataset (owner, source, partitions, etc.)

## Catalog

- Metadata
- Schema

## Processing

- Clean
- Validate
- Standardize
- Remove PII information

- Model
- Apply further transformations



Landing / raw

1			
2			
3			

Cleaned / Transformed  
Storage

1			
2			
3			

Curated / enriched



Analytics

Machine Learning

# Separate Data Lakes and Data Warehouses



Data Lake

Low-Cost Storage

Store large amounts of data

and

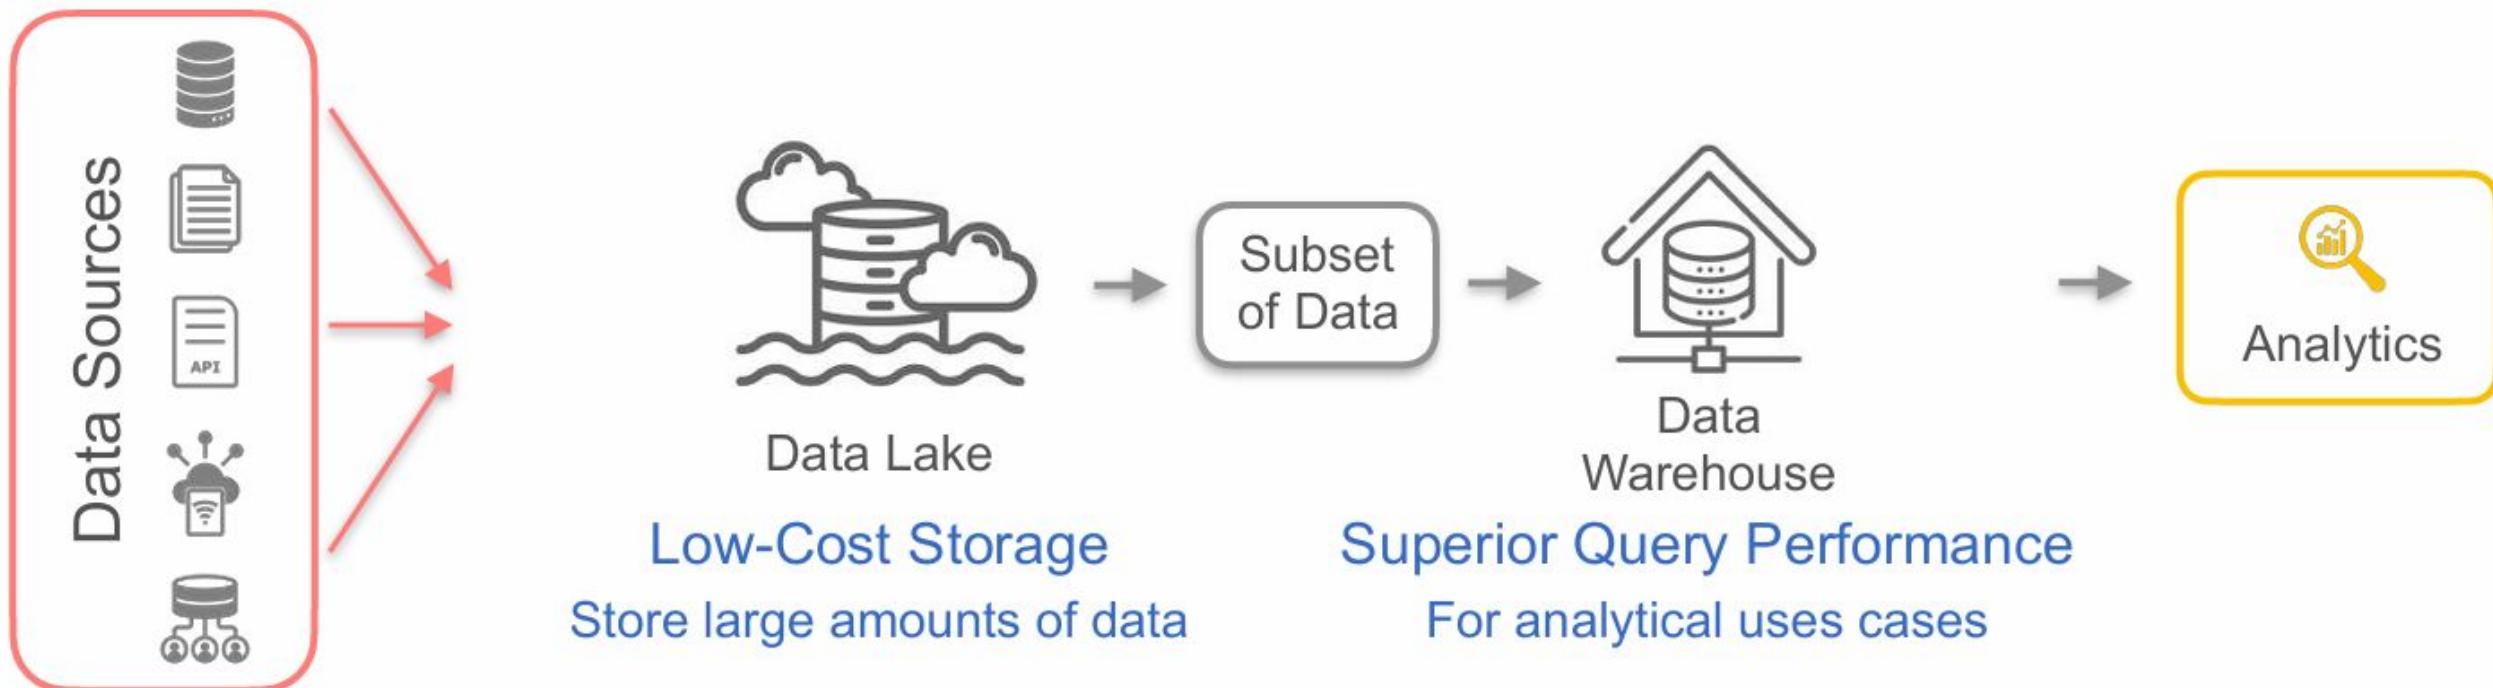


Data  
Warehouse

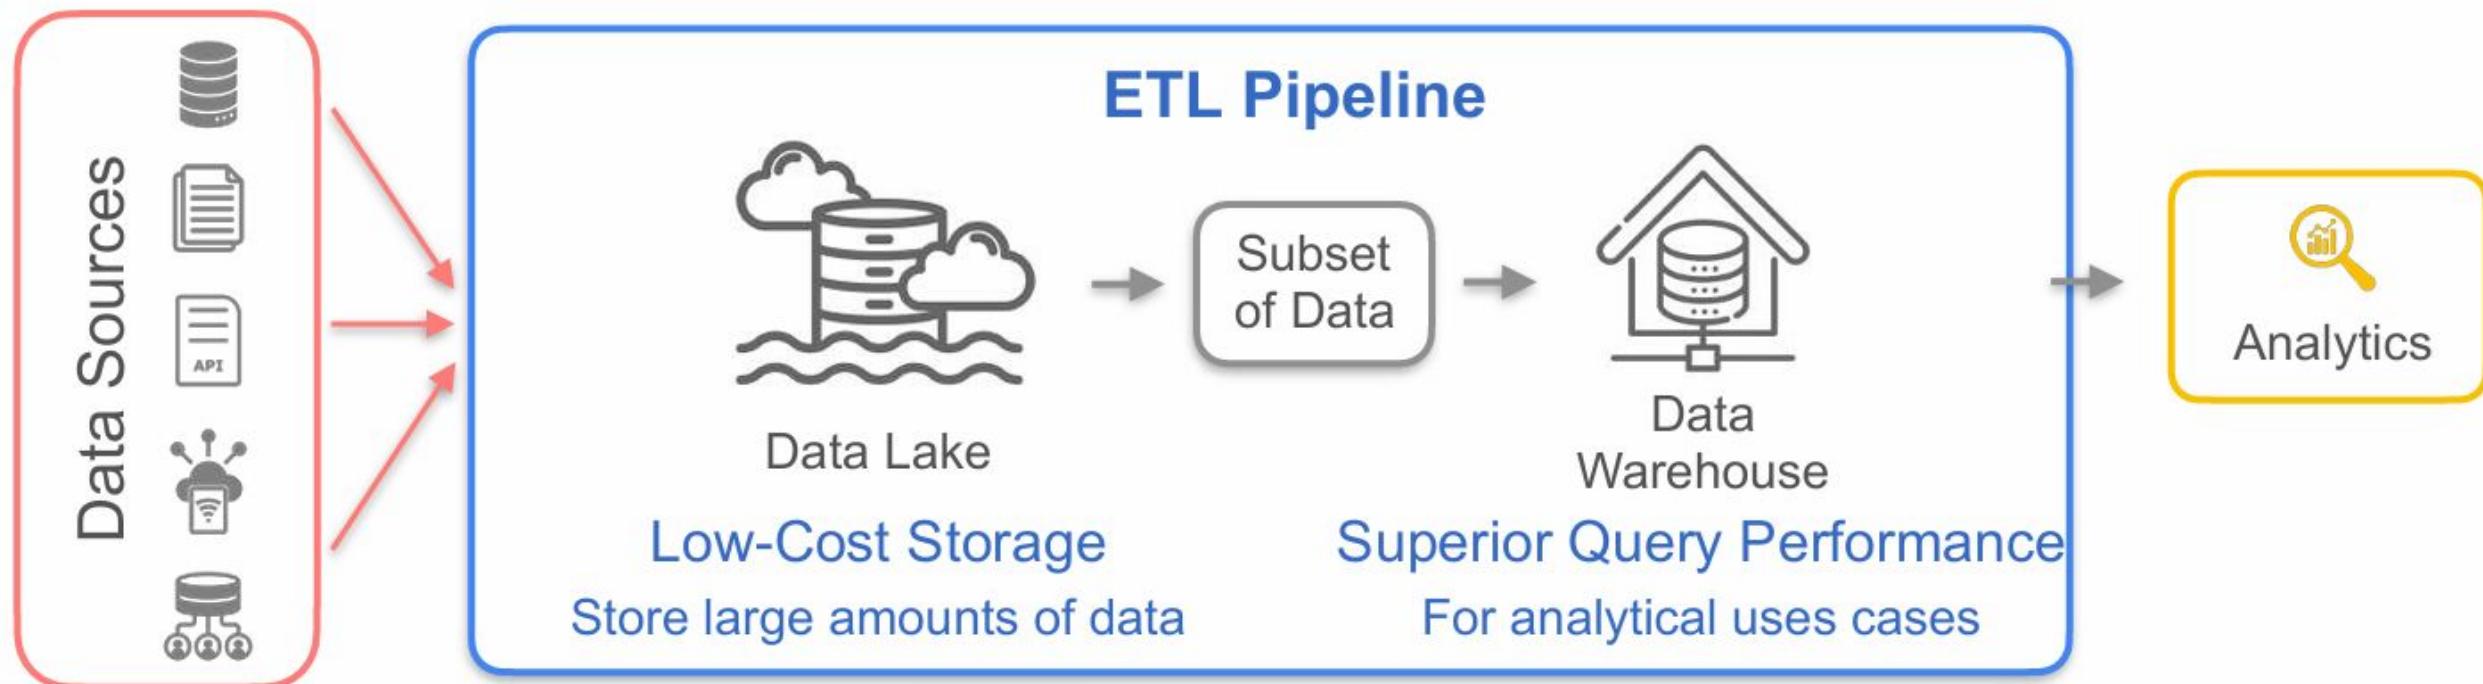
Superior Query Performance

For analytical use cases

# Separate Data Lakes and Data Warehouses

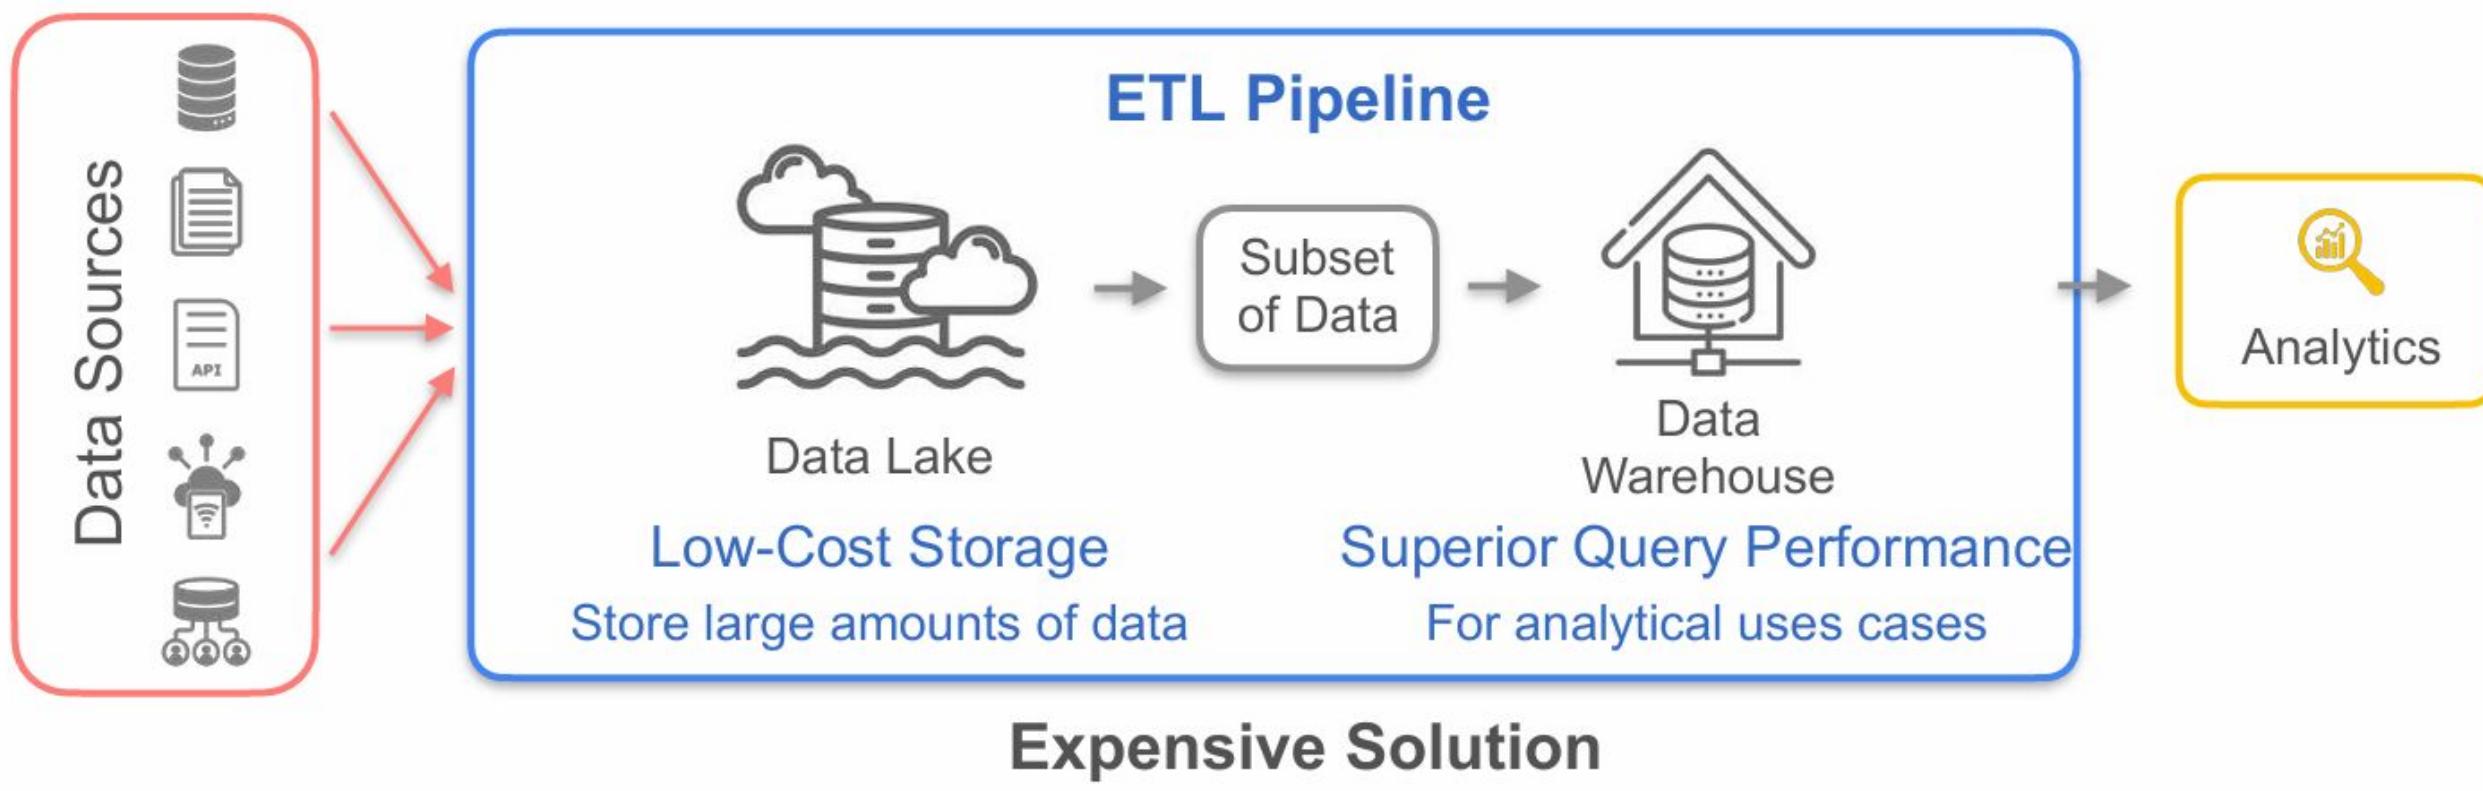


# Separate Data Lakes and Data Warehouses



**Expensive Solution**

# Separate Data Lakes and Data Warehouses



## Expensive Solution

- Can introduce bugs/failures
- Can cause issues with data quality, duplication, consistency

# Data Lakehouse

# Data Lakehouse



**Data Lake**

- Flexibility
- Low-Cost Storage

+



**Data Warehouse**

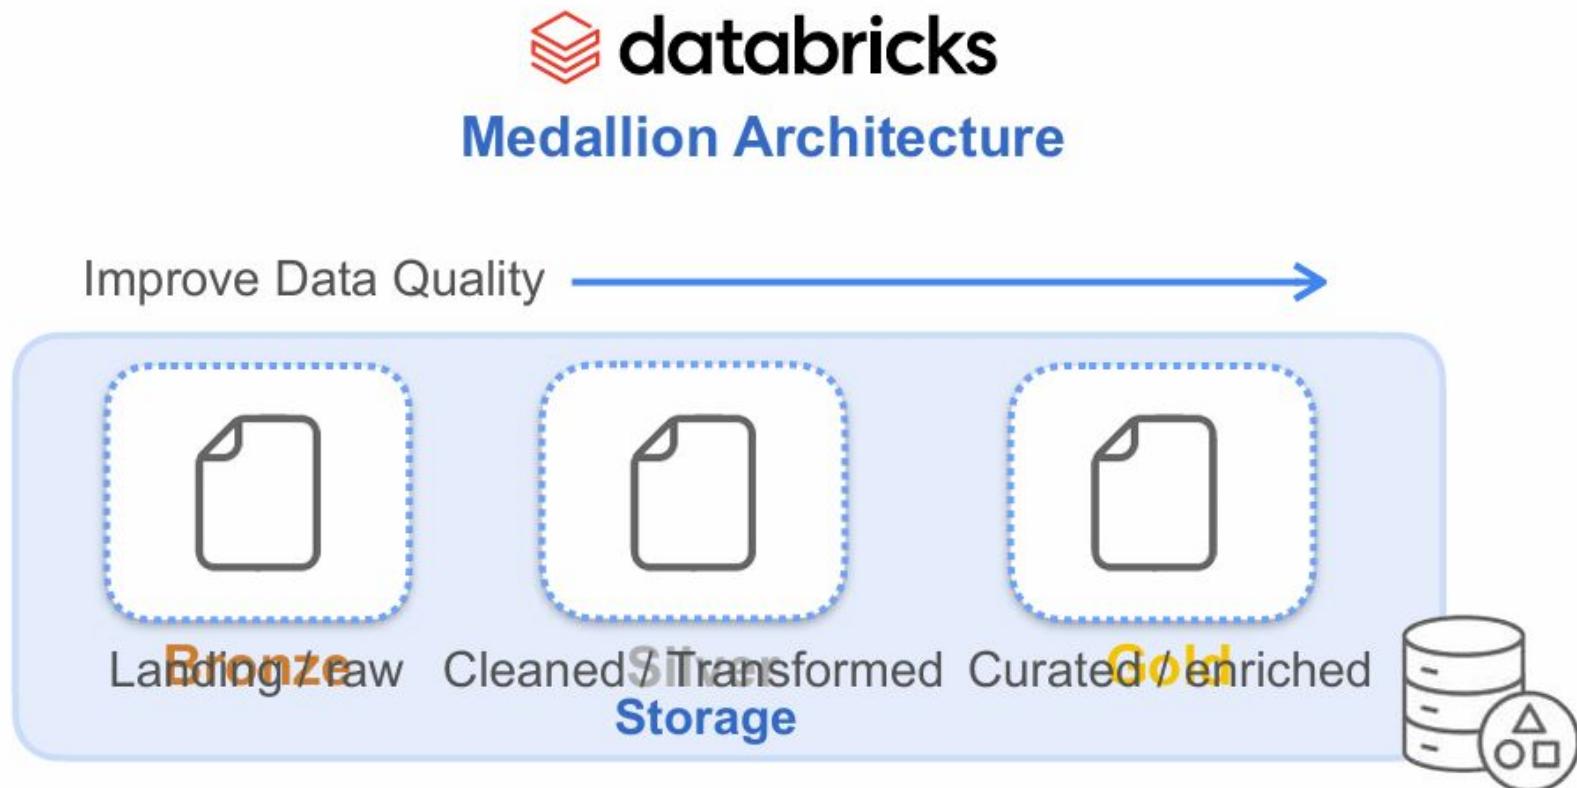
- Superior query performance
- Robust data management

=

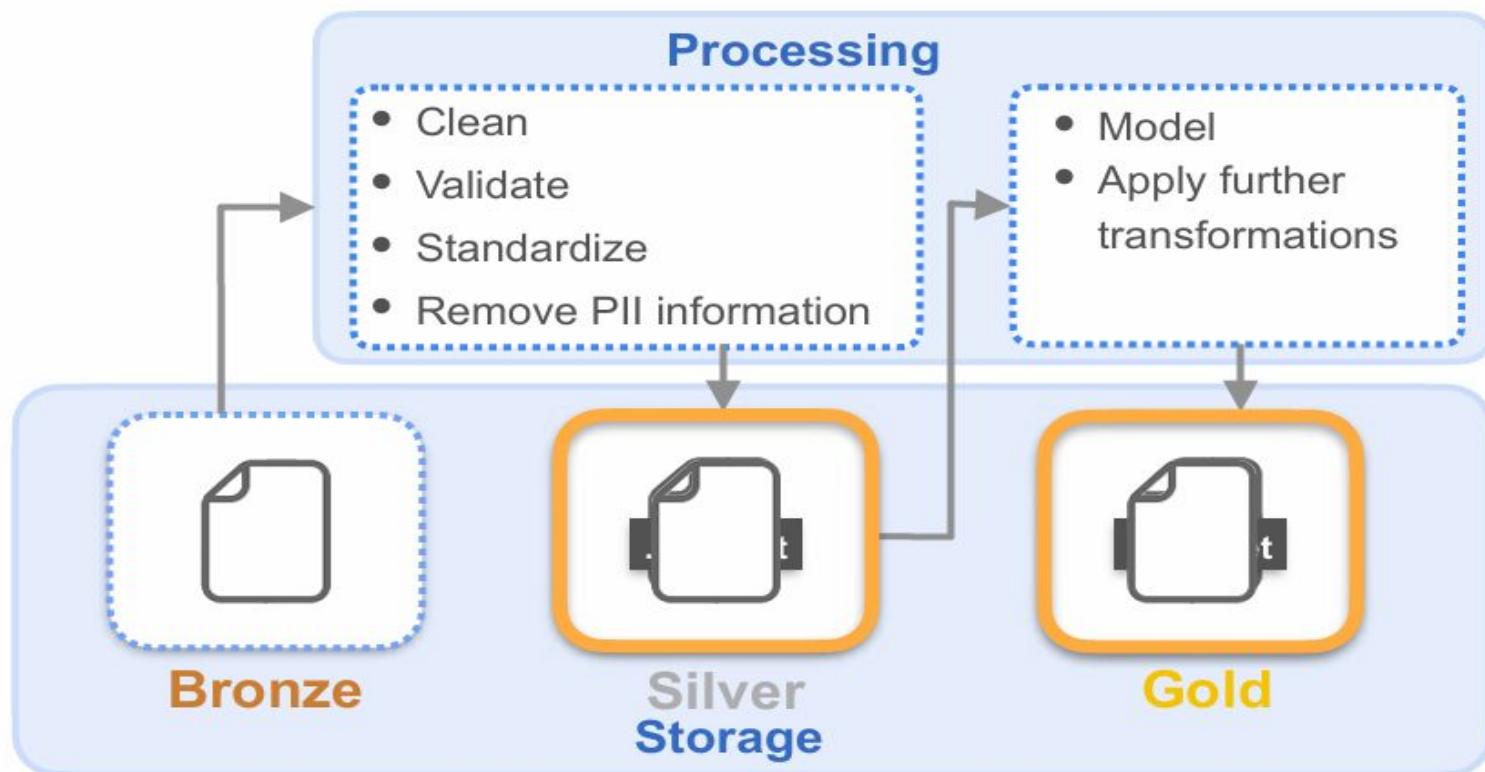


**Data Lakehouse**

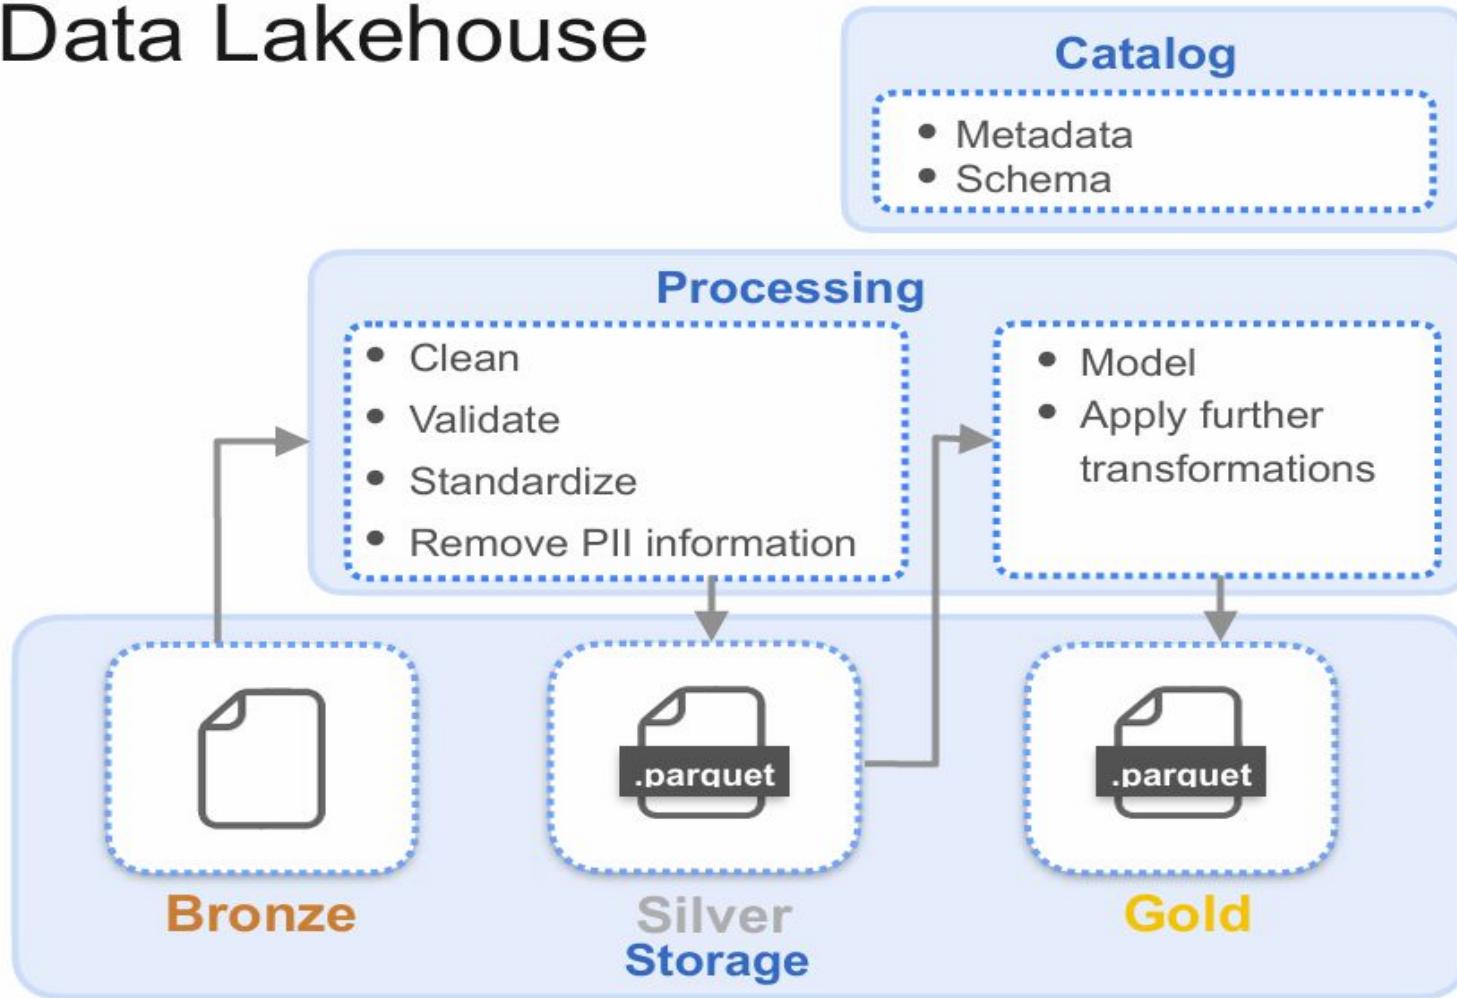
# Data Lake Features



# Data Lake Features

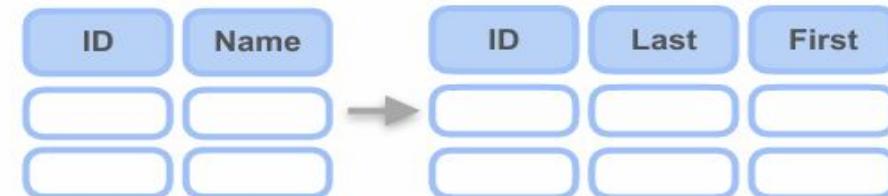


# Data Lakehouse

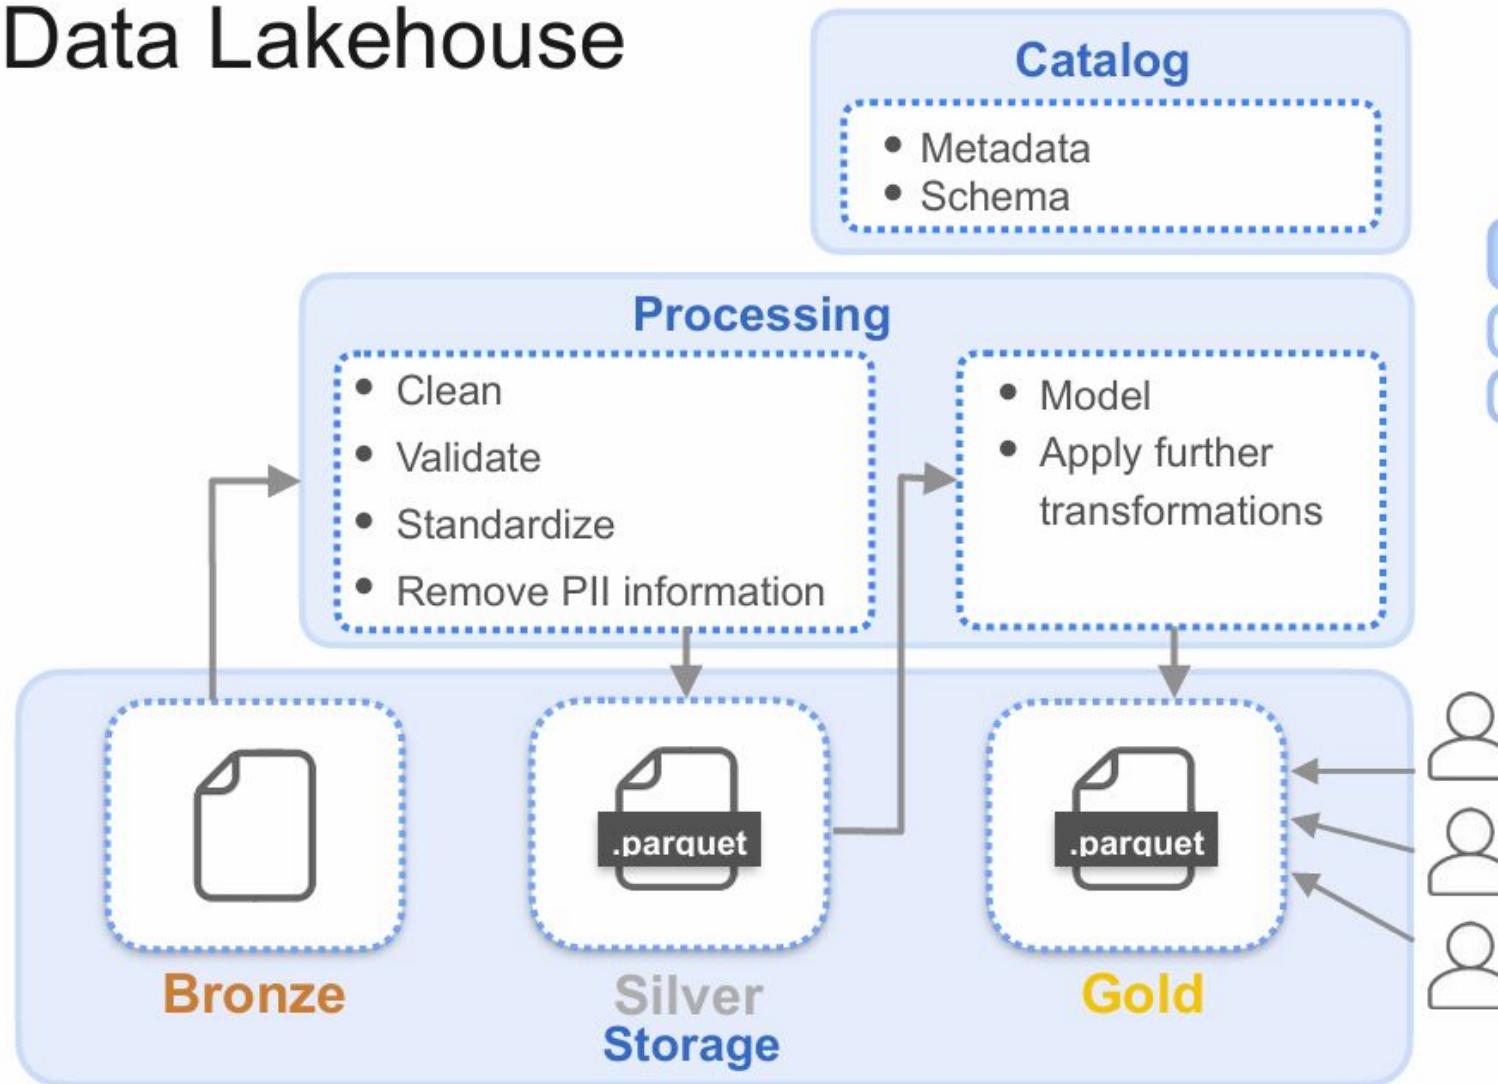


## Data Management from Data Warehouses

### Schema Enforcement

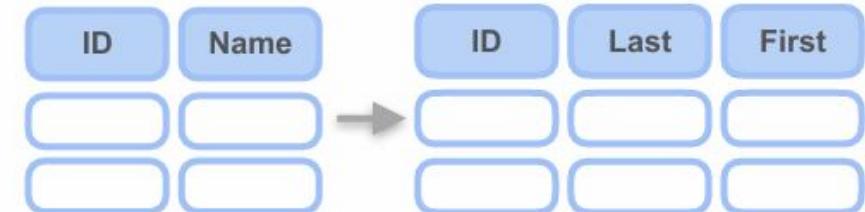


# Data Lakehouse



Data Management from Data Warehouses

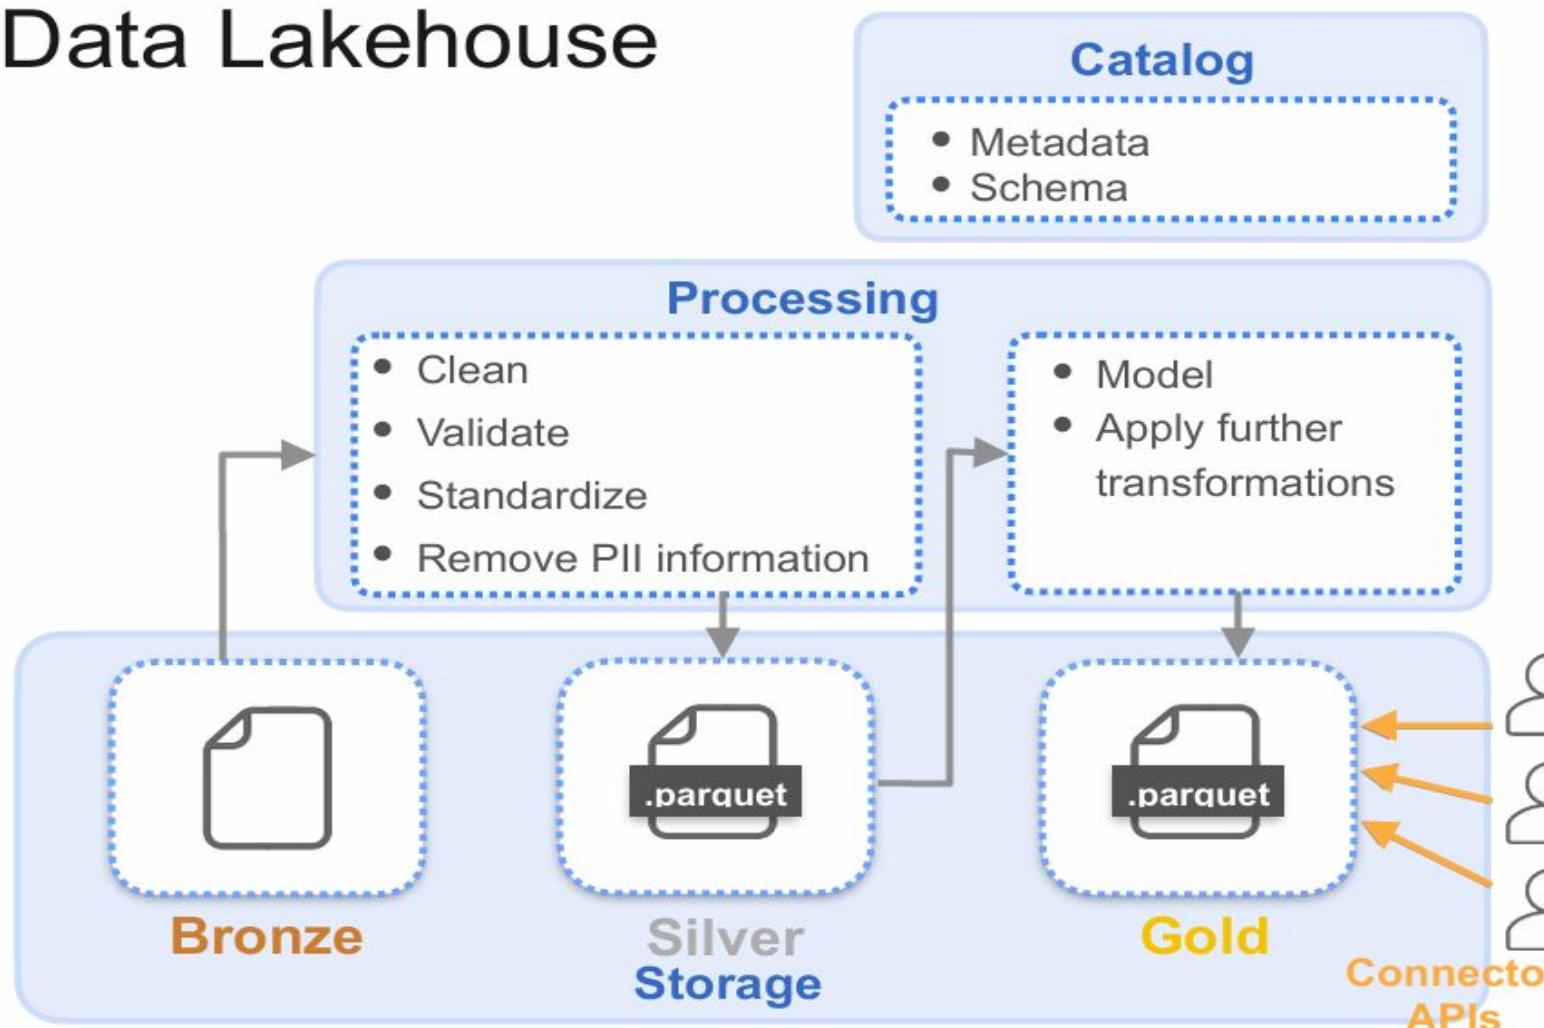
## Schema Enforcement



## ACID

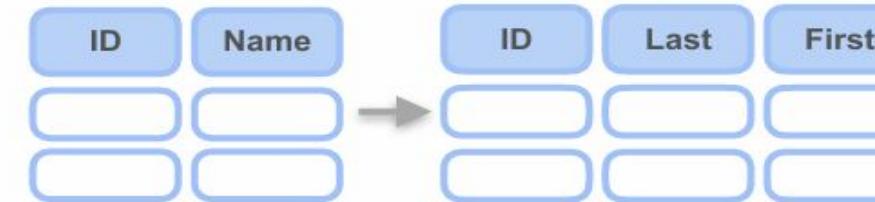
- Atomic, Consistent, Isolated, Durable
- Concurrent read, insert, update, delete

# Data Lakehouse



## Data Management from Data Warehouses

### Schema Enforcement



### ACID

- **A**tomic, **C**onsistent, **I**solated, **D**urable
- Concurrent read, insert, update, delete

### Data Governance & Security

- Robust access controls, data auditing capabilities, data lineage
- Incremental updates & deletions
- Rollback to or access any version of your historical data

# Date Lakehouse Implementation

# Open Table Formats

## Open Table Formats

Specialized storage formats that add transactional features to your data lakehouse

- Allows you to update and delete records
- Supports ACID principles

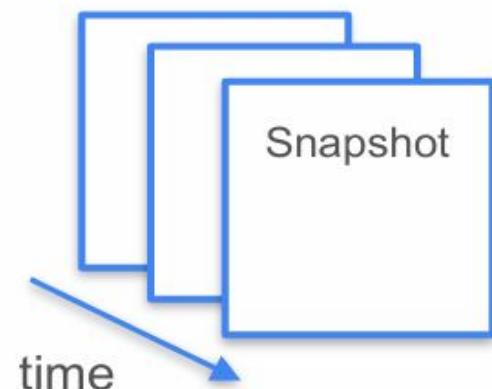


Hadoop  
Update  
Delete  
Incremental

# Open Table Formats

## Open Table Format

*Track changes in data*



## Data

*Insert, Update, Delete*

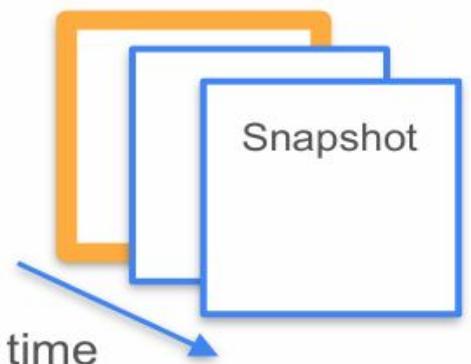
### Snapshot:

Reflects the state of the data at a given time

# Open Table Formats

## Open Table Format

*Track changes in data*



## Data

*Insert, Update, Delete*

### **Snapshot:**

Reflects the state of the data at a given time

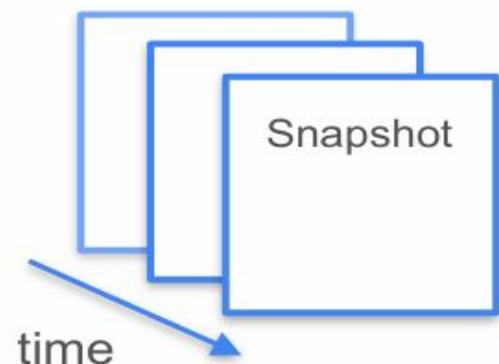
### **Time Travel:**

Query any previous version of a table

# Open Table Formats

## Open Table Format

*Track changes in data*



## Data

*Insert, Update, Delete*

### **Snapshot:**

Reflects the state of the data at a given time

### **Time Travel:**

Query any previous version of a table

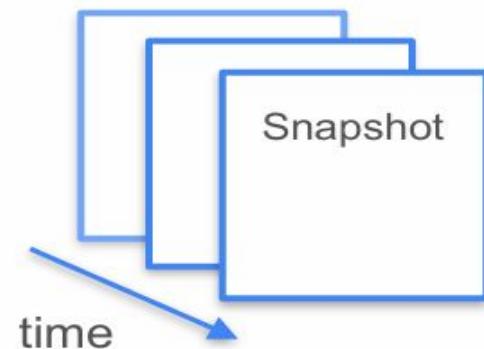
### **Schema & Partition Evolution:**

Ability to query the data even if you make changes to the schema or partitioning

# Open Table Formats

## Open Table Format

*Track changes in data*



## Data

*Insert, Update, Delete*

### Snapshot:

Reflects the state of the data at a given time

### Time Travel:

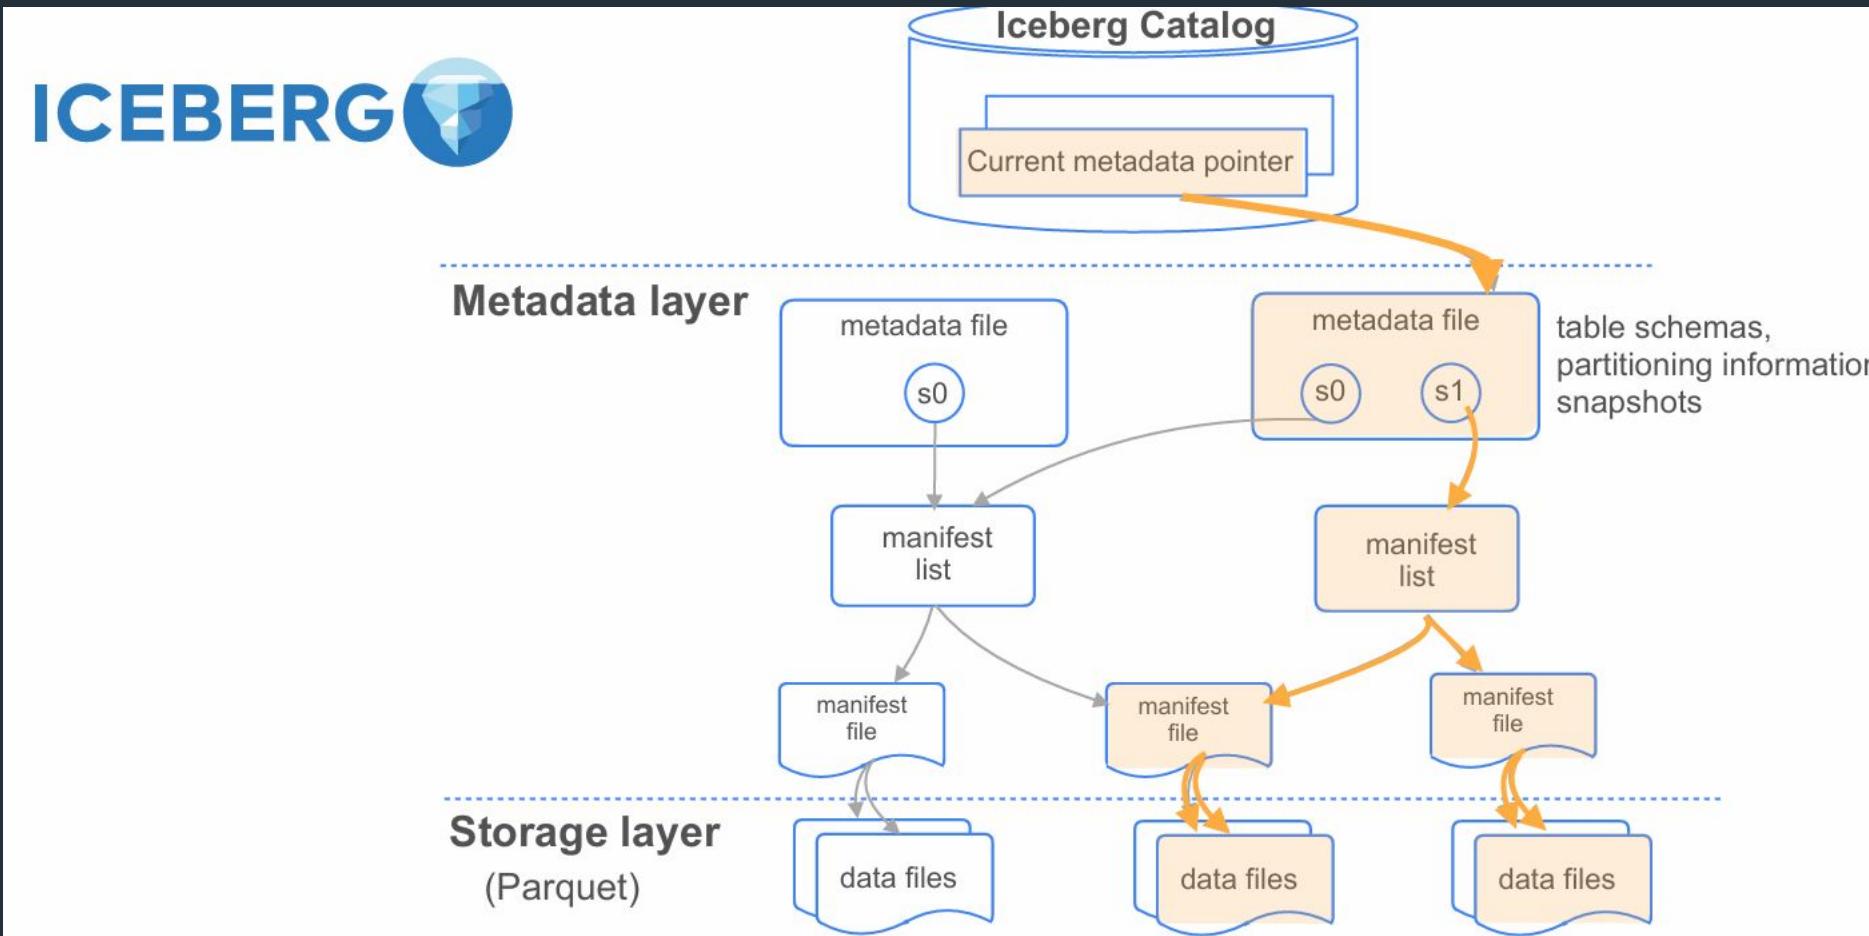
Query any previous version of a table

### Schema & Partition Evolution:

Ability to query the data even if you make changes to the schema or partitioning

### Open Source

Different query engines can access the data



# Storage Options

Production Database	Data Warehouse		
<ul style="list-style-type: none"> <li>Process small amounts of structured data</li> </ul>	<ul style="list-style-type: none"> <li>Bring together large volumes of structured / semi-structured data</li> <li>Query current and historical data</li> </ul>		
<b>Use case:</b> analytics, reporting	<b>Use case:</b> analytics, reporting		

# Storage Options

Production Database	Data Warehouse	Data Lake	
<ul style="list-style-type: none"> <li>Process small amounts of structured data</li> </ul>	<ul style="list-style-type: none"> <li>Bring together large volumes of structured / semi-structured data</li> <li>Query current and historical data</li> </ul>	<ul style="list-style-type: none"> <li>Process large volumes of structured, semi-structured and unstructured data</li> <li>Save on storage cost</li> </ul>	
<b>Use case:</b> reporting, analytics	<b>Use case:</b> reporting, analytics	<b>Use case:</b> machine learning	

# Storage Options



<b>Production Database</b>	<b>Data Warehouse</b>	<b>Data Lake</b>	<b>Data Lakehouse</b>
<ul style="list-style-type: none"> <li>• Process small amounts of structured data</li> </ul>	<ul style="list-style-type: none"> <li>• Bring together large volumes of structured / semi-structured data</li> <li>• Query current and historical data</li> </ul>	<ul style="list-style-type: none"> <li>• Process large volumes of structured, semi-structured and unstructured data</li> <li>• Save on storage cost</li> </ul>	<ul style="list-style-type: none"> <li>• Data management &amp; discoverability features</li> <li>• Low latency queries</li> </ul>
<b>Use case:</b> reporting, analytics	<b>Use case:</b> reporting, analytics	<b>Use case:</b> machine learning	<b>Use case:</b> machine learning, analytics, reporting

# Lakehouse Architecture on AWS

# AWS Lake Formation

